

Indexed Bellman Information Complexity

Yunbei Xu

National University of Singapore

yunbei@nus.edu.sg

Abstract

We develop *indexed Bellman information complexity*, a representation-level framework for reinforcement learning and interactive decision making. The primitive object is an indexed sequential learning with a full environment Ω and admissible nonanticipating policies; priors, posteriors, and ghost histories are used as reference or certificate devices for the fixed frequentist truth. The proof object is a Bellman-sufficient state together with an information index $Y = \chi(\Omega)$, often the optimal decision or value object rather than the full environment, thereby replacing syntactic model realizability with representation-level sufficiency. On the upper-bound side, learning is organized as a dynamic program on the sufficient state with a logarithmic information potential for the index. In fixed-truth analysis this potential is represented by the coordinate log loss $\gamma \log(1/q_t(\chi(\omega^*)))$; in the indexed Algorithmic Information Ratio (AIR) regret identities it gives rise to the log-posterior telescope, and after Bayesian posterior averaging it corresponds to an entropy term. On the lower side, a Bellman-Fano certificate uses the same state and index to compare the indexed information telescope with the ghost-good mass of low-regret reference trajectories. The central matching statement is therefore a conditional Bellman information-risk sandwich when the log-penalized Bellman upper value and the ghost-quantile lower certificate close on the same representation and at the same radius. UCB, E2D/DEC, and AMS/EBO then appear as tractable certificates or relaxations of this same log-potential Bellman program, rather than as separate notions of information complexity.

Contents

1	Introduction	2
1.1	Main contributions.	4
1.2	Organization.	5
2	Indexed sequential learning and Bellman-sufficient representation	5
2.1	Sequential learning through the dynamic programming lens	5
2.2	Latent environment and information index	7
2.3	Bellman-sufficient representation	8
3	Illustrations of Bellman-sufficient representations	9
3.1	Structured Bandits and experimental design	10
3.2	Full environment posterior is fallback sufficient state	11
3.3	Contextual bandits and reinforcement learning	12
4	Indexed information and exact AIR/MAIR identities	14
4.1	Posterior-reference histories and index information	14
4.2	Regret identity for indexed sequential learning	15
4.3	Exact AIR and MAIR regret identities	16

5	Upper bounds: one identity and four algorithm families	20
5.1	Information-potential Bellman programming	22
5.2	UCB families: calibration plus optimism	25
5.3	E2D/DEC: robust one-step offset optimization	26
5.4	AMS/EBO: robust convex belief optimization	26
6	Lower bounds: reference histories and quantile indices	28
6.1	Ghost probability and true good probability	28
6.2	Reference-history quantile theorem	28
6.3	Bellman-Fano lower certificates	30
6.4	DEC as one-step relaxation	31
7	Applications and extentions	33
7.1	Kernel bandits: indexed Bellman DP, GP-UCB, GP-E2D, and AMS/EBO	33
7.2	MAB lower bound via optimal-action index and reference history	39
7.3	Linear bandits lower bound via optimal-action index and reference history	42
7.4	Bellman-rank as indexed representation and lower-bound discussion	46
7.5	Frequentist estimators as effective Bellman-sufficient states	47
8	Conclusion	48

1 Introduction

Information-theoretic minimax theory balances two quantities: the information revealed by an experiment and the logarithmic mass of the set on which an estimator can be successful. In ordinary, noninteractive statistical estimation the experiment is fixed in advance, so this balance can be expressed through local entropy, prior mass, and KL information; this is the perspective behind the information-risk upper and lower bounds of [Zhang \(2006\)](#) and the entropy characterization of minimax rates in [Yang and Barron \(1999\)](#), together with classical Fano and Assouad methods ([Yu, 1997](#)). In interactive decision making the learner designs the experiment while learning. The action law at round t is a statistic of the past observations, so it may itself reveal information about the unknown environment. A fixed-design information calculation is therefore not enough. This adaptivity is exactly what makes a sharp theory difficult.

A powerful response is algorithmic Fano’s method, meaning a Fano-type inequality in which the comparison data set is generated from an algorithm-dependent reference or “ghost” experiment rather than from a fixed nonadaptive sampling scheme. This perspective appears in recent interactive Fano frameworks for general decision making ([Chen et al., 2024](#)) and in algorithmic lower bounds for statistical learning ([Xu, 2026](#)); it is also closely related to the decision-estimation coefficient (DEC) approach of [Foster et al. \(2021\)](#) and its constrained/localized refinements ([Foster et al., 2023](#)). The central obstacle in the algorithmic Fano method is what comes next: collapsing a T -round adaptive experiment into a single one-round action distribution generally requires a localized change-of-measure argument, because the action law under a hard environment need not coincide with the action law under the reference experiment. Such a reduction can obscure or discard the geometry of the posterior trajectory, making the theory more rigid than necessary.

The thesis of this paper is that the right invariant object is not a one-shot coefficient by itself, nor the full environment specification by itself, but a Bellman-sufficient representation together with an index. The full environment Ω remains the object that generates rewards and observations. The state records the part of the history needed to continue the Bellman recursion. The index $Y = \chi(\Omega)$

records the information whose acquisition is charged, such as the optimal action, optimal policy, value function, or full environment. In favorable problems the state is a minimal or nearly minimal sufficient statistic, such as a posterior, a finite marginal posterior, a calibrated confidence state, or a Bellman-error witness representation. When such a state is found, learning becomes a special dynamic program whose generic value-like potential is logarithmic information about the index. This differs from ordinary computational dynamic programming, where the value function and algorithm are usually identified case by case for a particular model; here the same entropy, relative-entropy, and log-partition potentials generate broad algorithmic families.

The basic upper-side identity is fixed-truth and coordinatewise. Fix a truth ω^* , let $y^* = \chi(\omega^*)$, and let q_t be the retained reference marginal on the index. The specialized logarithmic potential

$$\phi_t(s; y^*) = \eta^{-1}[-\log q_t(y^*)]$$

turns the one-step fixed-truth log gain into a Bellman potential drop. Thus the AIR/MAIR bracket is

$$\ell_{\omega^*}(S_t, p_t) + \mathbb{E}_{\omega^*}[\phi_{t+1}(S_{t+1}; y^*) \mid S_t, p_t] - \phi_t(S_t; y^*),$$

and its sum telescopes to a regret bound of order $\eta^{-1} \log(1/q_1(y^*))$ plus the accumulated bracket errors. Bayesian posterior averaging of the same coordinate identity gives the chain rule

$$I_\mu(Y; H_T) = \mathbb{E}_{H'_T} \sum_{t=1}^T \mathcal{I}_\chi(S'_t, p'_t),$$

where S'_t is the reference posterior state and p'_t is the algorithm's action distribution at that reference history.

The same reference history also defines the ghost probability that an algorithm has small target-wise average regret. The lower-bound theorem takes the form

$$\text{if } T\bar{C}_\chi^{\text{Alg}}(\mu) \leq \text{kl}\left(\frac{1}{2}, p_r^{\text{Alg}}(\mu, \chi)\right), \quad \text{then } \mathbb{E}_\mu^{\text{Alg}} \text{Reg}_T \geq \frac{Tr}{2}.$$

Here r is an average-regret threshold, p_r^{Alg} is the posterior-reference ghost probability of average regret at most r , and $\bar{C}_\chi^{\text{Alg}} = T^{-1} \mathbb{E} \sum_t \mathcal{I}_\chi(S'_t, p'_t)$ is average indexed information. The critical balance is

$$T\bar{C}_\chi \asymp \log \frac{1}{p_r}.$$

When the ghost-good mass is roughly $\exp(-d_{\text{eff}})$, the admissible total information is d_{eff} -level. This is the Fano/local-entropy radius, not the constant two-point radius.

Theorem 5.4 turns this balance into a frequentist minimax statement. After a state, index, reference update, and calibration mechanism have been fixed, the upper side is a log-penalized Bellman supersolution with coordinate cost $\gamma \log(1/q_1(\chi(\omega)))$; the lower side is a Bellman-Fano ghost certificate with entropy $-\log p_r$. If these two quantities meet at the same radius, the statistical part of the interactive decision problem is solved at that radius. This is the precise interactive analogue of noninteractive information-risk matching: the fixed experiment is replaced by a controlled Bellman recursion, and local prior mass is replaced by ghost-good mass along the reference history.

After fixing a prior μ , an index χ , and a state map $S_t = \phi_t(H_{t-1}, b_t)$, the two sides are comparable only if they use the same representation. In exact Bayes analyses the state contains the true posterior b_t ; in frequentist or UCB-style analyses it may contain an algorithmic belief or confidence object, but then a separate calibration step must relate that object to the fixed truth. The

upper side supplies either the fixed-truth log-potential certificate of Theorem 5.1, the calibrated frequentist log-penalized Bellman program of Theorem 5.2, or the Bayesian information-potential Bellman program of Theorem 5.3. The lower side supplies a Bellman-Fano certificate: an information supersolution and a ghost-good-mass certificate, as in Theorem 6.3, built on that same chosen state process and the same index. We use “indexed Bellman information complexity” for this compatible state/index accounting. It becomes a matching complexity statement only after the upper and lower certificates are separately verified and shown to agree up to the stated constants or logarithmic factors.

The word “Bellman” in the title is deliberate. The exact value of an interactive problem is a dynamic program, following Bellman’s principle of optimality (Bellman, 1957). An upper bound is a Bellman supersolution, or admissible relaxation in the sense of Rakhlin et al. (2012): it lies above the dynamic value and gives an algorithm; . A lower bound is a Bellman subsolution or, more flexibly, a Bellman-Fano certificate that combines an information supersolution with a ghost-good-mass supersolution. The canonical information-theoretic upper algorithm is therefore the AIR/MAIR information-potential Bellman program itself: it keeps the continuation value on the indexed state and optimizes the Bellman bracket before any one-step relaxation is imposed. UCB, DEC/E2D, and AMS/EBO are useful because they certify, relax, or robustify this Bellman bracket; they should not be confused with the exact Bellman program itself.

1.1 Main contributions.

1. We formulate *indexed sequential learning*, a representation-level framework for reinforcement learning and interactive decision making in which the problem is to choose an admissible policy satisfying Bellman optimality, while a chosen index $Y = \chi(\Omega)$ records the information whose acquisition is charged. The state must be sufficient for losses, observation kernels, posterior predictive laws, posterior updates, and indexed information increments. Several bandit and reinforcement learning examples illustrate the framework with states given by sufficient statistics, model posteriors, frequentist estimators, or other Bellman-sufficient representations.
2. We prove the indexed posterior-reference chain rule and the generic indexed AIR regret identity. AIR is recovered by taking the decision index $Y = \pi^*(\Omega)$, while MAIR is recovered by taking the model or environment index $Y = \Omega$. We then record simplified statements and alternative proofs, show how the gradient bracket unifies different update rules, and explain the role of Danskin’s theorem in the original AIR proof of Xu and Zeevi (2025).
3. We give a unified “one identity, four algorithm families” account centered on a log-penalized Bellman upper theorem and an information-risk sandwich. The first family is the exact AIR/MAIR information-potential Bellman program, which optimizes the dynamic bracket on the chosen state and is the upper object that can match the Bellman-Fano lower certificate. UCB methods (Auer et al., 2002; Abbasi-Yadkori et al., 2011; Srinivas et al., 2010) certify the fixed-truth log-potential bracket through self-normalized calibration and optimism. E2D/DEC methods (Foster et al., 2021, 2023) optimize a one-step offset obtained by replacing the continuation value by a local statistical separation penalty. AMS/EBO methods (Lattimore and György, 2021; Xu and Zeevi, 2025) optimize KL-regularized robust belief relaxations, equivalently dual log-partition relaxations, of the same indexed Bellman bracket.
4. We prove the reference-history quantile lower theorem and Bellman lower certificates. As an indexed posterior-reference specialization of the interactive Fano method, the lower side uses the same stepwise indexed information as the upper side, but compares it to ghost-good

entropy rather than to a one-round coefficient. Through bandit examples we show DEC/E2D is best viewed as a one-step relaxation of indexed Bellman information complexity, not as a universally tight conversion coefficient.

5. We develop several representative examples, including multi-armed bandits, linear bandits, kernel bandits with an indexable action set, and Bellman-rank embeddings. The kernel-bandit example highlights a regime in which the relevant active-action index is finite, even though the unknown RKHS function is infinite-dimensional. In such cases, optimism and its indexed AIR reinterpretation provide a natural mechanism for obtaining sharp complexity bounds, whereas global DEC optimization may be uninformative without appropriate localization or representation restriction.

1.2 Organization.

Section 2 defines indexed sequential learning and Bellman-sufficient representation. Section 3 illustrates the general framework through several examples from bandits and reinforcement learning. Section 4 gives the indexed information telescope and the exact AIR/MAIR regret identities. Section 5 the log-penalized Bellman upper theorem, the information-risk matching principle, and upper algorithmic relaxations including UCB families, E2D/DEC, and AMS/EBO. Section 6 gives the quantile lower theorem, Bellman-Fano certificates, and the status of one-round DEC. Section 7 provide applications to kernel bandits, multi-armed bandits, linear bandits, and Bellman-rank embeddings, and shows how frequentist estimators can serve as effective Bellman-sufficient states.

2 Indexed sequential learning and Bellman-sufficient representation

2.1 Sequential learning through the dynamic programming lens

The primitive problem is sequential learning with general across-time dependence. This formulation is broad enough to cover sequential risk minimization under severe misspecification or adversarial perturbations, as well as computational dynamic programming. Our main emphasis, however, is on the information-theoretic structure of sequential risk minimization. The sufficient-state viewpoint blurs, at the framework level, the boundary between agnostic learning and planning or reasoning; this raises further modeling and philosophical questions that we do not pursue here (Silver and Sutton, 2025). The key information-theoretic content comes instead from the specialization to posterior processes and logarithmic index potentials.

There is an environment class Ω , and performance is required for each fixed truth $\omega^* \in \Omega$. A prior μ is introduced only when we form a Bayes value, a Yao lower bound, a posterior-reference trajectory, or an algorithmic belief. At time t , after the history

$$H_{t-1} = (A_1, O_1, \dots, A_{t-1}, O_{t-1}),$$

the learner chooses a distribution $p_t(\cdot \mid H_{t-1})$ over available actions and samples A_t from it. Under truth ω^* , the observation is drawn from

$$O_t \sim P_{\omega^*, t}(\cdot \mid H_{t-1}, A_t),$$

and the one-step loss or regret is $\ell_t(\omega^*, H_{t-1}, A_t)$. The action A_t may itself be a composite object, such as a within-episode action sequence or a policy to be executed over an episode. The observation

O_t denotes all new information revealed before the next decision. Thus, in a contextual bandit, O_t may include both the realized bandit feedback and the next context; in an episodic MDP, A_t may be the policy or action sequence executed during the episode, and O_t may be the resulting trajectory-level feedback. The framework is general enough to include computational dynamic programming (Bertsekas, 2017; Cormen et al., 2001), although our present focus is on the information-theoretic structure of agnostic sequential learning.

The exact frequentist objective is

$$\mathfrak{R}_T^*(\Omega) := \inf_{\text{Alg}} \sup_{\omega \in \Omega} \mathbb{E}_{\omega}^{\text{Alg}} \sum_{t=1}^T \ell_t(\omega, H_{t-1}, A_t).$$

This is the reinforcement-learning and interactive-decision-making problem in its most general sequential risk-minimization form: choose an admissible nonanticipating policy whose Bellman value is small for every environment.

The full history is always a valid Bellman state. As illustrated in Section 3.2, the full environment posterior also provides a fallback sufficient state for episodic model learning. A compressed state is legitimate only when it preserves the objects needed to close the Bellman recursion. Let S_t be such a state, with action set $\mathcal{A}(S_t)$, transition/update map $S_{t+1} = \tau_t(S_t, A_t, O_t)$, and state-dependent kernels P_{ω, S_t, A_t} . If a Bayesian mixture law with posterior $b_t = \mathcal{L}_{\mu}(\Omega \mid H_{t-1})$ is used as a reference, the exact Bayes-regret dynamic program on an exact sufficient state is

$$V_{T+1}(s) = 0, \quad V_t(s) = \inf_{p \in \Delta(\mathcal{A}(s))} \left\{ \ell(s, p) + \mathbb{E}_{a \sim p, o \sim P_{s,a}} V_{t+1}(\tau(s, a, o)) \right\}, \quad (1)$$

where

$$\ell(s, p) = \mathbb{E}_{a \sim p} \int \ell_{\omega}(s, a) b_s(d\omega), \quad P_{s,a} = \int P_{\omega, s, a} b_s(d\omega).$$

Equation (1) is not the definition of the frequentist problem; it is the Bayes/reference Bellman recursion induced by the chosen representation. Minimax lower bounds are connected to it through Yao's principle (Yao, 1977),

$$\mathfrak{R}_T^*(\Omega) \geq \sup_{\mu \in \Delta(\Omega)} \inf_{\text{Alg}} \mathbb{E}_{\Omega \sim \mu} \mathbb{E}_{\Omega}^{\text{Alg}} L_{\Omega}(H_T).$$

The information-complexity layer asks whether the same frequentist problem can be controlled after compressing the history to a Bellman-sufficient state and charging information only about an index $Y = \chi(\Omega)$. The state should be as small as the proof permits—minimal when a minimal sufficient statistic exists, otherwise simply sufficient. The index is the final coordinate whose logarithmic mass or ghost entropy controls the statistical price of learning. This is the sense in which learning is a special dynamic program: unlike ordinary computational dynamic programming, where the value function is problem-specific, the statistical potential is generically logarithmic in the index belief.

Definition 2.1 (Bellman supersolution and subsolution). *Let T_t be the Bellman operator in (1). A sequence $\overline{V}_t$ is a regret supersolution if*

$$\overline{V}_{T+1} \geq 0, \quad \overline{V}_t(s) \geq (\mathsf{T}_t \overline{V}_{t+1})(s) \quad \text{for all } t, s.$$

A sequence $\underline{V}_t$ is a regret subsolution if

$$\underline{V}_{T+1} \leq 0, \quad \underline{V}_t(s) \leq (\mathsf{T}_t \underline{V}_{t+1})(s) \quad \text{for all } t, s.$$

Proposition 2.2 (Bellman sandwich). *If $\bar{V}$ is a supersolution, then any policy that attains the right-hand side of the supersolution inequality up to errors ε_t has Bayes regret at most $\bar{V}_1(s_1) + \sum_t \varepsilon_t$. If $\underline{V}$ is a subsolution, then every algorithm has Bayes regret at least $\underline{V}_1(s_1)$. Consequently,*

$$\underline{V}_1(s_1) \leq V_1(s_1) \leq \bar{V}_1(s_1).$$

Proof. For the upper statement, if the selected decision kernel satisfies

$$\ell(s, p_t) + \mathbb{E}\bar{V}_{t+1}(S_{t+1}) \leq \bar{V}_t(s) + \varepsilon_t,$$

then summing over t telescopes the potential and yields the claimed bound. The lower statement is the same induction with the inequality reversed: for any decision kernel p_t ,

$$\ell(s, p_t) + \mathbb{E}\underline{V}_{t+1}(S_{t+1}) \geq \underline{V}_t(s).$$

Summing gives expected loss at least $\underline{V}_1(s_1) - \mathbb{E}\underline{V}_{T+1}(S_{T+1}) \geq \underline{V}_1(s_1)$. $\square$

This proposition is the organizing principle of the paper. The upper theory constructs supersolutions by paying the Bellman bracket with information. The lower theory constructs subsolutions indirectly, by showing that any algorithm with too little information has too small a ghost-good probability to achieve low regret. The representation required by both sides is exactly the representation in which these Bellman backups are well-defined.

2.2 Latent environment and information index

The latent environment space is denoted by Ω . An information index is a measurable map

$$Y = \chi(\Omega) \in \mathcal{Y}.$$

The model-index choice $Y = \Omega$ gives MAIR. The action-index choice $Y = A^*(\Omega)$, where A^* is an optimal action under the environment, gives AIR. The loss may depend on more than Y . For index-based lower bounds we therefore use the conditional loss

$$\ell_t^\chi(y, H_{t-1}, a) := \mathbb{E}_\mu[\ell_t(\Omega, H_{t-1}, a) \mid \chi(\Omega) = y, H_{t-1}], \quad (2)$$

assuming the relevant regular conditional distributions exist. Define

$$L_\omega(H_T) = \sum_{t=1}^T \ell_t(\omega, H_{t-1}, A_t), \quad \bar{L}_\omega(H_T) = T^{-1} L_\omega(H_T),$$

and

$$L_\chi(y, H_T) = \sum_{t=1}^T \ell_t^\chi(y, H_{t-1}, A_t), \quad \bar{L}_\chi(y, H_T) = T^{-1} L_\chi(y, H_T).$$

Under the true Bayesian experiment,

$$\mathbb{E}_{\Omega \sim \mu, H_T \sim \mathbb{P}_\Omega^{\text{Alg}}} L_\Omega(H_T) = \mathbb{E}_{Y, H_T} L_\chi(Y, H_T). \quad (3)$$

The choice of information index is fairly flexible. If a pathwise, model-wise event is needed, one may take $Y = \Omega$. If the objective is decision identification, one may take $Y = A^*(\Omega)$ and pay only for information about the optimal decision. If a value-function class $\mathcal{F}$ provides Bellman sufficiency under the chosen state representation, one may take $Y = f^*(\Omega)$ and pay only for information about the optimal value function.

2.3 Bellman-sufficient representation

We work with a general sequential learning problem of Section 2.1 with an information index of Section 2.2. For minimax or frequentist risk, the environment $\omega^* \in \Omega$ is fixed. A prior or reference law μ on Ω is auxiliary: it may define an algorithmic belief, a Bayes value, a posterior-reference trajectory, or a Yao/Fano certificate, but it is not itself the frequentist truth.

Given an indexed Bellman-sufficient representation (formalized in Definition 2.3 below), the general frequentist sequential learning problem of Section 2.1 can be written in the following state-sufficient Markov form: at round t , the learner observes a state S_t , chooses $A_t \in \mathcal{A}(S_t)$, observes

$$O_t \sim P_{\omega^*,t}(\cdot \mid S_t, A_t),$$

incurs a nonnegative regret or surrogate loss

$$\ell_{\omega^*}(S_t, A_t) \in [0, L],$$

and updates

$$S_{t+1} = \tau_t(S_t, A_t, O_t).$$

Throughout the lower-bound analyses, we assume that the per-round loss or regret is nonnegative. For upper-bound analyses under misspecification or adversarial losses, one may instead work with a nonnegative surrogate loss or surrogate regret.

The Bellman-sufficient state is the operational substitute for carrying the full environment label through the analysis. One need not estimate all of ω^* if a smaller representation closes the prediction, regret, update, and indexed-information recursions. We formalize the definition of an indexed Bellman-sufficient representation below.

Definition 2.3 (Indexed Bellman-sufficient representation). *Fix a prior μ and an index map χ . Let $b_t = \mathcal{L}_\mu(\Omega \mid H_{t-1})$ be the posterior under the Bayesian mixture law. A representation is a process*

$$S_t = \phi_t(H_{t-1}, b_t).$$

The time index may be included in S_t ; equivalently, the kernels and updates may be time-inhomogeneous, and we suppress the subscript t when it is part of the state. It is exact Bellman-sufficient if, for every state s , action $a \in \mathcal{A}(s)$, and environment ω , there exist objects

$$P_{\omega,s,a} \in \Delta(\mathcal{O}), \quad \ell_\omega(s, a) \in [0, L], \quad \tau(s, a, o),$$

such that the following four conditions hold: three state-sufficient Markov conditions and one separate index-sufficiency condition.

- (i) Predictive sufficiency: $P_{\omega,t}(\cdot \mid H_{t-1}, a) = P_{\omega,S_t,a}$.
- (ii) Loss sufficiency: $\ell_t(\omega, H_{t-1}, a) = \ell_\omega(S_t, a)$.
- (iii) Update sufficiency: under the posterior-reference law, $S_{t+1} = \tau(S_t, A_t, O_t)$.
- (iv) Index sufficiency: for $q_s = \mathcal{L}(\chi(\Omega) \mid S_t = s)$ and $b_s^y = \mathcal{L}(\Omega \mid S_t = s, \chi(\Omega) = y)$ defined for q_s -almost every y , the conditional predictive laws

$$P_{s,a}^y := \int P_{\omega,s,a} b_s^y(d\omega)$$

and conditional losses

$$\ell_y^x(s, a) := \int \ell_\omega(s, a) b_s^y(d\omega)$$

are determined by s rather than the full history.

Note that the index sufficiency condition is necessary only for indexed information accounting.

The exact Bayesian information-gain formula in Definition 4.1 requires one additional condition. Namely, in the Fano-type lower-bound analysis and in the Bayesian upper-bound specialization, we require posterior predictive sufficiency: for each state-action pair (s, a) , the posterior predictive laws $P_{s,a} := \int P_{\omega,s,a} b_s(d\omega)$ where $b_s := \mathcal{L}_\mu(\Omega \mid S_t = s)$ are determined by s rather than the full history. This condition is not required for the frequentist upper-bound analyses.

The Bellman-sufficient state should be understood as a controlled sufficient statistic for both information accounting and algorithm design. Classical sufficiency starts from a statistic T , a random variable constructed from the data X alone and not depending on the unknown parameter θ . The statistic T is sufficient for the model $\{P_\theta : \theta \in \Theta\}$ if, after conditioning on T , the remaining conditional distribution of the data X is independent of θ (Fisher, 1922). Here the analogue is dynamic and decision-theoretic: a state is sufficient if it preserves everything needed for future prediction, loss, information, and updating under admissible controls.

The full history is always an admissible state. As illustrated in Section 3.2, for episodic model learning problems, the full environment posterior also provides a fallback sufficient state, offering a practical bridge between the full-history state and more problem-specific compressed states. Compressed states, such as frequentist estimators, Bayesian posteriors, or representation-dependent factorizations, are admissible when they preserve the information needed for prediction, updating, loss evaluation, and information accounting. Approximate representations may also be used, provided that their approximation errors are explicitly incorporated into the Bellman inequalities.

The index $Y = \chi(\Omega)$ is not required to identify the whole environment. It is the coordinate whose posterior log mass is charged by the information potential and whose ghost-good mass appears in the Fano certificate.

The formalism of information indices is aligned with the representation-level philosophy underlying low Bellman rank for contextual decision processes (Jiang et al., 2017). Contextual decision processes with low Bellman rank provide an important example of structure at the level of Bellman errors, rather than at the level of posterior beliefs. Bellman rank is not a posterior over models; it is a factorization of Bellman errors relative to a chosen function class, roll-in policy family, and sampling scheme. Thus, bounded Bellman rank is not by itself an information-theoretic complexity. It becomes an indexed Bellman sequential-learning representation only after one specifies the latent environment, the information index, the observable trajectory law, the loss, and the Bellman-error witness class whose estimation cost is being charged. Under this additional structure, a Bellman-rank representation can be viewed as a compressed Bellman-sufficient state, analogous to but distinct from a Bayesian posterior state.

3 Illustrations of Bellman-sufficient representations

This section provides verification templates for Bellman-sufficient representations. As these examples illustrate, structured bandit problems fit naturally within the Bellman-sufficiency framework. The full environment posterior provides a literal fallback sufficient statistic, offering a practical intermediate representation between the full-history state and more problem-specific compressed states. Finally, contextual bandits and linear MDPs can be incorporated in two clean ways: through

the exact posterior/reference-law formulation in Definition 2.3 and Section 3.2, or through the calibrated fixed-truth formulation in Definition 7.10.

3.1 Structured Bandits and experimental design

Structured bandit problems offer particularly transparent examples of the Bellman-sufficiency framework for indexed sequential learning, which is guided by the classical intuition of sufficient statistics.

Example 3.1 (Classical sufficient statistics as one-step Bellman states). Consider a nonadaptive dominated experiment with observations $X_1, \dots, X_n$ from P_θ , and no control. Let $T_n = T(X_1, \dots, X_n)$ satisfy the Fisher–Neyman factorization. If the loss and index depend on θ only through quantities whose posterior distribution is determined by T_n , then the process $S_t = (t, T_{t-1})$ is Bellman sufficient. Predictive sufficiency is the usual conditional factorization; update sufficiency is $T_t = u(T_{t-1}, X_t)$; posterior predictive and index sufficiency follow because the posterior or conditional reference law depends on the history only through T_t . In a regular exponential family,

$$p_\theta(x) = h(x) \exp\{\eta(\theta)^\top T(x) - A(\theta)\},$$

with conjugate or otherwise statistic-measurable reference law, the cumulative statistic $\sum_{i < t} T(X_i)$ gives the corresponding state. This is the static prototype for all examples below.

Example 3.2 (Finite stochastic multi-armed bandits). Let $\Omega \subseteq \Theta^K$ and let arm a produce an observation from P_{θ_a} . For unit-variance Gaussian arms, $P_{\theta_a} = N(\theta_a, 1)$. A reference Gaussian prior with independent coordinates gives the posterior state

$$S_t = (t, (n_{t,a}, \bar{X}_{t,a}, v_{t,a})_{a=1}^K),$$

where $n_{t,a}$ is the number of previous pulls of arm a , $\bar{X}_{t,a}$ is the empirical mean when $n_{t,a} > 0$, and $v_{t,a}$ is the posterior variance. Equivalently one may store the posterior hyperparameters. For a fixed truth $\omega = \theta$,

$$P_{\omega,s,a} = N(\theta_a, 1), \quad \ell_\omega(s, a) = \max_b \theta_b - \theta_a,$$

and the update of $(n, \bar{X}, v)$ after observing arm a is a function of (s, a, o) . The posterior predictive law is $N(m_{t,a}, 1 + v_{t,a})$, and for the action index $Y = A^*(\theta)$ the conditional predictive law $P_{s,a}^y$ and conditional loss $\ell_y^x(s, a)$ are obtained by integrating the same Gaussian posterior over the event $A^*(\theta) = y$. Hence Definition 2.3 is satisfied. The empirical mean alone is not sufficient: the count or posterior variance is needed for prediction, confidence, and information gain.

Example 3.3 (Finite-dimensional Gaussian linear bandits). Let $\Omega \subseteq \mathbb{R}^d$, actions satisfy $\|a\|_2 \leq 1$, and

$$O_t = \langle \theta, A_t \rangle + \xi_t, \quad \xi_t \sim N(0, 1).$$

Under a Gaussian reference prior, the exact Bayesian state is $S_t = (t, m_t, \Sigma_t)$, where

$$\Sigma_{t+1}^{-1} = \Sigma_t^{-1} + A_t A_t^\top, \quad m_{t+1} = \Sigma_{t+1}^{-1} \{\Sigma_t^{-1} m_t + A_t O_t\}.$$

For fixed θ , $P_{\theta,s,a} = N(\langle \theta, a \rangle, 1)$ and the regret loss is $\sup_{\|b\|_2 \leq 1} \langle \theta, b \rangle - \langle \theta, a \rangle$. The posterior predictive law is $N(\langle m_t, a \rangle, 1 + a^\top \Sigma_t a)$, and conditioning the Gaussian posterior on an index event such as $\chi(\theta) = \theta$, $\chi(\theta) = \theta / \|\theta\|_2$, or $\chi(\theta) = A^*(\theta)$ determines $P_{s,a}^y$ and $\ell_y^x(s, a)$. Therefore the Gaussian mean–covariance pair is an exact Bellman-sufficient state. A frequentist UCB state $(\hat{\theta}_t, V_t, \beta_t)$ is different: it can certify a fixed-truth bracket on the calibration event $\|\hat{\theta}_t - \theta^*\|_{V_t} \leq \beta_t$, but it is not an exact posterior-reference state unless a reference belief or confidence-to-belief map is added.

Example 3.4 (Finite active marginals in kernel bandits). Let $f^\star$ belong to an RKHS on a possibly infinite domain, but suppose the learner is evaluated on a finite active set $\mathcal{X} = \{x_1, \dots, x_n\}$. Under the GP reference law on the active reward vector $F = (f(x_1), \dots, f(x_n))$, the finite marginal posterior $(m_t, \Sigma_t) \in \mathbb{R}^n \times \mathbb{R}^{n \times n}$ is Bellman sufficient for the finite experiment. For fixed $f^\star$, pulling x_i has law $N(f^\star(x_i), \lambda)$ in the Gaussian reference calculation; the posterior predictive law is $N(m_{t,i}, \lambda + \Sigma_{t,ii})$; and the usual rank-one Gaussian update is a function of (m_t, Σ_t, i, O_t) . If the index is the optimal active action $Y = \arg \max_i F_i$, then index sufficiency follows by conditioning the finite Gaussian posterior on $\arg \max_i F_i = y$. The infinite-dimensional RKHS function is still the frequentist truth; the sufficient state is finite only because the decisions and observations use the active marginal.

3.2 Full environment posterior is fallback sufficient state

For indexed sequential learning, a natural Bellman-sufficient representation is obtained by retaining the full environment posterior. For suitably defined Markovian model classes over the state space, which include standard episodic model-based reinforcement learning settings, this posterior serves as a fallback sufficient state.

Example 3.5 (Full environment posterior as fallback sufficient state). *Consider a realizable model class in which each environment $\omega \in \Omega$ specifies, for every retained physical state or context z , action a , and stage t , an observation kernel $P_{\omega,t}(\cdot | z, a)$, a loss $\ell_{\omega,t}(z, a)$, and the physical-state transition rule. Let Z_t denote the physical state, context, stage, or episode coordinate needed to make these kernels Markov. Fix a reference prior μ and set*

$$\Pi_t = \mathcal{L}_\mu(\Omega | H_{t-1}), \quad S_t = (Z_t, \Pi_t).$$

Assume the observation laws are dominated on the relevant support. Then S_t gives an exact Bellman-sufficient state. Indeed, for each fixed truth ω ,

$$P_{\omega, S_t, a} = P_{\omega, t}(\cdot | Z_t, a), \quad \ell_\omega(S_t, a) = \ell_{\omega, t}(Z_t, a),$$

and the next retained state is obtained by updating the physical coordinate and the posterior

$$\Pi_{t+1}(d\omega) = \frac{p_{\omega, t}(O_t | Z_t, A_t) \Pi_t(d\omega)}{\int p_{\omega', t}(O_t | Z_t, A_t) \Pi_t(d\omega')}.$$

For an index $\chi : \Omega \rightarrow \mathcal{Y}$, the indexed marginal and conditional predictive components are

$$q_t = \chi_\# \Pi_t, \quad \Pi_t^y = \mathcal{L}_{\Pi_t}(\Omega | \chi(\Omega) = y),$$

$$P_{S_t, a}^y = \int P_{\omega, t}(\cdot | Z_t, a) \Pi_t^y(d\omega), \quad \ell_y^x(S_t, a) = \int \ell_{\omega, t}(Z_t, a) \Pi_t^y(d\omega).$$

Thus prediction, loss evaluation, posterior/reference updating, and indexed information accounting all close on S_t . This verifies Definition 2.3.

In this sense, the indexed sequential learning framework with Bellman-sufficient representations can be viewed as containing the Decision Making with Structured Observations (DMSO) framework (Foster et al., 2021) as a special case, when DMSO is interpreted as an independent episodic model learning formalism and the full environment posterior is taken as the Bellman-sufficient state. At the same time, indexed sequential learning also accommodates Markovian dependence across states and time, as in computational dynamic programming. For well-specified environments, however, the more important point is not this formal inclusion, which is essentially straightforward, but the possibility of dynamic-programming algorithms over sufficient states and tighter lower-bound certificates based on the indexed representation, even within the DMSO setting.

3.3 Contextual bandits and reinforcement learning

There are two clean ways to include contextual bandits and linear MDPs. The first is the exact belief-state lift: augment the physical or estimator state by a coherent reference belief, posterior, or indexed marginal whose update is closed. This verifies the exact clauses of Definition 2.3 for the reference experiment, and gives a fixed-truth upper bound only after calibration. The second is the calibrated-certificate route of Definition 7.10 (deferred to Section 7.5): keep the estimator, confidence set, or relaxation state used by the frequentist algorithm, and prove directly that it upper-bounds the fixed-truth Bellman bracket. The latter formulation naturally accommodates IPW estimators, least-squares states, and adaptively optimized algorithmic posteriors.

Example 3.6 (Adversarial contextual bandits and IPW estimator states). Consider an oblivious adversarial contextual bandit with finite action set $[K]$, policy class Π , contexts x_t , and losses $\ell_t(a) \in [0, 1]$. The fixed truth is the whole sequence

$$\omega^* = (x_t, \ell_t(1), \dots, \ell_t(K))_{t \leq T}.$$

At the decision time for round t , the current context x_t is part of the state. The observation after acting may be written as $O_t = (\ell_t(A_t), x_{t+1})$, with the next context omitted at $t = T$. This convention makes the transition $S_{t+1} = \tau_t(S_t, A_t, O_t)$ closed. Given an action distribution $p_t(\cdot | x_t)$, the inverse-probability-weighted policy-loss increment

$$\hat{\ell}_t(\pi) = \frac{\ell_t(A_t) \mathbb{1}\{\pi(x_t) = A_t\}}{p_t(A_t | x_t)}$$

satisfies $\mathbb{E}[\hat{\ell}_t(\pi) | H_{t-1}, x_t, \omega^*] = \ell_t(\pi(x_t))$ whenever the denominator is positive. Thus the cumulative vector $\hat{L}_t(\pi) = \sum_{s < t} \hat{\ell}_s(\pi)$ is a valid loss-estimation coordinate for relaxation-based contextual-bandit algorithms such as BISTRO (Rakhlin and Sridharan, 2016) and for mirror-descent information-ratio analyses (Lattimore and György, 2021). For an infinite policy class, $\hat{L}_t$ should be understood as the exact functional $\pi \mapsto \hat{L}_t(\pi)$, or as an oracle representation sufficient to evaluate the relaxation; a finite-dimensional sketch is sufficient only when the relaxation can be recovered from that sketch up to charged error.

It is not, however, an exact indexed Bellman state by itself. There are two rigorous augmentations. In the exact belief-state lift, one uses

$$S_t = (t, x_t, \hat{L}_t, p_t, b_t),$$

where b_t is a reference law over the remaining context/loss sequence or over a parametrized environment class. Then the reference predictive law, the conditional index predictive laws, and the indexed marginal $q_t = \chi_{\#} b_t$ are determined by S_t , so Definition 2.3 applies to the reference experiment. In the frequentist certificate route, one uses

$$S_t = (t, x_t, \hat{L}_t, p_t, C_t),$$

where C_t is the relaxation, confidence, or domination object proving that the IPW loss estimates upper-bound the fixed-truth Bellman bracket. This could verify the route of Definition 7.10 (see Section 7.5 for details), though not the straightforward posterior clauses of Definition 2.3. In both routes the IPW vector is an update coordinate; it is not the whole state. Unbiasedness identifies conditional expected losses, while Bellman sufficiency additionally requires closed prediction or certified domination, closed updating, and compatible index accounting.

Example 3.7 (Linear MDPs and least-squares Bellman states). Consider an episodic finite-horizon MDP. The fixed truth ω^* consists of the reward and transition kernels at every stage. At stage h , the physical state X_h is part of the Bellman state, the action is A_h , and the observation is $O_h = (R_h, X_{h+1})$. Suppose that there are known features $\phi_h(x, a) \in \mathbb{R}^d$ such that rewards and transition expectations are linear in ϕ_h . For a candidate next-value function V , the Bellman target has the form

$$r_h(x, a) + \mathbb{E}[V(X_{h+1}) \mid x, a] = \phi_h(x, a)^\top w_h(V).$$

The least-squares value-iteration state used in optimistic linear-MDP algorithms (Jin et al., 2019) is Bellman sufficient for a specified value class only when it contains enough information to compute $w_h(V)$ for every future candidate V in that class. A concise exact data state for stage h is

$$\Lambda_{h,t} = \lambda I + \sum_{i < t} \phi_{h,i} \phi_{h,i}^\top, \quad u_{h,t} = \sum_{i < t} \phi_{h,i} R_{h,i}, \quad \mathcal{M}_{h,t} : V \mapsto \sum_{i < t} \phi_{h,i} V(X_{h+1,i}),$$

so that

$$\hat{w}_{h,t}(V) = \Lambda_{h,t}^{-1} \{u_{h,t} + \mathcal{M}_{h,t} V\}.$$

The full Bellman state also includes (h, X_h) , the value class or admissible relaxation being optimized, and either a reference belief over the linear MDP parameters or a calibrated confidence set for the Bellman targets. If the candidate value class is finite-dimensional, $\mathcal{M}_{h,t}$ can be stored as finitely many sufficient coordinates; otherwise the full transition sample or an exact evaluation oracle is part of the state. Storing only a single fitted vector $\hat{w}_h$ is sufficient for one fixed backup, but not for the Bellman recursion if later backups may use different value functions.

Thus linear-MDP least squares fits the framework in the same two ways as contextual IPW. With a coherent algorithmic posterior over the reward and transition parameters, the augmented state verifies Definition 2.3 for the reference experiment and may exploit structure beyond a worst-case confidence set. With the usual optimistic LSVI analysis, the state could verify the route of Definition 7.10 (see Section 7.5 for details): uniform self-normalized inequalities certify that, for every fixed truth in the linear-MDP class, the least-squares bonus upper-bounds the Bellman target error for all candidate values in the maintained class. Least squares is therefore not sufficient because it is efficient; it is sufficient only when the stored operator and calibration event close all future Bellman backups.

Example 3.8 (Bellman rank and bilinear Bellman-error representations). Low Bellman rank provides a finite-dimensional factorization of Bellman errors, not automatically a full state. In a contextual decision process, suppose that for the relevant hypothesis, value-function, and roll-in classes the Bellman error can be written as

$$\mathcal{E}(h, f, \pi) = \langle u(h, \pi), v(f) \rangle, \quad u(h, \pi), v(f) \in \mathbb{R}^d,$$

as in low-Bellman-rank and bilinear-class analyses (Jiang et al., 2017; Du et al., 2021). These coordinates are sufficient for the scalar Bellman-error tests covered by the factorization. They verify Definition 2.3 only after augmentation by the planning variables and reference or confidence object that determine the actual observation kernels, losses, posterior/reference updates, and index-conditional predictive laws. Hence the valid state is schematically

$$S_t = (\text{roll-in/planning state, witness coordinates, value-class coordinates, } b_t \text{ or } C_t),$$

not the rank number d alone. When these additional sufficiency conditions hold, the information potential and ghost-good recursions may be run on the witness-coordinate state; when they do not, Bellman rank remains an upper-bound factorization rather than an indexed information-complexity representation.

The examples have two consequences for later applications. First, exact Bayesian posterior states, Gaussian marginal states, finite conjugate statistics, and coherent algorithmic-belief lifts directly satisfy Definition 2.3 or their reference experiments. In particular, Section 3.2 shows full environment posterior is a fallback sufficient state for episodic model learning. Second, estimator-based frequentist states such as IPW, matrix inverse IPW, UCB confidence ellipsoids, and least-squares Bellman fits are certificate states: they become Bellman sufficient only together with the calibration, reference belief, or domination inequalities that make the fixed-truth Bellman bracket valid.

4 Indexed information and exact AIR/MAIR identities

4.1 Posterior-reference histories and index information

The mutual-information identities in this subsection use the exact posterior-reference version of Definition 2.3. Fixed-truth log-gain regret identity in the next subsection should not be substituted into the posterior chain rule unless they define a coherent reference law.

For a prior μ and algorithm Alg, define the Bayesian posterior-reference trajectory law

$$\bar{\mathbb{P}}_\mu^{\text{Alg}} := \int \mathbb{P}_\omega^{\text{Alg}} \mu(d\omega).$$

A reference history is denoted

$$H'_T \sim \bar{\mathbb{P}}_\mu^{\text{Alg}}, \quad S'_t = \phi_t(H'_{t-1}, b'_t), \quad b'_t = \mathcal{L}_\mu(\Omega \mid H'_{t-1}),$$

and the algorithm's reference decision kernel is

$$p'_t = p_t(\cdot \mid H'_{t-1}).$$

Definition 4.1 (Indexed information gain). *At state s , let q_s be the posterior law of $Y = \chi(\Omega)$ and let $P_{s,a}^y$ and $P_{s,a}$ be the conditional and marginal predictive laws from Definition 2.3. For $p \in \Delta(\mathcal{A}(s))$, define*

$$\mathcal{I}_\chi(s, p) := \mathbb{E}_{a \sim p} \int D_{\text{KL}}(P_{s,a}^y \parallel P_{s,a}) q_s(dy). \quad (4)$$

When $Y = \Omega$, this is model-index information. When $Y = \pi^(\Omega)$, this is action-index information. In the action-index case, the belief is still a belief over environments; only the information target is the optimal action.*

Define the cumulative and average indexed information of an algorithm by

$$C_\chi^{\text{Alg}}(\mu) := \mathbb{E}_{H'_T \sim \bar{\mathbb{P}}_\mu^{\text{Alg}}} \sum_{t=1}^T \mathcal{I}_\chi(S'_t, p'_t), \quad \bar{C}_\chi^{\text{Alg}}(\mu) := \frac{1}{T} C_\chi^{\text{Alg}}(\mu). \quad (5)$$

Proposition 4.2 (Reference-history indexed chain rule). *For every prior μ , index map χ , and algorithm Alg,*

$$I_\mu(Y; H_T) = \mathbb{E}_{H'_T \sim \bar{\mathbb{P}}_\mu^{\text{Alg}}} \sum_{t=1}^T \mathcal{I}_\chi(S'_t, p'_t) = T \bar{C}_\chi^{\text{Alg}}(\mu). \quad (6)$$

Proof. The mutual-information chain rule gives

$$I_\mu(Y; H_T) = \sum_{t=1}^T I_\mu(Y; A_t, O_t \mid H_{t-1}).$$

Given H_{t-1} , the action A_t is sampled by the algorithm using only randomization independent of the environment and is conditionally independent of Y . Therefore

$$I_\mu(Y; A_t, O_t \mid H_{t-1}) = I_\mu(Y; O_t \mid H_{t-1}, A_t).$$

Conditioning further on $A_t = a$, the law of O_t given $Y = y$ is $P_{S_t, a}^y$, while the posterior predictive law is $P_{S_t, a}$. Hence

$$I_\mu(Y; O_t \mid H_{t-1}, A_t = a) = \int D_{\text{KL}}(P_{S_t, a}^y \parallel P_{S_t, a}) q_{S_t}(dy).$$

Averaging over $a \sim p_t(\cdot \mid H_{t-1})$ and over the marginal law of the history under the Bayesian mixture law proves the result. Writing this marginal law as the reference history law gives the displayed expression. $\square$

4.2 Regret identity for indexed sequential learning

The information gain in (4) is posterior-averaged. The regret identity used in frequentist upper bounds is slightly finer: it fixes the true environment and tracks the posterior log ratio of its index.

Fix a state s , an environment ω , and let $y = \chi(\omega)$. Assume domination and $q_s(y) > 0$. Define

$$\Lambda_\chi(s, a, o; y) := \log \frac{dP_{s, a}^y}{dP_{s, a}}(o).$$

For $p \in \Delta(\mathcal{A})$, define the fixed-truth index log gain

$$J_\chi(s, p; \omega) := \mathbb{E}_{a \sim p, O \sim P_{\omega, s, a}} \Lambda_\chi(s, a, O; \chi(\omega)). \quad (7)$$

Equivalently,

$$J_\chi(s, p; \omega) = \mathbb{E}_{a \sim p} \left[D_{\text{KL}}(P_{\omega, s, a} \parallel P_{s, a}) - D_{\text{KL}}(P_{\omega, s, a} \parallel P_{s, a}^{\chi(\omega)}) \right]. \quad (8)$$

Thus the fixed-truth log gain is a pointwise-in- ω quantity, whereas $\mathcal{I}_\chi(s, p)$ is a posterior-averaged index information. More explicitly, for each index value y with positive posterior mass,

$$\mathbb{E}_{\omega \sim \mu_s(\cdot \mid \chi(\omega)=y)} J_\chi(s, p; \omega) = \mathbb{E}_{a \sim p} D_{\text{KL}}(P_{s, a}^y \parallel P_{s, a}),$$

and averaging this display over $y \sim q_s$ gives $\mathcal{I}_\chi(s, p)$. Besides this conditional posterior averaging, a special model-index $\chi(\omega) = \omega$ case gives $P_{\omega, s, a} = P_{s, a}^{\chi(\omega)}$ and

$$J^{\text{MAIR}}(s, p; \omega) = \mathbb{E}_{a \sim p} D_{\text{KL}}(P_{\omega, s, a} \parallel P_{s, a}),$$

In this sense, $J_\chi(s, p; \omega)$ serves as a common fixed-truth log-gain object for stating both algorithmic information gain and fixed-truth estimation error.

For $\eta > 0$, define the fixed-truth indexed bracket

$$B_{\chi, \eta}(s, p; \omega) := \ell_\omega(s, p) - \frac{1}{\eta} J_\chi(s, p; \omega), \quad \ell_\omega(s, p) := \mathbb{E}_{a \sim p} \ell_\omega(s, a). \quad (9)$$

Theorem 4.3 (Regret identity for indexed sequential learning). *Suppose the algorithm chooses an H_{t-1} -measurable decision distribution $p_t \in \Delta(\mathcal{A}(S_t))$ and then samples $A_t \sim p_t$. Fix ω^* and $y^* = \chi(\omega^*)$. If the reference index posterior is updated by Bayes' rule after the realized action and observation, then, conditionally on H_{t-1} ,*

$$\ell_{\omega^*}(S_t, p_t) = B_{\chi, \eta}(S_t, p_t; \omega^*) + \frac{1}{\eta} \mathbb{E}_{A_t \sim p_t, O_t \sim P_{\omega^*, S_t, A_t}} \left[\log \frac{q_{t+1}(y^*)}{q_t(y^*)} \middle| H_{t-1} \right]. \quad (10)$$

Consequently,

$$\mathbb{E}_{\omega^*}^{\text{Alg}} L_{\omega^*}(H_T) = \frac{1}{\eta} \mathbb{E}_{\omega^*}^{\text{Alg}} \log \frac{q_{T+1}(y^*)}{q_1(y^*)} + \mathbb{E}_{\omega^*}^{\text{Alg}} \sum_{t=1}^T B_{\chi, \eta}(S_t, p_t; \omega^*). \quad (11)$$

The expectation is over the action randomization, the observation process, and any exogenous randomness of the algorithm. Thus the identity is an equality in expectation; a pathwise equality is available only after replacing the randomized loss by its realized counterpart and keeping the corresponding martingale difference.

Proof. For the realized action $A_t = a$, Bayes' rule for the reference index posterior gives

$$\log \frac{q_{t+1}(y^*)}{q_t(y^*)} = \log \frac{dP_{S_t, a}^{y^*}}{dP_{S_t, a}}(O_t).$$

Taking conditional expectation under the fixed truth ω^* , first over $O_t \sim P_{\omega^*, S_t, a}$ and then over $a \sim p_t$, yields

$$\mathbb{E}_{A_t \sim p_t, O_t \sim P_{\omega^*, S_t, A_t}} \left[\log \frac{q_{t+1}(y^*)}{q_t(y^*)} \middle| H_{t-1} \right] = J_{\chi}(S_t, p_t; \omega^*).$$

Rearranging the definition of $B_{\chi, \eta}$ proves the one-step identity. Summing over t and taking expectation telescopes the posterior log ratio. The equality between expected realized loss and $\mathbb{E} \sum_t \ell_{\omega^*}(S_t, p_t)$ follows from conditional expectation over $A_t \sim p_t$. $\square$

The following corollary gives the deterministic specialization of Theorem 4.3, which will be useful in the subsequent UCB analyses. It follows directly from conditional determinism and the conditional-expectation step in (10). In this case the outer expectation over randomized decision laws drops out, leaving only the posterior-ratio telescope and the fixed-truth gradient bracket.

Corollary 4.4 (Deterministic-action indexed regret identity). *Suppose the algorithm is deterministic after conditioning on H_{t-1} , so $p_t = \delta_{a_t}$. Fix ω^* and $y^* = \chi(\omega^*)$. If the reference index posterior is updated by Bayes' rule, then*

$$\mathbb{E}_{\omega^*}^{\text{Alg}} L_{\omega^*}(H_T) = \frac{1}{\eta} \mathbb{E}_{\omega^*}^{\text{Alg}} \log \frac{q_{T+1}(y^*)}{q_1(y^*)} + \mathbb{E}_{\omega^*}^{\text{Alg}} \sum_{t=1}^T B_{\chi, \eta}(S_t, \delta_{a_t}; \omega^*).$$

The outer expectation is only over the observation process and any exogenous randomness. In a deterministic noiseless experiment the same identity is pathwise.

4.3 Exact AIR and MAIR regret identities

We now present simplified statements and alternative proofs of several results that are central to the AIR/MAIR literature (Xu and Zeevi, 2025). These results are best understood as regret identities, rather than merely as upper bounds. The results in this subsection can be understood

as specializations and consequences of Theorem 4.3. In principle, the AIR/MAIR identities require only the Bellman-sufficient representation of Definition 2.3. We nevertheless formulate them in the DMSO language, both to align with earlier work and to simplify notation. Under this specialization, the index may be an optimal-action index or a model index. The state can be taken to be the full model posterior as illustrated in Section 3.2.

Let $\mathcal{M}$ be the model class and $\mathcal{A}$ be the decision space. Model M has reward mean $r_M(a)$ and observation law $P_{M,a}$. For a candidate optimal action a^* , write

$$\Delta_{M,a^*}(p) := r_M(a^*) - \mathbb{E}_{a \sim p} r_M(a).$$

AIR. An AIR belief is a pair belief $\nu \in \Delta(\mathcal{M} \times \mathcal{A})$ over (M, A^*) , while the reference is an action marginal $q \in \text{int } \Delta(\mathcal{A})$. Define

$$\text{AIR}_{q,\eta}(p, \nu) := \mathbb{E}_{(M,A^*) \sim \nu} \Delta_{M,A^*}(p) - \frac{1}{\eta} \mathbb{E}_{a \sim p} I_\nu(A^*; O | a) - \frac{1}{\eta} D_{\text{KL}}(q_\nu \| q), \quad (12)$$

where q_ν is the A^* -marginal of ν . Equivalently,

$$\text{AIR}_{q,\eta}(p, \nu) = \mathbb{E}_{(M,A^*),a,O} \left[\Delta_{M,A^*}(\delta_a) - \frac{1}{\eta} \log \frac{\nu(A^* | a, O)}{q(A^*)} \right].$$

Lemma 4.5 (AIR gradient bracket). *Assume finite spaces, or dominated spaces where the displayed Gateaux derivatives exist. Up to an additive constant on the simplex,*

$$\frac{\partial}{\partial \nu(M, a^*)} \text{AIR}_{q,\eta}(p, \nu) = \Delta_{M,a^*}(p) - \frac{1}{\eta} \mathbb{E}_{a \sim p, O \sim P_{M,a}} \log \frac{\nu(A^* = a^* | a, O)}{q(a^*)}. \quad (13)$$

Consequently, for every fixed pair (M, a^*) ,

$$\begin{aligned} B_{q,\eta}^{\text{AIR}}(p, \nu; M, a^*) &:= \text{AIR}_{q,\eta}(p, \nu) + \langle \nabla_\nu \text{AIR}_{q,\eta}(p, \nu), \delta_{(M,a^*)} - \nu \rangle \\ &= \Delta_{M,a^*}(p) - \frac{1}{\eta} \mathbb{E}_{a \sim p, O \sim P_{M,a}} \log \frac{\nu(A^* = a^* | a, O)}{q(a^*)}. \end{aligned} \quad (14)$$

Proof. For fixed a , set

$$w_{a^*,o} = \sum_M \nu(M, a^*) P_{M,a}(o), \quad w_o = \sum_{a^*} w_{a^*,o}.$$

Then

$$I_\nu(A^*; O | a) + D_{\text{KL}}(q_\nu \| q) = \sum_{a^*,o} w_{a^*,o} \log \frac{w_{a^*,o}}{w_o q(a^*)}.$$

Differentiating with respect to $\nu(M, a^*)$ gives

$$\mathbb{E}_{O \sim P_{M,a}} \log \frac{\nu(A^* = a^* | a, O)}{q(a^*)},$$

because the $+1$ terms from differentiating $w_{a^*,o} \log w_{a^*,o}$ and $w_o \log w_o$ cancel. Adding the derivative of the linear regret term proves (13). Pairing the gradient with $\delta_{(M,a^*)} - \nu$ cancels the posterior-averaged terms and gives (14). $\square$

Theorem 4.6 (Exact AIR regret identity). *Let an algorithm generate decision laws p_t , pair beliefs ν_t , and reference action beliefs $q_t \in \text{int } \Delta(\mathcal{A})$. After action a_t and observation O_t , set*

$$q_{t+1}(\cdot) = \nu_t(A^* \in \cdot \mid a_t, O_t).$$

Fix a model M^ and a comparator a^* with $q_t(a^*) > 0$ for all relevant t ; in the standard stochastic case, take $a^* \in \arg \max_a r_{M^*}(a)$. Let*

$$\text{Reg}_T(\text{Alg}, M^*, a^*) := \mathbb{E} \sum_{t=1}^T (r_{M^*}(a^*) - r_{M^*}(A_t)).$$

Then, for every $\eta > 0$,

$$\text{Reg}_T(\text{Alg}, M^*, a^*) = \frac{1}{\eta} \mathbb{E} \log \frac{q_{T+1}(a^*)}{q_1(a^*)} + \mathbb{E} \sum_{t=1}^T B_{q_t, \eta}^{\text{AIR}}(p_t, \nu_t; M^*, a^*). \quad (15)$$

Remark 4.7 (Adversarial model sequence). *Since Lemma 4.5 holds pointwise for every model–action pair, and does not require the true model to be stochastic, Theorem 4.6 extends directly to an adversarial model sequence $\mathbf{M} = (M_1, \dots, M_T) \in \mathcal{M}^T$. In this case, the comparator is the best-in-hindsight action a^* , and the regret identity remains exact. We omit this extension only to avoid additional notation.*

Proof. Condition on H_{t-1} . By the definition of the action-marginal posterior,

$$\mathbb{E} \left[\log \frac{q_{t+1}(a^*)}{q_t(a^*)} \mid H_{t-1} \right] = \mathbb{E}_{a \sim p_t, O \sim P_{M^*, a}} \log \frac{\nu_t(A^* = a^* \mid a, O)}{q_t(a^*)}.$$

Combining this equality with (14) gives a one-step equality between regret, bracket, and posterior-ratio increment. Taking expectations and summing over time telescopes the logarithm. $\square$

MAIR. For MAIR the belief is $\mu \in \Delta(\mathcal{M})$ and the reference is $\rho \in \text{int } \Delta(\mathcal{M})$. Write $\mu_{O|a} = \int P_{M, a} \mu(dM)$ and

$$\Delta_M(p) = r_M(a_M^*) - \mathbb{E}_{a \sim p} r_M(a), \quad a_M^* \in \arg \max_a r_M(a).$$

Define

$$\text{MAIR}_{\rho, \eta}(p, \mu) := \mathbb{E}_{M \sim \mu} \Delta_M(p) - \frac{1}{\eta} \mathbb{E}_{a \sim p} I_\mu(M; O \mid a) - \frac{1}{\eta} D_{\text{KL}}(\mu \parallel \rho). \quad (16)$$

The same differentiation gives

$$\begin{aligned} B_{\rho, \eta}^{\text{MAIR}}(p, \mu; M^*) &:= \text{MAIR}_{\rho, \eta}(p, \mu) + \langle \nabla_\mu \text{MAIR}_{\rho, \eta}(p, \mu), \delta_{M^*} - \mu \rangle \\ &= \Delta_{M^*}(p) - \frac{1}{\eta} \mathbb{E}_{a \sim p} D_{\text{KL}}(P_{M^*, a} \parallel \mu_{O|a}) - \frac{1}{\eta} \log \frac{\mu(M^*)}{\rho(M^*)}. \end{aligned} \quad (17)$$

If $\rho_{t+1}(\cdot) = \mu_t(\cdot \mid a_t, O_t)$, then

$$\text{Reg}_T(\text{Alg}, M^*) = \frac{1}{\eta} \mathbb{E} \log \frac{\rho_{T+1}(M^*)}{\rho_1(M^*)} + \mathbb{E} \sum_{t=1}^T B_{\rho_t, \eta}^{\text{MAIR}}(p_t, \mu_t; M^*). \quad (18)$$

The proof is the same posterior-ratio telescope as for AIR. Thus AIR and MAIR are one indexed identity with different choices of Y .

Remark 4.8 (Danskin, Nash equilibrium, and why the bracket is not accidental). *The direct proofs above are intentionally short. The original AIR proof of Xu and Zeevi (2025, Section 5) gives the variational reason behind the same equations. Introduce an auxiliary prediction rule $Q : (a, o) \mapsto \Delta(\mathcal{A})$ and the posterior-scoring objective*

$$\mathbb{E}_{(M, A^*), a, O} \left[\Delta_{M, A^*}(\delta_a) - \frac{1}{\eta} \log \frac{Q[a, O](A^*)}{q(A^*)} \right].$$

For fixed ν , the logarithmic score is strongly convex in the predictive distribution and is uniquely minimized by the Bayesian posterior $Q_\nu[a, o] = \nu(A^ \in \cdot \mid a, o)$. The optimized value is the AIR objective; the unoptimized objective is linear, hence concave, in ν . Danskin's theorem identifies the AIR gradient with the derivative at the optimizer Q_ν . This is the conceptual content of Lemma 5.3 in Xu and Zeevi (2025). The same concavity plus strong convexity also gives the Nash equilibrium statement of Lemma 5.2 by first-order/Danskin reasoning, without invoking a separate fixed-point theorem. This does not necessarily simplify the full proof, because compactness, differentiability, and boundary conditions still have to be checked. It does explain why the gradient bracket is not an accidental cancellation: posterior scoring produces the minimizer, Danskin produces the gradient, and the regret identity telescopes the posterior reference.*

The perspective developed in Xu and Zeevi (2025) is that frequentist sequential learning can be organized around algorithmic beliefs μ , interpreted as predictive laws, and decision rules induced by reference or posterior beliefs ρ , either through posterior sampling or through robust minimax optimization. We use MAIR to illustrate how a single functional captures both the confidence multiplier in UCB-type algorithms and the decision-estimation coefficient in one-round minimax rules.

Lemma 4.9 (MAIR gradient bracket and three canonical specializations). *Assume the observation laws are dominated and the displayed derivatives exist. For fixed p, ρ, μ , up to an additive constant on the probability simplex,*

$$\frac{\partial}{\partial \mu(M)} \text{MAIR}_{\rho, \eta}(p, \mu) = \Delta_M(p) - \frac{1}{\eta} \log \frac{\mu(M)}{\rho(M)} - \frac{1}{\eta} \mathbb{E}_{\pi \sim p} D_{\text{KL}}(P_{M, \pi} \parallel \mu_{o|\pi}).$$

Consequently, for every M^ ,*

$$\begin{aligned} & \text{MAIR}_{\rho, \eta}(p, \mu) + \langle \nabla_\mu \text{MAIR}_{\rho, \eta}(p, \mu), \delta_{M^*} - \mu \rangle \\ &= \Delta_{M^*}(p) - \frac{1}{\eta} \mathbb{E}_{\pi \sim p} D_{\text{KL}}(P_{M^*, \pi} \parallel \mu_{o|\pi}) - \frac{1}{\eta} \log \frac{\mu(M^*)}{\rho(M^*)}. \end{aligned} \quad (19)$$

This bracket has three important specializations.

1. *If $\mu = \rho$, the logarithmic correction vanishes and the bracket becomes the KL-DEC offset*

$$\Delta_{M^*}(p) - \frac{1}{\eta} \mathbb{E}_{\pi \sim p} D_{\text{KL}}(P_{M^*, \pi} \parallel \mu_{o|\pi}). \quad (20)$$

In Section 5 we will show that for GP-UCB,

$$\text{self-normalized calibration} + \text{UCB optimism} \Rightarrow \text{MAIR gradient-error condition}.$$

2. *If $\bar{\mu} \in \arg \max_{\mu \in \Delta(\mathcal{M})} \text{MAIR}_{\rho, \eta}(p, \mu)$ is an interior maximizer, then*

$$\text{MAIR}_{\rho, \eta}(p, \bar{\mu}) + \langle \nabla_\mu \text{MAIR}_{\rho, \eta}(p, \bar{\mu}), \delta_{M^*} - \bar{\mu} \rangle \leq \text{MAIR}_{\rho, \eta}(p, \bar{\mu}). \quad (21)$$

This is the MAMS/MEBO bracket control.

3. Let $p_{M,\pi} = dP_{M,\pi}/d\nu_\pi$ and define, for an auxiliary pair $(\bar{\pi}, \bar{o})$,

$$\frac{d\mu_{\bar{\pi},\bar{o}}^{1/2}}{d\rho}(M) := \frac{\sqrt{p_{M,\bar{\pi}}(\bar{o})}}{\int \sqrt{p_{N,\bar{\pi}}(\bar{o})} \rho(dN)}. \quad (22)$$

Let $h^2(P, Q) := 1 - \int \sqrt{dP dQ}$. If $\bar{\pi} \sim p$ and $\bar{o} \sim P_{M^*,\bar{\pi}}$, then

$$\begin{aligned} \mathbb{E}_{\bar{\pi},\bar{o}} \left[\text{MAIR}_{\rho,\eta}(p, \mu_{\bar{\pi},\bar{o}}^{1/2}) + \langle \nabla_\mu \text{MAIR}_{\rho,\eta}(p, \mu_{\bar{\pi},\bar{o}}^{1/2}), \delta_{M^*} - \mu_{\bar{\pi},\bar{o}}^{1/2} \rangle \right] \\ \leq \Delta_{M^*}(p) - \frac{1}{\eta} \mathbb{E}_{\pi \sim p} \mathbb{E}_{M \sim \rho} h^2(P_{M^*,\pi}, P_{M,\pi}). \end{aligned} \quad (23)$$

Thus the square-root posterior turns the MAIR bracket into a Hellinger E2D/DEC offset.

Proof. The derivative calculation is direct. The derivative of $\mathbb{E}_\mu \Delta_M(p)$ is $\Delta_M(p)$; the derivative of $D_{\text{KL}}(\mu \parallel \rho)$ is $\log(\mu(M)/\rho(M)) + 1$; and differentiating the predictive mixture in the mutual-information term gives $D_{\text{KL}}(P_{M,\pi} \parallel \mu_{o|\pi})$, up to the same simplex constant. Pairing the derivative with $\delta_{M^*} - \mu$ cancels the averaged terms in (16) and proves (19).

The case $\mu = \rho$ is immediate from the logarithmic cancellation.

The case $\bar{\mu} \in \arg \max \text{MAIR}$ follows from first-order optimality.

It remains to justify the square-root claim; this is a standard fact for the squared Hellinger distance. For fixed $\bar{\pi}$, let

$$Z_{\bar{\pi}}(\bar{o}) := \int \sqrt{p_{N,\bar{\pi}}(\bar{o})} \rho(dN).$$

Then

$$\log \frac{d\mu_{\bar{\pi},\bar{o}}^{1/2}}{d\rho}(M^*) = \frac{1}{2} \log p_{M^*,\bar{\pi}}(\bar{o}) - \log Z_{\bar{\pi}}(\bar{o}).$$

Under $\bar{o} \sim P_{M^*,\bar{\pi}}$, Jensen's inequality gives

$$\begin{aligned} \mathbb{E} \log \frac{d\mu_{\bar{\pi},\bar{o}}^{1/2}}{d\rho}(M^*) &= -\mathbb{E} \log \frac{Z_{\bar{\pi}}(\bar{o})}{\sqrt{p_{M^*,\bar{\pi}}(\bar{o})}} \\ &\geq -\log \int \sqrt{p_{M^*,\bar{\pi}}(o)} Z_{\bar{\pi}}(o) d\nu_{\bar{\pi}}(o) \\ &= -\log (1 - \mathbb{E}_{M \sim \rho} h^2(P_{M^*,\bar{\pi}}, P_{M,\bar{\pi}})) \\ &\geq \mathbb{E}_{M \sim \rho} h^2(P_{M^*,\bar{\pi}}, P_{M,\bar{\pi}}). \end{aligned}$$

Averaging over $\bar{\pi} \sim p$, substituting into (19), and dropping the nonnegative KL term proves (23). $\square$

5 Upper bounds: one identity and four algorithm families

We first isolate the fixed-truth log-potential telescope, because it is the pathwise algebra underneath AIR/MAIR upper bounds. Bayesian mutual information, UCB self-normalized gain, and EBO log-partition relaxations are different ways of controlling this telescope; they should not be conflated. Fix a truth ω^* , write $y^* = \chi(\omega^*)$, and let q_s be the maintained reference law of the index at state s . This reference law may be an exact posterior marginal, an algorithmic belief marginal, a confidence-to-belief construction, or an exponential-weights marginal. After action a and observation o , write $q^+(\cdot \mid s, a, o)$ for the next reference marginal. The fixed-truth reference log gain is

$$\mathbf{G}_\chi(s, p; \omega^*) := \mathbb{E}_{a \sim p, O \sim P_{\omega^*, s, a}} \log \frac{q^+(y^* \mid s, a, O)}{q_s(y^*)}. \quad (24)$$

When the reference update is exact indexed Bayes, this is $J_\chi(s, p; \omega^\star)$ from (7). When the reference is an AIR pair belief ν with marginal q , it is the AIR posterior-coordinate log ratio

$$\mathbb{E}_{a \sim p, O \sim P_{\omega^\star, a}} \log \frac{\nu(Y = y^\star \mid a, O)}{q(y^\star)}.$$

For any predictable continuation potential V_t on states and any coefficient $\gamma > 0$, define the specialized fixed-truth logarithmic potential

$$\Phi_t^{V, \gamma}(s; y^\star) := V_t(s) + \gamma \log \frac{1}{q_s(y^\star)}. \quad (25)$$

The corresponding AIR/MAIR Bellman bracket is

$$\mathfrak{B}_\chi^{V, \gamma}(s, p; \omega^\star) := \ell_{\omega^\star}(s, p) + \mathbb{E}_{\omega^\star}[V_{t+1}(S^+) \mid s, p] - V_t(s) - \gamma \mathbf{G}_\chi(s, p; \omega^\star), \quad (26)$$

where $S^+ = \tau(s, A, O)$ with $A \sim p$ and $O \sim P_{\omega^\star, s, A}$. The identity behind the whole upper section is the telescope

$$\begin{aligned} \mathbb{E}_{\omega^\star} L_{\omega^\star}(H_T) &= V_1(s_1) - \mathbb{E}_{\omega^\star} V_{T+1}(S_{T+1}) + \gamma \log \frac{1}{q_1(y^\star)} - \gamma \mathbb{E}_{\omega^\star} \log \frac{1}{q_{T+1}(y^\star)} \\ &\quad + \mathbb{E}_{\omega^\star} \sum_{t=1}^T \mathfrak{B}_\chi^{V, \gamma}(S_t, p_t; \omega^\star). \end{aligned} \quad (27)$$

Thus an upper bound is a Bellman supersolution for this log-potential bracket. The four algorithm families below differ only in how they make the bracket small: the exact log-penalized Bellman program optimizes it dynamically; UCB certifies it by calibration and optimism; E2D/DEC drops or bounds the continuation to obtain a one-step offset; and AMS/EBO optimizes a robust KL-dual belief relaxation whose first-order condition certifies it.

Theorem 5.1 (Generic fixed-truth indexed AIR/MAIR upper certificate). *Fix $\omega^\star$ and $y^\star = \chi(\omega^\star)$. Assume $q_1(y^\star) > 0$ and that the reference update remains positive on $y^\star$ along the analyzed trajectory. If an algorithm chooses p_t so that, for every reached state,*

$$\mathfrak{B}_\chi^{V, \gamma}(S_t, p_t; \omega^\star) \leq \varepsilon_t,$$

then

$$\mathbb{E}_{\omega^\star}^{\text{Alg}} L_{\omega^\star}(H_T) \leq V_1(s_1) + \gamma \log \frac{1}{q_1(y^\star)} - \mathbb{E}_{\omega^\star}^{\text{Alg}} V_{T+1}(S_{T+1}) - \gamma \mathbb{E}_{\omega^\star}^{\text{Alg}} \log \frac{1}{q_{T+1}(y^\star)} + \mathbb{E}_{\omega^\star}^{\text{Alg}} \sum_{t=1}^T \varepsilon_t. \quad (28)$$

Proof. Taking the conditional expectation of $\gamma \log(1/q_{t+1}(y^\star)) - \gamma \log(1/q_t(y^\star))$ under $P_{\omega^\star, s, a}$ gives $-\gamma \mathbf{G}_\chi(s, p; \omega^\star)$. Substituting the bracket inequality into the telescoping identity (27) proves (28). $\square$

It is useful to keep one concrete state specialization in mind. A fixed estimation procedure maps histories to algorithmic beliefs ν_t over environments, a decision procedure maps the retained state to p_t , and the index marginal is

$$q_t = \chi_\# \nu_t. \quad (29)$$

In a well-specified Bayes analysis, ν_t is the exact posterior. In a frequentist analysis, ν_t may be only an algorithmic belief or confidence object; then a calibration event must put the fixed truth

into the comparison set used by the robust Bellman program below. Posterior averaging of the coordinate potential gives

$$\mathbb{E}_{Y \sim \nu_\chi} \gamma \log \frac{1}{q_s(Y)} = \gamma H(\nu_\chi) + \gamma D_{\text{KL}}(\nu_\chi \| q_s), \quad (30)$$

and for a finite index set the same logarithmic scoring rule has the KL-dual log-partition form

$$\Psi_\gamma(q, z) := \gamma \log \sum_{y \in \mathcal{Y}} q(y) e^{zy/\gamma} = \sup_{r \in \Delta(\mathcal{Y})} \{ \langle r, z \rangle - \gamma D_{\text{KL}}(r \| q) \}. \quad (31)$$

The coordinate log loss, its posterior average, and the dual log partition are three representations of the same KL/logarithmic scoring rule. They are not interchangeable in a proof: fixed-truth guarantees use (25), posterior-averaged Bayes identities use (30), and AMS/EBO relaxations use (31).

In the model-index specialization, write the posterior/reference belief as μ_t . If the maintained reference update is exact Bayes, then $\mathbf{G}_\chi = J_\chi$ and the fixed-truth MAIR information increment is

$$J_t^{\text{MAIR}}(p; \omega^*) := \mathbb{E}_{a \sim p} D_{\text{KL}}(P_{\omega^*, t}(\cdot | S_t, a) \| P_{\mu_t, t}(\cdot | S_t, a)). \quad (32)$$

The action-index AIR version is identical with the index $Y = A^*$, the reference marginal q_t , and the AIR posterior coordinate $q^+(a^* | a, o)$ in place of the model posterior.

5.1 Information-potential Bellman programming

The exact frequentist upper algorithm is a log-penalized Bellman program on calibrated state/index representations. Let $\mathfrak{C}_t(s)$ be the set of environments that the algorithmic state regards as admissible at state s . The set may be a confidence set, a localized posterior support, or a robust belief support. For $\gamma > 0$, define

$$W_{T+1}^\gamma(s) = 0, \quad W_t^\gamma(s) := \inf_{p \in \Delta(\mathcal{A}(s))} \sup_{\omega \in \mathfrak{C}_t(s)} \{ \ell_\omega(s, p) + \mathbb{E}_\omega[W_{t+1}^\gamma(S^+) | s, p] - \gamma \mathbf{G}_\chi(s, p; \omega) \}. \quad (33)$$

A policy attaining the infimum in (33) is the information-penalized Bellman dynamic-programming algorithm. This construction builds on the principle of admissible relaxations (Rakhlin et al., 2012), which organize learning through dynamic programming over loss states, and is particularly close to the partial-information relaxation viewpoint of Rakhlin and Sridharan (2016), where estimator-like state variables and potential terms quantify uncertainty. It can also be viewed as a sequential analogue of classical minimum-complexity, information-risk, and Gibbs-type estimation in fixed experiments (Barron and Cover, 1991; Yang and Barron, 1999; Zhang, 2006). The key new feature is that a generic logarithmic mass penalty is propagated through a controlled Bellman recursion, rather than applied only once to a terminal estimator, and the recursion is carried by a Bellman-sufficient state together with an explicit information index.

Theorem 5.2 (Frequentist log-penalized Bellman upper bound). *Fix $\Omega_0 \subseteq \Omega$ and $\gamma > 0$. Suppose the reference update and the comparison sets $\mathfrak{C}_t$ are calibrated for Ω_0 in the following sense: for every $\omega^* \in \Omega_0$, on the calibration event $\mathcal{E}$, $\omega^* \in \mathfrak{C}_t(S_t)$ whenever round t is reached. Suppose the algorithm chooses p_t so that*

$$\sup_{\omega \in \mathfrak{C}_t(S_t)} \{ \ell_\omega(S_t, p_t) + \mathbb{E}_\omega[W_{t+1}^\gamma(S^+) | S_t, p_t] - \gamma \mathbf{G}_\chi(S_t, p_t; \omega) \} \leq W_t^\gamma(S_t) + \varepsilon_t.$$

Then, for every fixed truth $\omega^* \in \Omega_0$ with $y^* = \chi(\omega^*)$ and $q_1(y^*) > 0$,

$$\mathbb{E}_{\omega^*}^{\text{Alg}} L_{\omega^*}(H_T) \leq W_1^\gamma(s_1) + \gamma \log \frac{1}{q_1(y^*)} + \mathbb{E}_{\omega^*}^{\text{Alg}} \sum_{t=1}^T \varepsilon_t + \text{cal}_T(\omega^*), \quad (34)$$

provided the index is finite or countable, where

$$\text{cal}_T(\omega^*) := \mathbb{E}_{\omega^*}[L_{\omega^*}(H_T) \mathbb{1}\{\mathcal{E}^c\}].$$

If calibration holds surely, then $\text{cal}_T(\omega^*) = 0$; if losses are bounded by L , then $\text{cal}_T(\omega^*) \leq LT\mathbb{P}_{\omega^*}(\mathcal{E}^c)$.

Proof. On the calibration event, the fixed truth is one of the environments in the robust supremum. Therefore the displayed approximate minimization implies

$$\ell_{\omega^*}(S_t, p_t) + \mathbb{E}_{\omega^*}[W_{t+1}^\gamma(S^+) \mid S_t, p_t] - W_t^\gamma(S_t) - \gamma G_\chi(S_t, p_t; \omega^*) \leq \varepsilon_t.$$

Apply Theorem 5.1 with $V_t = W_t^\gamma$. Since $W_{T+1}^\gamma = 0$ and the terminal index log loss is nonnegative for a finite or countable index, the result follows on $\mathcal{E}$. Splitting the expectation according to $\mathcal{E}$ and $\mathcal{E}^c$ gives the calibration term. $\square$

The posterior-averaged Bayesian recursion is a specialization rather than the primitive fixed-truth statement. If the robust comparison set is replaced by the current posterior law and one averages the coordinate identity over $Y \sim q_s$, then

$$\mathbb{E}[H_\chi(S_{t+1}) \mid s, p] - H_\chi(s) = -\mathcal{I}_\chi(s, p).$$

For a fixed multiplier $\gamma > 0$, define the Bayes information-regularized value

$$U_{T+1}^\gamma \equiv 0, \quad U_t^\gamma(s) := \inf_{p \in \Delta(\mathcal{A}(s))} \{ \ell(s, p) + \mathbb{E}[U_{t+1}^\gamma(S_{t+1}) \mid s, p] - \gamma \mathcal{I}_\chi(s, p) \}. \quad (35)$$

Equivalently,

$$\Phi_t^\gamma(s) := U_t^\gamma(s) + \gamma H_\chi(s) \quad (36)$$

solves the admissible-relaxation recursion

$$\Phi_t^\gamma(s) = \inf_{p \in \Delta(\mathcal{A}(s))} \{ \ell(s, p) + \mathbb{E}[\Phi_{t+1}^\gamma(S_{t+1}) \mid s, p] \}, \quad \Phi_{T+1}^\gamma(s) = \gamma H_\chi(s). \quad (37)$$

Theorem 5.3 (Posterior-averaged information-potential upper bound). *Assume the indexed Bellman representation is exact under the Bayesian mixture law and $H_\mu(Y) < \infty$. If a policy chooses p_t satisfying*

$$\ell(S_t, p_t) + \mathbb{E}[U_{t+1}^\gamma(S_{t+1}) \mid S_t, p_t] - \gamma \mathcal{I}_\chi(S_t, p_t) \leq U_t^\gamma(S_t) + \varepsilon_t,$$

then

$$\mathbb{E} L_\Omega(H_T) \leq U_1^\gamma(s_1) + \gamma I_\mu(Y; H_T) + \sum_{t=1}^T \varepsilon_t \leq U_1^\gamma(s_1) + \gamma H_\mu(Y) + \sum_{t=1}^T \varepsilon_t. \quad (38)$$

This is the posterior average of the fixed-truth coordinate telescope, not a pointwise replacement for Theorem 5.2.

Proof. Rearrange the displayed inequality and telescope U_t^γ . Proposition 4.2 identifies $\mathbb{E} \sum_t \mathcal{I}_\chi(S_t, p_t)$ with $I_\mu(Y; H_T)$, and $I_\mu(Y; H_T) \leq H_\mu(Y)$ gives the final bound. The equivalence with (37) follows from the posterior entropy-drop identity. $\square$

Theorem 5.4 (Frequentist Bellman information-risk matching principle). *Fix a problem class Ω_0 , a horizon T , an index χ , and a Bellman-sufficient state/reference update. Suppose the upper side has a calibrated log-penalized Bellman program as in Theorem 5.2. Define its worst-case upper value*

$$\mathfrak{U}_T(\Omega_0) := \inf_{\gamma > 0, q_1, \mathfrak{C}, \text{Alg}} \sup_{\omega \in \Omega_0} \left\{ W_1^\gamma(s_1) + \gamma \log \frac{1}{q_1(\chi(\omega))} + \text{cal}_T(\omega) + \text{opt}_T(\omega) \right\},$$

where the infimum is over initial reference laws, calibrated comparison mechanisms, and algorithms that approximately solve the displayed Bellman recursion, and where

$$\text{opt}_T(\omega) := \mathbb{E}_\omega^{\text{Alg}} \sum_{t=1}^T \varepsilon_t$$

records optimization or relaxation error. For the exact log-penalized Bellman dynamic program, $\text{opt}_T \equiv 0$. Then

$$\mathfrak{R}_T^*(\Omega_0) \leq \mathfrak{U}_T(\Omega_0).$$

Conversely, suppose there are a prior μ supported on Ω_0 , the same index χ , the same state process, and Bellman-Fano lower certificates $\overline{C}, \overline{\Gamma}^r$ satisfying the critical condition (58). Then

$$\mathfrak{R}_T^*(\Omega_0) \geq \frac{Tr}{2}.$$

Consequently, if for the same representation and some radius r one has

$$\mathfrak{U}_T(\Omega_0) \leq C_{\text{up}} Tr \quad \text{and} \quad \overline{C}_1(s_1) + \log 2 \leq \frac{1}{2} \mathfrak{E}_1^r(\mu, \chi),$$

then the frequentist minimax regret is sandwiched as

$$\frac{Tr}{2} \leq \mathfrak{R}_T^*(\Omega_0) \leq C_{\text{up}} Tr.$$

This is the precise sense in which the log-penalized Bellman dynamic program matches the Bellman-Fano lower bound. The theorem is conditional on representational sufficiency and calibration; it is not a claim that an arbitrary index, belief, or one-step relaxation is sharp.

Proof. The upper inequality is Theorem 5.2 followed by taking the supremum over $\omega \in \Omega_0$ and then optimizing over the displayed design choices. The lower inequality is Theorem 6.3 and Yao's principle, since μ is supported on Ω_0 . The sandwich conclusion is the conjunction of the two bounds. The ghost history enters only through the lower certificate; the upper side may use the exact Bellman recursion or any admissible relaxation, including random-payout or sequential symmetrization, that upper bounds W_1^γ . $\square$

Theorem 5.4 is the interactive analogue of noninteractive information-risk minimization. In a fixed experiment, the statistical radius is determined by balancing a loss scale against a log prior mass or metric entropy. Here the loss scale is propagated by Bellman recursion, the upper log mass is the coordinate cost $\log(1/q_1(\chi(\omega)))$, and the lower entropy is the ghost-good quantile $-\log p_r$. The exact recursion is often computationally difficult, but it is the benchmark that the following algorithm families approximate or certify.

5.2 UCB families: calibration plus optimism

An upper-confidence-bound (UCB) algorithm maintains a prediction and an uncertainty scale, builds a confidence band from them, and chooses an action that is optimistic within that band (Auer et al., 2002; Abbasi-Yadkori et al., 2011; Srinivas et al., 2010). We use the following notation throughout. The unmultiplied uncertainty scale is denoted by $\sigma_t(a)$ in the generic discussion and by $s_t(x)$ in the GP specialization below. The confidence half-width is

$$w_t(a) = \beta_t \sigma_t(a),$$

where β_t is the confidence multiplier. Calibration says that the fixed truth is contained in the band of radius $w_t(a)$; optimism says that the selected action maximizes the upper envelope. Together, in reward-maximization problems, these two facts give

$$r_{\omega^*}(s, a^*) \leq m_t(a^*) + w_t(a^*) \leq m_t(a_t) + w_t(a_t) \leq r_{\omega^*}(s, a_t) + 2w_t(a_t),$$

and hence the instantaneous regret is at most $2\beta_t \sigma_t(a_t)$. An elliptical-potential/log-determinant lemma then controls the realized sum of $\sigma_t^2(a_t)$. Thus UCB does not solve the AIR/MAIR bracket optimization directly; it certifies the relevant fixed-truth bracket after the optimistic action has been chosen.

Lemma 5.5 (Calibration and optimism imply bracket control). *Fix a truth ω^* and condition on an event $\mathcal{E}$ on which the algorithm's state is calibrated for that truth. Suppose that, for every reached round t , the action a_t chosen by the algorithm and a predictable uncertainty scale $\sigma_t(a) \geq 0$ satisfy*

$$\ell_{\omega^*}(S_t, a_t) \leq c_{\text{opt}} \beta_t \sigma_t(a_t), \quad (39)$$

where c_{opt} is a numerical optimism constant; in the usual two-sided reward-confidence proof above, $c_{\text{opt}} = 2$. Then for every $\eta > 0$,

$$\ell_{\omega^*}(S_t, a_t) - \frac{1}{\eta} \sigma_t^2(a_t) \leq \frac{\eta c_{\text{opt}}^2 \beta_t^2}{4}. \quad (40)$$

Consequently, for the deterministic-action fixed-truth bracket in (9),

$$B_{\chi, \eta}(S_t, \delta_{a_t}; \omega^*) \leq \frac{\eta c_{\text{opt}}^2 \beta_t^2}{4} + \frac{1}{\eta} (\sigma_t^2(a_t) - J_{\chi}(S_t, \delta_{a_t}; \omega^*)), \quad (41)$$

where the fixed-truth index log gain is defined in (7). If, in addition, the log-determinant/elliptical-potential argument gives a realized bound

$$\sum_{t=1}^T \sigma_t^2(a_t) \leq c_{\text{sn}} G_T,$$

for an information telescope G_T , then summing (41) reduces UCB regret control to the comparison between G_T and the fixed-truth log-gain telescope.

Proof. The optimism/calibration hypothesis (39) and Young's inequality give

$$c_{\text{opt}} \beta_t \sigma_t(a_t) \leq \frac{\eta c_{\text{opt}}^2 \beta_t^2}{4} + \frac{\sigma_t^2(a_t)}{\eta}.$$

This proves (40). Subtracting $\eta^{-1} J_{\chi}(S_t, \delta_{a_t}; \omega^*)$ from both sides and using the definition of $B_{\chi, \eta}$ proves (41). The final sentence is the result of summing the displayed inequality over time. $\square$

The correction $\sigma_t^2(a_t) - J_\chi(S_t, \delta_{a_t}; \omega^*)$ is not a second estimation cost. It is the algebraic difference between the tractable self-normalized variance proxy and the fixed-truth posterior log gain appearing in the bracket. In linear-UCB and GP-UCB proofs, the log-determinant or elliptical-potential argument is precisely the certificate that the realized sum of $\sigma_t^2(a_t)$ is controlled by an information telescope. The confidence multiplier β_t is separate: it is the price of converting calibrated uncertainty into regret through optimism. Therefore

$$\text{self-normalized calibration} + \text{optimism} \implies \text{fixed-truth AIR/MAIR bracket control.}$$

In this organizing perspective, the round-heterogeneous summation of unmultiplied uncertainty scales, usually controlled by an elliptical-potential or log-determinant argument, plays the role of a realized estimation-complexity telescope. The confidence multiplier enters instead through the regret-to-uncertainty conversion, and hence through the coefficient in the fixed-truth bracket. This separates two mechanisms that are often blurred in informal UCB discussions: posterior or confidence widths quantify the “estimation complexity” of the underlying parameter, as explained by [Lattimore \(2023\)](#), while optimism and calibration determine how that geometry pays for regret. The resulting unification reveals a common posterior-telescoping structure while keeping distinct the estimation telescope and the coefficient mechanisms behind UCB, E2D, and AMS/EBO.

5.3 E2D/DEC: robust one-step offset optimization

The decision-estimation coefficient (DEC) ([Foster et al., 2021](#)) and estimation-to-decision (E2D) algorithms choose p by solving a one-round robust offset. In model-index notation the schematic KL form is

$$\inf_{p \in \Delta(\mathcal{A})} \sup_{\omega \in \mathcal{N}(s)} \{ \ell_\omega(s, p) - \gamma \mathbb{E}_{a \sim p} D_{\text{KL}}(P_{\omega, s, a} \| P_{\rho, s, a}) \}, \quad (42)$$

where ρ is a reference belief or reference model and $\mathcal{N}(s)$ is a localized comparison set that contains the fixed truth on the calibration event. The same one-round minimax principle applies to the original Hellinger form and constrained variants ([Foster et al., 2023](#)).

What E2D/DEC keeps is the immediate regret and a one-step statistical-separation penalty. What it drops, freezes, or upper bounds is the Bellman continuation term

$$\mathbb{E}_\omega[V_{t+1}(S^+) \mid s, p] - V_t(s),$$

together with the exact coordinate reference log gain G_χ in (26). Thus E2D/DEC is not an optimizer of the full Bellman potential unless one proves that the omitted continuation is dominated by the displayed one-step penalty. In an exact Bayesian calculation, posterior averaging can turn the KL penalty into a one-step mutual-information term. In a fixed-truth or frequentist calculation, the corresponding statement is the reference posterior-ratio identity for G_χ , or its exact-Bayes specialization J_χ , plus whatever calibration or localization is needed to compare $P_{\rho, s, a}$ with the reference predictive law in that identity. Without this compatibility, one may pay for estimation difficulty once through the one-step DEC and again through a separate log-determinant or posterior-ratio telescope. The safe interpretation is therefore: DEC/E2D is a one-step relaxation, re-solved at each state, of the fixed-truth log-posterior Bellman bracket.

5.4 AMS/EBO: robust convex belief optimization

AMS/EBO should be read as a tractable relaxation of the robust log-penalized Bellman program. Work with an action or policy index $Y = \chi(\Omega)$. At state s , let $q \in \text{int } \Delta(\mathcal{Y})$ be the current reference

marginal and let $\mathfrak{B}_t(s)$ be a calibrated convex set of pair beliefs ν on (ω, y) , supported on $y = \chi(\omega)$. For each ν , write ν_χ for its index marginal and define

$$\ell_\nu(s, p) = \mathbb{E}_{(\omega, y) \sim \nu} \ell_\omega(s, p), \quad \mathcal{I}_\chi(s, p; \nu) = \mathbb{E}_{a \sim p} I_\nu(Y; O \mid s, a).$$

With coefficient $\gamma > 0$, the one-step robust AIR/EBO objective is

$$\mathfrak{A}_{q, \gamma}(s, p, \nu) := \ell_\nu(s, p) - \gamma \mathcal{I}_\chi(s, p; \nu) - \gamma D_{\text{KL}}(\nu_\chi \| q). \quad (43)$$

The robust maximization over beliefs is essential:

$$p_t \in \arg \min_{p \in \Delta(\mathcal{A}(s))} \max_{\nu \in \mathfrak{B}_t(s)} \mathfrak{A}_{q_t, \gamma}(s, p, \nu), \quad \bar{\nu}_t \in \arg \max_{\nu \in \mathfrak{B}_t(s)} \mathfrak{A}_{q_t, \gamma}(s, p_t, \nu). \quad (44)$$

Without the inner $\max_\nu$, the display is only a posterior AIR calculation for a chosen belief, not a robust AMS/EBO rule.

The convex-analytic interpretation is the following. The fixed-truth proof uses the coordinate log loss $\gamma \log(1/q_t(y^*))$. Averaging this coordinate under a candidate belief gives $\gamma H(\nu_\chi) + \gamma D_{\text{KL}}(\nu_\chi \| q_t)$. When the next reference marginal is the posterior coordinate induced by ν , the expected change of this averaged potential is $-\gamma \mathcal{I}_\chi(s, p; \nu) - \gamma D_{\text{KL}}(\nu_\chi \| q_t)$. Equivalently, optimizing over possible next index laws yields the log-partition dual (31). Hence the EBO potential is the KL-dual log partition, while the fixed-truth AIR certificate is the coordinate posterior log loss; they coincide only through KL duality and the chosen update, not as identical scalar functions.

For a genuine convex belief optimization one must verify the convexity/concavity assumptions used by the saddle argument. In finite exponential-weights or admissible-relaxation settings this follows from the log-partition dual. In general model classes, mutual information need not be concave in an arbitrary belief when both the index marginal and the observation kernels vary. One must either verify concavity of $\nu \mapsto \mathfrak{A}_{q, \gamma}(s, p, \nu)$ on $\mathfrak{B}_t(s)$, or replace it by a certified concave upper relaxation.

Lemma 5.6 (AMS/EBO saddle certifies the AIR fixed-truth bracket). *Assume the finite or dominated differentiability setting of Lemma 4.5. Fix (ω^*, y^*) with $y^* = \chi(\omega^*)$. Suppose $\nu \mapsto \mathfrak{A}_{q_t, \gamma}(s, p_t, \nu)$ is concave on the convex set $\mathfrak{B}_t(s)$, $\bar{\nu}_t$ maximizes it as in (44), and the comparison direction $\delta_{(\omega^*, y^*)} - \bar{\nu}_t$ is admissible for $\mathfrak{B}_t(s)$; any failure of this containment is charged as a calibration error. If*

$$\max_{\nu \in \mathfrak{B}_t(s)} \mathfrak{A}_{q_t, \gamma}(s, p_t, \nu) \leq \varepsilon_t,$$

then the fixed-truth AIR bracket with coefficient γ satisfies

$$\ell_{\omega^*}(s, p_t) - \gamma \mathbb{E}_{a \sim p_t, O \sim P_{\omega^*, s, a}} \log \frac{\bar{\nu}_t(Y = y^* \mid s, a, O)}{q_t(y^*)} \leq \varepsilon_t. \quad (45)$$

If the reference update is $q_{t+1}(\cdot) = \bar{\nu}_t(Y \in \cdot \mid s, A_t, O_t)$ and the calibration assumptions hold along the trajectory, then Theorem 5.1 with $V_t \equiv 0$ gives

$$\mathbb{E}_{\omega^*} \text{Reg}_T \leq \gamma \log \frac{1}{q_1(y^*)} + \sum_{t=1}^T \varepsilon_t + \text{cal}_T \quad (46)$$

for finite or countable index sets. Averaging this coordinate bound under a Bayesian prior replaces $\log(1/q_1(Y))$ by the corresponding cross-entropy or by $H_\mu(Y)$ when $q_1 = \mu_\chi$.

Proof. The AIR gradient calculation in Lemma 4.5, with $\gamma = 1/\eta$, identifies the fixed-truth bracket as the first-order value of $\mathfrak{A}_{q_t, \gamma}$ in the direction $\delta_{(\omega^*, y^*)} - \bar{\nu}_t$. Concavity and optimality of $\bar{\nu}_t$ make that first-order correction nonpositive for every admissible comparison direction. Hence the bracket is no larger than $\mathfrak{A}_{q_t, \gamma}(s, p_t, \bar{\nu}_t)$, which is at most the robust value and therefore at most ε_t . The regret bound is the fixed-truth log-potential telescope. $\square$

A dynamic AMS/EBO relaxation is obtained by adding a continuation term before the robust maximization:

$$\mathfrak{A}_{t, q, \gamma}^V(s, p, \nu) = \ell_\nu(s, p) + \mathbb{E}_\nu V_{t+1}(S^+) - V_t(s) - \gamma \mathcal{I}_\chi(s, p; \nu) - \gamma D_{\text{KL}}(\nu_\chi \| q).$$

This is a valid relaxation of the robust Bellman recursion (33) only after the same first-order or direct-bracket certificate is proved. The standard one-step EBO display is the case $V_t \equiv 0$; the dynamic version is an admissible relaxation of the central log-penalized Bellman program, not an automatic consequence of the one-step objective.

Remark 5.7 (The one identity and the four families). *The identity is the fixed-truth log-potential telescope (27). The exact information-risk Bellman program is (33). Posterior-averaged Bayes analysis is obtained by averaging the same coordinate identity, giving (35). UCB certifies the bracket through confidence and optimism. E2D/DEC replaces the continuation by a one-step separation penalty. AMS/EBO optimizes a KL-dual robust belief relaxation, with the required $\max_\nu$ appearing explicitly. Thus the common object is not a universal one-step coefficient but the logarithmic Bellman bracket on a sufficient state and a chosen index.*

6 Lower bounds: reference histories and quantile indices

6.1 Ghost probability and true good probability

Fix an average-regret threshold $r \geq 0$. The posterior-reference ghost-good probability is

$$p_r^{\text{Alg}}(\mu, \chi) := \mathbb{P}_{Y \sim \mu_\chi, H'_T \sim \bar{\mathbb{P}}_\mu^{\text{Alg}}} (\bar{L}_\chi(Y, H'_T) \leq r), \quad (47)$$

where Y and H'_T are independent under the ghost law and μ_χ is the marginal law of Y . The true good probability is

$$q_r^{\text{Alg}}(\mu, \chi) := \mathbb{P}_{Y, H_T} (\bar{L}_\chi(Y, H_T) \leq r), \quad (48)$$

where (Y, H_T) are generated by $\Omega \sim \mu$ and $H_T \sim \mathbb{P}_\Omega^{\text{Alg}}$.

The ghost probability is algorithm-dependent. That dependence is intentional: it preserves the reference trajectory geometry instead of replacing it with a worst-case static small ball.

6.2 Reference-history quantile theorem

Let

$$\text{kl}(q, p) = q \log \frac{q}{p} + (1 - q) \log \frac{1 - q}{1 - p}$$

be binary KL divergence. For $p \in [0, 1]$ and $c \geq 0$, define

$$\psi(p, c) := \sup\{q \in [0, 1] : \text{kl}(q, p) \leq c\}.$$

Theorem 6.1 (Reference-history quantile indexed-information lower bound). *For every prior μ , index map χ , threshold $r \geq 0$, and algorithm Alg ,*

$$\text{kl}(q_r^{\text{Alg}}(\mu, \chi), p_r^{\text{Alg}}(\mu, \chi)) \leq C_\chi^{\text{Alg}}(\mu) = T\bar{C}_\chi^{\text{Alg}}(\mu). \quad (49)$$

Consequently,

$$\mathbb{E}_{\Omega \sim \mu} \mathbb{E}_{\mathbb{P}_\Omega^{\text{Alg}}} L_\Omega(H_T) \geq Tr \left[1 - \psi(p_r^{\text{Alg}}(\mu, \chi), T\bar{C}_\chi^{\text{Alg}}(\mu)) \right]. \quad (50)$$

In particular, if

$$T\bar{C}_\chi^{\text{Alg}}(\mu) \leq \text{kl}\left(\frac{1}{2}, p_r^{\text{Alg}}(\mu, \chi)\right), \quad (51)$$

then

$$\mathbb{E}_{\Omega \sim \mu} \mathbb{E}_{\mathbb{P}_\Omega^{\text{Alg}}} \bar{L}_\Omega(H_T) \geq \frac{r}{2}, \quad \mathbb{E}_{\Omega \sim \mu} \mathbb{E}_{\mathbb{P}_\Omega^{\text{Alg}}} L_\Omega(H_T) \geq \frac{Tr}{2}. \quad (52)$$

A convenient sufficient condition is

$$T\bar{C}_\chi^{\text{Alg}}(\mu) + \log 2 \leq \frac{1}{2} \log \frac{1}{p_r^{\text{Alg}}(\mu, \chi)}. \quad (53)$$

Proof. Consider the true joint law of (Y, H_T) and the ghost law of (Y, H'_T) . Both have the same marginal law of Y , while Y and H'_T are independent under the ghost law. Hence

$$D_{\text{KL}}(\mathcal{L}(Y, H_T) \| \mathcal{L}(Y, H'_T)) = I_\mu(Y; H_T).$$

Data processing applied to the indicator of the event $\{\bar{L}_\chi(Y, \cdot) \leq r\}$ gives

$$\text{kl}(q_r^{\text{Alg}}, p_r^{\text{Alg}}) \leq I_\mu(Y; H_T).$$

Proposition 4.2 identifies the right-hand side with $C_\chi^{\text{Alg}}(\mu) = T\bar{C}_\chi^{\text{Alg}}(\mu)$, proving (49). The definition of ψ gives $q_r^{\text{Alg}} \leq \psi(p_r^{\text{Alg}}, T\bar{C}_\chi^{\text{Alg}})$. By nonnegativity of loss and by (3),

$$\begin{aligned} \mathbb{E}_{\Omega, H_T} L_\Omega(H_T) &= \mathbb{E}_{Y, H_T} L_\chi(Y, H_T) \\ &\geq Tr \mathbb{P}\{\bar{L}_\chi(Y, H_T) > r\} \\ &= Tr(1 - q_r^{\text{Alg}}), \end{aligned}$$

which proves (50). If (51) holds, then $q_r^{\text{Alg}} \leq 1/2$, giving (52). Finally,

$$\text{kl}\left(\frac{1}{2}, p\right) = \frac{1}{2} \log \frac{1}{4p(1-p)} \geq \frac{1}{2} \log \frac{1}{p} - \log 2,$$

so (53) implies (51). □

Relation to interactive Fano. Theorem 6.1 should be viewed as an indexed, posterior-reference specialization of the interactive Fano method of Chen et al. (2024). The binary-KL step is the same ghost-data argument: one compares the true experiment with a reference experiment, applies data processing to the small-regret event, and obtains a Bayes risk lower bound when the trajectory information is smaller than the logarithmic inverse ghost-good probability. The additional point here is representational. We choose the reference experiment to be the Bayesian posterior-predictive

trajectory and choose an information index $Y = \chi(\Omega)$. With this choice, the trajectory information admits the exact stepwise identity

$$I_\mu(Y; H_T) = \mathbb{E}_{H'_T \sim \mathbb{P}_\mu^{\text{Alg}}} \sum_{t=1}^T \mathcal{I}_\chi(S'_t, p'_t).$$

Thus the same reference history H'_T generates both the ghost quantile and the indexed information telescope. This is the lower-bound counterpart of the AIR/MAIR upper-bound identity. In particular, $\chi(\Omega) = \Omega$ gives the model-index MAIR form, while $\chi(\Omega) = \pi^*(\Omega)$ gives the action-index AIR form.

The novelty is not a new Fano inequality, but a new representation of interactive Fano in posterior-reference indexed-information coordinates. This avoids prematurely collapsing the adaptive experiment to a one-round decision distribution, the step that creates the localized change-of-measure difficulty in one-round DEC reductions. It also makes the lower-bound critical radius directly comparable to AIR/MAIR supersolution upper bounds, since both sides are expressed as sums of the same indexed information along posterior-reference histories.

Remark 6.2 (Why average notation is useful). *The theorem could be written entirely with cumulative quantities. Writing $\bar{C}_\chi = T^{-1} \mathbb{E} \sum_t \mathcal{I}_\chi$ and $\bar{L} = T^{-1} L$ makes the scaling visible:*

$$\text{information side} = T \bar{C}_\chi, \quad \text{regret side} = Tr.$$

Thus T appears exactly as in one-round reductions, but no posterior trajectory has been collapsed into a static action distribution.

6.3 Bellman-Fano lower certificates

Theorem 6.1 is algorithm-specific. We now give algorithm-uniform certificates. The first bounds the information of any algorithm; the second bounds the ghost probability of success of any algorithm.

Information supersolution. For a function F on states, define

$$(\mathbf{C}_t F)(s) := \sup_{p \in \Delta(\mathcal{A}(s))} \left\{ \mathcal{I}_\chi(s, p) + \mathbb{E}_{a \sim p, o \sim P_{s,a}} F(\tau(s, a, o)) \right\}.$$

A sequence $\bar{C}_t$ is an information supersolution if

$$\bar{C}_{T+1} \geq 0, \quad \bar{C}_t(s) \geq (\mathbf{C}_t \bar{C}_{t+1})(s) \quad \forall t, s.$$

Then, for every algorithm starting from s_1 ,

$$C_\chi^{\text{Alg}}(\mu) \leq \bar{C}_1(s_1). \tag{54}$$

The proof is the same backward induction as Proposition 2.2.

Ghost-good-mass supersolution. Let $m \in \Delta(\mathcal{Y})$ be the target-index distribution used to evaluate success, let s be the posterior-reference state, and let $g : \mathcal{Y} \rightarrow [0, L]$ be an accumulated average-loss profile. For threshold r , define the exact recursion

$$\Gamma_{T+1}^r(m, s, g) = m\{y : g(y) \leq r\},$$

$$\Gamma_t^r(m, s, g) = \sup_{p \in \Delta(\mathcal{A}(s))} \mathbb{E}_{a \sim p, o \sim P_{s,a}} \Gamma_{t+1}^r \left(m, \tau(s, a, o), g + \frac{1}{T} \ell^X(s, a) \right),$$

where $\ell^X(s, a)$ is the function $y \mapsto \ell_y^X(s, a)$. If $\bar{\Gamma}_t^r$ is a supersolution of this recursion, then for every algorithm,

$$p_r^{\text{Alg}}(\mu, \chi) \leq \bar{\Gamma}_1^r(\mu_\chi, s_1, 0). \quad (55)$$

Define the ghost entropy lower certificate

$$\underline{\mathcal{E}}_1^r(\mu, \chi) := -\log \bar{\Gamma}_1^r(\mu_\chi, s_1, 0). \quad (56)$$

Theorem 6.3 (Bellman quantile-index lower certificate). *Suppose $\bar{C}_t$ is an information supersolution and $\bar{\Gamma}_t^r$ is a ghost-good-mass supersolution. Then*

$$\inf_{\text{Alg}} \mathbb{E}_{\Omega \sim \mu} \mathbb{E}_{\mathbb{P}_\Omega^{\text{Alg}}} L_\Omega(H_T) \geq Tr \left[1 - \psi(e^{-\underline{\mathcal{E}}_1^r(\mu, \chi)}, \bar{C}_1(s_1)) \right]. \quad (57)$$

In particular, if

$$\bar{C}_1(s_1) + \log 2 \leq \frac{1}{2} \underline{\mathcal{E}}_1^r(\mu, \chi), \quad (58)$$

then

$$\inf_{\text{Alg}} \mathbb{E}_{\Omega \sim \mu} \mathbb{E}_{\mathbb{P}_\Omega^{\text{Alg}}} L_\Omega(H_T) \geq \frac{Tr}{2}. \quad (59)$$

Consequently,

$$\mathfrak{R}_T^* \geq \sup_{\mu, \chi, r} \left\{ \frac{Tr}{2} : \bar{C}_1(s_1) + \log 2 \leq \frac{1}{2} \underline{\mathcal{E}}_1^r(\mu, \chi) \right\}.$$

Proof. For every algorithm, (54) gives $T\bar{C}_\chi^{\text{Alg}}(\mu) = C_\chi^{\text{Alg}}(\mu) \leq \bar{C}_1(s_1)$, and (55) gives $p_r^{\text{Alg}}(\mu, \chi) \leq e^{-\underline{\mathcal{E}}_1^r(\mu, \chi)}$. Substitute these two inequalities into Theorem 6.1, using the monotonicity of ψ in both arguments. The minimax statement follows from Yao's principle. $\square$

The critical average-regret radius is

$$r_*(\mu, \chi) := \sup \left\{ r : \bar{C}_1(s_1) + \log 2 \leq \frac{1}{2} \underline{\mathcal{E}}_1^r(\mu, \chi) \right\}. \quad (60)$$

The lower bound is order Tr_* . A two-point Le Cam proof corresponds to the special case in which $\underline{\mathcal{E}} = O(1)$. A Fano or local-prior-mass lower bound has $\underline{\mathcal{E}}$ of order dimension or effective dimension.

6.4 DEC as one-step relaxation

As discussed in Section 5, offset DEC is obtained from the fixed-truth log-potential bracket (26) by making a one-step relaxation: the coordinate potential drop, or its posterior-averaged Bellman continuation counterpart, is replaced by a static one-step separation penalty, and the posterior-reference law is replaced by a fixed one-step reference predictive law.

In the model-index case $\chi(\omega) = \omega$, let $\bar{\nu}$ be a reference belief at state s , and write

$$P_{\bar{\nu}, s, a} := \int P_{\omega, s, a} \bar{\nu}(d\omega)$$

for its predictive law. The corresponding pointwise KL offset is

$$\text{dec}_\gamma^{\text{KL}}(s; \bar{\nu}) := \inf_{p \in \Delta(\mathcal{A}_s)} \sup_{\omega \in \Omega_s} \{ \mathbb{E}_{a \sim p} \ell_\omega(s, a) - \gamma \mathbb{E}_{a \sim p} D_{\text{KL}}(P_{\omega, s, a} \| P_{\bar{\nu}, s, a}) \}. \quad (\text{one-step model-index offset})$$

Thus the usual offset DEC is a worst-case, static-reference relaxation of the Bellman program: the learner chooses the one-step decision distribution p , and the local adversary then chooses the hardest model relative to the fixed reference predictive law.

The same one-step logic also appears on the lower-bound side, but only after an additional reduction from a dynamic posterior-reference experiment to a one-step reference experiment. Theorem 6.1 and Theorem 6.3 retain the dynamic reference posteriors. A one-step quantile DEC is obtained only by relaxing this dynamic proof: one fixes a reference model, or more generally a reference predictive law, decomposes the trajectory divergence into one-step divergences, and replaces the algorithm-dependent ghost history by a one-round decision distribution. In Chen et al. (2024, Section 3.2.2), this idea is made explicit through the quantile PAC-DEC. A reference model $\bar{M}$ generates the ghost experiment, an adversarial model M supplies the hard alternative, and the loss certificate is a quantile of the one-step risk rather than its expectation. The further reduction to constrained regret-DEC in that paper is precisely where localization, regularity, and problem-specific assumptions enter.

The passage from the dynamic theorem to a one-step coefficient is therefore not merely notational. On both the upper- and lower-bound sides, existing DEC measures can be viewed as one-step relaxations obtained by freezing the posterior-reference process into reference models or predictive laws. This interpretation is primarily conceptual: even the distinction between quantile and expected loss raises nontrivial technical questions when deriving regret-DEC lower bounds under regularity conditions. Nevertheless, the present paper adopts this viewpoint because it clarifies how one-step DEC certificates arise from the dynamic indexed Bellman picture.

Thus DEC/E2D should be read as a one-step relaxation of indexed Bellman information complexity, rather than as a universally tight conversion coefficient from information to regret. When the posterior geometry is simple, this relaxation can be sharp. In more structured problems, however, a one-round coefficient can obscure the dynamic object in two distinct ways: it may charge an estimation or entropy cost that the posterior-reference lower bound has already paid, or it may discard the geometry of the reference trajectory that enables localization. The examples in Section 7 illustrate indexed-Bellman certificates and diagnostics; the following informal diagnostics explain why the corresponding one-round DEC accounting should be interpreted with care.

- *Vacuity for RKHS classes without localization.* In infinite-dimensional kernel bandits, a global one-round DEC over an RKHS ball can be unbounded or vacuous unless the comparison is localized to the finite active experiment actually visited by the algorithm; see the end of Section 7.1 for details. Known routes to matching DEC-style lower bounds introduce this localization explicitly or impose regularity conditions that enforce it implicitly through a critical-radius or sub-root fixed point (Foster et al., 2023). Indexed Bellman information complexity keeps this localization inside the posterior-reference process itself.
- *Double counting in multi-armed bandits.* The hard prior and the ghost entropy already contain the cost of identifying one arm among K . Therefore, after a Fano or posterior-reference calculation has produced the logarithmic entropy term, appending an independent DEC “estimation” penalty charges the same information budget a second time. This is not a statement that DEC-based algorithms are necessarily suboptimal for multi-armed bandits; it is a statement that the entropy appearing in the lower certificate and the estimation term appearing in a one-step upper relaxation should not be added as independent costs.
- *Collapsed lower bounds in linear bandits.* A one-round coefficient obtained by collapsing the adaptive trajectory to a single action distribution can leave a localization or logarithmic gap in the lower bound. The gap is not intrinsic to the linear-bandit experiment: it is introduced

by replacing the posterior trajectory by a static action law. The indexed Bellman formulation keeps the posterior covariance, the information telescope, and the ghost quantile on the same reference path.

- *Representational ambiguity for Bellman rank.* Bellman rank is a representation-relative algebraic certificate: it depends on a hypothesis class, value-function class, roll-in distribution, discrepancy, and Bellman-error witness (Jiang et al., 2017). Without fixing these observable witnesses, bounded Bellman rank is not itself a DMSO model class and does not determine an intrinsic DEC. Once a bilinear-class witness and its estimation complexity are specified, DEC-style questions become meaningful (Du et al., 2021). Thus, a lower bound requires an explicit hard family, observable witnesses, and a ghost-entropy calculation within the specified class.

Definition 6.4 (Tight indexed Bellman representation). *Fix a problem family, a horizon sequence, and a class of admissible indices χ . An indexed Bellman representation is tight if, for the relevant initial states s_1 and for some prior-index pair (μ, χ) , there exist a regret supersolution $\bar{V}$ and Bellman-Fano lower certificates $\bar{C}, \underline{\mathcal{E}}$ on the same state space such that, up to the stated universal constants or logarithmic factors,*

$$\bar{V}_1(s_1) \lesssim Tr_\star(\mu, \chi), \quad \mathfrak{R}_T^\star \gtrsim Tr_\star(\mu, \chi),$$

where $r_\star(\mu, \chi)$ is the critical average-regret radius in (60).

This definition is the representation-level replacement for asking whether a one-round DEC coefficient is universally tight. The question is not whether a single static coefficient is always exact, but whether the same Bellman-sufficient state supports both an upper supersolution and a posterior-reference quantile lower certificate at the same critical radius.

7 Applications and extentions

7.1 Kernel bandits: indexed Bellman DP, GP-UCB, GP-E2D, and AMS/EBO

This subsection specializes the indexed Bellman language to kernel bandits on a finite active action set. The point is not to make the unknown function finite-dimensional. The truth remains an element of an RKHS. What becomes finite is the marginal experiment observed through the active actions. This finite marginal is enough to define the GP posterior update, the model-index MAIR information, and the action-index AIR belief over the optimal active action. We present the exact finite-marginal indexed Bellman program first, and then compare GP-UCB, GP-E2D, and AMS/EBO as tractable relaxations or certificates of its bracket.

Problem and finite active marginal. Let $\mathcal{H}_k$ be a reproducing kernel Hilbert space (RKHS) on a possibly infinite domain $\mathcal{X}_{\text{all}}$, with kernel k satisfying $k(x, x) \leq \kappa^2$. The learner is evaluated on a finite active set

$$\mathcal{X} = \{x_1, \dots, x_n\} \subset \mathcal{X}_{\text{all}}.$$

The unknown reward function $f^\star \in \mathcal{H}_k$ satisfies $\|f^\star\|_{\mathcal{H}_k} \leq B$, and at round t the learner chooses $X_t \in \mathcal{X}$ and observes

$$Y_t = f^\star(X_t) + \varepsilon_t, \quad \mathbb{E}[\varepsilon_t \mid H_{t-1}, X_t] = 0, \quad (61)$$

where the noise is conditionally R -sub-Gaussian. For the Gaussian calculations below we use the algorithmic reference model with variance parameter $\lambda > 0$; the frequentist confidence event is obtained by the usual self-normalized argument.

Let $F = (f(x_1), \dots, f(x_n)) \in \mathbb{R}^n$ be the active reward vector and let $K_{\mathcal{X}}$ be the $n \times n$ kernel matrix. The GP reference prior is

$$F \sim N(0, K_{\mathcal{X}}), \quad Y_t \mid F, X_t = x_i \sim N(F_i, \lambda).$$

After history H_{t-1} , the algorithmic posterior on F is Gaussian,

$$F \mid H_{t-1} \sim N(m_t, \Sigma_t).$$

For $x_i \in \mathcal{X}$, write $m_t(x_i) = e_i^\top m_t$ and $s_t^2(x_i) = e_i^\top \Sigma_t e_i$. If $X_t = x_i$, the exact posterior update is

$$m_{t+1} = m_t + \frac{\Sigma_t e_i}{\lambda + s_t^2(x_i)} \{Y_t - m_t(x_i)\}, \quad (62)$$

$$\Sigma_{t+1} = \Sigma_t - \frac{\Sigma_t e_i e_i^\top \Sigma_t}{\lambda + s_t^2(x_i)}. \quad (63)$$

This is a finite marginal posterior update. It should not be read as an assumption that f^* has n parameters. The RKHS function may be infinite-dimensional; only the observations and decisions in this finite experiment depend on the vector F .

Model-index information and action-index information. For the model-index specialization, take the index to be the active reward vector $Y_{\text{MAIR}} = F$. The one-step GP/MAIR information at state H_{t-1} is

$$\mathcal{I}_t^{\text{GP}}(p) := \mathbb{E}_{x \sim p} \frac{1}{2} \log \left(1 + \frac{s_t^2(x)}{\lambda} \right). \quad (64)$$

Along a realized action sequence,

$$C_T^{\text{GP}} := \sum_{t=1}^T \mathcal{I}_t^{\text{GP}}(\delta_{X_t}) = \frac{1}{2} \log \det(I + \lambda^{-1} K_{X_{1:T}, X_{1:T}}), \quad (65)$$

where repeated actions are included in the Gram matrix. This is the realized information gain of the algorithmic trajectory. It can be much smaller than the maximal information gain that maximizes over all length- T designs.

For the action-index specialization, assume ties are broken by a fixed rule and set

$$A^*(F) \in \arg \max_{x \in \mathcal{X}} F_x, \quad q_t(a) := \mathbb{P}(A^* = a \mid H_{t-1}).$$

Let ν_t^a be the exact conditional law of F given H_{t-1} and $A^* = a$. It is generally a truncated Gaussian on the cone where a is optimal. Define the conditional predictive law

$$P_{t,a,x}(\cdot) := \int N(F_x, \lambda)(\cdot) \nu_t^a(dF),$$

for the observation at action x given the event $A^* = a$. The marginal predictive is

$$P_{t,x} = \sum_{a \in \mathcal{X}} q_t(a) P_{t,a,x} = N(m_t(x), \lambda + s_t^2(x)).$$

The AIR information increment is

$$\mathcal{I}_t^{\text{AIR}}(p) := \mathbb{E}_{x \sim p} \sum_{a \in \mathcal{X}} q_t(a) D_{\text{KL}}(P_{t,a,x} \| P_{t,x}). \quad (66)$$

The conditional AIR regret is

$$\Delta_t^{\text{AIR}}(p) := \sum_{a \in \mathcal{X}} q_t(a) \mathbb{E}_{F \sim \nu_t^a} \mathbb{E}_{x \sim p} [F_a - F_x]. \quad (67)$$

The exact indexed chain rule gives

$$\sum_{t=1}^T \mathbb{E} \mathcal{I}_t^{\text{AIR}}(p_t) = I(A^*; H_T) \leq H(A^*) \leq \log n.$$

Thus AIR places the entropy telescope on the decision index rather than on the ambient RKHS. The price is that the coefficient converting action-index information into regret must be controlled by an algorithmic bracket.

Algorithm 1: finite-marginal AIR/MAIR Bellman DP. On the finite active marginal, the exact frequentist benchmark is the log-penalized Bellman program of Theorem 5.2. The state contains (m_t, Σ_t, q_t) , where (m_t, Σ_t) is the algorithmic GP marginal on F and q_t is the maintained reference law of the index, either $A^*(F)$ for AIR or F for MAIR. Let $\mathfrak{C}_t(m, \Sigma, q)$ be a calibrated local set of candidate active reward vectors containing the fixed truth F^* on the calibration event. In the action-index version define

$$y(f) = A^*(f), \quad \mathbf{G}_t^{\text{AIR}}(p; f) := \mathbb{E}_{x \sim p, O \sim N(f_x, \lambda)} \log \frac{q^+(y(f) \mid m, \Sigma, q, x, O)}{q(y(f))},$$

where q^+ is the chosen reference update for the optimal-action marginal. The robust finite-marginal Bellman recursion is

$$W_t^\gamma(m, \Sigma, q) = \inf_{p \in \Delta(\mathcal{X})} \sup_{f \in \mathfrak{C}_t(m, \Sigma, q)} \{ \Delta_f(p) + \mathbb{E}_f W_{t+1}^\gamma(m^+, \Sigma^+, q^+) - \gamma \mathbf{G}_t^{\text{AIR}}(p; f) \}.$$

The model-index MAIR version replaces $y(f)$ by the finite reward vector f and uses the corresponding model-coordinate posterior log gain. By Theorem 5.2, an approximate minimizer of this recursion has regret at most

$$W_1^\gamma(m_1, \Sigma_1, q_1) + \gamma \log \frac{1}{q_1(A^*(F^*))} + \text{optimization and calibration errors.}$$

If the maintained belief is the exact Bayesian posterior and one averages this coordinate recursion over F , the robust supremum is replaced by posterior expectation and the dynamic program reduces to the posterior-averaged AIR recursion

$$U_t^\gamma(m, \Sigma, q) = \inf_{p \in \Delta(\mathcal{X})} \{ \Delta_t^{\text{AIR}}(p) + \mathbb{E} U_{t+1}^\gamma(m^+, \Sigma^+, q^+) - \gamma \mathcal{I}_t^{\text{AIR}}(p) \}.$$

Thus the finite-marginal Bellman DP is the clean upper object for comparison with the lower theorem: the frequentist version uses the coordinate log potential and calibrated candidate truths, while the Bayesian version is its posterior average. The following three rules are tractable certificates or relaxations of this benchmark.

Algorithm 2: GP-UCB. GP-UCB uses the Gaussian marginal posterior (m_t, Σ_t) on the finite active marginal. In this paragraph, e_x denotes the coordinate vector associated with $x \in \mathcal{X}$, and $s_t(x) = (e_x^\top \Sigma_t e_x)^{1/2}$ denotes the unmultiplied posterior standard deviation, i.e., the generic uncertainty scale $\sigma_t(x)$ from Lemma 5.5. The confidence half-width is

$$w_t(x) := \beta_t s_t(x),$$

where β_t is the confidence multiplier. The algorithm chooses

$$X_t \in \arg \max_{x \in \mathcal{X}} \{m_t(x) + \beta_t s_t(x)\}. \quad (68)$$

The following event is the frequentist calibration certificate:

$$\mathcal{E}_{\text{GP}} := \{|f^*(x) - m_t(x)| \leq w_t(x) = \beta_t s_t(x) \text{ for all } t \leq T, x \in \mathcal{X}\}. \quad (69)$$

For finite $\mathcal{X}$, standard RKHS self-normalized concentration gives $\mathbb{P}(\mathcal{E}_{\text{GP}}) \geq 1 - \delta$ for the usual choice of β_t . For example, up to universal constants one may take

$$\beta_t \asymp B + R \sqrt{\Gamma_{t-1}^{\text{GP}} + \log(1/\delta)}, \quad \Gamma_m^{\text{GP}} := \sup_{x_{1:m} \in \mathcal{X}^m} \frac{1}{2} \log \det(I + \lambda^{-1} K_{x_{1:m}, x_{1:m}}),$$

with the precise form depending on the normalization of k and the ridge parameter. Here $K_{x_{1:m}, x_{1:m}}$ is the Gram matrix of the possibly repeated design sequence $(x_1, \dots, x_m)$. The quantity Γ_m^{GP} is a worst-case calibration device for the confidence event; it is different from the realized information gain C_T^{GP} paid in the regret bound. The next theorem isolates the deterministic part of the GP-UCB proof.

Theorem 7.1 (GP-UCB realized-information bound). *On the event $\mathcal{E}_{\text{GP}}$, the GP-UCB rule (68) satisfies*

$$\text{Reg}_T(f^*) \leq 2\beta_T \sum_{t=1}^T s_t(X_t) \leq 2\beta_T \sqrt{T c_{\kappa, \lambda} C_T^{\text{GP}}}, \quad (70)$$

where

$$c_{\kappa, \lambda} := \frac{2\kappa^2}{\log(1 + \kappa^2/\lambda)}$$

and C_T^{GP} is the realized information gain in (65). Consequently, if rewards are bounded in $[0, 1]$ and $\mathbb{P}(\mathcal{E}_{\text{GP}}) \geq 1 - \delta$, then

$$\mathbb{E} \text{Reg}_T(f^*) \leq 2\beta_T \sqrt{T c_{\kappa, \lambda} \mathbb{E} C_T^{\text{GP}}} + T\delta.$$

Proof. Let $x^* \in \arg \max_{x \in \mathcal{X}} f^*(x)$. On $\mathcal{E}_{\text{GP}}$, calibration at x^* and the optimistic choice of X_t give

$$f^*(x^*) \leq m_t(x^*) + w_t(x^*) \leq m_t(X_t) + w_t(X_t).$$

Calibration at X_t gives $m_t(X_t) \leq f^*(X_t) + w_t(X_t)$. Hence

$$f^*(x^*) - f^*(X_t) \leq 2w_t(X_t) = 2\beta_t s_t(X_t) \leq 2\beta_T s_t(X_t).$$

Summing and applying Cauchy-Schwarz yields

$$\text{Reg}_T(f^*) \leq 2\beta_T \sqrt{T \sum_{t=1}^T s_t^2(X_t)}.$$

For $u \in [0, \kappa^2/\lambda]$, monotonicity of $u/\log(1+u)$ on a bounded interval gives

$$\lambda u \leq \frac{\kappa^2}{\log(1 + \kappa^2/\lambda)} \log(1 + u).$$

Applying this with $u = s_t^2(X_t)/\lambda$ and using (65) proves (70). The expectation bound adds the trivial regret bound T on $\mathcal{E}_{\text{GP}}^c$ and uses Jensen's inequality. $\square$

GP–UCB as a MAIR decomposition. The theorem is Lemma 5.5 with the generic uncertainty scale $\sigma_t(X_t) = s_t(X_t)$ and the realized-information log determinant telescope

$$\sum_t \frac{1}{2} \log(1 + s_t^2(X_t)/\lambda) = C_T^{\text{GP}}.$$

The confidence half-width is $w_t(X_t) = \beta_t s_t(X_t)$, but the multiplier β_t is not part of the information telescope; it is the coefficient paid for calibration and optimism in the fixed-truth bracket. This separation is the reason the proof uses both s_t and w_t : s_t^2 is summed by the log-determinant/elliptical-potential certificate, whereas w_t is used only to upper bound instantaneous regret. GP–UCB is often practically strong because it follows the realized trajectory and pays the realized information gain C_T^{GP} , not the worst possible design information Γ_T^{GP} , even though its confidence multiplier may be conservative in step-uniform minimax analyses.

Algorithm 3: GP–E2D. A GP–E2D rule keeps the same GP posterior reference but chooses a distribution over actions by solving a localized one-round MAIR/DEC offset. Given a localized comparison set $\mathcal{F}_t$ containing plausible functions, define

$$\Delta_f(p) := \max_{x \in \mathcal{X}} f(x) - \mathbb{E}_{x \sim p} f(x),$$

and let $P_{f,x} = N(f(x), \lambda)$ while $P_{t,x} = N(m_t(x), \lambda + s_t^2(x))$ is the posterior predictive. A schematic GP–E2D decision is

$$p_t \in \arg \min_{p \in \Delta(\mathcal{X})} \sup_{f \in \mathcal{F}_t} \{\Delta_f(p) - \gamma_t \mathbb{E}_{x \sim p} D_{\text{KL}}(P_{f,x} \| P_{t,x})\}. \quad (71)$$

If $f^* \in \mathcal{F}_t$ and the optimized value in (71) is at most ε_t , then

$$\Delta_{f^*}(p_t) \leq \gamma_t \mathbb{E}_{x \sim p_t} D_{\text{KL}}(P_{f^*,x} \| P_{t,x}) + \varepsilon_t. \quad (72)$$

Summing (72) gives a regret bound once the fixed-truth KL terms are related to a posterior-ratio or information telescope. Under a well-specified Bayesian GP reference this relation holds in expectation by the MAIR chain rule. Under a fixed frequentist truth it requires calibration/localization. This is the technical point that a GP–E2D analysis must address: the one-round DEC is optimized, but the compatibility between that coefficient and the realized GP log determinant is not automatic. An unlocalized offset over an infinite-dimensional RKHS ball can be unbounded or vacuous; meaningful GP–E2D guarantees must localize the comparison set to the posterior-reference experiment.

Algorithm 4: AMS/EBO on the action index. AMS/EBO attacks the same specialized log-potential bracket from the robust posterior side. In the finite-action kernel marginal, the natural AIR index is A^* . Let q_t be the current reference distribution on A^* , let $\mathfrak{B}_t(q_t)$ be a calibrated local set of candidate beliefs ν on the active reward vector F , and fix a constant coefficient $\gamma > 0$ (equivalently $\eta = 1/\gamma$). For each ν , define $\Delta_\nu^{\text{AIR}}(p)$ and $\mathcal{I}_\nu^{\text{AIR}}(p)$ by the analogues of (67) and (66), using the conditional laws induced by ν . A robust AIR/EBO decision is

$$p_t \in \arg \min_{p \in \Delta(\mathcal{X})} \sup_{\nu \in \mathfrak{B}_t(q_t)} \left\{ \Delta_\nu^{\text{AIR}}(p) - \gamma \mathcal{I}_\nu^{\text{AIR}}(p) - \gamma D_{\text{KL}}(\nu_{A^*} \| q_t) \right\}. \quad (73)$$

The maximization over ν is essential: without it the display is only the posterior AIR bracket for a chosen belief, not the robust AMS/EBO relaxation. In a finite-index abstraction, the last two terms are the KL-dual form of the log-partition potential (31). If the robust value in (73) is at most ε_t , the required concavity/first-order conditions in Lemma 5.6 hold, and the belief set is calibrated so that the fixed truth f^* is an admissible comparison direction up to error cal_T , then Theorem 5.1 gives

$$\mathbb{E}_{f^*} \text{Reg}_T \leq \gamma \log \frac{1}{q_1(a^*)} + \mathbb{E}_{f^*} \sum_{t=1}^T \varepsilon_t + \text{cal}_T, \quad (74)$$

where $a^* = A^*(f^*)$. The constant coefficient avoids the false shortcut of replacing a weighted fixed-truth log-gain sum by a single telescope. If a varying coefficient is used, the corresponding weighted telescope must be proved separately. Under a Bayesian reference prior, posterior averaging of the same coordinate bound gives

$$\mathbb{E} \text{Reg}_T \leq \gamma H(A^*) + \mathbb{E} \sum_{t=1}^T \varepsilon_t + \mathbb{E} \text{cal}_T.$$

A ratio form gives the familiar Cauchy–Schwarz variant after replacing the fixed coefficient by the appropriate sequential AIR ratio and applying the action-index chain rule. Since $H(A^*) \leq \log n$, this route scales with the finite decision index and the sequential AIR coefficient. It is different from GP–E2D: AMS/EBO uses a robust action-index belief trajectory, whereas GP–E2D optimizes a one-round model-index offset whose relationship to the realized log determinant or posterior-ratio telescope must be justified separately.

Position relative to the GP–UCB optimality debate. The original GP–UCB analysis is practical and trajectory-sensitive because it pays the realized information gain (65), or a worst-case relaxation only at the final step (Srinivas et al., 2010). Later work asks whether the remaining suboptimality comes from the algorithm or from the confidence analysis; in particular, Vakili et al. (2021b) ask for sharper online RKHS confidence intervals and sharper finite-action kernel-bandit guarantees. The indexed view separates the issue into four mechanisms. The finite-marginal indexed Bellman DP is the exact information-potential benchmark, but may be computationally demanding. GP–UCB has the strongest realized-information geometry but may have a conservative MAIR coefficient. AMS/EBO can keep the action-index AIR telescope $I(A^*; H_T) \leq \log n$ and therefore has a principled route to finite-decision complexity, provided the robust bracket can be controlled. GP–E2D directly optimizes a one-round DEC coefficient, but an unlocalized DEC can be vacuous and a localized DEC must be matched to the same posterior-reference trajectory. Thus these algorithms should not be compared by a single universal coefficient; they are one exact Bellman benchmark and three relaxations or certificates of the same indexed Bellman identity, with different estimation and coefficient mechanisms.

Why global RKHS DEC can be vacuous. Consider a Gaussian RKHS bandit with noise variance σ^2 , feature map $\phi(x) \in \mathcal{H}$ satisfying $\|\phi(x)\|_{\mathcal{H}} \leq 1$, and suppose the action set contains points $\{x_j : j \geq 1\}$ with $\phi(x_j) = e_j$, where $\{e_j\}_{j \geq 1}$ is orthonormal in $\mathcal{H}$. Let the reference model be $f_0 \equiv 0$, and let $\mathcal{F}_B = \{f \in \mathcal{H} : \|f\|_{\mathcal{H}} \leq B\}$. For any one-round decision distribution p , define

$$q_j := \mathbb{E}_{x \sim p} \langle e_j, \phi(x) \rangle_{\mathcal{H}}^2.$$

Since

$$\sum_{j \geq 1} q_j \leq \mathbb{E}_{x \sim p} \|\phi(x)\|_{\mathcal{H}}^2 \leq 1,$$

there are directions with arbitrarily small q_j . For the alternative $f_j = B e_j$, the optimal action is x_j , while

$$\mathbb{E}_{x \sim p} \Delta_{f_j}(x) = B - B \mathbb{E}_{x \sim p} \langle e_j, \phi(x) \rangle_{\mathcal{H}} \geq B(1 - \sqrt{q_j}),$$

and the one-round KL information against f_0 is

$$\mathbb{E}_{x \sim p} D_{\text{KL}}(N(f_j(x), \sigma^2) \parallel N(0, \sigma^2)) = \frac{B^2}{2\sigma^2} q_j.$$

Hence, for every finite offset coefficient γ , one may choose j with q_j so small that

$$\mathbb{E}_{x \sim p} \Delta_{f_j}(x) - \gamma \mathbb{E}_{x \sim p} D_{\text{KL}}(N(f_j(x), \sigma^2) \parallel N(0, \sigma^2)) \geq cB$$

for a numerical constant $c > 0$. The constrained version has the same pathology: for any fixed information radius, a sufficiently rare direction has information inside the radius while retaining regret of order B . Thus an unlocalized one-round DEC over the entire RKHS ball need not reflect the effective difficulty of the finite experiment being run. The meaningful object is the finite active marginal together with its posterior-reference trajectory. Indexed Bellman information complexity keeps this localization inside the state process, rather than appending a separate DEC localization after the fact.

GP–E2D offset optimization should be localized. The model-index offset in (71) is meaningful only relative to a comparison set and a reference law. If the comparison set is the entire infinite-dimensional RKHS ball, then a one-round adversary can place mass in directions receiving arbitrarily small action probability. The resulting regret remains order one while the one-step information is arbitrarily small, making the global offset vacuous. If the comparison set is instead the finite active posterior neighborhood, the offset may be finite and useful. The point is that the localization is part of the representation and posterior-reference trajectory; it should not be appended as an independent estimation term after a separate log-determinant analysis.

7.2 MAB lower bound via optimal-action index and reference history

In this subsection, we apply our lower-bound approach to multi-armed bandits (MAB). Consider $K \geq 3$ Gaussian arms with unit noise variance. Let the ambient model space be a class $\Omega_{\text{MAB}} \subseteq \mathbb{R}^K$ of mean vectors, and write a model environment as $\omega = (\theta_{\omega}(1), \dots, \theta_{\omega}(K))$. Pulling arm a in environment ω produces an observation distributed as $N(\theta_{\omega}(a), 1)$. Let

$$A^*(\omega) \in \arg \max_{a \in [K]} \theta_{\omega}(a)$$

with an arbitrary fixed tie-breaking rule, and define the information index

$$\chi(\omega) = A^*(\omega) \in [K].$$

Thus the random environment is the full mean vector Ω , while the information target used in the lower certificate is only

$$Y = \chi(\Omega) = A^*(\Omega).$$

For the lower bound, fix $\Delta > 0$ and use the K -point prior μ_Δ supported on the spike environments

$$\omega^j = \Delta e_j, \quad j \in [K],$$

where e_j is the j th coordinate vector. Equivalently, draw $J \sim \text{Unif}([K])$ and set $\Omega = \omega^J$. On this prior support the optimal arm is unique and

$$Y = \chi(\Omega) = J.$$

The equality $Y = J$ is only a property of this least-favorable prior; it should not be read as identifying the information target with the full model environment. The one-step regret in environment ω^j is

$$\ell_{\omega^j}(a) = \max_b \theta_{\omega^j}(b) - \theta_{\omega^j}(a) = \Delta \mathbb{1}\{a \neq j\}.$$

For brevity write $\ell_j = \ell_{\omega^j}$ on this subfamily.

Lemma 7.2 (MAB ghost entropy). *Let $Y = \chi(\Omega)$ under the prior $\Omega \sim \mu_\Delta$. For every algorithm and every reference history H'_T independent of $Y \sim \text{Unif}([K])$,*

$$\mathbb{P}\left(\bar{L}_Y(H'_T) \leq \frac{\Delta}{2}\right) \leq \frac{2}{K},$$

where $\bar{L}_y(H'_T) = T^{-1} \sum_{t=1}^T \ell_y(A'_t)$ and A'_t denotes the arm selected in the reference history at time t .

Proof. For a fixed reference action sequence $a'_1, \dots, a'_T$, let

$$N'_y = \sum_{t=1}^T \mathbb{1}\{a'_t = y\}.$$

The average regret against the spike environment ω^y is

$$\bar{L}_y(H'_T) = \Delta(1 - N'_y/T).$$

The event $\bar{L}_y(H'_T) \leq \Delta/2$ implies $N'_y \geq T/2$. Since Y is uniform and independent of the reference sequence,

$$\mathbb{E}[N'_Y | H'_T] = \frac{1}{K} \sum_{y=1}^K N'_y = \frac{T}{K}.$$

Markov's inequality gives

$$\mathbb{P}(N'_Y \geq T/2 | H'_T) \leq \frac{2}{K}.$$

Averaging over H'_T proves the claim. $\square$

Lemma 7.3 (MAB information bound). *There is a universal constant $c_0 > 0$ such that if $\Delta^2 T/K \leq c_0$, then every adaptive algorithm satisfies*

$$I(Y; H_T) \leq \frac{1}{4} \log(K/2),$$

where $Y = \chi(\Omega)$ and $\Omega \sim \mu_\Delta$.

Proof. Let $\mathbb{P}_0$ be the law of the history under the zero-mean reference model, let $\mathbb{P}_j$ be the law under the spike environment ω^j , and let

$$\bar{\mathbb{P}} = \frac{1}{K} \sum_{j=1}^K \mathbb{P}_j$$

be the marginal law of H_T under the prior μ_Δ . Under this prior, $Y = j$ is equivalent to $\Omega = \omega^j$, but the mutual information is still the action-index information $I(Y; H_T)$, not information about an ambient full mean vector outside the prior support. Since $\bar{\mathbb{P}} \geq K^{-1} \mathbb{P}_j$, the likelihood ratio $d\mathbb{P}_j/d\bar{\mathbb{P}}$ is bounded by K . The inequality

$$D_{\text{KL}}(P\|Q) \leq (2 + \log K)H^2(P, Q) \quad \text{whenever } Q \geq K^{-1}P,$$

together with the triangle inequality and convexity for squared Hellinger distance, gives

$$I(Y; H_T) = \frac{1}{K} \sum_{j=1}^K D_{\text{KL}}(\mathbb{P}_j\|\bar{\mathbb{P}}) \leq 4(2 + \log K) \frac{1}{K} \sum_{j=1}^K H^2(\mathbb{P}_j, \mathbb{P}_0).$$

Using $H^2(P, Q) \leq D_{\text{KL}}(Q\|P)$ and the adaptive Gaussian KL chain rule under the zero model,

$$\frac{1}{K} \sum_{j=1}^K H^2(\mathbb{P}_j, \mathbb{P}_0) \leq \frac{1}{K} \sum_{j=1}^K D_{\text{KL}}(\mathbb{P}_0\|\mathbb{P}_j) = \frac{\Delta^2}{2K} \mathbb{E}_0 \sum_{t=1}^T \sum_{j=1}^K \mathbb{1}\{A_t = j\} = \frac{\Delta^2 T}{2K}.$$

Therefore

$$I(Y; H_T) \leq 2(2 + \log K) \frac{\Delta^2 T}{K}.$$

Choosing, for example,

$$c_0 \leq \inf_{K \geq 3} \frac{\log(K/2)}{8(2 + \log K)}$$

which is a strictly positive universal constant, proves the displayed bound. $\square$

Theorem 7.4 (MAB lower bound). *For Gaussian K -armed bandits with unit noise variance, over any ambient mean-vector class Ω_{MAB} that contains the spike instances $\{\Delta e_j : j \in [K]\}$ used below,*

$$\mathfrak{R}_T^* \geq c\sqrt{KT}$$

for a universal constant $c > 0$ whenever $T \geq K \geq 3$. In particular, this applies to any bounded class such as $[0, 1]^K$ once the universal constant in the choice of Δ is small enough.

Proof. Choose

$$\Delta = c_1 \sqrt{K/T}$$

with c_1 small enough that Lemma 7.3 applies. Let $r = \Delta/2$. Lemma 7.2 gives $p_r \leq 2/K$, while Lemma 7.3 gives

$$T\bar{C} = I(Y; H_T) \leq \frac{1}{4} \log(K/2).$$

For c_1 small enough, condition (53) holds. Theorem 6.1, applied uniformly over algorithms, gives a Bayes regret lower bound under the prior μ_Δ supported on the spike subfamily. Since this prior is supported inside the ambient model class Ω_{MAB} , the minimax regret over Ω_{MAB} is at least this Bayes risk. Hence

$$\mathfrak{R}_T^* \geq \frac{Tr}{2} = \frac{T\Delta}{4} = c\sqrt{KT}$$

for a universal constant $c > 0$. $\square$

On the upper-bound side, the same distinction between the full model environment and the action index is important. An action-index AIR analysis uses $Y = \chi(\Omega) = A^*(\Omega)$ as the information target, rather than the full environment Ω . This avoids charging information for distinctions between environments that do not change the optimal action. A model-index E2D relaxation can pay a K -factor both through model estimation and through the DEC term, whereas the action-index AIR identity in Theorem 4.6 incurs only a $\log K$ information-gain term and pays the K -factor through the AIR coefficient (Xu and Zeevi, 2025). This nearly matches the lower bound.

7.3 Linear bandits lower bound via optimal-action index and reference history

Consider stochastic linear bandits with action set

$$\mathcal{A} = \{a \in \mathbb{R}^d : \|a\|_2 \leq 1\}.$$

Let the ambient model class be the Euclidean unit parameter ball

$$\Omega_{\text{lin}} = \{\omega \in \mathbb{R}^d : \|\omega\|_2 \leq 1\},$$

and regard an environment $\omega \in \Omega_{\text{lin}}$ as the full mean parameter vector. At time t , after choosing $A_t \in \mathcal{A}$, the learner observes

$$O_t = \langle \omega, A_t \rangle + \xi_t, \quad \xi_t \sim N(0, 1),$$

where the noises are independent over time and independent of the learner's external randomization. For $\omega \neq 0$, the unique optimal action over $\mathcal{A}$ is

$$A^*(\omega) = \frac{\omega}{\|\omega\|_2}.$$

At $\omega = 0$, fix any deterministic tie-breaking action; this case is not used by the lower-bound prior. The information index is the optimal action

$$\chi(\omega) = A^*(\omega), \quad Y = \chi(\Omega),$$

not the full environment Ω . Thus two environments that have the same optimal action are not distinguished by the index unless the prior support itself makes them distinguishable.

We prove the standard $\Omega(d\sqrt{T})$ minimax lower bound through the quantile indexed information theorem. The proof uses a hypercube prior together with an adaptive information bound. The role of the hypercube variable below is only to code the finitely many optimal actions in the prior support. It should not be read as replacing the ambient environment Ω by a model label. Note that for a non-Gaussian prior, one cannot condition on the realized adaptive design matrix and then apply the fixed-design Gaussian-channel capacity formula, because the design itself is a statistic of the past observations and may carry information about the unknown environment and hence about Y . Chen et al. (2024, Section 3.3) prove a closely related lower bound by a different route: they use the interactive Fano method with a Gaussian prior and a unit-sphere calculation. Both approaches show that action-indexed information arguments can recover the sharp linear-bandit order, whereas model-indexed relaxations such as DEC can lose factors when they charge for estimating more of the model than is needed to identify the optimal action.

Let

$$V \sim \text{Unif}(\{\pm 1\}^d), \quad \Omega = \Delta V.$$

Assume $\Delta\sqrt{d} \leq 1$, so that this prior is supported on Ω_{lin} . On this prior support,

$$Y = \chi(\Omega) = A^*(\Omega) = \frac{V}{\sqrt{d}}. \quad (75)$$

The map $v \mapsto v/\sqrt{d}$ is one-to-one on $\{\pm 1\}^d$. Hence, for this prior only, V may be used as a code for the action index Y , and

$$I(Y; H_T) = I(V; H_T).$$

For an index value $y = v/\sqrt{d}$, write ℓ_v for the regret under the corresponding environment $\omega = \Delta v$. The one-step regret of action $a \in \mathbb{R}^d$, $\|a\|_2 \leq 1$, is

$$\ell_v(a) = \Delta\sqrt{d} - \Delta\langle v, a \rangle.$$

Assume $T \geq d^2$ and choose $\Delta \leq 1/\sqrt{d}$, so that $\|\Omega\|_2 = \Delta\sqrt{d} \leq 1$.

Lemma 7.5 (Linear ghost entropy). *There is a universal constant $c_{\text{gh}} > 0$ such that every deterministic reference action sequence $a'_1, \dots, a'_T$, with $\|a'_t\|_2 \leq 1$, satisfies*

$$\mathbb{P}_{V \sim \text{Unif}(\{\pm 1\}^d)} \left(\frac{1}{T} \sum_{t=1}^T (\Delta\sqrt{d} - \Delta\langle V, a'_t \rangle) \leq \frac{\Delta\sqrt{d}}{4} \right) \leq \exp(-c_{\text{gh}}d).$$

Consequently, the same bound holds after averaging over any random reference history H'_T whose induced action sequence is independent of the action index Y , equivalently independent of V under the hypercube prior.

Proof. Let

$$\bar{a} := \frac{1}{T} \sum_{t=1}^T a'_t.$$

Since $\|\bar{a}\|_2 \leq T^{-1} \sum_t \|a'_t\|_2 \leq 1$, the event in the display implies

$$\Delta\sqrt{d} - \Delta\langle V, \bar{a} \rangle \leq \frac{\Delta\sqrt{d}}{4},$$

or equivalently,

$$\langle V, \bar{a} \rangle \geq \frac{3\sqrt{d}}{4}.$$

The random variable $\langle V, \bar{a} \rangle = \sum_{i=1}^d V_i \bar{a}_i$ is a centered Rademacher sum with variance proxy $\|\bar{a}\|_2^2 \leq 1$. Hoeffding's inequality therefore gives

$$\mathbb{P} \left(\langle V, \bar{a} \rangle \geq \frac{3\sqrt{d}}{4} \right) \leq \exp \left(-\frac{9d}{32} \right).$$

Thus the first claim holds, for example with $c_{\text{gh}} = 9/32$. If the reference action sequence is random but independent of Y , then because $Y = V/\sqrt{d}$ is a bijective function of V on the prior support, the reference sequence is independent of V . Conditioning on the reference history and applying the deterministic bound gives the second claim. $\square$

Lemma 7.6 (Adaptive hypercube information capacity). *For the hypercube prior above, every possibly randomized adaptive algorithm satisfies*

$$I(Y; H_T) = I(V; H_T) \leq \Delta^2 T,$$

where $H_T = (A_1, O_1, \dots, A_T, O_T)$.

Proof. Since $Y = V/\sqrt{d}$ is a bijective function of V on the hypercube support, $I(Y; H_T) = I(V; H_T)$. It remains to bound $I(V; H_T)$. Let U denote the learner's internal random seed, independent of V . Then

$$I(V; H_T) \leq I(V; H_T, U) = \mathbb{E}_U I(V; H_T \mid U).$$

It is therefore enough to prove the claim after conditioning on U , so that the algorithm is deterministic.

For $v \in \{\pm 1\}^d$, let $\mathbb{P}_v$ denote the trajectory law when $\Omega = \Delta v$. We use the chain rule over coordinates:

$$I(V; H_T) = \sum_{i=1}^d I(V_i; H_T \mid V_1, \dots, V_{i-1}).$$

Fix a coordinate i and a prefix $u = (v_1, \dots, v_{i-1})$. Let $\mathbb{P}_{u,+}$ and $\mathbb{P}_{u,-}$ be the trajectory laws after averaging uniformly over the remaining coordinates $V_{i+1}, \dots, V_d$, conditional respectively on $V_i = +1$ and $V_i = -1$. Then

$$I(V_i; H_T \mid V_{<i} = u) = D_{\text{JS}}(\mathbb{P}_{u,+}, \mathbb{P}_{u,-}),$$

where D_{JS} denotes Jensen–Shannon divergence with equal weights. For any two probability measures P, Q ,

$$D_{\text{JS}}(P, Q) \leq \frac{1}{4} \left(D_{\text{KL}}(P \parallel Q) + D_{\text{KL}}(Q \parallel P) \right).$$

Next couple the two mixtures by using the same suffix. If $w \in \{\pm 1\}^{d-i}$ and

$$v^+(u, w) = (u, +1, w), \quad v^-(u, w) = (u, -1, w),$$

then joint convexity of relative entropy gives

$$D_{\text{KL}}(\mathbb{P}_{u,+} \parallel \mathbb{P}_{u,-}) \leq \mathbb{E}_w D_{\text{KL}}(\mathbb{P}_{v^+(u,w)} \parallel \mathbb{P}_{v^-(u,w)}),$$

and the same argument gives the analogous reverse inequality.

For two fixed vertices v^+ and v^- that differ only in coordinate i , the adaptive KL chain rule gives

$$\begin{aligned} D_{\text{KL}}(\mathbb{P}_{v^+} \parallel \mathbb{P}_{v^-}) &= \frac{1}{2} \mathbb{E}_{v^+} \sum_{t=1}^T \left(\Delta \langle v^+ - v^-, A_t \rangle \right)^2 \\ &= 2\Delta^2 \mathbb{E}_{v^+} \sum_{t=1}^T A_{t,i}^2. \end{aligned}$$

Here the action kernel contributes no KL because, after conditioning on the learner's internal randomness, the algorithm uses the same decision rule under both environments; only the Gaussian observation kernel changes. Similarly,

$$D_{\text{KL}}(\mathbb{P}_{v^-} \parallel \mathbb{P}_{v^+}) = 2\Delta^2 \mathbb{E}_{v^-} \sum_{t=1}^T A_{t,i}^2.$$

Combining the preceding displays yields

$$\begin{aligned} I(V_i; H_T \mid V_{<i} = u) &\leq \frac{\Delta^2}{2} \mathbb{E}_w \left[\mathbb{E}_{v^+(u,w)} \sum_{t=1}^T A_{t,i}^2 + \mathbb{E}_{v^-(u,w)} \sum_{t=1}^T A_{t,i}^2 \right] \\ &= \Delta^2 \mathbb{E} \left[\sum_{t=1}^T A_{t,i}^2 \mid V_{<i} = u \right], \end{aligned}$$

where the final expectation is under the original hypercube prior and the algorithm's trajectory law. Averaging over $V_{<i}$ and summing over i gives

$$\begin{aligned} I(V; H_T) &\leq \Delta^2 \mathbb{E} \sum_{t=1}^T \sum_{i=1}^d A_{t,i}^2 \\ &= \Delta^2 \mathbb{E} \sum_{t=1}^T \|A_t\|_2^2 \\ &\leq \Delta^2 T. \end{aligned}$$

This proves the claim. $\square$

Remark 7.7 (Why the fixed-design log-determinant argument is not used). *For a Gaussian prior on a full parameter vector, posterior conjugacy gives the adaptive information identity*

$$I(\Theta; H_T) = \frac{1}{2} \mathbb{E} \log \det \left(I + \Sigma_0^{1/2} \left(\sum_{t=1}^T A_t A_t^\top \right) \Sigma_0^{1/2} \right),$$

and this is bounded by a trace-constrained log determinant. For the hypercube prior, however, the posterior is not Gaussian and the realized design matrix is itself a statistic of the observations. Conditioning only on that realized design matrix controls the conditional observation information given the design; it does not by itself control the information carried by the adaptive design. Lemma 7.6 therefore uses a coordinatewise adaptive KL argument to control $I(Y; H_T)$ directly.

Theorem 7.8 (Linear bandit lower bound). *For stochastic linear bandits with action set $\{a \in \mathbb{R}^d : \|a\|_2 \leq 1\}$, unit Gaussian noise, and unit-ball parameter class $\Omega_{\text{lin}} = \{\omega \in \mathbb{R}^d : \|\omega\|_2 \leq 1\}$,*

$$\mathfrak{R}_T^* \geq c d \sqrt{T}$$

for a universal constant $c > 0$, whenever $T \geq d^2$.

Proof. Assume first that d is larger than a sufficiently large universal constant d_0 to be fixed below. Choose

$$\Delta = c_1 \sqrt{\frac{d}{T}},$$

where $c_1 > 0$ is a sufficiently small universal constant. Since $T \geq d^2$, choosing $c_1 \leq 1$ ensures

$$\Delta \sqrt{d} = c_1 \frac{d}{\sqrt{T}} \leq c_1 \leq 1,$$

so the hypercube prior is supported on the unit parameter ball.

Let

$$r = \frac{\Delta\sqrt{d}}{4}.$$

Lemma 7.5 gives the reference ghost-good probability bound

$$p_r \leq \exp(-c_{\text{gh}}d).$$

Lemma 7.6 gives the adaptive action-index information bound

$$T\bar{C} = I(Y; H_T) = I(V; H_T) \leq \Delta^2 T = c_1^2 d.$$

Choose c_1 so small that $c_1^2 \leq c_{\text{gh}}/4$, and then choose d_0 so large that $\log 2 \leq c_{\text{gh}}d/4$ for all $d \geq d_0$. Then, for all $d \geq d_0$,

$$T\bar{C} + \log 2 \leq \frac{1}{2} \log \frac{1}{p_r}.$$

Theorem 6.1, applied to the action index $Y = \chi(\Omega) = A^*(\Omega)$, gives a Bayes regret lower bound under the hypercube prior of at least

$$\frac{Tr}{2} = \frac{T\Delta\sqrt{d}}{8} = \frac{c_1}{8}d\sqrt{T}.$$

Since this prior is supported on Ω_{lin} , the minimax regret over Ω_{lin} is at least this Bayes risk.

It remains to handle the finitely many dimensions $1 \leq d < d_0$. For these dimensions, use the two-point prior $\Omega = \sigma\delta e_1$, where $\sigma \sim \text{Unif}(\{\pm 1\})$, e_1 is the first coordinate vector, and $\delta = (2\sqrt{T})^{-1}$. This prior is also supported on the unit parameter ball. Under $\sigma = +1$, the optimal action is e_1 ; under $\sigma = -1$, the optimal action is $-e_1$. Let $\mathbb{P}_+$ and $\mathbb{P}_-$ be the laws of the history under these two environments, and let $\mathbb{P}_\pm^{t-1}$ denote the law of the history before action t . For any algorithm,

$$\begin{aligned} \frac{1}{2}\mathbb{E}_+[\delta(1 - A_{t,1})] + \frac{1}{2}\mathbb{E}_-[\delta(1 + A_{t,1})] &= \delta \left(1 - \frac{1}{2}(\mathbb{E}_+ A_{t,1} - \mathbb{E}_- A_{t,1}) \right) \\ &\geq \delta (1 - \text{TV}(\mathbb{P}_+^{t-1}, \mathbb{P}_-^{t-1})). \end{aligned}$$

By Pinsker's inequality and the adaptive Gaussian KL chain rule,

$$\text{TV}(\mathbb{P}_+^{t-1}, \mathbb{P}_-^{t-1}) \leq \sqrt{\frac{1}{2}D_{\text{KL}}(\mathbb{P}_+^{t-1} \parallel \mathbb{P}_-^{t-1})} \leq \delta\sqrt{t-1} \leq \frac{1}{2}.$$

Therefore each round has Bayes regret at least $\delta/2$, and the total Bayes regret is at least $T\delta/2 = \sqrt{T}/4$. Since $d < d_0$, this is at least $(4d_0)^{-1}d\sqrt{T}$. Reducing the universal constant c , if necessary, completes the proof for all d . $\square$

7.4 Bellman-rank as indexed representation and lower-bound discussion

Low Bellman rank was introduced by Jiang et al. (2017) as a sufficient structural condition under which contextual decision processes with suitable bilinear Bellman-error structure are PAC-learnable. In the present language, a Bellman-rank factorization can help build a low-dimensional Bellman-sufficient representation when it is supplemented by predictive, update, loss, and index sufficiency; the factorization alone controls Bellman-error evaluation and does not, by itself, determine the full observation process. The following proposition is structural; it is not a new sharp lower bound for every Bellman-rank class.

Proposition 7.9 (Bellman rank with predictive sufficiency gives an indexed representation). *Suppose a contextual decision process admits an exact Bellman-error factorization of rank d in the sense of Jiang et al. (2017): for the stages and hypothesis-policy pairs in the planning class, the relevant Bellman error can be written as an inner product of two d -dimensional witness-feature vectors. Fix a prior μ and an index map χ . If the observation model is dominated by a common σ -finite measure and if, conditional on the current posterior or calibrated confidence object, the predictive law, the posterior-reference update, the relevant loss or benchmark terms, and the conditional χ -indexed predictive laws and losses in Definition 2.3 are all determined by these Bellman-error features together with the retained planning variables, then the process admits an indexed Bellman sufficient representation with state that may be chosen as a posterior or calibrated confidence object over the d -dimensional witness coordinates, augmented by the planning variables needed for Bellman updates. Consequently, whenever one constructs a valid information supersolution and a valid ghost-good-mass supersolution for the induced feature experiment, for example in a linear-Gaussian or uniformly bounded-likelihood setting, Theorem 6.3 yields a Bellman quantile-index lower certificate.*

Proof. The Bellman-error factorization supplies d -dimensional coordinates sufficient to evaluate the one-step Bellman-error terms covered by the factorization. It does not by itself imply predictive sufficiency, loss sufficiency, update sufficiency, or index sufficiency. These are precisely the additional assumptions imposed in the proposition. If the predictive observation law and posterior update are functions of this statistic and the current posterior, and if the relevant losses, benchmarks, and χ -conditional predictive objects are functions of the same retained state, then Definition 2.3 is satisfied with S_t equal to the posterior or calibrated confidence object over these features together with the accumulated planning state. The information and ghost-good Bellman recursions are then well-defined on the same state. Applying Theorem 6.3 to any valid pair of supersolutions for this state space gives the claimed certificate. The proposition is intentionally conditional because Bellman-rank learnability is primarily a structural upper-bound statement; sharp lower bounds require a packing or ghost-entropy calculation inside the particular rank- d class. $\square$

This application clarifies the connection between the present framework and the broader perspective on representation learning developed in Li and Xu (2026); Xu (2026). The point is not that every low-Bellman-rank problem has the same minimax rate or that Bellman rank alone determines a lower bound. The point is that a Bellman-rank factorization can provide a candidate state on which the upper supersolution and the lower quantile-information certificate should be compared, after the necessary predictive, update, loss, and index-sufficiency conditions have been verified.

7.5 Frequentist estimators as effective Bellman-sufficient states

The exact Bellman-sufficiency definition in Definition 2.3 corresponds to the posterior-process/reference-law formulation. This formulation is designed to state the fundamental specification conditions and to support lower-bound analysis. In this formulation, posteriors induced by fixed algorithmic beliefs may be used as frequentist estimators and may qualify as Bellman-sufficient states when they close the required Bellman recursions. By contrast, estimator-based frequentist algorithms, such as IPW contextual bandits and least-squares value iteration, typically verify a calibrated certificate rather than exact posterior-predictive sufficiency. The same is true for the adaptively optimized algorithmic posteriors of Xu and Zeevi (2025), where the posterior process is chosen for algorithmic or certification purposes rather than inherited from a fixed Bayesian belief. The following more flexible certificate is based on the upper-bound theory and notation of Section 5, and extends the spirit of Definition 2.3 to common frequentist estimators beyond algorithmic posteriors.

Definition 7.10 (Calibrated fixed-truth Bellman certificate). *Fix a class $\Omega_0 \subseteq \Omega$, an index χ , an indexed marginal q_s whenever a logarithmic information potential is used, and a potential sequence Ψ_t . A state process S_t is a calibrated fixed-truth Bellman certificate on Ω_0 if it satisfies the closed action/update requirements and, for each state s , determines a comparison set $\mathfrak{C}_t(s) \subseteq \Omega_0$ or a reference belief b_s , together with state-measurable quantities sufficient to evaluate or upper-bound the robust fixed-truth bracket*

$$\mathcal{B}_t^{\Psi, \gamma}(s, p) := \sup_{\omega \in \mathfrak{C}_t(s)} \left\{ \ell_\omega(s, p) + \mathbb{E}_{a \sim p, O \sim P_{\omega, s, a}} \Psi_{t+1}(\tau(s, a, O)) - \Psi_t(s) - \gamma G_\chi(s, p; \omega) \right\},$$

where

$$G_\chi(s, p; \omega) = \mathbb{E}_{a \sim p, O \sim P_{\omega, s, a}} \log \frac{q_{\tau(s, a, O)}(\chi(\omega))}{q_s(\chi(\omega))}$$

when an indexed belief q_s is maintained; if no such belief is used, the last term is replaced by the state-measurable uncertainty or relaxation increment appearing in the chosen supersolution. Calibration means that, for every fixed truth $\omega^* \in \Omega_0$, either $\omega^* \in \mathfrak{C}_t(S_t)$ on the relevant good event or the failure probability and approximation error are explicitly charged, and the actual conditional Bellman bracket under ω^* is bounded by $\mathcal{B}_t^{\Psi, \gamma}(S_t, p_t)$ plus those charged errors.

Thus a sufficient state for calibrated Bellman certificate need not be an exact posterior state, but it must contain enough information to make the fixed-truth Bellman inequality a function of the state. Unbiasedness, least-squares normal equations, or efficiency of an estimator are useful only insofar as they establish this calibration condition.

8 Conclusion

The paper identifies the statistical core of interactive decision making as an information-risk Bellman problem. A frequentist environment remains fixed; the proof chooses a Bellman-sufficient state, an information index, and a calibrated reference update. On that representation, the upper side is a log-penalized Bellman supersolution with coordinate potential $\gamma \log(1/q_t(\chi(\omega^*)))$, and the lower side is a Bellman-Fano ghost certificate whose entropy is $-\log p_r$. When these two logarithmic quantities close at the same radius, Theorem 5.4 gives a matching minimax rate. Equivalently, the critical balance is

$$T\bar{C}_\chi \asymp \log \frac{1}{p_r},$$

now interpreted on a controlled reference trajectory rather than on a fixed nonadaptive experiment.

AIR and MAIR are two choices of index within the same log-gain Bellman supersolution. Choosing $\chi(\Omega) = A^*(\Omega)$ gives action-index AIR, which is often the right description of decision difficulty. Choosing $\chi(\Omega) = \Omega$, or the active reward vector in a bandit problem, gives model-index MAIR, which is often the convenient description of estimation geometry. With the fixed-truth coordinate bracket, the Bayesian posterior-averaged bracket, and the robust belief relaxation kept separate, the exact AIR/MAIR Bellman program, UCB, E2D/DEC, and AMS/EBO become four algorithmic forms of the same indexed log-potential identity: the exact program optimizes the dynamic information-potential Bellman recursion; UCB certifies the fixed-truth bracket through self-normalized calibration and optimism; E2D/DEC optimizes a one-step frequentist offset that drops or freezes the Bellman continuation; and AMS/EBO optimizes a KL-dual log-partition relaxation over candidate beliefs.

The kernel-bandit application illustrates why this separation matters. The finite-marginal indexed Bellman program is the exact benchmark on the sufficient state. GP-UCB is practical because it keeps the realized log-determinant information gain of the actual trajectory. AMS/EBO, in an AIR form, can scale with the entropy of the finite or low-dimensional optimal-action index and the sequential AIR bracket, provided the robust belief set is calibrated to the fixed truth. GP-E2D optimizes a one-round DEC offset, but its interaction with realized information gain is a theorem to be proved rather than a formal consequence of the offset definition; without localization a global RKHS DEC may be vacuous. Thus the right comparison is not a universal conversion coefficient. It is whether a single Bellman-sufficient state supports matching upper supersolutions and lower quantile certificates without charging estimation difficulty twice.

Several directions remain. The first is to identify representations for which the exact information-potential Bellman upper certificate and the Bellman-Fano lower certificate match up to constants or logarithms. The second is to replace one-step DEC calculations by dynamically consistent Bellman relaxations when posterior evolution matters. The third is to sharpen the finite-action kernel theory by proving which of the indexed Bellman DP, GP-UCB, AMS/EBO, or GP-E2D attains the finite-marginal information scale under the weakest calibration assumptions. These are representation-level questions. The goal is not a universal one-shot notion of dimension, but a distribution-dependent dynamic theory of state spaces that accounts for estimation, decision, and lower-bound entropy costs once, rather than repeatedly.

References

- Yasin Abbasi-Yadkori, Dávid Pál, and Csaba Szepesvári. Improved algorithms for linear stochastic bandits. In *Advances in Neural Information Processing Systems*, 2011.
- Peter Auer, Nicolo Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47:235–256, 2002.
- Andrew R. Barron and Thomas M. Cover. Minimum complexity density estimation. *IEEE Transactions on Information Theory*, 37(4):1034–1054, 1991.
- Richard Bellman. *Dynamic Programming*. Princeton University Press, 1957.
- Dimitri P. Bertsekas. *Dynamic Programming and Optimal Control*, Volumes I and II. Athena Scientific, 4th edition, 2017.
- Fan Chen, Dylan J. Foster, Yanjun Han, Jian Qian, Alexander Rakhlin, and Yunbei Xu. Assouad, Fano, and Le Cam with interaction: A unifying framework for lower bounds. arXiv:2410.05117, 2024.
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms*. MIT Press and McGraw–Hill, 2nd edition, 2001.
- Simon S. Du, Sham M. Kakade, Jason D. Lee, Shachar Lovett, Gaurav Mahajan, Wen Sun, and Ruosong Wang. Bilinear classes: A structural framework for provable generalization in RL. In *Proceedings of the 38th International Conference on Machine Learning*, PMLR 139:2826–2836, 2021.
- Ronald A. Fisher. On the mathematical foundations of theoretical statistics. *Philosophical Transactions of the Royal Society of London. Series A*, 222:309–368, 1922.

- Dylan J. Foster, Sham M. Kakade, Jian Qian, and Alexander Rakhlin. The statistical complexity of interactive decision making. *arXiv:2112.13487*, 2021.
- Dylan J. Foster, Noah Golowich, and Yanjun Han. Tight guarantees for interactive decision making with the decision-estimation coefficient. *arXiv:2301.08215*, 2023.
- Nan Jiang, Akshay Krishnamurthy, Alekh Agarwal, John Langford, and Robert E. Schapire. Contextual decision processes with low Bellman rank are PAC-learnable. In *Proceedings of the 34th International Conference on Machine Learning*, 2017.
- Chi Jin, Zhuoran Yang, Zhaoran Wang, and Michael I. Jordan. Provably efficient reinforcement learning with linear function approximation. *arXiv:1907.05388*, 2019.
- Tor Lattimore and András Györfy. Mirror descent and the information ratio. In *Proceedings of the 34th Annual Conference on Learning Theory*, PMLR 134:2965–2992, 2021.
- Tor Lattimore. A lower bound for adaptive linear regression. In *Proceedings of the 34th International Conference on Algorithmic Learning Theory*, 2023.
- Shaojie Li and Yunbei Xu. Pointwise generalization in deep neural networks. *arXiv:2605.18598*, 2026.
- Alexander Rakhlin, Ohad Shamir, and Karthik Sridharan. Relax and randomize: From value to algorithms. In *Advances in Neural Information Processing Systems*, 2012.
- Alexander Rakhlin and Karthik Sridharan. BISTRO: An efficient relaxation-based method for contextual bandits. In *Proceedings of the 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pages 1977–1985. PMLR, 2016.
- David Silver and Richard S. Sutton. Welcome to the era of experience. Preprint of a chapter to appear in *Designing an Intelligence*, MIT Press, 2025.
- Niranjan Srinivas, Andreas Krause, Sham M. Kakade, and Matthias Seeger. Gaussian process optimization in the bandit setting: No regret and experimental design. In *International Conference on Machine Learning*, 2010.
- Sattar Vakili, Jonathan Scarlett, and Tara Javidi. Open problem: Tight online confidence intervals for RKHS elements. In *Proceedings of the 34th Annual Conference on Learning Theory*, PMLR 134:280–289, 2021.
- Yunbei Xu. Pointwise complexity for Gaussian fields: Upper envelopes, algorithmic lower bounds, and separation. *arXiv:2606.07931*, 2026.
- Yunbei Xu and Assaf Zeevi. Bayesian design principles for frequentist sequential learning. *Journal of the ACM*, 72(5), Article 37, 2025.
- Yuhong Yang and Andrew R. Barron. Information-theoretic determination of minimax rates of convergence. *The Annals of Statistics*, 27(5):1564–1599, 1999.
- Andrew C. Yao. Probabilistic computations: Toward a unified measure of complexity. In *Proceedings of the 18th Annual Symposium on Foundations of Computer Science*, pages 222–227, 1977.
- Bin Yu. Assouad, Fano, and Le Cam. In *Festschrift for Lucien Le Cam*, pages 423–435. Springer, New York, 1997.

Tong Zhang. Information-theoretic upper and lower bounds for statistical estimation. *IEEE Transactions on Information Theory*, 52(4):1307–1321, 2006.