

The Dimension of Nonterminating Resampling Computations

Yunbei Xu
National University of Singapore
yunbei@nus.edu.sg

Abstract

How much randomness can remain in a random tape on which a computation runs forever? For a resampling computation over finitely many independent variables, reversing a finite run yields an exact likelihood identity: the self-information of the repaired events minus the description length of the repair sequence under a comparison law equals the log-likelihood ratio between a backward record law and the source law of the tape prefix consumed by the run. If this difference is Δ_T , then

$$\mathbb{P}\{\Delta_T \geq u\} \leq 2^{-u}.$$

A linear lower bound on Δ_T for every execution lasting T repairs therefore gives an exponential running-time tail and, under computability, an upper bound on the Kolmogorov complexity of every surviving tape prefix. For rejection sampling from a computable full-support source π on a finite alphabet of at least two symbols and a nonempty rejection set U , let $\delta \in [0, 1]$ be the unique solution of $\sum_{x \in U} \pi(x)^\delta = 1$. Every tape rejected forever has effective strong dimension relative to the source at most δ , the bound is attained, and the nonterminating set has source-relative Hausdorff dimension δ . Finally, two whole-block resampling systems are constructed with the same event-variable incidence graph and the same bad-event probabilities but different exact survival exponents and, consequently in these uniform examples, different maximal effective strong dimensions and Hausdorff dimensions. Thus these standard local-lemma data do not determine the geometry or the exact exponential decay of nontermination, even when every repair redraws the entire random block.

1 Introduction

A randomized algorithm maps its random tape into a trajectory. Once the tape is fixed, the computation is deterministic: it either stops or it does not. Let $\mathbb{P}$ denote probability under the source law of the random tape. Let

$$\mathcal{N} = \{\omega : \tau(\omega) = \infty\},$$

where τ is the number of rejection or repair steps before stopping and $\tau = \infty$ on a nonterminating tape. The tail probability $\mathbb{P}\{\tau \geq n\}$ measures how often the computation survives n steps. When the following extended limit exists in $[0, \infty]$, define the *survival exponent* (the base-two escape rate)

$$\alpha = - \lim_{n \rightarrow \infty} \frac{1}{n} \log_2 \mathbb{P}\{\tau \geq n\}.$$

The survival exponent is an ensemble quantity: it measures the exponential decay of the nested survival events $\{\tau \geq n\}$. Even when the algorithm terminates almost surely and $\mathbb{P}(\mathcal{N}) = 0$, the set $\mathcal{N}$ may be nonempty. The additional question is pointwise. Once a tape $\omega \in \mathcal{N}$ is fixed, the

run is deterministic, and effective dimension measures the incompressibility of the prefixes of that individual tape. Source-relative Hausdorff dimension instead measures $\mathcal{N}$ as a set.

The paper makes two contributions. First, an exact backward likelihood identity for finite variable-model resampling computations yields both running-time tail bounds and Kolmogorov-complexity bounds for individual surviving tape prefixes. Second, a nested system and a disjoint system with the same event-variable structure and bad-event probabilities show that these data need not determine the exact survival exponent or either dimension quantity. Before the general identity is developed, rejection sampling serves as the basic case: its nonterminating tapes form an exactly solvable symbolic fractal.

1.1 Rejection sampling: the basic case

A familiar example is the simplest form of rejection sampling. Draw $X_1, X_2, \dots$ independently from a full-support probability law π on a finite set Ω with $|\Omega| \geq 2$, reject a draw in $U \subseteq \Omega$, and return the first draw in $G = \Omega \setminus U$. This is the special case of the classical accept-reject method in which each proposed value is accepted or rejected deterministically (von Neumann, 1951; Devroye, 1986).

Write $p = \pi(U)$ and $a = \pi(G) = 1 - p$. If $0 < a < 1$, then

$$\mathbb{P}\{\tau \geq n\} = p^n, \quad \mathbb{P}\{\tau = t\} = p^t a, \quad \mathbb{E}\tau = \frac{p}{a}, \quad \mathbb{E}(\tau + 1) = \frac{1}{a} \quad (n, t \in \mathbb{Z}_{\geq 0}). \quad (1)$$

Here τ is the number of rejected draws and $\tau + 1$ is the total number of draws. For $a > 0$, the returned value has the conditional law

$$\mathbb{P}\{X_{\tau+1} \in B\} = \frac{\pi(B \cap G)}{\pi(G)} \quad (B \subseteq \Omega). \quad (2)$$

If $a = 1$, the first draw is accepted; if $a = 0$, the algorithm never stops. When $0 < a < 1$, the algorithm stops almost surely although the set $U^{\mathbb{N}}$ of tapes on which it never stops is nonempty. Probability assigns this set mass zero. Effective dimension asks how much algorithmic randomness can nevertheless remain in one of its individual elements.

Now suppose that the source is uniform and $U \neq \emptyset$. Write $Q = |\Omega| \geq 2$ and $M = |U|$, so $1 \leq M \leq Q$. If E_n is the event that the first n blocks are rejected, then

$$\mathbb{P}(E_n) = \left(\frac{M}{Q}\right)^n,$$

and there are exactly M^n possible surviving prefixes. Each such prefix can be described using $n \log_2 M + o(n)$ bits, whereas the source supplies $n \log_2 Q$ bits. Their ratio is

$$\frac{\log_2 M}{\log_2 Q}.$$

In the two-bit example

$$\Omega = \{00, 01, 10, 11\}, \quad U = \{00, 01\}.$$

Every surviving block has its first bit fixed and its second bit free. Exactly one half of the source information can remain.

For a computable full-support source $P = \pi^{\mathbb{N}}$, let $\omega_{1:n}$ denote the first n source symbols (or blocks). For a finite word w , let $[w]$ be the cylinder of tapes beginning with w , and write $P[w] = P([w])$. Define their self-information by

$$I_P(\omega_{1:n}) = -\log_2 P[\omega_{1:n}].$$

Let $K(w)$ be the length of the shortest program that prints w on a fixed universal prefix-free, or self-delimiting, machine. More generally, $K(w \mid z)$ is the corresponding program length when z is supplied as auxiliary input. The effective strong P -dimension of a tape ω is

$$\text{Dim}_{\text{eff},P}(\omega) = \limsup_{n \rightarrow \infty} \frac{K(\omega_{1:n})}{I_P(\omega_{1:n})}. \quad (3)$$

Replacing the limsup by a liminf defines the effective P -dimension $\text{dim}_{\text{eff},P}(\omega)$. For a fair binary source, write

$$\text{Dim}_{\text{eff}}(\omega) = \limsup_{n \rightarrow \infty} \frac{K(\omega_{1:n})}{n}.$$

This is the constructive strong dimension relative to a source measure (Athreya et al., 2007; Lutz, 2011). It is recalled here for rejection sampling; later sections relate it to nontermination in general resampling computations. A value near one means that the tape is nearly as incompressible as an ordinary draw from P . A smaller value means that continued execution has imposed structure on the tape.

For a set E of tapes, $\text{dim}_H^P E$ denotes its Billingsley, or source-relative Hausdorff, dimension: in the Hausdorff covering construction, the cylinder $[w]$ is assigned diameter $P[w]$ (Billingsley, 1965; Lutz, 2011). Under a uniform q -ary source, this is ordinary Hausdorff dimension in the symbolic metric $d(\omega, \eta) = q^{-k}$, where k is the length of the longest common prefix of ω and η .

There are three different questions:

- (i) $\mathbb{P}\{\tau \geq n\}$ is the probability of performing at least n rejection or repair steps;
- (ii) $\text{Dim}_{\text{eff},P}(\omega)$ measures the randomness of one tape; and
- (iii) $\text{dim}_H^P \mathcal{N}$ measures the size of the set of all nonterminating tapes.

Of the three questions, (i) is an ensemble or distributional quantity and (iii) is a set-level geometric quantity, whereas (ii) is a pointwise quantity for one fixed tape without averaging over tapes. In general, effective strong dimension, which uses a limsup, is associated with packing dimension rather than Hausdorff dimension (Athreya et al., 2007). The equality between maximal effective strong dimension and Hausdorff dimension proved below is a feature of the regular product and finite-state sets considered here, not a general identity. For rejection sampling, the tail in (i) is $\pi(U)^n$. The two dimension questions are governed by the number δ that solves

$$\sum_{x \in U} \pi(x)^\delta = 1. \quad (4)$$

Then every tape rejected forever satisfies

$$\text{Dim}_{\text{eff},P}(\omega) \leq \delta \quad (\omega \in U^\mathbb{N}). \quad (5)$$

Moreover, write $\mathbf{1}\{\cdot\}$ for the indicator of the stated condition, and set

$$Q_\delta(x) = \pi(x)^\delta \mathbf{1}\{x \in U\}.$$

For a computable law Q , a tape is Martin–Löf random if it avoids every effective Q -null test; equivalently, the Levin–Schnorr theorem gives $K(\omega_{1:n}) \geq -\log_2 Q[\omega_{1:n}] - O(1)$ for every n . The displayed Q_δ is a probability law, and a tape that is Martin–Löf random under $Q_\delta^\mathbb{N}$ attains equality. The classical set calculation gives $\text{dim}_H^P(U^\mathbb{N}) = \delta$. Under the uniform law, $\pi(U) = Q^{\delta-1}$, so the tail exponent and the dimension determine one another. For a nonuniform law, they need not.

The law Q_δ produces tapes attaining the pointwise bound; the source-relative Hausdorff dimension of the whole set equals the same value δ because this symbolic set is regular.

1.2 Reconstructing the random tape

Rejection sampling is the simplest resampling computation. The same question can be posed for finite variable-model resampling algorithms. At each step, such an algorithm selects a bad event and redraws the variables on which that event depends. A finite run can be read backward from its final assignment, its ordered list of repairs, and the values erased by those repairs. Restoring the erased values in reverse order recovers every random choice used by the run.

This reconstruction gives a simple comparison:

- information revealed by the repaired events
- information needed to describe the repair history = a log-likelihood ratio.

Because the backward descriptions have total probability at most one, this ratio has an exponential upper tail. If the description is computable, the random choices used by the run admit a description shorter than their source self-information by at least the same number of bits. Theorem 2.1 formalizes both conclusions. Event probabilities alone are not enough; the description must include the repair order and any control information needed to reverse the run.

1.3 Nested rejection rules

The event–variable incidence graph is the bipartite graph joining each bad event to every random variable on which it depends. Joining two event vertices whenever they share a variable gives the canonical dependency graph. A dependency graph has one vertex for each rejection rule and joins two rules that may interact. In the variable model, two rules are joined when they depend on a common random variable. The Lovász local lemma and its variants use this graph to give conditions under which all bad events can be avoided (Erdős and Lovász, 1975; Shearer, 1985; Scott and Sokal, 2005). Algorithmic analyses use the same graph to bound the number and weight of repair histories (Moser and Tardos, 2010; Haeupler and Harris, 2017; Fialho et al., 2022). The graph does not record how the rejected sets intersect or whether one contains another.

Suppose a finite family $\mathcal{A}$ of nested rules can reject the same random block. The algorithm records the smallest rule that applies and then redraws the whole block. Its true one-step survival probability is

$$\pi \left(\bigcup_{A \in \mathcal{A}} A \right),$$

whereas the complete-graph history enumeration considered below assigns the one-step weight $\sum_{A \in \mathcal{A}} \pi(A)$, counting an outcome once for every event that contains it. Dividing the union into disjoint cases preserves the recorded rule and recovers the exact survival probability. Section 3 gives nested Hamming balls for which the graph enumeration yields only the trivial dimension bound, followed by two systems with the same graph and marginals but different survival exponents and dimensions.

Because every rejection redraws the whole block, this example has no persistent state. Partially overlapping local repairs, which retain variables relevant to later repairs, are the harder case.

1.4 Contribution and limits

The first contribution is Theorem 2.1, an exact likelihood-ratio identity for every finite run of a finite variable-model resampling computation. Its probabilistic interpretation controls the survival tail, while a code built from the same backward law controls the Kolmogorov complexity of each

surviving tape prefix. The running-time and pointwise bounds are therefore two consequences of one equality.

The second contribution is a quantitative separation. Theorem 3.1 computes the survival exponent and dimension from the union of the rejected sets. Example 3.5 gives two whole-block resampling computations with the same event-variable incidence graph and the same bad-event probabilities but different exact survival exponents and, consequently for these uniform systems, different maximal effective strong dimensions and Hausdorff dimensions. The contribution is not the known observation that dependency graphs omit intersection information: intersection-sensitive refinements of graph-based local-lemma bounds already exploit such information (Shearer, 1985; He et al., 2017, 2023). The point here is its exact consequence for the survival rate and tape geometry of fixed resampling computations.

The present analysis does not yet give matching upper and lower bounds for partially overlapping local repairs with persistent state. When one repair resamples only part of the variables involved in another, the retained values may affect the next repair. The repair sequence alone may therefore be insufficient to predict or reconstruct the run. The open problem is to find an intermediate state description that still determines the survival exponent and permits backward reconstruction.

1.5 Related work

The Lovász local lemma asks whether all bad events can be avoided; its algorithmic forms ask whether resampling can find such an outcome efficiently (Erdős and Lovász, 1975; Moser and Tardos, 2010). The present work instead studies the exponential decay rate of $\mathbb{P}\{\tau \geq n\}$, the maximal effective strong dimension of an individual tape in $\mathcal{N}$, and the source-relative Hausdorff dimension of $\mathcal{N}$.

The survival exponent uses the logarithmic scaling of large-deviation theory, but it is not a large-deviation rate function: a rate function assigns costs to an entire space of limiting observables (Dembo and Zeitouni, 2010). More formally, suppose Z_n takes values in a topological space $\mathbf{X}$, $\{\tau \geq n\} = \{Z_n \in F\}$ for a fixed $F \subseteq \mathbf{X}$, and (Z_n) satisfies a large-deviation principle with speed n and rate function $I : \mathbf{X} \rightarrow [0, \infty]$ under the natural-log convention. If

$$\inf_{z \in F^\circ} I(z) = \inf_{z \in \overline{F}} I(z),$$

where F° and $\overline{F}$ denote the interior and closure of F , so that the large-deviation lower and upper bounds agree for F , then

$$\alpha = \frac{1}{\ln 2} \inf_{z \in F} I(z).$$

No such principle is assumed here. Theorem 2.1 gives only the one-sided bound

$$\limsup_{T \rightarrow \infty} \frac{1}{T} \log_2 \mathbb{P} \left\{ \frac{\Delta_T}{T} \geq a \right\} \leq -a \quad (a \geq 0),$$

not a matching lower bound or a full large-deviation principle. For rejection sampling, α nevertheless equals the endpoint Bernoulli cost $D_2(1 \parallel \pi(U)) = \log_2(1/\pi(U))$, where

$$D_2(r \parallel p) = r \log_2 \frac{r}{p} + (1 - r) \log_2 \frac{1 - r}{1 - p}$$

is the Bernoulli Kullback–Leibler divergence in bits, with $0 \log_2 0 = 0$ and the usual continuous extension at the boundary. In the finite-state model of Appendix A, let B_1 be the matrix of one-step

transition probabilities between nonterminal states, omitting the probability of termination. Its entries are nonnegative and its row sums are at most one; such a matrix is called substochastic. The appendix proves

$$\alpha = -\log_2 \rho(B_1).$$

Here, for any square matrix C , $\rho(C) = \max\{|\lambda| : \lambda \text{ is an eigenvalue of } C\}$ denotes its spectral radius. Thus α is the base-two escape rate of B_1 (Keller and Liverani, 2009).

The analysis builds on two established lines of work. For finite-state systems, the relation among weighted path growth, Hausdorff dimension, and the Markov law obtained by reweighting paths is classical (Parry, 1964; Mauldin and Williams, 1988; Bowen, 1975; Simpson, 2015). Entropy-compression arguments encode repair histories reversibly, and their interpretation through Kolmogorov complexity is also known (Moser, 2009; Moser and Tardos, 2010; Fortnow, 2009; Tao, 2009). These results are used as tools rather than claimed as new.

Recent work assigns pointwise complexity to a realized neural network or to a location in a Gaussian field (Li and Xu, 2026; Xu, 2026b). Here the realized object is a random tape. A related question asks how much state information is needed for a Bellman recursion (Xu, 2026a); in the present setting, that question reappears for overlapping local repairs. These analogies motivate the pointwise viewpoint, but they are not used in the proofs and do not yield a learning-theoretic generalization bound.

2 Reversing a resampling computation

Let $X_1, \dots, X_n$ be independent finite-valued variables with product law $\pi = \bigotimes_i \pi_i$. For $J \subseteq \{1, \dots, n\}$, write $\pi_J = \bigotimes_{i \in J} \pi_i$ for the marginal product law on the coordinates in J . Each nonempty bad event A is determined by the coordinates $v(A)$ and has probability $p_A > 0$. Fix a deterministic rule for selecting a currently true bad event.

List the random choices in the order in which the algorithm uses them. First draw $S_0 \sim \pi$. Whenever the current assignment contains a bad event, select A_t , draw $Z_t \sim \pi_{v(A_t)}$ independently of the past, and write Z_t on $v(A_t)$. Let S_t denote the resulting assignment after the t th repair. Let E_T be the event that at least T resamplings occur, and put

$$U_T = (S_0, Z_1, \dots, Z_T).$$

Let P_{tape} denote the law generated by the initial draw and these successive independent resampling draws. Then $\mathbb{P}(E) = P_{\text{tape}}(E)$ for every tape event E ; for a finite used-tape prefix w (including U_T), $P_{\text{tape}}(w)$ denotes the probability of its cylinder. Thus $E_T = \{\tau \geq T\}$. For each fixed T , the minimal tape prefixes that realize T resamplings form a prefix-free family: no member is a proper prefix of another. Their cylinders are therefore disjoint, and their probabilities sum to $\mathbb{P}(E_T)$.

Let $Y_t = S_{t-1}|_{v(A_t)}$ be the configuration overwritten at time t , and let $L_T = (A_1, \dots, A_T)$. Write $\pi_{v(A)}(\cdot \mid A)$ for the marginal law $\pi_{v(A)}$ conditioned on the local configurations satisfying A . Coordinatewise, the entries of the used tape are partitioned into the final assignment and the overwritten configurations. Therefore

$$P_{\text{tape}}(U_T) = \pi(S_T) \prod_{t=1}^T p_{A_t} \pi_{v(A_t)}(Y_t \mid A_t). \quad (6)$$

Let Q_T be a subprobability mass function—a nonnegative mass function of total mass at most one—on lists $\ell = (A_1, \dots, A_T)$, positive on every feasible T -repair list. Define a comparison

subprobability law on complete records by

$$G_T(s, \ell, y_{1:T}) = \pi(s) Q_T(\ell) \prod_{t=1}^T \pi_{v(A_t)}(y_t \mid A_t). \quad (7)$$

The map

$$U_T \mapsto (S_T, L_T, Y_{1:T})$$

is injective. Indeed, reverse the repairs from T to 1: the current values on $v(A_t)$ recover Z_t , and restoring Y_t recovers the preceding assignment.

Theorem 2.1 (Backward likelihood identity). *Every execution in E_T satisfies*

$$\log_2 \frac{G_T(S_T, L_T, Y_{1:T})}{P_{\text{tape}}(U_T)} = \sum_{t=1}^T \log_2 \frac{1}{p_{A_t}} - \log_2 \frac{1}{Q_T(L_T)}. \quad (8)$$

Call the right side Δ_T on E_T and set $\Delta_T = -\infty$ otherwise. Then, for every $u \geq 0$,

$$\mathbb{P}\{\Delta_T \geq u\} \leq 2^{-u}. \quad (9)$$

If the instance and the family $(Q_T)_{T \geq 1}$ are uniformly computable, then

$$-\log_2 P_{\text{tape}}(U_T) - K(U_T \mid T, \text{instance}) \geq \Delta_T - O(1). \quad (10)$$

Here instance denotes a fixed effective description of the source laws, bad events, and deterministic selection rule. Uniform computability means that one algorithm can approximate $Q_T(\ell)$ to any requested precision from T and ℓ , with the instance data encoded effectively in the same sense.

Proof. Dividing (7) by (6) cancels the final-state factor and all conditional overwritten-value factors. The remaining terms give

$$\frac{G_T(S_T, L_T, Y_{1:T})}{P_{\text{tape}}(U_T)} = \frac{Q_T(L_T)}{\prod_{t=1}^T p_{A_t}} = 2^{\Delta_T},$$

which proves (8).

It remains to use this likelihood ratio. Let $\mathcal{W}_T$ be the prefix-free set of minimal tape prefixes that produce at least T repairs, and let $\mathcal{R}_T$ be the set of formal records on which G_T is defined. For $w \in \mathcal{W}_T$, write

$$R_T(w) = (S_T(w), L_T(w), Y_{1:T}(w)).$$

The reconstruction argument shows that R_T is injective.

First observe that G_T has total mass at most one. Indeed, for a fixed list $\ell = (A_1, \dots, A_T)$, the final-state law and each conditional overwritten-value law have total mass one. Consequently,

$$\sum_{r \in \mathcal{R}_T} G_T(r) = \sum_{\ell} Q_T(\ell) \leq 1.$$

Now pull G_T back to live tape prefixes by defining

$$\mu_T(w) = G_T(R_T(w)), \quad w \in \mathcal{W}_T.$$

Injectivity of R_T gives

$$\sum_{w \in \mathcal{W}_T} \mu_T(w) = \sum_{r \in R_T(\mathcal{W}_T)} G_T(r) \leq \sum_{r \in \mathcal{R}_T} G_T(r) \leq 1.$$

Thus μ_T is a subprobability mass function. Its total mass can be strictly less than one because Q_T may itself have mass less than one, or because G_T may assign mass to formal records that are not produced by the algorithm.

For every $w \in \mathcal{W}_T$, the identity already proved gives

$$2^{\Delta_T(w)} = \frac{G_T(R_T(w))}{P_{\text{tape}}(w)} = \frac{\mu_T(w)}{P_{\text{tape}}(w)}.$$

Because $\Delta_T = -\infty$ outside E_T and the cylinders determined by $\mathcal{W}_T$ are disjoint, for every $u \geq 0$,

$$\begin{aligned} \mathbb{P}\{\Delta_T \geq u\} &= \sum_{\substack{w \in \mathcal{W}_T \\ \Delta_T(w) \geq u}} P_{\text{tape}}(w) \\ &= \sum_{\substack{w \in \mathcal{W}_T \\ \Delta_T(w) \geq u}} 2^{-\Delta_T(w)} \mu_T(w) \\ &\leq 2^{-u} \sum_{\substack{w \in \mathcal{W}_T \\ \Delta_T(w) \geq u}} \mu_T(w) \\ &\leq 2^{-u}. \end{aligned}$$

This proves (9); equivalently, it is Markov's inequality applied to a subprobability likelihood ratio.

Finally, uniform computability of the instance and of $(Q_T)_{T \geq 1}$ makes $(\mu_T)_{T \geq 1}$ a uniformly computable family of subprobability mass functions. Shannon–Fano coding, or equivalently the coding theorem for computable subprobability masses (Shannon, 1948; Fano, 1949), therefore gives

$$K(U_T \mid T, \text{instance}) \leq -\log_2 \mu_T(U_T) + O(1).$$

Since

$$-\log_2 \mu_T(U_T) = -\log_2 P_{\text{tape}}(U_T) - \Delta_T,$$

rearranging proves (10). $\square$

The two terms in (8) have different origins. The first is the information revealed by the events that were repaired. The second is the number of bits needed to describe which repairs occurred. Neither term alone gives a bound; their difference does, because it is a likelihood ratio.

For example, color the vertices of a graph independently with q colors and repeatedly recolor the endpoints of the first monochromatic edge. A fixed edge is monochromatic with probability $1/q$, so each repair contributes $\log_2 q$ bits to the first term. The second term pays for the repaired edges and their order. The comparison is therefore between information ruled out by the violations and information used by the algorithm's history.

Corollary 2.2 (A tail bound and a pointwise dimension bound). *Suppose every T -repair execution satisfies*

$$\sum_{t=1}^T \log_2 \frac{1}{p_{A_t}} \geq \beta T - C_0, \quad -\log_2 Q_T(L_T) \leq \gamma T + C_1,$$

where $\beta > \gamma$. Then

$$\mathbb{P}(E_T) \leq 2^{C_0 + C_1 - (\beta - \gamma)T}. \quad (11)$$

For the pointwise conclusion, assume in addition that

- (i) the instance and the family $(Q_T)_{T \geq 1}$ are uniformly computable;
- (ii) the computation reads a fixed fair binary tape sequentially, the initial draw consumes $c_0 = O(1)$ bits, and every repair consumes exactly c bits, so that U_T is uniformly computably interconvertible with $\omega_{1:c_0+cT}$.

Then every nonterminating tape satisfies

$$\text{Dim}_{\text{eff}}(\omega) \leq \max \left\{ 0, 1 - \frac{\beta - \gamma}{c} \right\}. \quad (12)$$

Proof. The assumptions give $\Delta_T \geq (\beta - \gamma)T - (C_0 + C_1)$. Equation (11) follows from (9). Equation (10), together with the $O(\log_2 T)$ bits needed to specify T , gives $K(U_T) \leq cT - (\beta - \gamma)T + o(T)$ along an infinite execution. Divide by cT ; the gap between consecutive repair prefixes is the fixed number c of bits. $\square$

Remark 2.3 (What the record must contain). For a recursive Moser-style implementation, L_T must be enlarged to include the recursion tree or return structure unless those data are recoverable by a separate prefix-free decoding rule. In Theorem 2.1, everything needed for reversal contributes to the probability $Q_T(L_T)$.

3 What a dependency graph forgets

To use Theorem 2.1, one must describe the possible lists of repairs. The Lovász local lemma summarizes event interactions by a dependency graph, and its algorithmic proofs use that graph to organize repair histories.

3.1 Why dependency graphs enter

Let $\mathcal{A}$ be a finite family of bad events. The Lovász local lemma asks when there is a positive probability of avoiding all of them:

$$\mathbb{P} \left(\bigcap_{A \in \mathcal{A}} A^c \right) > 0. \quad (13)$$

In the variable model, where independent variables generate the events, the standard graph D joins A and B whenever their variable sets overlap. For example, in random vertex coloring a bad event says that one edge is monochromatic, and two such events are adjacent in D when their edges share a vertex. More generally, write $\Gamma(A)$ for the neighbors of A . A loopless graph on $\mathcal{A}$ is a dependency graph if every event A is independent of the sigma-field generated by its nonneighbors; equivalently, A is independent of every Boolean combination of events indexed by $\mathcal{A} \setminus (\{A\} \cup \Gamma(A))$. The vertices of D are events; this is different from the bipartite variable–event incidence graph used in Section 3.5. If $\mathbb{P}(A) \leq p$ for every A and D has maximum degree d , the symmetric local lemma gives the familiar sufficient condition

$$ep(d+1) \leq 1. \quad (14)$$

Here e is Euler’s number. For a fixed graph D , Shearer’s region is the set of probability vectors $(p_A)_{A \in \mathcal{A}}$ for which every event family having dependency graph D and satisfying $\mathbb{P}(A) \leq p_A$ has positive probability of avoiding all events. It is the sharp worst-case criterion based only on D and the marginal event probabilities (Erdős and Lovász, 1975; Shearer, 1985; Scott and Sokal, 2005).

Moser and Tardos turn this existence argument into a random process. Draw all variables independently; while some bad event is true, choose one and redraw the variables on which it depends. If there are numbers $x_A \in (0, 1)$ such that

$$\mathbb{P}(A) \leq x_A \prod_{B \in \Gamma(A)} (1 - x_B) \quad (A \in \mathcal{A}), \quad (15)$$

then the expected number of times A is redrawn is at most $x_A/(1 - x_A)$ (Moser and Tardos, 2010). Consequently,

$$\mathbb{E}\tau \leq \sum_{A \in \mathcal{A}} \frac{x_A}{1 - x_A} < \infty. \quad (16)$$

The algorithm therefore stops almost surely and returns an assignment in which no bad event occurs. Bounds of this kind are proved by encoding repair histories as trees, directed acyclic graphs (DAGs), or entropy-compression records whose weights are determined by event probabilities and adjacency in D (Moser and Tardos, 2010; Haeupler and Harris, 2017; Fialho et al., 2022).

The graph summary omits intersection data. The graph records which repairs may interact, but it does not determine $\mathbb{P}(A \cap B)$, set inclusion, or $\mathbb{P}(\bigcup_A A)$. Two event systems can therefore have the same graph and the same marginal probabilities but different survival exponents and different nonterminating sets. The next three subsections compare the exact state-space description, the graph bound, and a nested example where the loss is explicit.

3.2 The full state space

Every finite variable-model resampling algorithm is a finite-state randomized computation on assignments. Set $P = P_{\text{tape}}$, the sequential source law of the initial assignment and subsequent resampling draws; finite initial factors do not affect the dimensions below. Delete zero-probability transitions and assume that every surviving transition has probability strictly less than one. Fix a deterministic rule for selecting a currently true bad event. On reachable nonterminal assignments and for $\theta \geq 0$, define

$$(B_\theta)_{xy} = \sum_{z: F(x,z)=y} \pi_{v(A(x))}(z)^\theta, \quad (17)$$

where $A(x)$ is the selected event and $F(x, z)$ is the updated assignment. Decompose the nonterminal transition graph into reachable strongly connected components C containing a directed cycle, where a strongly connected component is a maximal set of states that are mutually reachable by directed paths. Let $B_{\theta,C}$ be the restriction of B_θ to C . Proposition A.1 gives the exact source-relative dimension

$$\dim_H^P \mathcal{N} = \max_C \delta_C, \quad \rho(B_{\delta_C,C}) = 1. \quad (18)$$

Here $\delta_C \in [0, 1]$ is the unique solution of the displayed spectral-radius equation for component C . This formula is exact, but its matrix usually has exponentially many rows and columns.

3.3 The graph bound

Classical entropy-compression and Shearer arguments use the dependency graph and the probabilities p_A to bound the total weight of possible repair lists (Cartier and Foata, 1969; Scott and Sokal, 2005). Only the complete-graph case is needed here. If $D = K_m$ is complete on $m = |\mathcal{A}|$ event vertices, the

graph-only enumeration allows every word $(A_1, \dots, A_t) \in \mathcal{A}^t$ and assigns it weight $\prod_{i=1}^t p_{A_i}$. Its total weight is

$$\left(\sum_{A \in \mathcal{A}} p_A \right)^t.$$

This is an upper bound obtained by enumeration from D and the marginals; it need not equal the probability mass of the histories produced by a fixed selection rule. For uniform q -ary blocks of size b , this gives the dimension bound in (29). The bound is classical. A short graph description also does not make exact computation easy: evaluating the relevant independence polynomial and testing membership in the general Shearer region are #P-hard in general (Harvey et al., 2018).

3.4 Nested rejection rules

Let Ω be a finite block space with source law π satisfying $0 < \pi(x) < 1$ for every x , and let $\mathcal{A}$ be a laminar family of bad subsets: for $A, B \in \mathcal{A}$, either $A \cap B = \emptyset$, $A \subseteq B$, or $B \subseteq A$. Every repair resamples the whole block from π . If the current block x lies in the bad union $U = \bigcup_{A \in \mathcal{A}} A$, select the inclusion-minimal event containing x . Write $\phi(x)$ for this event, which is unique by laminarity.

Let $\text{ch}(A)$ be the maximal proper members of $\mathcal{A}$ contained in A , and define the disjoint region

$$C_A = A \setminus \bigcup_{B \in \text{ch}(A)} B. \quad (19)$$

Let $\mathcal{I} = \{A \in \mathcal{A} : C_A \neq \emptyset\}$ index the nonempty regions. They are disjoint, their union is U , and the algorithm records A exactly on C_A . Define

$$w_A(\theta) = \sum_{x \in C_A} \pi(x)^\theta, \quad W(\theta) = \sum_{A \in \mathcal{I}} w_A(\theta) = \sum_{x \in U} \pi(x)^\theta. \quad (20)$$

Write $w(\theta) = (w_A(\theta))_{A \in \mathcal{I}}$ for the corresponding column vector.

Theorem 3.1 (Laminar whole-block rejection). *Assume $U \neq \emptyset$, and let $P = \pi^{\mathbb{N}}$. There is a unique $\delta \in [0, 1]$ such that*

$$W(\delta) = \sum_{x \in U} \pi(x)^\delta = 1. \quad (21)$$

(a) *If the alphabet, π , U , and the law Q_δ below are computable, every nonterminating tape satisfies*

$$\text{Dim}_{\text{eff}, P}(\omega) \leq \delta \quad (\omega \in \mathcal{N} = U^{\mathbb{N}}). \quad (22)$$

Every tape that is Martin–Löf random for $Q_\delta^{\mathbb{N}}$ satisfies $\dim_{\text{eff}, P}(\omega) = \text{Dim}_{\text{eff}, P}(\omega) = \delta$. Moreover,

$$\dim_H^P \mathcal{N} = \delta, \quad \mathbb{P}\{\tau \geq t\} = \pi(U)^t. \quad (23)$$

Under the uniform law on $\Omega = [q]^b$, writing $M = |U|$,

$$\dim_H \mathcal{N} = \delta = \frac{\log_2 M}{b \log_2 q}, \quad \mathbb{P}\{\tau \geq t\} = \left(\frac{M}{q^b} \right)^t. \quad (24)$$

In particular, under a uniform source, a tape of effective strong dimension greater than δ must make the algorithm terminate.

(b) Group rejected blocks according to the region C_A that contains them. For every $\theta \geq 0$, the resulting weighted transition matrix is

$$\widehat{B}_\theta(A, B) = w_B(\theta) \quad (A, B \in \mathcal{I}). \quad (25)$$

Its spectral radius is $W(\theta)$, the same as that of the full transition matrix, and the all-ones vector $\mathbf{1} \in \mathbb{R}^{|\mathcal{I}|}$ is a right eigenvector.

(c) The dimension and survival bounds are attained by the following reversible record laws. The dimension-attaining law is

$$Q_\delta(x) = \pi(x)^\delta \mathbf{1}\{x \in U\}. \quad (26)$$

For every T -repair execution, there is a probability record law $G_{\delta, T}$ satisfying

$$\log_2 \frac{G_{\delta, T}(S_T, Y_{1:T})}{P_{\text{tape}}(U_T)} = (1 - \delta) \sum_{t=1}^T \log_2 \frac{1}{\pi(Y_t)}. \quad (27)$$

The second is the source conditioned on another rejection, $Q_{\text{surv}}(x) = \pi(x) \mathbf{1}\{x \in U\} / \pi(U)$. The corresponding record law $G_{\text{surv}, T}$ satisfies

$$\log_2 \frac{G_{\text{surv}, T}(S_T, Y_{1:T})}{P_{\text{tape}}(U_T)} = T \log_2 \frac{1}{\pi(U)}. \quad (28)$$

Proof. The function W is continuous and strictly decreasing because $0 < \pi(x) < 1$. Moreover, $W(0) = |U| \geq 1$ and $W(1) = \pi(U) \leq 1$, so (21) has a unique solution in $[0, 1]$.

For full nonterminal states $x, y \in U$, resampling the common scope gives $B_\theta(x, y) = \pi(y)^\theta$. Hence, for every $x \in C_A$,

$$\sum_{y \in C_B} B_\theta(x, y) = \sum_{y \in C_B} \pi(y)^\theta = w_B(\theta),$$

independently of x . Thus grouping states by the sets C_A leaves the weighted transition calculation unchanged and proves (25). The reduced matrix is $\mathbf{1} w(\theta)^\top$, so its only nonzero eigenvalue is $W(\theta)$ and $\mathbf{1}$ is a right eigenvector. Proposition A.1 now gives $\dim_H^P \mathcal{N} = \delta$. Each repair draws a fresh block from π and continues precisely on U , proving the survival formula and its uniform specialization.

For the dimension bound, observe that Q_δ is computable and supported on U . Every $\omega \in U^\mathbb{N}$ therefore satisfies

$$K(\omega_{1:n}) \leq -\log_2 Q_\delta^\mathbb{N}[\omega_{1:n}] + O(\log_2 n) = \delta(-\log_2 P[\omega_{1:n}]) + O(\log_2 n).$$

Division by the source information proves (22). If ω is Martin–Löf random for $Q_\delta^\mathbb{N}$, Levin–Schnorr gives the reverse inequality up to $O(1)$, and hence equality. This also proves (5).

It remains to construct the record laws. The repair list is determined by the overwritten states because $A_t = \phi(Y_t)$. Reversal recovers the used tape from $(S_T, Y_{1:T})$. The source factorization is

$$P_{\text{tape}}(U_T) = \pi(S_T) \prod_{t=1}^T \pi(Y_t).$$

Since (21) holds, the record law

$$G_{\delta, T}(s, y_{1:T}) = \pi(s) \prod_{t=1}^T \left[\pi(y_t)^\delta \mathbf{1}\{y_t \in U\} \right]$$

has total mass one. Dividing by the source factorization proves (27). Similarly, define

$$G_{\text{surv},T}(s, y_{1:T}) = \pi(s) \prod_{t=1}^T \frac{\pi(y_t) \mathbf{1}\{y_t \in U\}}{\pi(U)}.$$

This is a probability law, and division by the source factorization proves (28). The region containing y_t determines the recorded rejection label $A_t = \phi(y_t)$. $\square$

Remark 3.2 (Beyond laminar families). Laminarity is used only to define the selector cells canonically through set inclusion. For an arbitrary family of whole-block rejection rules with a deterministic selector ϕ , the same conclusions hold with $C_A = \{x \in U : \phi(x) = A\}$.

Remark 3.3 (The one-state reduction). Every repair redraws the whole block. Consequently, the next step is independent of the current rejected block, and the one-state matrix $[W(\theta)]$ already gives both $W(1) = \pi(U)$ and the dimension root. The disjoint regions retain the recorded reason for rejection, but do not change the stopping dynamics. Partially overlapping local repairs, which leave relevant state unchanged, are the harder case.

The survival law is also a direct application of Theorem 2.1. Relabel each repair by the disjoint region C_A containing the overwritten block; this does not change the run or its resampling rule. Put $r_A = \pi(C_A)$ and choose

$$Q_T(A_1, \dots, A_T) = \prod_{t=1}^T \frac{r_{A_t}}{\pi(U)}.$$

For the relabeled event C_{A_t} , the event contribution in (8) is $\log_2(1/r_{A_t})$, whereas encoding the region label under Q_T costs $\log_2(\pi(U)/r_{A_t})$ bits. Their difference is $\log_2(1/\pi(U))$ at every repair, giving (28). Conditioning only on the original event A_t would count overlapping mass more than once.

Comparison with the dependency graph. Under the standard variable-overlap definition, all common-block events are pairwise adjacent. The graph calculation therefore uses the one-step weight $\sum_{A \in \mathcal{A}} \pi(A)$, while the true rejection probability is $\pi(U)$. Under the uniform source on $[q]^b$, the resulting upper bound and the true dimension are

$$1 + \frac{\log_2 \sum_{A \in \mathcal{A}} \pi(A)}{b \log_2 q} \quad \text{and} \quad 1 + \frac{\log_2 \pi(U)}{b \log_2 q}, \quad (29)$$

again truncated to $[0, 1]$. Before truncation, the two expressions agree when the events are disjoint, whereas overlap produces overcounting. Truncation can conceal this difference when both values are at least one.

In the variable model, a nonedge certifies disjoint scopes and hence commuting repairs; an edge marks only possible interaction. A complete graph therefore contains no information about how the corresponding bad events intersect. The following standard family makes the loss explicit.

Example 3.4 (Nested Hamming balls). Let $b \geq 1$ and $0 \leq r \leq b$ be integers, and let $\Omega = \{0, 1\}^b$ carry the uniform law. Reject a block when its Hamming weight is at most r . Writing $|x|$ for the Hamming weight of x , let

$$U_r = \{x : |x| \leq r\}, \quad M_r = |U_r| = \sum_{j=0}^r \binom{b}{j},$$

Theorem 3.1 gives, for every nonterminating tape,

$$\text{Dim}_{\text{eff}}(\omega) \leq \frac{1}{b} \log_2 M_r \quad \text{for every } \omega \in U_r^{\mathbb{N}}, \quad \dim_H(U_r^{\mathbb{N}}) = \frac{1}{b} \log_2 M_r. \quad (30)$$

Now add the stricter rule $U_s \subset U_r$, where $0 \leq s < r$, and record U_s whenever it applies. The standard variable-overlap dependency graph on the two original rules is K_2 and the graph-only sum is $(M_s + M_r)/2^b$, although the computation rejects on a set of mass only $M_r/2^b$.

For example, take $b \geq 3$, $r = b - 1$, and $s = b - 2$. The graph sum is $2 - (b + 2)/2^b > 1$, so the graph bound gives only the trivial upper bound 1, whereas the true dimension is $b^{-1} \log_2(2^b - 1) < 1$. Splitting U_r into U_s and $U_r \setminus U_s$ removes this double counting.

Example 3.5 (Same local-lemma data, different survival exponents). The graph and bad-event probabilities also fail to determine the exact survival exponent and dimensions. Fix $b \geq 3$, and let π be the uniform law on $\Omega = \{0, 1\}^b$. Let

$$B = U_{b-2}, \quad A = \{0^b\}, \quad A' = \{1^b\}.$$

Compare the two event systems $\{A, B\}$ and $\{A', B\}$. Every event depends on all b block coordinates, so both systems have the same event-variable incidence graph $K_{2,b}$, the complete bipartite graph joining the two event vertices to all b variable vertices, and hence the same canonical dependency graph K_2 . With the singleton event listed first, they also have the same ordered vector of bad-event probabilities:

$$(\pi(A), \pi(B)) = (\pi(A'), \pi(B)) = \left(2^{-b}, \frac{2^b - b - 1}{2^b}\right).$$

Their intersections are different: $A \subset B$, while $A' \cap B = \emptyset$. Consequently,

$$|A \cup B| = 2^b - b - 1, \quad |A' \cup B| = 2^b - b.$$

After t redraws, the two survival probabilities are therefore

$$\left(\frac{2^b - b - 1}{2^b}\right)^t \quad \text{and} \quad \left(\frac{2^b - b}{2^b}\right)^t,$$

so their survival exponents are

$$\alpha_1 = b - \log_2(2^b - b - 1), \quad \alpha_2 = b - \log_2(2^b - b).$$

Let $\mathcal{N}_1 = (A \cup B)^{\mathbb{N}}$ and $\mathcal{N}_2 = (A' \cup B)^{\mathbb{N}}$. Theorem 3.1 gives

$$\max_{\omega \in \mathcal{N}_1} \text{Dim}_{\text{eff}}(\omega) = \dim_H \mathcal{N}_1 = \frac{1}{b} \log_2(2^b - b - 1), \quad \max_{\omega \in \mathcal{N}_2} \text{Dim}_{\text{eff}}(\omega) = \dim_H \mathcal{N}_2 = \frac{1}{b} \log_2(2^b - b).$$

Thus the common graph and bad-event probabilities do not determine the exact survival exponent and, consequently in this uniform whole-block setting, do not determine the maximal effective strong dimension or the Hausdorff dimension.

Relation to the variable local lemma. The separation does not contradict the tightness of the Shearer region, which is a worst-case statement over event systems with fixed graph and event probabilities. In fact, the common marginal vector in Example 3.5 satisfies

$$\pi(A) + \pi(B) = 1 - \frac{b}{2^b} < 1,$$

and therefore lies strictly inside the Shearer region for K_2 . The variable local lemma gives existence or convergence guarantees from event probabilities and event–variable structure (He et al., 2017); refinements that also use intersection information are known (He et al., 2023). The backward likelihood identity instead holds for every finite resampling run, whether or not a local-lemma criterion is satisfied. For the instances that can be solved exactly, it determines the survival exponent and the geometry of the nonterminating tapes. This is a tape-level analysis of fixed resampling computations, not a stronger universal local-lemma criterion.

3.5 Partially overlapping local repairs

The preceding examples are memoryless because every action redraws the same block. A partially overlapping local repair leaves values relevant to later repairs unchanged, and those retained values can affect which event is selected next.

Existing representations of repair histories retain some of this persistent state information (Achlioptas and Iliopoulos, 2016; Haeupler and Harris, 2017; Fialho et al., 2022). A path decomposition is a sequence of vertex bags such that every edge occurs in some bag and the bags containing each vertex are consecutive; its width is one less than the largest bag size. A concrete test case is bounded pathwidth. Consider a q -ary resampling instance of size n whose variable–event incidence graph has a path decomposition of width w and whose deterministic selection rule can be updated from the current bag using $\text{poly}(n)$ control states. Can the boundary assignment, which contains at most $w + 1$ active variables, together with this control state define a matrix of size $q^{O(w)} \text{poly}(n)$ whose spectral-radius equation gives the exact dimension? Can its dimension-attaining change of measure be implemented by reversible execution records, giving matching upper and lower bounds for the effective dimension of a tape? Forbidden-pattern generation and bounded-pathwidth k -SAT provide two test cases.

The pointwise result requires both an exact state-space calculation and a computable encoding of actual runs. A spectral calculation alone provides no such code, while an encoding that represents one run more than once need not recover the Hausdorff dimension.

4 Conclusion

This paper connects three views of nontermination: the decay of the survival probability, the incompressibility of one fixed tape, and the geometry of the set of all nonterminating tapes. Its central result is an exact backward likelihood identity for finite variable-model resampling computations. The self-information revealed by repaired events, minus the description length of the repair sequence under a comparison law, is a likelihood ratio whose tail controls $\mathbb{P}\{\tau \geq T\}$. Under computability, the same ratio bounds the Kolmogorov complexity of every surviving tape prefix. The probabilistic and pointwise conclusions therefore come from the same equality.

For rejection sampling and whole-block resampling, the calculation is exact. The critical law Q_δ produces individual tapes attaining the effective dimension bound, and the source-relative Hausdorff dimension measures the entire nonterminating set. Disjoint selector cells recover the exact survival exponent when rejection rules overlap. Example 3.5 shows that the event–variable incidence graph and bad-event probabilities can be identical while the exact survival exponents, and consequently the dimensions in these uniform systems, differ.

The ultimate goal is an equally exact theory for overlapping local repairs with persistent state. A repair label or dependency graph may omit retained values needed to predict and reverse the next step; the full assignment is exact but generally exponential in the number of variables. The main open problem is to find a substantially smaller state representation that still supports both an exact

survival calculation and reversible coding of individual runs. Such a representation would provide a route to matching bounds for the survival exponent, individual-tape dimension, and Hausdorff dimension under genuinely local overlap.

A Finite-state computations

Let $\mathcal{G}$ be a finite directed multigraph of nonterminal states with fixed initial state s_0 . Each outgoing edge $e : s \rightarrow t$ represents a distinct next source outcome and has probability $p(e) > 0$. The outgoing probabilities from each state sum to at most one; the remaining probability terminates the computation. A surviving path $w = e_1 \cdots e_n$, traversing states $s_0, s_1, \dots, s_n$, determines a tape cylinder of probability

$$P[w] = \prod_{i=1}^n p(e_i).$$

Let $\mathcal{W}_n$ be the set of surviving paths of length n from s_0 . Assume that the reachable nonterminal graph is strongly connected, contains a directed cycle, and satisfies $\max_e p(e) < 1$. For $\theta \geq 0$, define

$$(B_\theta)_{st} = \sum_{e:s \rightarrow t} p(e)^\theta, \quad \Lambda(\theta) = \log_2 \rho(B_\theta). \quad (31)$$

For any square matrix C , its spectral radius is $\rho(C) = \max\{|\lambda| : \lambda \text{ is an eigenvalue of } C\}$. At $\theta = 1$, B_1 is the substochastic transition matrix among nonterminal states: each row sum is at most one, and the missing mass is the one-step termination probability. The quantity $\Lambda(\theta)$ is the base-two pressure, the exponential growth rate of the total θ -weight of surviving paths. Graphs that are not strongly connected are handled by applying the result to each reachable strongly connected component containing a directed cycle and taking the largest dimension.

Use $P[w]$ as the diameter of the cylinder $[w]$. The resulting Hausdorff dimension is the source-relative dimension $\dim_H^P$ defined in the introduction; for a uniform q -ary source it is ordinary Hausdorff dimension in the q -adic metric.

Proposition A.1 (The finite-state formula). *There is a unique $\delta \in [0, 1]$ satisfying $\rho(B_\delta) = 1$. Moreover,*

$$\lim_{n \rightarrow \infty} \frac{1}{n} \log_2 \sum_{w \in \mathcal{W}_n} P[w]^\theta = \Lambda(\theta), \quad (32)$$

$$\lim_{n \rightarrow \infty} -\frac{1}{n} \log_2 \mathbb{P}\{\tau \geq n\} = -\Lambda(1), \quad (33)$$

$$\dim_H^P \mathcal{N} = \delta. \quad (34)$$

If $B_\delta h = h$ with $h_s > 0$ for every state s , then

$$Q_\delta(e \mid s) = p(e)^\delta \frac{h_t}{h_s} \quad (35)$$

is a probability law on surviving paths. For every $w = e_1 \cdots e_n \in \mathcal{W}_n$, define its probability under this law by

$$Q_\delta[w] = \prod_{i=1}^n Q_\delta(e_i \mid s_{i-1}).$$

Then it satisfies

$$\log_2 \frac{Q_\delta[w]}{P[w]} = (1 - \delta) \log_2 \frac{1}{P[w]} + \log_2 \frac{h_{s_n}}{h_{s_0}}. \quad (36)$$

For a computable system with computable Q_δ , every nonterminating tape satisfies $\text{Dim}_{\text{eff},P}(\omega) \leq \delta$. A tape that is Martin–Löf random under Q_δ satisfies $\dim_{\text{eff},P}(\omega) = \text{Dim}_{\text{eff},P}(\omega) = \delta$.

Proof. Since the graph contains a directed cycle, $\rho(B_0) \geq 1$. Since B_1 is substochastic, $\rho(B_1) \leq 1$. If $\theta' > \theta$, then every nonzero entry contributed by an edge is smaller in $B_{\theta'}$ than in B_θ , because $p(e) < 1$. Strict Perron comparison on the common strongly connected support therefore gives $\rho(B_{\theta'}) < \rho(B_\theta)$. Together with continuity, this proves existence and uniqueness of δ . The matrix identity

$$\sum_{w \in \mathcal{W}_n} P[w]^\theta = e_{s_0}^\top B_\theta^n \mathbf{1},$$

where e_{s_0} is the coordinate vector of the initial state and $\mathbf{1}$ is the all-ones column vector. Perron–Frobenius theory proves (32); setting $\theta = 1$ gives (33).

If $\theta > \delta$, then $\rho(B_\theta) < 1$, so the total θ -weight of the length- n surviving cylinders tends to zero. These cylinders cover $\mathcal{N}$, giving $\dim_H^P \mathcal{N} \leq \delta$. At $\theta = \delta$, (35) has row sums one, and telescoping gives

$$Q_\delta[w] = P[w]^\delta \frac{h_{s_n}}{h_{s_0}}.$$

Because h is positive on a finite state space, there is a constant $C < \infty$, independent of w , such that $Q_\delta[w] \leq CP[w]^\delta$. The mass-distribution principle states that a probability measure satisfying this cylinder bound forces $\dim_H^P \mathcal{N} \geq \delta$. It therefore gives the required lower bound, proving (34) and (36) (Bowen, 1975; Mauldin and Williams, 1988).

For a computable system, coding under Q_δ gives, for every surviving prefix w of length n ,

$$K(w) \leq -\log_2 Q_\delta[w] + O(\log_2 n) = \delta \log_2 \frac{1}{P[w]} + O(\log_2 n).$$

Division by source self-information proves the upper effective-dimension bound. Levin–Schnorr gives the reverse inequality for a Q_δ -Martin–Löf-random path (Lutz, 2003; Athreya et al., 2007; Simpson, 2015). $\square$

Under a uniform q -ary source, let $M = (M_{st})$, where M_{st} counts the source symbols that move s to t . Then $B_\theta = q^{-\theta} M$. If

$$\alpha = -\lim_{n \rightarrow \infty} \frac{1}{n} \log_2 \mathbb{P}\{\tau \geq n\}$$

is the survival exponent, then

$$\dim_H \mathcal{N} = \log_q \rho(M) = 1 - \frac{\alpha}{\log_2 q}.$$

Writing $h_{\text{surv}} = \log_2 \rho(M)$ for the topological entropy, in bits per step, of the surviving path space gives

$$\alpha = \log_2 q - h_{\text{surv}}, \quad \dim_H \mathcal{N} = \frac{h_{\text{surv}}}{\log_2 q}.$$

Thus α is the entropy deficit imposed by survival, not an entropy function. Equation (36) is the finite-state form of the comparison used in Theorem 2.1.

Acknowledgements

The author used ChatGPT as a research and editorial aid. The author assumes full responsibility for the statements and proofs in the paper.

References

- D. Achlioptas and F. Iliopoulos. Random walks that find perfect objects and the Lovász Local Lemma. *Journal of the ACM*, 63(3):22:1–22:29, 2016.
- K. B. Athreya, J. M. Hitchcock, J. H. Lutz, and E. Mayordomo. Effective strong dimension, algorithmic information, and computational complexity. *SIAM Journal on Computing*, 37(3):671–705, 2007. doi: 10.1137/S0097539703446912.
- P. Billingsley. *Ergodic Theory and Information*. John Wiley & Sons, 1965.
- R. Bowen. *Equilibrium States and the Ergodic Theory of Anosov Diffeomorphisms*, volume 470 of *Lecture Notes in Mathematics*. Springer, 1975. doi: 10.1007/BFb0081029.
- P. Cartier and D. Foata. *Problèmes combinatoires de commutation et réarrangements*, volume 85 of *Lecture Notes in Mathematics*. Springer, 1969.
- A. Dembo and O. Zeitouni. *Large Deviations Techniques and Applications*, volume 38 of *Stochastic Modelling and Applied Probability*. Springer, Berlin, Heidelberg, second edition, 2010. doi: 10.1007/978-3-642-03311-7.
- L. Devroye. *Non-Uniform Random Variate Generation*. Springer, New York, 1986. doi: 10.1007/978-1-4613-8643-8. See Chapter II, Section 3.
- P. Erdős and L. Lovász. Problems and results on 3-chromatic hypergraphs and some related questions. In A. Hajnal, R. Rado, and V. T. Sós, editors, *Infinite and Finite Sets, Vol. II*, volume 10 of *Colloquia Mathematica Societatis János Bolyai*, pages 609–627. North-Holland, Amsterdam, 1975.
- R. M. Fano. The transmission of information. Technical Report 65, Research Laboratory of Electronics, Massachusetts Institute of Technology, Cambridge, Massachusetts, 1949.
- P. M. S. Fialho, B. N. B. de Lima, and A. Procacci. Moser–Tardos resampling algorithm, entropy compression method and the subset gas. *Annales de l’Institut Henri Poincaré D*, 9(2):435–471, 2022. doi: 10.4171/AIHPD/122.
- L. Fortnow. A Kolmogorov complexity proof of the Lovász Local Lemma. [Computational Complexity blog](#), 2009.
- B. Haeupler and D. G. Harris. Parallel algorithms and concentration bounds for the Lovász Local Lemma via witness DAGs. *ACM Transactions on Algorithms*, 13(4):53:1–53:40, 2017.
- N. J. A. Harvey, P. Srivastava, and J. Vondrák. Computing the independence polynomial: From the tree threshold down to the roots. In *Proceedings of the Twenty-Ninth Annual ACM–SIAM Symposium on Discrete Algorithms*, SODA ’18, pages 1557–1576. Society for Industrial and Applied Mathematics, 2018. doi: 10.1137/1.9781611975031.102.
- K. He, L. Li, X. Liu, Y. Wang, and M. Xia. Variable version Lovász Local Lemma: Beyond Shearer’s bound. In *2017 IEEE 58th Annual Symposium on Foundations of Computer Science*, pages 451–462, 2017. doi: 10.1109/FOCS.2017.48.
- K. He, Q. Li, and X. Sun. Moser–Tardos algorithm: Beyond Shearer’s bound. In *Proceedings of the 2023 Annual ACM–SIAM Symposium on Discrete Algorithms*, SODA ’23, pages 3362–3387. Society for Industrial and Applied Mathematics, 2023. doi: 10.1137/1.9781611977554.ch129.

- G. Keller and C. Liverani. Rare events, escape rates and quasistationarity: Some exact formulae. *Journal of Statistical Physics*, 135(3):519–534, 2009. doi: 10.1007/s10955-009-9747-8.
- S. Li and Y. Xu. Pointwise generalization in deep neural networks. [arXiv preprint arXiv:2605.18598](#), 2026.
- J. H. Lutz. The dimensions of individual strings and sequences. *Information and Computation*, 187(1):49–79, 2003. doi: 10.1016/S0890-5401(03)00187-0.
- J. H. Lutz. A divergence formula for randomness and dimension. *Theoretical Computer Science*, 412(1–2): 166–177, 2011. doi: 10.1016/j.tcs.2010.09.005.
- R. D. Mauldin and S. C. Williams. Hausdorff dimension in graph directed constructions. *Transactions of the American Mathematical Society*, 309(2):811–829, 1988. doi: 10.1090/S0002-9947-1988-0961615-4.
- R. A. Moser. A constructive proof of the Lovász Local Lemma. In *Proceedings of the 41st Annual ACM Symposium on Theory of Computing*, pages 343–350, 2009. doi: 10.1145/1536414.1536462.
- R. A. Moser and G. Tardos. A constructive proof of the general Lovász Local Lemma. *Journal of the ACM*, 57(2):11:1–11:15, 2010. doi: 10.1145/1667053.1667060.
- W. Parry. Intrinsic Markov chains. *Transactions of the American Mathematical Society*, 112:55–66, 1964. doi: 10.1090/S0002-9947-1964-0161372-1.
- A. D. Scott and A. D. Sokal. The repulsive lattice gas, the independent-set polynomial, and the Lovász Local Lemma. *Journal of Statistical Physics*, 118:1151–1261, 2005. doi: 10.1007/s10955-004-2055-4.
- C. E. Shannon. A mathematical theory of communication. *The Bell System Technical Journal*, 27(3):379–423, July 1948. doi: 10.1002/j.1538-7305.1948.tb01338.x.
- J. B. Shearer. On a problem of Spencer. *Combinatorica*, 5(3):241–245, 1985. doi: 10.1007/BF02579368.
- S. G. Simpson. Symbolic dynamics: Entropy = dimension = complexity. *Theory of Computing Systems*, 56(3):527–543, 2015. doi: 10.1007/s00224-014-9546-8.
- T. Tao. Moser’s entropy compression argument. [What’s New](#), 2009.
- J. von Neumann. Various techniques used in connection with random digits. In A. S. Householder, G. E. Forsythe, and H. H. Germond, editors, *Monte Carlo Method*, volume 12 of *National Bureau of Standards Applied Mathematics Series*, pages 36–38. U.S. Government Printing Office, Washington, DC, 1951.
- Y. Xu. Bellman-sufficient information complexity. [arXiv preprint arXiv:2606.11171](#), 2026a.
- Y. Xu. Pointwise complexity for Gaussian fields: Upper envelopes, algorithmic lower bounds, and separation. [arXiv preprint arXiv:2606.07931](#), 2026b.