

Pointwise Generalization in Deep Neural Networks

Shaojie Li Yunbei Xu*
National University of Singapore
{li_sj,yunbei}@nus.edu.sg

Abstract

We address the fundamental question of why deep neural networks generalize by establishing a pointwise generalization theory for fully connected networks. This framework resolves long-standing barriers to characterizing the rich nonlinear feature-learning regime and builds a new statistical foundation for representation learning. For each trained model, we characterize the hypothesis via a pointwise Riemannian Dimension, derived from the eigenvalues of the learned feature representations across layers. This establishes a principled framework for deriving hypothesis-dependent, representation-aware generalization bounds. These bounds offer a systematic upgrade over approaches based on model size, products of norms, and infinite-width linearizations, yielding guarantees that are orders of magnitude tighter in both theory and experiment. Analytically, we identify the structural properties and mathematical principles that explain the tractability of deep networks. Empirically, the pointwise Riemannian Dimension exhibits substantial feature compression, decreases with increased over-parameterization, and captures the implicit bias of optimizers. Taken together, our results indicate that deep networks are mathematically tractable in practical regimes and that their generalization is sharply explained by pointwise, feature-spectrum-aware complexity.

Contents

1	Introduction	3
1.1	Organization and Contributions	5
2	The Nature of Pointwise Generalization	6
2.1	Pointwise Dimension Strengthens PAC-Bayes and Generic Chaining	6
2.2	Necessity of Finite-Scale Pointwise Geometry and Structural Analysis	11
3	Deep Neural Networks and Pointwise Riemannian Dimension	12
3.1	Neural Network Setup	12
3.2	Non-Perturbative Expansion and Layer-wise Correlation	13
3.3	Hierarchical Covering from Local Chart to Global Atlas	14
3.3.1	Gold Standard: Effective Dimension	14
3.3.2	Key Challenge: Prior Independence from W	16
3.3.3	Hierarchical covering (mixture prior over subspaces).	16

*In keeping with standard practice in mathematics and theory, authors are listed alphabetically. Yunbei Xu (yunbei@nus.edu.sg) is the corresponding author.

4	Generalization Bounds and Implications	19
4.1	Unconditional Generalization Bound for DNN	19
4.2	Fully Observable Generalization Bound for DNN	21
4.3	Comparison with Norm Bounds, VC, and NTK	23
4.4	Algorithmic Implications and Excess Risk Bound	24
5	Experiments	25
5.1	Riemannian Dimension Explains Overparameterization	26
5.2	Feature Learning Compresses Effective Rank	27
5.3	SGD Finds Low Riemannian Dimension Point	28
6	Conclusion	29
A	Related Works and Experimental Setup	30
A.1	Related Works	30
A.2	Experimental Setup	33
B	Proofs for Pointwise Generalization Framework (Section 2)	34
B.1	The “Uniform Pointwise Convergence” Principle	34
B.1.1	Necessary and Sufficient Conditions for Pointwise Generalization	35
B.2	The PAC-Bayes Optimization Problem	36
B.2.1	Pointwise Dimension Bound via Metric Ball	38
B.2.2	Lower Bound and Optimality of PAC-Bayes Optimization	39
B.3	Subset Homogeneity of Pointwise Dimension	40
B.4	Pointwise Generalization Bound via Ghost Sample	43
B.5	Proof of Theorem 2	48
B.6	Proof of Theorem 3	53
B.7	Background on Gaussian and Empirical Processes	55
C	Proofs for Deep Neural Networks and Riemannian Dimension (Section 3)	60
C.1	Proof of Lemma 1 (Non-Perturbative Feature Expansion)	60
C.2	Metric Domination Lemma	60
C.3	Pointwise Dimension Bound with Reference Subspace	61
C.4	Proof of Riemannian Dimension Bound for DNN (Theorem 4)	66
C.4.1	Decomposition Properties of NN-surrogate Metric Tensor	66
C.4.2	Proof of Theorem 4	67
D	Ellipsoidal Covering of the Grassmannian (Lemma 3)	73
D.1	Grassmannian Manifold, Stiefel Parameterization, and Orthogonal Groups	75
D.2	Principal Angles between Subspaces	77
D.3	Local Charts of the Grassmannian	78
D.4	Global Atlas of Graph Charts	81
D.5	Decomposition and Lipschitz Properties inside Graph Chart	83
D.6	Proof of the Main Result	86

E	Proofs for Generalization Bounds and Comparison (Section 4)	92
E.1	Proof of Theorem 5	92
E.2	Proof of Theorem 6	92
E.3	Ghost formulation and proof of the DNN bound	93
E.4	Empirical-Ghost Subspace Isomorphism for the Pointwise Ellipsoidal Metric	95
E.5	How feature regularity can imply the subspace isomorphism	100
E.6	Proof for Regularized ERM in Section 4.4	102
E.7	Improvement over Norm Bounds in Section 4.3	103
E.7.1	Exponential Improvement to a Norm Bound and Comparison	103
E.7.2	Proof of Corollary 2	106

1 Introduction

Deep learning has ushered in a new era of AI, delivering striking generalization across scientific tasks [LeCun et al., 2015, Bengio et al., 2013]. Yet, a fundamental paradox remains: while classical theory predicts severe overfitting for massive models, practice exhibits strong generalization. This gap has fueled a prevailing view that neural networks are opaque “black boxes” resistant to principled explanation [Goodfellow et al., 2016]. We narrow this gap by addressing the generalization problem for the canonical fully connected Deep Neural Networks (DNN). We demonstrate that, under verifiable spectral conditions on the *learned feature representations*, deep neural networks fall into a tractable regime with tight generalization guarantees. Crucially, these conditions impose no constraints on parameter count or weight sparsity, standing in sharp contrast to prior statistical conventions based on pure weight-space compression. Methodologically, our characterization leverages a pointwise generalization paradigm that fundamentally transcends classical uniform-convergence and covering-number approaches. We believe this is an especially notable conclusion: deep neural networks can admit sharp, hypothesis-dependent finite-sample guarantees with a degree of tractability often thought possible only for linearized models, *throughout the rich nonlinear feature-learning regime*, without relying on infinite-width limits, linearized approximations around initialization, or exponential dependence on key problem parameters. We hope that this framework helps demystify generalization and provides a systematic methodology for analyzing nonlinear, overparameterized models in representation learning.

We study standard fully connected (feed-forward) networks on a dataset $X = [x_1, \dots, x_n] \in \mathbb{R}^{d_0 \times n}$, where each column is one input example. The network has widths $d_1, \dots, d_L$, and weight matrices $W_l \in \mathbb{R}^{d_l \times d_{l-1}}$ for $l = 1, \dots, L$. We define the *feature matrix* at layer l by the recursion

$$F_l(W, X) := \sigma_l(W_l F_{l-1}(W, X)) \in \mathbb{R}^{d_l \times n}, \quad l = 1, \dots, L, \quad (1.1)$$

where $F_0 := X$ and the nonlinear activation σ_l acts columnwise. Each *column* of F_l is the feature vector of one data point at layer l ; each *row* of F_l is the activation of one neuron across the dataset.

Our focus is the *generalization gap*—the difference between test and training loss at the learned weights W . Informally—up to universal constants, optimistic logarithmic factors¹, and reasonable simplification (made precise in Theorem 5 and Theorem 6 with discussion on the feasibility of these

¹Throughout the paper, the phrase “optimistic logarithmic factors” is used in the following precise sense illustrated by Example 1: under common rapid spectral-decay or explicit low-rank assumptions on learned-feature compression, the relevant scaling is logarithmic. Whenever absolute logarithmic terms appear, they are stated explicitly.

simplifications)—we prove that this gap is controlled by the *effective dimension* of the learned features: uniformly over every $W \in \mathbb{R}^{\sum_l d_l \cdot d_{l-1}}$,

$$\mathcal{L}_{\text{test}}(W) - \mathcal{L}_{\text{train}}(W) \lesssim \sqrt{\frac{1}{n} \sum_{l=1}^L (d_l + d_{l-1}) d_{\text{eff}} \left(F_{l-1}(W, X) F_{l-1}(W, X)^\top \right)}. \quad (1.2)$$

Here $d_{\text{eff}}(\cdot)$ denotes the (layerwise) *effective dimension*—a smoothed, spectrum-aware notion of rank—of the feature Gram matrix $F_{l-1}(W, X) F_{l-1}(W, X)^\top$, i.e., the number of meaningful directions the feature data actually occupies at that layer. Intuitively, each layer contributes a term proportional to its size $(d_l + d_{l-1})$ multiplied by how many directions its features $F_{l-1}(W, X)$ truly use, d_{eff} . When features are correlated, low rank, or exhibit a rapidly decaying spectrum (a few large eigenvalues dominating many small ones), d_{eff} is small, so the bound remains tight even for very wide/deep networks. Such “feature compression” are widely observed in modern deep learning [Huh et al., 2021, Wang et al., 2025, Parker et al., 2023]. Strikingly, in our experiments, increasing overparameterization often induces pronounced *feature-rank compression*: the bound (1.2) decreases as model size grows (Section 5); for example, in ResNet trained on CIFAR-10, a majority of layers compress to (near-)zero effective rank.

Inequality (1.2) yields a strong *uniform, hypothesis- and data-dependent* guarantee, which we term *pointwise generalization*. It tracks how features evolve across layers of the *trained* model and explains overparameterization in practice. Moreover, the right-hand side of (1.2) can be used directly as a *regularizer*, leading to algorithms that adapt to the effective ranks around a benchmark W^* (Section 4.4). The spectrum-aware effective-dimension notion we adopt is standard and minimax-sharp in linear and kernel settings [Even and Massoulié, 2021]. In contrast, existing bounds either (i) rely on infinite-width linearizations (the NTK line of work, e.g., Jacot et al. [2018]), (ii) blow up exponentially with products of norms (e.g., Neyshabur et al. [2018], Bartlett et al. [2017]), or (iii) scale with model size (e.g., VC dimension [Bartlett et al., 2019]). Our bounds avoid these pathologies, providing a pointwise, spectrum-aware account grounded in unifying structural principles, clear links to prior frameworks, and qualified tightness discussions. By confronting the long-standing challenge of obtaining sharp, *nonlinear* guarantees for representation learning beyond uniform convergence—emphasized in Bartlett et al. [2021], Zhang et al. [2021], Neyshabur et al. [2017], Nagarajan and Kolter [2019], Wilson [2025]—we aim to show that generalization in deep neural networks can be mathematically tractable in the above sense.

Fully observable bounds and an open problem. On the technical issue of full observability from the empirical sample and learned features, it is important to clarify the role of the ghost sample and the unknown population distribution in our analysis. We first prove an unconditional compressed pointwise theorem under a mixed empirical-population feature metric. We then provide a fully observed-sample theorem that isolates the additional finite-resolution subspace observability needed to compute the same compressed spectrum from the training sample alone. Verifying this condition from intrinsic feature regularity is a substantive open problem, rather than a routine technical step. Algorithm design offers another potential route to resolving the observability problem, for example through sample splitting and cross-validation.

1.1 Organization and Contributions

The paper is organized into three parts: (i) a pointwise generalization framework (Section 2); (ii) structural principles of deep networks (Sections 3 to 4); and (iii) empirical validation (Section 5). Related work appears in Appendix A.1, and all proofs are in Appendices B, C, D, E. Below we summarize the main novelties in each part.

Pointwise Generalization and Finite-Scale Geometry. We develop a pointwise framework that analyzes the *trained* hypothesis and yields generalization bounds with (qualified) matching upper and lower rates via a finite-scale notion of *pointwise dimension*. This fundamentally upgrades generic chaining and all covering-number approaches by assigning each hypothesis its *own complexity* that directly controls its error. The bounds can also be read as an optimally tuned PAC–Bayes objective specialized to deterministic predictors. Taken together, this framework reframes generalization as a study of pointwise geometry governed by dimension reduction at finite precision, clarifying why nonlinear models can generalize even in the absence of uniform convergence.

Structural Principles and Tight Bounds for Neural Networks. We develop a *non-perturbative* analysis based on exact telescoping identities (rather than Taylor linearizations) that preserves the finite-scale geometry of deep networks. This yields our first structural principle: *cross-layer correlations factor through the feature matrices, approximately preserving a pointwise-linear structure* (Section 3.2). Next, we show that bounding the pointwise dimension reduces, on each local chart, to the gold-standard *effective dimension*, and we lift this to a global statement by constructing an *ellipsoidal* covering of the Grassmannian of subspaces. This extension—beyond classical differential-geometric/Lie-algebraic treatments—establishes our second structural principle: *the complexity of the global atlas (covering reference eigenspaces) is commensurate with that of the local charts*. Building on these principles, we introduce the *Riemannian Dimension*—a spectrum-aware, pointwise effective complexity—which governs generalization at the trained model and yields tight, analyzable bounds (Section 3.3). We provide explicit generalization guarantees for deep neural networks and argue that they are tight in a qualified sense. We establish a principled *exponential* improvement over spectral-norm bounds, provide an explicit separation, and discuss implications of our pointwise generalization theory for algorithm design (Section 4).

Unconditional compression, observability, and an open problem. The paper deliberately separates two statements that are often conflated. Theorem 5 is unconditional and already gives a compressed pointwise bound through the mixed empirical-population Riemannian Dimension (Section 4.1). Theorem 6 gives the fully observable upgrade: under a finite-resolution subspace isomorphism (Section 4.2), it replaces the population feature contribution by the observed spectrum. This is a sufficient finite-resolution condition in our derivation. We further pose its verification from subspace-level sub-Gaussian or small-ball feature regularity as an open question, identifying a precise route toward fully observed, pointwise feature-compressed bounds.

Empirical Findings and Evidences. The experiments are designed to systematically examine three central questions in modern deep learning: (i) why does overparameterization often improve generalization? (ii) how does feature learning evolve during training? and (iii) what implicit regularization is encoded by the baseline optimizer? Across the experimental results, we observe that

(i) the overparameterization impressively leads to decreasing Riemannian Dimension; (ii) feature learning compresses the effective ranks of learned features during the training; and (iii) stochastic gradient descent (SGD) with momentum implicitly regularizes the Riemannian Dimension.

2 The Nature of Pointwise Generalization

In this section, we develop our pointwise framework for generalization analysis, which introduces a tight tool—pointwise dimension—to characterize generalization. We illustrate its advancement to existing methodologies and bring new understandings on the nature of generalization.

2.1 Pointwise Dimension Strengthens PAC-Bayes and Generic Chaining

Let $\mathcal{F}$ be a hypothesis class, let $z \in \mathcal{Z}$ be a random data (e.g., input-label pair $z = (x, y)$), and let $\ell : \mathcal{F} \times \mathcal{Z} \rightarrow \mathbb{R}$ be a real-valued loss. Denote by $\mathbb{P}$ the (unknown) population distribution, and by $\mathbb{P}_n$ the empirical distribution associated with an i.i.d. sample $S = \{Z_i\}_{i=1}^n \sim \mathbb{P}^{\otimes n}$. For any integrable function $g : \mathcal{Z} \rightarrow \mathbb{R}$, define the population and empirical averaging operators by

$$\mathbb{P}g := \mathbb{E}_{z \sim \mathbb{P}}[g(z)], \quad \mathbb{P}_ng := \frac{1}{n} \sum_{i=1}^n g(z_i).$$

For convenience, we will often write $g(z)$ inside these operators (e.g., $g(z) = \ell(f; z)$). Our goal is to control the *generalization gap* $(\mathbb{P} - \mathbb{P}_n)\ell(f; z)$ in the following *pointwise generalization* manner: for $\delta \in (0, 1)$, with probability at least $1 - \delta$, uniformly over every $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) := \mathbb{E}_{z \sim \mathbb{P}}[\ell(f; z)] - \frac{1}{n} \sum_{i=1}^n \ell(f; z_i) \leq C \sqrt{\frac{d(f) + \log \frac{1}{\delta}}{n}}, \quad (2.1)$$

where $d(f)$ is a hypothesis-dependent complexity measure that aims to characterize the intrinsic complexity of every *trained* hypothesis f , different from class-wide, uniformly defined complexity measures. In Appendix B.1, we state necessary and sufficient conditions for pointwise generalization through the “uniform pointwise convergence” principle proposed in Xu and Zeevi [2020, 2025]. This stands in opposition to the uniform-convergence paradigm [Vapnik, 2013], which seeks a single worst-case bound on $\sup_{f \in \mathcal{F}} (\mathbb{P} - \mathbb{P}_n)\ell(f; z)$ rather than the precise, hypothesis-dependent (pointwise) control we pursue here. The failure of class-wide uniform convergence to meaningfully capture generalization in deep networks—in sharp contrast to its success in classical linear and kernel models—motivates a fundamental departure toward tight, hypothesis-dependent complexity measures that can distinguish generalizable solutions from arbitrary interpolants [Zhang et al., 2021, Neyshabur et al., 2017, Belkin et al., 2019, Nagarajan and Kolter, 2019, Wilson, 2025].

In the spirit of (2.1), we introduce the central notion of this section, the *pointwise dimension*: a finite-scale analogue of asymptotic notions from fractal geometry [Falconer, 1997], and a measure of the local complexity encoded by the Fernique-Talagrand integral in generic chaining [Fernique, 1974, Talagrand, 2014]. In addition, Maurer [2010], Chu and Raginsky [2023] provide direct antecedents of our pointwise-dimension viewpoint: they establish simultaneous bounds in terms of point-dependent variants of the Fernique-Talagrand integral. We formalize the underlying ball-mass complexity as a finite-scale analogue of classical pointwise dimension in general metric spaces and derive simple, direct generalization bounds in terms of this quantity. Throughout the paper, “metric” ϱ means a *semi-metric*: all metric axioms hold except that $\varrho(f_1, f_2) = 0$ need not imply $f_1 = f_2$.

Definition 1 (Pointwise Dimension) Given a function class $\mathcal{F}$, a metric ϱ on $\mathcal{F}$, and a prior π over $\mathcal{F}$, the local dimension at f with scale ε is defined as the log inverse density of the ε -ball $B_\varrho(f, \varepsilon) = \{f' \in \mathcal{F} : \varrho(f, f') \leq \varepsilon\}$ centered at f :

$$\log \frac{1}{\pi(B_\varrho(f, \varepsilon))}. \quad (2.2)$$

We first state the one-shot bound suggested by PAC-Bayes optimization, and then the multiscale bound obtained by generic chaining. Write $\mathbb{P}_S = \mathbb{P}_n$ and define the mixed empirical-population metrics

$$\begin{aligned} \varrho_{(S, \mathbb{P}), \ell}^{(1)}(f, g) &:= (\mathbb{P}_S + \mathbb{P})|\ell(f; z) - \ell(g; z)|, \\ \varrho_{(S, \mathbb{P}), \ell}(f, g) &:= \sqrt{(\mathbb{P}_S + \mathbb{P})(\ell(f; z) - \ell(g; z))^2}. \end{aligned} \quad (2.3)$$

For the one-shot bound, we can use the smaller gap distance

$$\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}(f, g) := |(\mathbb{P} - \mathbb{P}_n)(\ell(f; z) - \ell(g; z))|, \quad (2.4)$$

$$\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}(f, g) \leq \varrho_{(S, \mathbb{P}), \ell}^{(1)}(f, g) \leq \sqrt{2} \varrho_{(S, \mathbb{P}), \ell}(f, g). \quad (2.5)$$

This is the smallest symmetric distance controlling differences of generalization gaps. It depends on the unknown population gaps; the mixed L_1 distance gives a useful relaxation in terms of loss differences, while the mixed L_2 metric is used by chaining and the subsequent DNN analysis.

Theorem 1 (One-Shot Pointwise Generalization Bound) Let $\ell(f; z) \in [0, 1]$. For any prior π on $\mathcal{F}$ independent of S and any $\delta \in (0, 1)$, with probability at least $1 - \delta$, uniformly over every $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq C \left(\inf_{0 \leq \varepsilon \leq 2} \left\{ \underbrace{\varepsilon}_{\text{bias (approximate } f)} + \underbrace{\sqrt{\frac{\log \frac{1}{\pi(B_{\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}(f, \varepsilon))}}{n}}}_{\text{variance (PAC-Bayes term)}} \right\} + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right), \quad (2.6)$$

where $C > 0$ is an absolute constant.

Here, the bias parameter ε specifies the one-shot approximation resolution in the posterior optimization problem underlying this closed-form objective and can be optimized pointwise for each realized hypothesis f . Theorem 7 in Appendix B.2 provides a direct PAC-Bayes proof of this theorem. Theorem 3 establishes the worst-case optimality of its complexity term.

The concept of pointwise dimension and Theorems 1 and 2 strengthen several established generalization methodologies, including PAC-Bayesian analysis, Kolmogorov complexity, generic chaining, and covering numbers. We elaborate below. The short, direct proof based on PAC-Bayes optimization originally led us to identify pointwise dimension as a quantity governing generalization. The chaining approach extends this core idea and requires developing several new technical ideas and arguments.

Theorem 1 sharpens PAC-Bayes posterior optimization. A key novelty is that, by measuring the prior mass accumulated in a *localized metric ball* centered at f , rather than the prior mass assigned to the singleton $\{f\}$, the pointwise dimension framework naturally accommodates general uncountable classes. This circumvents the discreteness limitations inherent in hypothesis-singleton bounds corresponding to singleton localization, such as Occam-type or Kolmogorov-complexity bounds (e.g., Lotfi et al. [2022], Sutskever [2023]), thereby yielding strictly stronger guarantees; interested readers are referred to Lutz [2016] for related connections between pointwise dimension and algorithmic complexity. Additionally, our perspective brings the best possible PAC-Bayesian mechanism under the selected metric: generalization is recast as a bias–variance tradeoff optimized over a user–chosen posterior, applies to *deterministic* hypotheses, and shows that the pointwise dimension optimally governs the complexity; see Appendix B.2 for this perspective. This clarifies and strengthens earlier PAC-Bayes optimization approaches, which typically restrict the posterior to a tractable Gaussian family in the Euclidean weight space, rather than optimizing over arbitrary posteriors on the nonlinear hypothesis class, in order to obtain computable and nonvacuous bounds for random outputs [Hinton et al., 2012, Dziugaite and Roy, 2017] and deterministic hypotheses [Neyshabur et al., 2017, 2018]; see the end of Section 2.2 for more details.

In addition, the scope of Theorems 1 and 2 goes well beyond traditional local Rademacher complexity analysis. The key distinction is that our approach localizes the dimension factor—that is, the model complexity itself—through the “right” notion of pointwise dimension, rather than merely localizing a norm radius as in standard localization arguments. While the latter is sufficient for obtaining fast and adaptive rates, it does not capture the overparameterization phenomena studied in this paper. See the framework in Xu and Zeevi [2020, 2025], where the current work is positioned as showing that pointwise dimension is the central localized complexity for this purpose.

Theorem 2 (Multiscale Pointwise Generalization Bound) *Let $\ell(f; z) \in [0, 1]$. For any prior π on $\mathcal{F}$ independent of S and any $\delta \in (0, 1)$, with probability at least $1 - \delta$, uniformly over every $f \in \mathcal{F}$,*

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq C \left(\inf_{0 \leq \alpha \leq \sqrt{2}} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho_{(S, \mathbb{P}), \ell}}(f, \varepsilon))}} d\varepsilon \right\} + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right),$$

where $C > 0$ is an absolute constant and $\varrho_{(S, \mathbb{P}), \ell}$ is defined in (2.3).

The proof introduces a ghost sample and then applies Lemma 9 to obtain the displayed empirical–population bound. A sharper version retaining the conditional ghost expectation appears in Lemma 8. The main challenge for empirical processes is to establish a multiscale bound governed by the complexity of the realized hypothesis, integrating over scales ε rather than using a single resolution, in the spirit of earlier Gaussian-process bounds [Maurer, 2010, Chu and Raginsky, 2023]. Here, the prior π must be selected independently of the sample, in contrast to class-uniform tools such as generic chaining and empirical covering that rely on sample-dependent measures or covers. The truncation level α can be optimized pointwise for each realized hypothesis f , a feature that is also new to the chaining and covering literature. We retain truncation for technical convenience, although it plays a secondary role in the theory of the multiscale integral.

Theorem 2 upgrades generic chaining and covering numbers to a pointwise form. Building on the core idea behind the PAC-Bayes optimization bound in Theorem 1, we develop

several new technical ideas to obtain *pointwise* refinements of classical generic chaining bounds based on the Fernique-Talagrand integral [Fernique, 1974, Talagrand, 1987, 2014, Block et al., 2021]. The resulting *pointwise* generalization bounds can be strictly stronger than entropy-integral bounds based on *uniform* covering numbers (e.g., Dudley’s integral), whose integrand takes a supremum over the entire class $\mathcal{F}$ rather than localizing at the realized hypothesis. In particular, (B.21) in Appendix B.3 demonstrates that the *class-wide* fractional covering number

$$\inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{F}} \frac{1}{\pi(B_\varrho(f, \varepsilon))} \quad (2.7)$$

is (up to absolute constants) equivalent to the *canonical covering number* of $\mathcal{F}$ with metric ϱ at scale ε . Consequently, Theorem 2 goes beyond classical covering analyses by (i) recasting covering-number complexity as the inverse-prior-density objective (2.7), and (ii) localizing this complexity *pointwise* in f . Through the results and tools developed in this paper, we advocate the “prior-density + localization” (i.e., pointwise dimension) viewpoint as a new paradigm for statistical complexity analysis—one that is markedly more flexible than the classical covering-number approaches prevalent in statistics and machine learning.

Corollary 1 (Common One-Shot L_2 Relaxation) *Each of Theorems 1 and 2 implies the following bound: under their common hypotheses, with probability at least $1 - \delta$, uniformly over every $f \in \mathcal{F}$,*

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq C \left(\inf_{0 \leq \varepsilon \leq \sqrt{2}} \left\{ \varepsilon + \sqrt{\frac{\log \frac{1}{\pi(B_{\varrho(S, \mathbb{P}), \ell}(f, \varepsilon))}}{n}} \right\} + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right). \quad (2.8)$$

From Theorem 1, use (2.5) and the inclusion $B_{\varrho(S, \mathbb{P}), \ell}(f, \varepsilon) \subseteq B_{\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}}(f, \sqrt{2}\varepsilon)$. From Theorem 2, use that its integrand is nonincreasing, so the integral from α to $\sqrt{2}$ is bounded by $\sqrt{2}$ times the integrand’s value at α . Our pointwise optimization allows us to set the truncation level α equal to the chosen approximation resolution ε . In either case, enlarge the absolute constant C as needed.

The gap-distance, L_1 , and L_2 bounds. Theorem 1 also holds with $\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}$ replaced by $\varrho_{(S, \mathbb{P}), \ell}^{(1)}$. By (2.5), the gap-distance bound is no larger than this L_1 relaxation, which in turn sharpens the L_2 one-shot bound up to absolute constants. At a fixed L_2 metric, chaining improves the one-shot relaxation, but there is no general dominance order between the L_1 one-shot bound and the L_2 chaining bound: the former exploits L_1 approximation, whereas the latter exploits all scales. The DNN analysis below uses L_2 geometry to express complexity through feature spectra; the experiments evaluate its one-shot relaxation using the observed feature geometry, subject to the observability distinction in Section 4.2.

The multiscale improvement can be substantial in the same L_2 metric. For example, a pointwise dimension estimate of order $d(f)\varepsilon^{-2}$ gives a truncated-integral bound of order $\sqrt{d(f)/n}$ up to logarithmic factors, whereas the corresponding one-shot estimate is of order $(d(f)/n)^{1/4}$.

The next theorem distinguishes worst-case sharpness from an instancewise characterization. Let $A_\pi(f)$ and $C_\pi(f)$ denote the infima in Theorems 1 and 2, respectively, excluding their common confidence remainder and outer constants.

Theorem 3 (Lower Bounds and Scope of Tightness) *Let $\ell(f; z) \in [0, 1]$ and $n \geq 2$. There are positive absolute constants c, C with the following properties.*

One-shot gap-distance lower bound. *For every sample and every prior π , we have*

$$\sup_{f \in \mathcal{F}} |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| \geq \frac{1}{2} \sup_{f \in \mathcal{F}} A_\pi(f). \quad (2.9)$$

Chaining integral lower bound. *With $\varrho_{n,\ell}(f, g)^2 = \mathbb{P}_n(\ell(f; z) - \ell(g; z))^2$,*

$$\mathbb{E} \sup_{f \in \mathcal{F}} |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| \geq \frac{c}{\sqrt{n \log n}} \mathbb{E} \inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{F}} \int_0^1 \sqrt{\log \frac{1}{\pi(B_{\varrho_{n,\ell}}(f, \varepsilon))}} d\varepsilon - \frac{C \sup_{f \in \mathcal{F}} \mathbb{P}\ell(f; z)}{\sqrt{n \log n}}.$$

The absolute value may be omitted if the loss class is closed under complementation.

Both upper bounds can be sharp. *For every pair of integers $2 \leq d \leq n$, there is a finite loss class and a distribution for which, under the uniform prior U , each of $\mathbb{E} \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)|$ and $\mathbb{E} \sup_f A_U(f)$ lies between $c\sqrt{d/n}$ and $C\sqrt{d/n}$, and*

$$c\sqrt{\frac{d}{n}} \leq \mathbb{E} \inf_{\pi} \sup_f C_\pi(f) \leq \mathbb{E} \sup_f C_U(f) \leq C\sqrt{\frac{d}{n}}. \quad (2.10)$$

The empirical and mixed chaining functionals are also equivalent in expectation on this family, both with and without pointwise optimized truncation.

Chaining need not be sharp on every problem. *For every $n \geq 2$ and $0 < p < 1$, there is another finite loss class and distribution for which each of $\mathbb{E} \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)|$ and $\mathbb{E} \sup_f A_U(f)$ lies between cp and Cp , but*

$$\frac{\sqrt{p}}{2} \leq \mathbb{E} \inf_{\pi} \sup_f C_\pi(f) \leq \mathbb{E} \sup_f C_U(f) \leq \sqrt{2p}. \quad (2.11)$$

In both constructions, U is independent of the sample. The chaining lower bounds hold even when the infimum over priors is taken after observing S .

Theorem 3 establishes worst-case sharpness of the complexity terms in our one-shot and multiscale pointwise generalization bounds (Theorems 1 and 2). For the one-shot bound, the deterministic comparison (2.9) unconditionally controls the class-wide complexity by twice the largest absolute generalization gap for every loss class and distribution, although it uses the unknown population gaps. For the multiscale bound, sharpness is established in the weaker existential sense through matching examples.

The worst-case statements are analogous in spirit to the minimax optimality in frequentist decision theory [Wald, 1945]. It does not preclude sharper guarantees for a fixed hypothesis f . A matching pointwise lower bound, without the outer supremum over $f \in \mathcal{F}$, cannot hold in general: a hypothesis may have zero generalization gap while its complexity term remains positive. Appendix B.6 provides the proofs and constructions illustrating both sharpness and limitations. These statements also highlight the importance of metric choice for both pointwise generalization and uniform convergence.

We defer technical innovations and connections to existing methodologies to the Appendix—most notably the unified pointwise-generalization framework of Xu and Zeevi [2020, 2025] which we build upon (Appendix B.1), and the alternative PAC-Bayesian perspective (Appendix B.2). We conclude that the proposed *pointwise dimension* yields hypothesis-specific generalization guarantees whose complexity terms can be sharp in the worst case.

2.2 Necessity of Finite-Scale Pointwise Geometry and Structural Analysis

The transition from uniform convergence to the “prior-density + localization” (pointwise dimension) perspective offers a fundamental tightening over standard covering number approaches. However, translating this theoretical advantage into a practical framework for deep learning requires addressing two distinct challenges. First, we will distinguish the *geometric nature* of generalization from classical infinitesimal geometry: relevance lies not in the limit $\varepsilon \rightarrow 0$, but motivates a new program of finite-scale geometric analysis. Second, we must overcome the *computational intractability* of evaluating the pointwise dimension directly, which necessitates a dedicated structural analysis for deep neural networks.

Asymptotic vs. Finite-Scale Dimension. Although powerful in mathematics, standard differential-geometric tools (e.g., pointwise metrics and subspace angles) have not been systematically used in generalization theory, largely because they define dimension in infinitesimal notions. For instance, the *asymptotic pointwise dimension*—central to fractal and Riemannian geometry [Falconer, 1997, Jost, 2008] and used to characterize Hausdorff and packing dimensions (e.g., Theorem 3 of Lutz [2016])—is defined via a limit:

$$\lim_{\varepsilon \rightarrow 0} \frac{\log \pi(B_\varrho(f, \varepsilon))}{\log \varepsilon}.$$

We argue that generalization is distinct from, and in some ways more challenging than, infinitesimal geometry: the nature of generalization in deep models lies in reducing geometric dimension at a *finite scale* of precision for each hypothesis. Crucially, the pointwise dimension $\log \frac{1}{\pi(B_\varrho(f, \varepsilon))}$ is monotonic: it naturally decreases as the resolution ε increases. Therefore, a finite-scale analysis reveals significant dimension reduction that infinitesimal analysis misses. In our one-shot bound (2.6), the objective is to identify the optimal finite scale ε^* where the trade-off between precision and pointwise complexity is minimized. At this scale, the effective dimension can be orders of magnitude smaller than the asymptotic dimension, explaining the tractability of overparameterized models. To the best of our knowledge, this distinction is novel; prior uses of geometric dimension in generalization (e.g., Birdal et al. [2021]) have largely emphasized globally uniform and infinitesimal notions. And the Neural Tangent Kernel (NTK) [Jacot et al., 2018] and Gaussian-process [Lee et al., 2018] viewpoints are valid only in an infinitesimal neighborhood of initialization (equivalently, in the infinite-width regime). A precise account of deep-model generalization thus calls for a shift from infinitesimal calculus to finite-scale, pointwise geometry.

Computation and the Necessity of Structural Analysis. Although worst-case sharp, Theorems 1 and 2—like many abstract bounds (PAC-Bayes, mutual-information, generic chaining)—is generally not computationally tractable on its own; practical use requires adapting it to the function class at hand and introducing suitable relaxations. If we denote an effective dimension by $d(f) = \log \frac{1}{\pi(B_\varrho(f, \varepsilon^*))}$ (ε^* tuned in one-shot optimization), a brute-force Monte Carlo estimator using i.i.d. draws $f' \sim \pi$ would require on the order of $e^{d(f)}$ samples to obtain a single hit $f' \in B_\varrho(f, \varepsilon)$ with constant probability. For high-dimensional deep networks, where $d(f)$ should be moderate to large, this is computationally prohibitive.

This intractability helps explain why much of the PAC-Bayes literature turns to tractable weight-space surrogates that are effectively linearized: instead of optimizing over arbitrary posteriors on the nonlinear hypothesis class $\mathcal{F}$, one restricts to Gaussian priors and posteriors over

the weights $W \in \mathbb{R}^p$, typically with isotropic or fixed covariance structure. This restriction yields closed-form KL terms and computable objectives [Hinton and Van Camp, 1993, Dziugaite and Roy, 2017]; see Sections 3 and 6 of Dziugaite and Roy [2017] for representative formulations. However, this strategy implicitly imposes a uniform linearization that discards the distinctive *pointwise* geometry of deep networks, effectively flattening a curved manifold. To retain the sharpness of pointwise dimension without incurring the simulation barrier, we therefore avoid black-box sampling and instead develop explicit *structural principles* of deep networks that allow analytic control of the pointwise dimension—yielding generalization guarantees that are both theoretically rigorous and practically computable.

3 Deep Neural Networks and Pointwise Riemannian Dimension

In this section we develop a systematic pointwise dimension analysis for deep neural networks, focusing on the empirical DNN geometry induced by the observed sample S ; the separate reduction from the mixed empirical-population metric in Theorem 2 to this observed-sample complexity is handled in Section 4. Section 3.1 formalizes the standard fully connected architecture and notation. Section 3.2 introduces a non-perturbative calculus (avoiding infinitesimal Taylor expansions) to analyze finite-scale behavior. Section 3.3 introduces a hierarchical covering scheme—our key technical innovation—that overcomes the well-known linear/kernel bottleneck in classical statistical learning and enables a principled treatment of genuinely nonlinear models.

3.1 Neural Network Setup

We consider fully connected (feed-forward) networks that map an input $x \in \mathbb{R}^{d_0}$ to an output $f_L(W, x) \in \mathbb{R}^{d_L}$. The architecture is specified by widths $d_0, \dots, d_L$ and weight matrices $W = \{W_1, \dots, W_L\}$ with $W_l \in \mathbb{R}^{d_l \times d_{l-1}}$ for $l = 1, \dots, L$. Let $\sigma_1, \dots, \sigma_L$ be nonlinear activations (e.g., ReLU), acting componentwise on column vectors, and each $\sigma_l : \mathbb{R}^{d_l} \rightarrow \mathbb{R}^{d_l}$ is assumed 1-Lipschitz. The network’s forward map is the composition

$$f_L(W, x) := \sigma_L \left(W_L \sigma_{L-1} (W_{L-1} \cdots \sigma_1 (W_1 x)) \right).$$

Let $X = [x_1, \dots, x_n] \in \mathbb{R}^{d_0 \times n}$ collect the n training inputs as columns. For each layer $l \in \{1, \dots, L\}$, define the depth- l map and the corresponding *feature matrix*

$$f_l(W, x) := \sigma_l \left(W_l \sigma_{l-1} (W_{l-1} \cdots \sigma_1 (W_1 x)) \right), F_l(W, X) := [f_l(W, x_1) \cdots f_l(W, x_n)] \in \mathbb{R}^{d_l \times n}.$$

Equivalently (full, non-recursive form consistent with (1.1)),

$$F_l(W, X) = \sigma_l \left(W_l \sigma_{l-1} (W_{l-1} \cdots \sigma_1 (W_1 X)) \right),$$

where for a matrix $A = [a_1, \dots, a_n]$ we write $\sigma_l(A) := [\sigma_l(a_1), \dots, \sigma_l(a_n)]$. Thus $F_L(W, X)$ collects the network outputs on the dataset X .

We denote $\|\cdot\|_{\mathbf{F}}$ for the Frobenius norm, $\|\cdot\|_{\text{op}}$ for the spectral norm, and $\|\cdot\|_2$ for the Euclidean norm on vectors. We abbreviate norm balls by $B_{\mathbf{F}}(R)$, $B_{\text{op}}(R)$, and $B_2(R)$ (all centered at 0; with

radius R). The empirical $L_2(\mathbb{P}_n)$ distance between two hypotheses W, W' is (a $1/\sqrt{n}$ scaling is used to keep consistency with Section 2)

$$\varrho_n(W, W') := \sqrt{\|F_L(W, X) - F_L(W', X)\|_{\mathbf{F}}^2/n}.$$

The function-level empirical metric and generalization statements in Section 2 for the loss $x \mapsto \ell(f_L(W, x), y)$ at data-label pairs $z = (x, y)$ specialize, on the dataset X , to the metric ϱ_n defined above. We assume the loss $\ell(\cdot, y)$ is β -Lipschitz in its first argument with respect to $f_L(W, x)$. This bridges the loss-induced metric on $\mathcal{F}$, studied in Section 2, with the weight-space metric used here.

3.2 Non-Perturbative Expansion and Layer-wise Correlation

Throughout, our finite-scale analysis relies on *non-perturbative* expansions. Borrowing terminology from theoretical physics, “non-perturbative” here means we avoid Taylor/derivative expansions and instead use exact, telescoping algebraic identities that hold at finite scale. For example,

$$W'_2 W'_1 - W_2 W_1 = W'_2(W'_1 - W_1) + (W'_2 - W_2)W_1, \quad \Sigma'^{-1} - \Sigma^{-1} = \Sigma'^{-1}(\Sigma - \Sigma')\Sigma^{-1},$$

with analogous decompositions used throughout. This viewpoint preserves the full finite-scale geometry of deep networks, rather than linearizing around an infinitesimal neighborhood.

To present our non-perturbative expansion for DNN, we define *local Lipschitz constant* as follows.

Definition 2 (Local Lipschitz Constant) For each $l = 1, \dots, L$, we define $M_{l \rightarrow L}(W, \varepsilon)$ as the (outer) local Lipschitz constant, which characterizes the sensitivity of the layer L output, F_L , to variations in layer l ’s output, within a neighborhood around F_l . Formally, we assume that for every $W' \in B_{\varrho_n}(W, \varepsilon)$

$$\|F_L(F_l(W', X), \{W'_i\}_{i=l+1}^L) - F_L(F_l(W, X), \{W_i\}_{i=l+1}^L)\|_{\mathbf{F}} \leq M_{l \rightarrow L}(W, \varepsilon) \|F_l(W', X) - F_l(W, X)\|_{\mathbf{F}}.$$

Local Lipschitz constants are typically much smaller than products of spectral norms and can be computed by formal-verification toolchains [Shi et al., 2022]. In our bounds, these constants enter only through *optimistic logarithmic factors* and therefore do not affect the leading rates; see Example 1 for the precise sense in which the relevant effective dimension is logarithmically sensitive to such multiplicative scale factors. We propose a telescoping decomposition to replace conventional Taylor expansion, where in each summand the only difference lies in W'_l and W_l .

$$\begin{aligned} & F_L(W', X) - F_L(W, X) \\ &= \sum_{l=1}^L \underbrace{[\sigma_L(W'_L \cdots W'_{l+1}) \underbrace{\sigma_l(W'_l F_{l-1}(W, X))}_{\text{learned feature}}] - \sigma_L(W'_L \cdots W'_{l+1}) \underbrace{\sigma_l(W_l F_{l-1}(W, X))}_{\text{learned feature}}]}_{\text{controlled by } M_{l \rightarrow L} \text{ by 1}}. \end{aligned} \quad (3.1)$$

Note that this is a *non-perturbative* expansion that holds unconditionally and does not rely on infinitesimal approximation, and crucially keeps the *learned* feature $F_{l-1}(W, X)$ at the *trained* weight W . From this decomposition and applying basic inequalities, we have the following key lemma.

Lemma 1 (Non-Perturbative Feature Expansion) For all $W' \in B_{\varrho_n}(W, \varepsilon)$,

$$\|F_L(W', X) - F_L(W, X)\|_{\mathbf{F}}^2 \leq \sum_{l=1}^L L \cdot M_{l \rightarrow L}[W, \varepsilon]^2 \cdot \|(W'_l - W_l) F_{l-1}(W, X)\|_{\mathbf{F}}^2.$$

The lemma captures the first structural principle of fully connected DNN: cross-layer correlations mostly pass through the feature matrices, preserving an approximate pointwise linear structure.

Since enlarging the metric only shrinks metric balls and hence *increases* the pointwise dimension (2.2) we analyze in Section 2 (formalized as Lemma 18; metric domination lemma), it suffices to analyze pointwise dimension under the *pointwise ellipsoidal metric* that appears on the right-hand side of Lemma 1. Concretely, $F_{l-1}(W, X) F_{l-1}(W, X)^\top$, the feature Gram matrix from layer $l-1$, faithfully encodes the spectral information induced by the network-data geometry at layer l . Working with the corresponding pointwise ellipsoidal metric yields sharp, *pointwise*, *spectrum-aware* bounds with the desired properties for deep networks, and underpins our tractability results (with the structural principles and technical innovations to be developed in the next subsection).

3.3 Hierarchical Covering from Local Chart to Global Atlas

Lemma 1 suggests that the following *pointwise ellipsoidal metric* dominates $n \cdot \varrho_n$ at every W (here, NP stands for “non-perturbative”):

$$\begin{aligned} G_{\text{NP}}(W) &= \text{blockdiag} \left(\cdots, LM_{l \rightarrow L}^2(W, \varepsilon) \cdot F_{l-1}(W, X) F_{l-1}^\top(W, X) \otimes I_{d_l}, \cdots \right) \\ \varrho_{G_{\text{NP}}(W)}(W, W')^2 &= \text{vec}(W' - W)^\top G_{\text{NP}}(W) \text{vec}(W' - W). \end{aligned} \quad (3.2)$$

We are therefore interested in bounding the enlarged pointwise dimension under the pointwise ellipsoidal metric $\varrho_{G_{\text{NP}}(W)}$:

$$\log \frac{1}{\pi(B_{\varrho_n}(f(W, \cdot), \varepsilon))} \leq \log \frac{1}{\pi(B_{\varrho_{G_{\text{NP}}(W)}}(W, \sqrt{n}\varepsilon))}.$$

This section offers a deep dive past classical effective dimension, shifting to hierarchical covering and a global geometric analysis.

3.3.1 Gold Standard: Effective Dimension

Classical studies of static ellipsoidal metrics suggest that if π is chosen to be uniformly constrained on the top- r eigenspace of a PSD matrix $G(W)$, and the vectorized weights $W \in \mathbb{R}^p$ are restricted to the Euclidean ball $B_2(R) := \{w \in \mathbb{R}^p : \|w\|_2 \leq R\}$, then one can achieve a tight effective dimension as follows: define the *effective rank*

$$r_{\text{eff}}(G(W), R, \varepsilon) := \max\{k : \lambda_k(G(W))R^2 \geq n\varepsilon^2/2\}, \quad (3.3)$$

where the eigenvalues $\{\lambda_k(G(W))\}$ are ordered nonincreasingly; and define the spectrum-aware *effective dimension*

$$d_{\text{eff}}(G(W), R, \varepsilon) := \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}(G(W), R, \varepsilon)} \log \left(\frac{8R^2 \lambda_k(G(W))}{n\varepsilon^2} \right). \quad (3.4)$$

This definition serves as a gold standard for static ellipsoidal metrics and is asymptotically tight, as established by the covering number of the unit ball with ellipsoids in Section 3.3 of [Even and Massoulié \[2021\]](#). For this static ellipsoidal metric, the same volume-ratio calculation also controls,

up to absolute constants, the local prior mass of ellipsoidal balls under the uniform measure on the top- r eigenspace. Thus d_{eff} is also the gold standard for pointwise dimension. For brevity, we write r for $r_{\text{eff}}(G(W), R, \varepsilon)$, and denote by $\mathcal{V} \subseteq \mathbb{R}^p$ the r -dimensional subspace corresponding to the top- r_{eff} eigenspace of $G(W)$.

Example 1 (Logarithmic Sensitivity to Scale) *It is useful to record two elementary regimes in which the dependence of (3.4) on the radius R , the precision ε , and any multiplicative rescaling of the metric tensor is only logarithmic, in the optimistic sense used throughout this paper. Let*

$$\tilde{G}(W) := aG(W), \quad a > 0,$$

so that a may represent, for instance, the layerwise multiplier $LM_{l \rightarrow L}^2(W, \varepsilon)$ appearing in the neural-network metric tensor.

Exponential spectral decay. *Suppose that for some $\lambda_0 > 0$ and $\gamma > 0$,*

$$\lambda_k(G(W)) \leq \lambda_0 e^{-\gamma(k-1)}, \quad k \geq 1.$$

Then by the definition of effective rank $r_{\text{eff}}(\tilde{G}(W), R, \varepsilon) = \max\{k : a\lambda_k(G(W))R^2 \geq n\varepsilon^2/2\}$,

$$r_{\text{eff}}(\tilde{G}(W), R, \varepsilon) \leq 1 + \frac{1}{\gamma} \left[\log \left(\frac{2a\lambda_0 R^2}{n\varepsilon^2} \right) \right]_+, \quad [t]_+ := \max\{t, 0\}.$$

Moreover, if the decay is exact, $\lambda_k(G(W)) = \lambda_0 e^{-\gamma(k-1)}$, and $r = r_{\text{eff}}(\tilde{G}(W), R, \varepsilon)$, then

$$d_{\text{eff}}(\tilde{G}(W), R, \varepsilon) = \frac{1}{2} \sum_{k=1}^r \log \left(\frac{8R^2 a \lambda_0 e^{-\gamma(k-1)}}{n\varepsilon^2} \right) = \frac{r}{2} \log \left(\frac{8a\lambda_0 R^2}{n\varepsilon^2} \right) - \frac{\gamma}{4} r(r-1).$$

In particular, in the general upper-bound case, combining the two bounds above yields

$$d_{\text{eff}}(\tilde{G}(W), R, \varepsilon) \leq C_\gamma \left[\log \left(e + \frac{a\lambda_0 R^2}{n\varepsilon^2} \right) \right]^2,$$

where C_γ depends only on the exponential-decay rate. Thus even a large multiplicative factor a changes the effective dimension only through a squared logarithm.

Strict low rank. *Suppose instead that $\text{rank}(G(W)) \leq q$ and $\lambda_1(G(W)) \leq \lambda_0$. Then, for the rescaled metric $\tilde{G}(W) = aG(W)$,*

$$r_{\text{eff}}(\tilde{G}(W), R, \varepsilon) \leq q,$$

and hence

$$d_{\text{eff}}(\tilde{G}(W), R, \varepsilon) = \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}} \log \left(\frac{8R^2 \lambda_k(\tilde{G}(W))}{n\varepsilon^2} \right) \leq \frac{r_{\text{eff}}}{2} \log \left(e + \frac{8aR^2 \lambda_0}{n\varepsilon^2} \right) \leq \frac{q}{2} \log \left(e + \frac{8a\lambda_0 R^2}{n\varepsilon^2} \right).$$

In this strict feature-compression regime, the leading dependence is the intrinsic rank q , not the ambient dimension p , while R , ε^{-1} , and the metric multiplier a enter only through a single logarithm. This is the precise sense in which the effective-dimension calculus demotes products of norms and local-Lipschitz multipliers from leading-order complexity to logarithmic scale factors whenever the learned feature spectrum is compressed.

3.3.2 Key Challenge: Prior Independence from W .

However, the main challenge is that the prior π must be chosen independently of the training data. This means that the construction of π cannot rely on knowledge of the learned weights W , including their top- r_{eff} eigenspace, yet still capture the underlying geometric structure. The next lemma extends classical results on static ellipsoidal metrics by showing that a uniform prior over a reference subspace $\bar{\mathcal{V}}$ suffices to bound the pointwise dimension for all W whose top- r eigenspace of $G(W)$ can be approximated by $\bar{\mathcal{V}}$.

Lemma 2 (Pointwise Dimension via Reference Subspace) *Consider the weight space $B_2(R) \subset \mathbb{R}^p$ for vectorized weights, and a pointwise ellipsoidal metric defined via PSD $G(W)$. Let $\bar{\mathcal{V}} \subseteq \mathbb{R}^p$ be a fixed r -dimensional subspace. Define the prior $\pi_{\bar{\mathcal{V}}} = \text{Unif}(B_2(1.58R) \cap \bar{\mathcal{V}})$. Then, uniformly over all (W, ε) such that the top- r eigenspace $\mathcal{V}$ of $G(W)$ can be approximated by $\bar{\mathcal{V}}$ to precision*

$$\varrho_{\text{proj}, G(W)}(\mathcal{V}, \bar{\mathcal{V}}) := \|G(W)^{1/2}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}} \leq \frac{\sqrt{n\varepsilon}}{4R}, \quad (3.5)$$

we have

$$\log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\varrho_{G(W)}}(W, \sqrt{n\varepsilon}))} \leq \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}(G(W), R, \varepsilon)} \log \left(\frac{40R^2 \lambda_k(G(W))}{n\varepsilon^2} \right) \leq d_{\text{eff}}(G(W), \sqrt{5}R, \varepsilon).$$

In (3.5), $\mathcal{P}_{\mathcal{V}}$ denotes the orthogonal projector onto the subspace $\mathcal{V}$, and $\varrho_{\text{proj}, G(W)}$ thus defines an ellipsoidal projection metric between subspaces. Further details are provided in the appendix.

3.3.3 Hierarchical covering (mixture prior over subspaces).

We introduce a hierarchical covering framework that pushes statistical learning beyond classical linear and kernel paradigms, providing a principled toolkit for genuinely nonlinear models—one of the central innovations of this work. It operates on two levels: a bottom-level local-chart covering that captures spectrum-aware behavior within a fixed subspace, and a top-level global geometric analysis over the Grassmannian.

(i) For each reference subspace $\bar{\mathcal{V}}$, placing a uniform prior on $\bar{\mathcal{V}}$ yields a tight pointwise-dimension bound for all “local” weights W whose top- r eigenspace of $G(W)$ is well approximated by $\bar{\mathcal{V}}$ (see Lemma 2).

(ii) At the top level, we place a prior over reference subspaces $\bar{\mathcal{V}}$ and average the local priors, producing a data-independent prior and the final bound.

By combining these two levels of priors, we obtain a pointwise dimension bound using a prior π that is completely blind to the choice of W . To formalize this, we introduce a top-level distribution μ over the Grassmannian

$$\text{Gr}(p, r) := \{r\text{-dimensional linear subspaces of } \mathbb{R}^p\}$$

the collection of all r -dimensional subspaces, and define

$$\pi(W) = \sum_{\mathcal{V}} \mu(\mathcal{V}) \pi_{\mathcal{V}}(W).$$

We refer to this two-stage construction as the hierarchical covering argument. Under the resulting prior π , the following bound holds uniformly over all (vectorized) $W \in B_2(R)$, the pointwise dimension $\log \frac{1}{\pi(B_{\ell_{G(W)}}(W, \sqrt{n\varepsilon}))}$ is bounded by two parts:

$$\underbrace{\log \frac{1}{\mu(B_{\ell_{\text{proj}, G(W)}}(\mathcal{V}, \sqrt{n\varepsilon/4R}))}}_{\text{covering Grassmannian (global atlas)}} + \underbrace{\sup_{\bar{\mathcal{V}} \in B_{\ell_{\text{proj}, G(W)}}(\mathcal{V}, \sqrt{n\varepsilon/4R})} \log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\ell_{G(W)}}(W, \sqrt{n\varepsilon}))}}_{\text{covering local charts}}, \quad (3.6)$$

In differential–geometric terms, our argument has two components.

- *Local (chart) analysis:* fixing a reference subspace $\bar{\mathcal{V}}$, we use effective dimension as the gold standard to determine the metric entropy of the corresponding local chart.
- *Global (atlas) covering:* we cover the Grassmannian by such reference subspaces, i.e., we bound the metric entropy of the global atlas and account for the cost of transitioning across local charts.

Lemma 2 controls the local part, while the following new result (Lemma 3) on the *ellipsoidal* covering of the Grassmannian controls the global part.²

Lemma 3 (Ellipsoidal Covering of the Grassmannian manifold) *Consider the Grassmannian $\text{Gr}(d, r)$. For uniform prior $\mu = \text{Unif}(\text{Gr}(d, r))$, we have that for every $\mathcal{V} \in \text{Gr}(d, r)$, every $\varepsilon > 0$ and every PSD matrix Σ with eigenvalues $\lambda_1 \geq \dots \geq \lambda_d$, we have the pointwise dimension bound*

$$\log \frac{1}{\mu(B_{\ell_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} \leq \frac{d-r}{2} \sum_{k=1}^r \log \frac{C \max\{\lambda_k, \varepsilon^2\}}{\varepsilon^2} + \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{C \max\{\lambda_k, \varepsilon^2\}}{\varepsilon^2},$$

where $C > 0$ is an absolute constant.

The result above is mathematically significant in its own right. It extends the classical metric-entropy (covering number) theory for the Grassmannian—where $\log$ covering number $\asymp r(d-r) \log(C/\varepsilon)$ under the *isotropic* projection metric—to an *ellipsoidal* (anisotropic) metric that captures feature- and model-induced geometry. This generalization translates the traditional differential-geometric and Lie-algebraic treatments (see Appendix D) and, we believe, illustrates a two-way exchange: deep mathematical structure is essential to understanding generalization in modern neural networks, and, conversely, generalization theory can motivate new questions and results in pure mathematics.

Leveraging the block-decomposable structure in (3.2), the l -th block is

$$G_l(W) = A_l(W) \otimes I_{d_l}, \text{ where } A_l(W) = LM_{l \rightarrow L}^2(W, \varepsilon) \cdot F_{l-1}(W, X) F_{l-1}(W, X)^\top \in \mathbb{R}^{d_{l-1} \times d_{l-1}}.$$

²Since the effective rank r of $\bar{\mathcal{V}}$ can take any value in $\{1, \dots, p\}$, the top-level Grassmannian covering must range over all $\text{Gr}(p, r)$. This adds only a negligible $O(\log p)$ overhead to the global-level cost. Accordingly, we construct a data-independent prior in a three-level prior hierarchy: (“global- r ”) choose the rank r (paying the $\log p$ overhead), (“global- $\bar{\mathcal{V}}$ ”) choose a reference subspace $\bar{\mathcal{V}} \in \text{Gr}(p, r)$, and (“local”) sample within $\bar{\mathcal{V}}$ using the local chart prior; see Figure 1 for an illustration. For conceptual clarity, Lemma 3 focuses on the Grassmannian covering cost at a fixed rank r ; and we defer the layer-specific specialization (to each $d_{l-1} \times d_{l-1}$ feature Gram-matrix block) to the calculation in (3.8).

Since the Kronecker factor is I_{d_l} , the spectrum of G_l consists of the eigenvalues of the feature Gram matrix $F_{l-1}F_{l-1}^\top$ (scaled by $LM_{l \rightarrow L}^2$), each repeated d_l times. For readability, in the following display we suppress the radius and precision parameters in d_{eff} ; the precise convention is given in Theorem 4 below. Consequently, the *local-chart* (within-subspace) covering cost at layer l scales as

$$d_l \cdot d_{\text{eff}}\left(LM_{l \rightarrow L}^2(W, \varepsilon) \cdot F_{l-1}(W, X)F_{l-1}(W, X)^\top\right), \quad (3.7)$$

while the *atlas* (subspace-selection) cost is the Grassmannian term over $\text{Gr}(d_{l-1}, r_{\text{eff}}[W, l])$, where $r_{\text{eff}}[W, l]$ is the effective rank of $A_l(W) \in \mathbb{R}^{d_{l-1} \times d_{l-1}}$. By Lemma 3 (and the footnote preceding it), the *global-atlas* (choosing-subspace) covering cost at layer ℓ scales as

$$d_{l-1} \cdot d_{\text{eff}}\left(LM_{l \rightarrow L}^2(W, \varepsilon) \cdot F_{l-1}(W, X)F_{l-1}(W, X)^\top\right) + \log(d_{l-1}). \quad (3.8)$$

Together, (3.7) and (3.8) yield a clean layerwise decomposition: the width d_l multiplies the spectral complexity of incoming features (local charts), whereas the input dimension d_{l-1} governs the Grassmannian covering (global atlas). This complementary, seemingly magical “duality” underlies the calculation below.

Theorem 4 (Riemannian Dimension for DNN) *Consider the weight space $B_{\mathbf{F}}(R)$, and a pointwise ellipsoidal metric defined via the ellipsoidal metric $G_{\text{NP}}(W)$ defined in (3.2). Define the pointwise Riemannian Dimension*

$$d_{\text{R}}(W, \varepsilon) = \sum_{l=1}^L \left(\underbrace{d_l \cdot d_{\text{eff}}(A_l(W))}_{\text{covering local charts}} + \underbrace{d_{l-1} \cdot d_{\text{eff}}(A_l(W))}_{\text{covering global atlas}} + \underbrace{\log(d_{l-1})}_{\text{covering discrete } r_{\text{eff}}} + \log n \right),$$

where $A_l(W)$ is the feature Gram matrix $LM_{l \rightarrow L}^2(W, \varepsilon) \cdot F_{l-1}(W, X)F_{l-1}^\top(W, X)$; and $d_{\text{eff}}(A_l(W))$ is abbreviation of $d_{\text{eff}}(A_l(W), C \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon)$ with $C > 0$ an absolute constant. Then we have the pointwise dimension bound: there exists a prior π such that uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\log \frac{1}{\pi(B_{\varrho_n}(f(W, \cdot), \varepsilon))} \leq d_{\text{R}}(W, \varepsilon).$$

This concludes our program for fully connected networks: we establish *Riemannian Dimension* as a principled complexity measure that explains—and sharply bounds—generalization. We summarize the *second structural principle of fully connected DNN*: The complexity of the *global atlas* (covering the space of reference top eigenspaces) remains commensurate with the layerwise, spectrum-aware complexity of covering the *local charts*. On closer inspection, the effect hinges on the block-decomposable structure in (3.2). This structure is intrinsic to layered neural networks and typically absent in generic nonlinear models, which helps explain why DNN are particularly amenable to sharp generalization analysis.

For intuition, we illustrate the construction of the prior π in the single-layer case—via the schematic in Figure 1. From a top-down view, the prior π can be generated by first sampling the effective rank r , then a subspace $\bar{\mathcal{V}}$ on the Grassmannian, and finally a weight W inside that subspace. The general L -layer setting is then obtained by applying the same construction independently to each layer and taking a product measure, which is enabled by the layer-wise decomposable structure of neural networks (a consequence of our non-perturbative analysis).




Figure 1: Hierarchical construction of the data-independent prior π and its role in the pointwise-dimension bound (one single-layer case).

4 Generalization Bounds and Implications

4.1 Unconditional Generalization Bound for DNN

Theorem 2 is stated in the mixed empirical-population metric. We first apply it to the learned feature geometry, and then give a fully observed-sample bound under a finite-resolution subspace condition. The first bound does not require that condition; the condition is used only to replace population feature information by the observed spectrum.

Write $F_l^S(W) := F_l(W, X_S)$ and define

$$\Sigma_l^{\mathbb{P}}(W) := \mathbb{E}_{X \sim \mathbb{P}}[f_l(W, X)f_l(W, X)^\top], \quad \Gamma_l^{S, \mathbb{P}}(W) := F_l^S(W)F_l^S(W)^\top + n\Sigma_l^{\mathbb{P}}(W).$$

Thus $n^{-1}\Gamma_l^{S, \mathbb{P}}$ is the feature second-moment matrix for $\mathbb{P}_S + \mathbb{P}$.

Let $d_R^{S, \mathbb{P}}(W, \varepsilon)$ be the Riemannian Dimension from Theorem 4 with $F_{l-1}F_{l-1}^\top$ replaced by $\Gamma_{l-1}^{S, \mathbb{P}}(W)$. The outer transfer constants $\bar{M}_{l \rightarrow L}(W, \varepsilon)$ are common bounds validating the non-perturbative expansion in $L_2(\mathbb{P}_S)$ and $L_2(\mathbb{P})$ on the same bounded weight charts used in Theorem 4. Uniform bounds over inputs on these charts are sufficient. These population transfer bounds, like the population feature moments, need not be observable from S ; the empirical result below uses the same valid bounds.

Combining Theorem 4 with Theorem 2 gives the following result.

Theorem 5 (Generalization Bound for DNN; mixed empirical-population form) *Assume the finite population feature moments and common transfer bounds specified above. Let the loss $\ell(f(W, x), y)$ be bounded in $[0, 1]$ and β -Lipschitz with respect to $f(W, x)$. For every $\delta \in (0, 1)$, with probability at least $1 - \delta$ over the observed sample S , uniformly over all $W \in B_{\mathbf{F}}(R)$,*

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C_1 \left(\beta \inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\mathbf{R}}^{S, \mathbb{P}}(W, c\varepsilon)} d\varepsilon \right\} + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right), \end{aligned}$$

where $c, C_1 > 0$ are absolute constants and

$$\begin{aligned} d_{\mathbf{R}}^{S, \mathbb{P}}(W, \varepsilon) = & \frac{1}{2} \sum_{l=1}^L \left((d_l + d_{l-1}) \underbrace{\sum_{k=1}^{r_{\text{eff}}^{S, \mathbb{P}}[W, l]} \log \frac{8C_2^2 \lambda_k(\Gamma_{l-1}^{S, \mathbb{P}}(W))}{n\varepsilon^2}}_{\text{spectrum of inner layers } 1:l-1} \right. \\ & \left. + (d_l + d_{l-1}) r_{\text{eff}}^{S, \mathbb{P}}[W, l] \cdot \underbrace{\log \left(\bar{M}_{l \rightarrow L}^2(W, \varepsilon) L \max\{\|W\|_{\mathbf{F}}^2, R^2/4^n\} \right)}_{\text{spectrum of outer layers } l+1:L} + 2 \log(d_{l-1}n) \right), \end{aligned} \quad (4.1)$$

with

$$r_{\text{eff}}^{S, \mathbb{P}}[W, l] := r_{\text{eff}} \left(L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) \Gamma_{l-1}^{S, \mathbb{P}}(W), C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon \right),$$

and $C_2 > 0$ is an absolute constant.

Tightness of Each Step and Resulting Theorem. We conclude by reviewing our comprehensive theory for generalization in fully connected networks and justifying the tightness of the resulting bound in Theorem 5.

1. **First**, in Section 2 we develop a framework based on *pointwise dimension*. The upper bounds are complemented by a general empirical-metric lower bound and a worst-case lower bound in the same mixed metric (Theorem 3), and the framework has a profound connection to finite-scale geometry—evidence that this is the right organizing principle.
2. **Second**, Section 3 introduces a *non-perturbative* expansion. Lemma 1 applies Cauchy-Schwarz layerwise (treating each layer as a block). While there may be room to improve depth dependence, the telescoping decomposition (3.1) is an exact *equality*, so the expansion is generally sharp (and fully avoid linearization).
3. **Third**, the hierarchical covering argument shows that the resulting *Riemannian Dimension* bound matches the gold standard of *effective dimension*. Thus our pointwise, spectrum-aware bounds achieve the optimal form dictated by static ellipsoid theory, now in strongly correlated deep networks.

4.2 Fully Observable Generalization Bound for DNN

Building on Theorem 5, we state a fully empirical bound under a finite-resolution *subspace* condition. We require population-to-empirical comparison only on the observed active eigenspaces and control the population energy on their inactive complement at the same finite resolution. The condition is stated separately on the observed active subspace and its inactive complement at each scale.

For the observed-sample theorem, write

$$\Gamma_l^S(W) := F_l^S(W) F_l^S(W)^\top, \quad l = 0, \dots, L-1, \quad (4.2)$$

and let $d_R^S(W, \varepsilon)$ denote the observed-sample Riemannian Dimension obtained from Theorem 4 after replacing each $F_{l-1}(W, X) F_{l-1}(W, X)^\top$ by $\Gamma_{l-1}^S(W)$. Equivalently, it has the same explicit $FF^\top$ eigenvalue form as the mixed quantity in Theorem 5:

$$\begin{aligned} d_R^S(W, \varepsilon) := & \frac{1}{2} \sum_{l=1}^L \left((d_l + d_{l-1}) \underbrace{\sum_{k=1}^{r_{\text{eff}}^S[W, l]} \log \frac{8C_2^2 \lambda_k(\Gamma_{l-1}^S(W))}{n\varepsilon^2}}_{\text{observed spectrum of inner layers } 1:l-1} \right. \\ & \left. + (d_l + d_{l-1}) r_{\text{eff}}^S[W, l] \cdot \underbrace{\log \left(\bar{M}_{l \rightarrow L}^2(W, \varepsilon) L \max\{\|W\|_{\mathbf{F}}^2, R^2/4^n\} \right)}_{\text{spectrum of outer layers } l+1:L} + 2 \log(d_{l-1}n) \right), \end{aligned} \quad (4.3)$$

where

$$r_{\text{eff}}^S[W, l] := r_{\text{eff}}(L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) \Gamma_{l-1}^S(W), C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon),$$

and $C_2 > 0$ is an absolute constant.

For later use, define the finite spectral resolution associated with feature layer j by

$$R_W := C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \quad \vartheta_j(W, \varepsilon) := \frac{n\varepsilon^2}{2L \bar{M}_{j+1 \rightarrow L}^2(W, \varepsilon) R_W^2}, \quad j = 0, \dots, L-1. \quad (4.4)$$

The resolution $\vartheta_j(W, \varepsilon)$ is the finite spectral resolution that determines which observed feature directions contribute to the effective rank at the next layer. Indeed, by the definition of $r_{\text{eff}}^S[W, j+1]$,

$$\lambda_k(L \bar{M}_{j+1 \rightarrow L}^2(W, \varepsilon) \Gamma_j^S(W)) R_W^2 \geq \frac{n\varepsilon^2}{2} \iff \lambda_k(\Gamma_j^S(W)) \geq \vartheta_j(W, \varepsilon).$$

Thus the eigendirections of $\Gamma_j^S(W)$ with eigenvalues at least $\vartheta_j(W, \varepsilon)$ are precisely the observed feature directions counted by $r_{\text{eff}}^S[W, j+1]$. Let $P_j^S(W, \varepsilon)$ be the spectral projector of $\Gamma_j^S(W)$ onto the span of those eigendirections; we call its image the observed active subspace. Set

$$Q_j^S(W, \varepsilon) := I_{d_j} - P_j^S(W, \varepsilon),$$

whose image is the corresponding inactive complement.

Definition 3 (Layer-wise finite-resolution subspace isomorphism) Fix constants $\kappa \geq 1$ and $b_{\text{sub}} \geq 1$. We say that the observed sample S satisfies the layer-wise finite-resolution population subspace isomorphism if the following two estimates hold simultaneously for every feature layer $j = 0, \dots, L-1$, every center $W \in B_{\mathbf{F}}(R)$, and every scale $\varepsilon \in (0, 1/\beta]$:

$$P_j^S(W, \varepsilon)(n\Sigma_j^{\mathbb{P}}(W))P_j^S(W, \varepsilon) \preceq \kappa P_j^S(W, \varepsilon)\Gamma_j^S(W)P_j^S(W, \varepsilon), \quad (4.5)$$

$$\|Q_j^S(W, \varepsilon)(n\Sigma_j^{\mathbb{P}}(W))Q_j^S(W, \varepsilon)\|_{\text{op}} \leq b_{\text{sub}} \vartheta_j(W, \varepsilon). \quad (4.6)$$

The first display is a multiplicative isomorphism only on the observed active subspace. The second display says that the population features may have energy outside the observed active subspace, but only below the finite resolution already discarded by the effective-rank truncation. This subspace isomorphism implies finite-scale domination of the mixed empirical-population metric by the empirical metric in the pointwise dimension, yielding the following theorem based only on the observed sample. Moreover, the squared ellipsoidal projection metric in (4.6) is precisely the metric used to define the finite-scale resolution in (3.5) of Lemma 2 and Lemma 3, at the same order in ε . It provides a sufficient finite-resolution condition for the empirical bound. We also note that (κ, b_{sub}) enter the Riemannian Dimension below only through optimistic logarithmic factors, so their effect can be mild under fast spectral decay.

Theorem 6 (Fully empirical DNN bound under subspace isomorphism) Assume the setting of Theorem 5, with valid $\bar{M}_{l \rightarrow L}$ bounds available from S and the fixed network constraints. Suppose that, with probability at least $1 - \delta_{\text{iso}}$ over S , the layer-wise finite-resolution population subspace isomorphism in Definition 3 holds. Then there is a constant $c_{\text{sub}} > 0$, depending only on (κ, b_{sub}) , such that for every $\delta \in (0, 1)$, with probability at least $1 - \delta - \delta_{\text{iso}}$ over S , uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C \left(\beta \inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\mathbf{R}}^S(W, c_{\text{sub}}\varepsilon)} d\varepsilon \right\} + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right), \end{aligned} \quad (4.7)$$

where $C > 0$ is an absolute constant. In particular, if $\kappa, b_{\text{sub}} = O(1)$, then c_{sub} is an absolute constant.

Open problem: subspace-level feature regularity. The preceding certificate leads to the following concrete open question.

Can Definition 3 be further reduced from a uniform-in-subspace condition to subspace-level tail regularity of the learned features, by leveraging the approximate pointwise linearity and subspace-decomposition structure of deep neural networks?

This question isolates the remaining gap between an unconditional mixed empirical-population compressed theorem and a fully observed-sample compressed theorem. It is not a limitation of the pointwise-compression principle; rather, it is the precise observability problem one must solve to certify, from the training sample alone, the feature subspaces that the unconditional theory already identifies. We defer further discussion of this open problem, including a concrete route based on the two-layer sub-Gaussian/small-ball special case, to Appendix E.5.

Interpreting Theorem 6 to the Informal Rate (1.2). In the expression for the observed Riemannian Dimension in (4.3), the quantity $r_{\text{eff}}^S[W, l]$ incorporates local Lipschitz factors. Specifically, the effective rank is computed for

$$L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) F_{l-1}(W, X_S) F_{l-1}(W, X_S)^\top$$

rather than $F_{l-1} F_{l-1}^\top$ alone. When $F_{l-1} F_{l-1}^\top$ has rapidly decaying eigenvalues this dependence is strongly suppressed, and it disappears entirely under strict low rank. Consequently, under mild low-rank or spectral-decay conditions, the bound aligns with the informal rate (1.2). For each layer l , the first term in (4.3) quantifies the contribution of the inner layers $1:(l-1)$ through the observed feature Gram, while the second term captures the influence of the outer layers $(l+1):L$ through $\bar{M}_{l \rightarrow L}$. Together, these terms provide a complete layerwise account of the effective dimension in the informal rate (1.2).

4.3 Comparison with Norm Bounds, VC, and NTK

We compare our generalization bound for fully connected DNN (Theorem 5) with three established lines of work: (i) bounds based on products of spectral norms, (ii) VC-dimension-type capacity bounds, and (iii) Neural Tangent Kernel (NTK) linearizations that are valid only in an infinitesimal neighborhood of initialization. Our framework yields *exponentially* tighter rates than norm-product bounds, refines VC-type statements into hypothesis- and data-dependent guarantees, and replaces infinitesimal linearization with a finite-scale, non-perturbative analysis that holds simultaneously for every trained hypothesis. For space, we defer the recovery of representative norm bounds to Appendix E.7.1 and a broader literature review to Appendix A.1.

Norm Bounds: The trace relaxation of Theorem 5 does not use Definition 3. Write $R_n(W) := \max\{\|W\|_{\mathbf{F}}, R/2^n\}$. The inequality $\log x \leq \log(1+x) \leq x$ gives, for each layer,

$$\sum_{k=1}^{r_{\text{eff}}^{S, \mathbb{P}}[W, l]} \log \left(\frac{C \lambda_k(\Gamma_{l-1}^{S, \mathbb{P}}(W)) L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) R_n(W)^2}{n \varepsilon^2} \right) \leq \frac{C \text{Tr}(\Gamma_{l-1}^{S, \mathbb{P}}(W)) L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) R_n(W)^2}{n \varepsilon^2}.$$

Here

$$\text{Tr}(\Gamma_{l-1}^{S, \mathbb{P}}(W)) = \|F_{l-1}^S(W)\|_{\mathbf{F}}^2 + n \mathbb{E} \|f_{l-1}(W, X)\|_2^2.$$

When the activations are 1-Lipschitz and satisfy $\sigma_l(0) = 0$, the inner feature norms are bounded by the corresponding products of spectral norms. If the common outer transfer bounds also satisfy $\bar{M}_{l \rightarrow L}(W, \varepsilon) \leq C \prod_{i>l} \|W_i\|_{\text{op}}$ on the relevant weight charts, uniformly over the scales used, the rank-free argument in Appendix E.7.1 yields

$$(\mathbb{P} - \mathbb{P}_n) \ell(f(W, x), y) \leq \tilde{O} \left(\frac{\beta \|W\|_{\mathbf{F}}}{n} \sqrt{\|X_S\|_{\mathbf{F}}^2 + n \mathbb{E} \|X\|_2^2} \sqrt{L \sum_{l=1}^L (d_l + d_{l-1}) \prod_{i \neq l} \|W_i\|_{\text{op}}^2} \right). \quad (4.8)$$

As in Corollary 2, the display suppresses the confidence term and the stated logarithmic and discretization terms. If $\|X\|_2 \leq B_x$ almost surely, the input factor is at most $\sqrt{2n} B_x$, giving the usual $n^{-1/2}$ scaling. Retaining the feature spectra and logarithmic ellipsoidal cost can be exponentially sharper than this representative spectral-norm relaxation. Appendix E.7.1 derives this relaxation directly from Theorem 5 and gives an explicit exponential separation.

VC Dimension: Let L be the number of layers and $P = \sum_{l=1}^L d_l d_{l-1}$ be the total number of weights, Bartlett et al. [2019] prove a nearly tight VC-dimension bound $\text{VCdim} \leq O(PL \log P)$, supported by a lower bound $\text{VCdim} \geq \Omega(PL \log(P/L))$. This VC dimension bound is roughly equivalent to be $L \sum_{l=1}^L d_l d_{l-1}$.³

Our Riemannian Dimension bound, by contrast, substantially sharpens this rate: it removes the explicit dependence on depth L and replaces the crude width factor with a (layerwise) effective-rank term.

Neural Tangent Kernel (NTK): Our approach uses an exact, non-perturbative expansion that preserves the finite-scale geometry of deep networks, thereby going beyond NTK-based Taylor linearizations. A major bottleneck of the NTK approach is that the Taylor linearization remains valid only in an infinitesimal neighborhood of initialization, or equivalently in the infinite-width “lazy” regime [Jacot et al., 2018, Arora et al., 2019]. Outside this regime, the NTK approximation typically breaks down, limiting its explanatory power for practical networks.

From a generalization perspective, this initialization-centric and infinitesimal view suppresses the feature learning that drives generalization in modern deep networks, and therefore cannot fully explain their empirical behavior. In contrast, our results provide a finite-scale, pointwise theory that operates directly in practical regimes and explicitly captures feature learning through the spectra of the *learned* feature matrices.

4.4 Algorithmic Implications and Excess Risk Bound

Pointwise Dimension as Regularization and Excess Risk Bound. Our bounds imply a natural regularization strategy for algorithm design. Given the pointwise generalization inequality (2.1) (e.g., the Riemannian Dimension bound in Theorem 5), we consider a regularized ERM objective that explicitly minimizes this complexity measure:

$$\hat{f} = \arg \min_{f \in \mathcal{F}} \left\{ \mathbb{P}_n \ell(f; z) + C \sqrt{\frac{d(f) + \log(2/\delta)}{n}} \right\}. \quad (4.9)$$

With probability at least $1 - \delta$, its excess risk is bounded by (compared to any benchmark $f^* \in \mathcal{F}$):

$$\begin{aligned} & \mathbb{P} \ell(\hat{f}; z) - \mathbb{P} \ell(f^*; z) \\ & \leq \inf_{f \in \mathcal{F}} \left\{ \mathbb{P}_n \ell(f; z) + C \sqrt{\frac{d(f) + \log(2/\delta)}{n}} \right\} - \mathbb{P} \ell(f^*; z) \\ & \leq (C + \sqrt{1/2}) \sqrt{\frac{d(f^*) + \log(2/\delta)}{n}}; \end{aligned} \quad (4.10)$$

see Appendix E.6 for full proof. Thus we obtain a problem-dependent oracle bound of order $\sqrt{d(f^*)/n}$ that adapts to the optimal hypothesis f^* .

³The extra factor L beyond parameter count in VCdim is essentially unavoidable: for nonlinear compositional models, VC/packing dimensions depend on the logarithm of a global worst-case Lipschitz constant, and in depth- L networks that constant grows multiplicatively across layers, yielding an additional linear dependence on L .

From Explicit Regularization to Implicit Bias of Practical Algorithms. Since modern optimizers like SGD routinely drive empirical risk to near-zero, convergence analysis alone offers limited insight into generalization. The central theoretical challenge is therefore not determining *whether* a minimum is reached, but identifying *which* of the infinite interpolating solutions the optimizer selects. Plain ERM is insufficient for this task: without constraints on pointwise dimension, an empirical risk minimizer yields no guarantee of controlled excess risk. In contrast, our RD-regularized objective (4.9) explicitly enforces the low-complexity structure required for the generalization bound in (4.10). Although this intuition is rooted in the earliest practices of deep learning, our pointwise theory rigorously articulates the underlying mathematical reasoning.

This motivates a concrete agenda for optimization in deep learning: characterize algorithms whose implicit bias drives iterates toward solutions with *low pointwise complexity*, in particular low *Riemannian Dimension* (RD). Analogous phenomena are well documented in linear and kernel settings: gradient descent converges to max-margin (logistic loss) or minimum-norm (least squares) solutions [Soudry et al., 2018, Gunasekar et al., 2018], iterate-averaged SGD behaves like ridge regression [Neu and Rosasco, 2018], and “ridgeless” kernel regression can generalize with an optimally zero ridge parameter [Liang and Rakhlin, 2020]; see Vardi [2023] for a survey. Our regularizer in (4.9), based on pointwise dimension and, in particular, the RD from Theorem 5, is strictly more informative than any single norm, making it a natural target for such analyses.

Empirically, we say an algorithm exhibits *Riemannian-Dimension implicit bias* if it preferentially returns solutions with small RD despite RD’s large dynamic range; in Section 5.3 we observe that SGD indeed finds low-RD solutions.

5 Experiments

We evaluate the one-shot L_2 relaxation of the Riemannian Dimension bound in Theorem 6 on two standard architectures—Fully Connected Networks (FCNs) and ResNets, using two benchmark datasets—MNIST [LeCun et al., 1998] and CIFAR-10 [Krizhevsky, 2009], respectively. We consider a 9-hidden-layer FCN architecture, where, except for the fixed layers, hidden layers share a common width h , with $h \in \{2^6, 2^7, 2^8, 2^9, 2^{10}, 2^{11}, 2^{12}\}$. Increasing h monotonically enlarges both layer widths and model sizes. We adopt canonical ResNet architectures—ResNet- $\{20, 32, 44, 56, 74, 110\}$ —which differ only in the number of residual blocks per stage while maintaining the same overall architecture (three-stage, basic-block design) as introduced by [He et al., 2016]. These ResNet architectures provides a clean capacity sweep via depth. In what follows, we organize experiments around the two complementary regimes—width scaling on FCNs and depth scaling on ResNets.

This design lets us systematically study three central questions in modern deep learning: (i) why does overparameterization often improve generalization? (ii) how does feature learning evolve during training? and (iii) what implicit regularization is encoded by the baseline optimizer? Detailed experimental setups are deferred to Appendix A.2. The code to reproduce all experiments is available at our GitHub repository [here](#).

In particular, we use the spectral-norm product $\prod_{i>l} \|W_i\|_{\text{op}}$ as a numerical surrogate for the local Lipschitz constant $\bar{M}_{l \rightarrow L}(W, \varepsilon)$; state-of-the-art formal-verification toolchains [Shi et al., 2022] can compute local Lipschitz constants much more sharply—with well-developed packages and rigorous numerical guarantees—than this spectral-product surrogate, and could therefore further strengthen all our empirical results (an active research area). This substitution gives a certified upper bound when the spectral product dominates the transfer constants throughout the rele-

Table 1: Final-epoch Metrics of FCNs on MNIST. Column descriptions are as follows: 1) Width-2^{*} denotes a hidden width of $h = 2^*$; 2) Train: training error; 3) Gen: generalization gap, defined as test error minus training error; 4) Spectral Norm: the spectrally normalized margin bound of [Bartlett et al., 2017], a tightness norm-based bound in the literature to our knowledge; we computed it with margin normalization; 5) # Parameters: parameter counts of the network; 6) VC dimension: we adopt a nearly tight VC-dimension bound from [Bartlett et al., 2019] and report $PL \log P$ for brevity (see Section 4.3); 7) R-D: the proposed Riemannian Dimension.

Model	Train	Gen	Spectral Norm	# Parameters	VC dimension	R-D
Width-2 ⁶	0.0002	0.0205	3.146×10^{15}	5.961×10^6	9.299×10^8	6.433×10^7
Width-2 ⁷	0.0002	0.0187	2.695×10^{15}	6.167×10^6	9.641×10^8	6.097×10^7
Width-2 ⁸	0.0000	0.0191	2.093×10^{15}	6.726×10^6	1.057×10^9	5.589×10^7
Width-2 ⁹	0.0000	0.0186	2.401×10^{15}	8.434×10^6	1.345×10^9	5.316×10^7
Width-2 ¹⁰	0.0000	0.0215	4.816×10^{15}	1.421×10^7	2.340×10^9	5.266×10^7
Width-2 ¹¹	0.0000	0.0160	1.001×10^{16}	3.520×10^7	6.116×10^9	4.972×10^7
Width-2 ¹²	0.0000	0.0210	1.466×10^{16}	1.149×10^8	2.133×10^{10}	4.803×10^7

Table 2: Final-Epoch Metrics of ResNets on CIFAR-10

Model	Train Error	Gen Gap	# Parameters	VC dimension	R-D
ResNet-20	0.0016	0.0752	2.690×10^5	6.727×10^7	8.801×10^6
ResNet-32	0.0003	0.0695	4.630×10^5	1.933×10^8	9.992×10^6
ResNet-44	0.0001	0.0627	6.570×10^5	3.872×10^8	6.339×10^6
ResNet-56	0.0000	0.0637	8.510×10^5	6.507×10^8	5.200×10^6
ResNet-74	0.0000	0.0615	1.142×10^6	1.179×10^9	3.237×10^6
ResNet-110	0.0000	0.0576	1.724×10^6	2.723×10^9	2.583×10^6

vant weight charts, as required in Corollary 2; the pointwise product alone does not establish this condition.

5.1 Riemannian Dimension Explains Overparameterization

This section studies why does overparameterization—despite exploding model capacity—often improve generalization. We investigate this paradox by tracking our Riemannian Dimension across models with varying parameter counts, asking whether more parameters truly enlarge capacity or instead deduce complexity.

Final-epoch metrics of FCNs on MNIST and ResNets on CIFAR-10 are reported in Table 1 and Table 2, respectively. In these Tables, the train error quickly collapses to zero for sufficiently large models, confirming their expressive capacity. Consistently, the generalization can continue to be improved as parameters increase, especially on ResNets (Table 2). This phenomenon means the overfitting does not appear and reflects a paradoxical truth of deep learning: over-parameterization is not a curse, but can benefit the generalization. However, classical complexity measures—e.g., the spectral norm and the VC dimension, often scale exponentially as the parameter count grows. Notably, the spectral norm is about 10^6 times larger than the VC dimension and seems to perform

Table 3: Final-epoch Effective Ranks for FCNs on MNIST, where Width-2^{*} means $h = 2^*$, and where for the form A/B, A represents the effective rank and B represents the original dimension, and where Layer-1 means the input layer.

Metric	Width-2 ⁶	Width-2 ⁷	Width-2 ⁸	Width-2 ⁹	Width-2 ¹⁰	Width-2 ¹¹	Width-2 ¹²
Layer-1	713/763	712/763	710/763	710/763	707/763	707/763	704/763
Layer-2	2048/2048	2044/2048	2042/2048	2048/2048	2047/2048	2048/2048	2048/2048
Layer-3	2048/2048	2045/2048	2037/2048	2019/2048	1925/2048	1460/2048	1009/2048
Layer-4	61/64	97/128	92/256	85/512	79/1024	79/2048	59/4096
Layer-5	23/64	43/128	34/256	33/512	28/1024	26/2048	22/4096
Layer-6	20/64	24/128	20/256	21/512	19/1024	18/2048	15/4096
Layer-7	15/64	18/128	17/256	15/512	15/1024	14/2048	13/4096
Layer-8	15/64	14/128	15/256	11/512	13/1024	13/2048	12/4096
Layer-9	14/64	14/128	15/256	13/512	13/1024	12/2048	12/4096
Layer-10	13/64	13/128	12/256	14/512	12/1024	13/2048	14/4096
Total	4970	5024	4994	4969	4858	4390	3908

Table 4: Final-epoch Effective Ranks for ResNets on CIFAR-10, where for the form A/B, A represents the effective rank and B represents the original dimension, and where Layer-0% means the input layer.

Metric	ResNet-20	ResNet-32	ResNet-44	ResNet-56	ResNet-74	ResNet-110
Layer-0%	384/3072	384/3072	17/3072	0/3072	0/3072	0/3072
Layer-25%	2048/16384	2048/16384	7/16384	1/16384	0/16384	0/16384
Layer-50%	1024/8192	1024/8192	1024/8192	227/8192	0/8192	0/8192
Layer-75%	512/4096	512/4096	512/4096	512/4096	58/4096	0/4096
Layer-100%	8/64	8/64	8/64	8/64	8/64	8/64
Total	23432	37768	27564	16294	11401	6925

worse (see Table 1). The two measures therefore struggle to explain the generalization of modern overparameterized networks. In contrast, Riemannian Dimension exhibits a consistent downward trend as model size grows—both under width scaling (last column of Table 1) and depth scaling (last column of Table 2), and it is about 10^3 times smaller than the VC dimension, suggesting that the effective dimension—not raw parameter count—is the more informative indicator of generalization in deep learning. In summary, increased parameterization is associated with reduced intrinsic model complexity, and Riemannian Dimension captures this phenomenon.

5.2 Feature Learning Compresses Effective Rank

We investigate the dynamics of feature learning by monitoring the effective rank of the feature Gram matrices $F_{l-1}F_{l-1}^\top$ scaled by $L\|W\|_{\mathbf{F}}^2 \prod_{i>l} \|W_i\|_{\text{op}}^2$ (i.e., $F_{l-1}F_{l-1}^\top \cdot L\|W\|_{\mathbf{F}}^2 \prod_{i>l} \|W_i\|_{\text{op}}^2$), as dictated by the theory. We report our empirical results in Tables 3, 4 and Figure 2.

Experimental results reveal some clear patterns. First, as training proceeds, the effective ranks of feature Gram matrices decrease sharply after a short transient phase; refer to Figure 2. Second,




Figure 2: Effective Rank evolutions of FCNs on MNIST (left) and ResNets on CIFAR-10 (right) across the training




Figure 3: Riemannian Dimension evolutions of FCNs on MNIST (left) and ResNets on CIFAR-10 (right) across the training

increasing the parameter count, either by width scaling in FCNs or by depth scaling in ResNets, accelerates and strengthens this effective-rank compression; refer to Figure 2. Third, on the largest FCN, the degree of effective rank compression can reach as much as $1/300$, which explains the significant improvement of the Riemannian Dimension relative to the VC dimension; refer to Table 3. On the largest ResNet, the effective ranks of the vast majority of layers compress to zero, thereby explaining why deeper networks can exhibit a smaller Riemannian Dimension; refer to Table 4. Together, these experimental results suggest that feature learning progressively reduces the intrinsic dimensionality of learned representations, and that overparameterization intensifies this dimension-reduction effect.

5.3 SGD Finds Low Riemannian Dimension Point

Prior work has shown that various norms are implicit bias of optimizers, but typically limited to linear models [Vardi, 2023]. This section studies whether SGD with momentum, in modern deep learning, implicitly regularized the Riemannian Dimension across training dynamics. We examine whether this optimizer preferentially converges to solutions with lower Riemannian Dimension, and the experimental results are presented in Figure 3.

Empirical results show a repeatable pattern across the architectures: SGD with momentum

drives the networks toward solutions with lower intrinsic Riemannian Dimension complexity, after an early transient; refer to Figure 3. We observe that the Riemannian Dimension drops by orders of magnitude, whereas the VC dimension remains essentially unchanged. The alignment between optimization dynamics and complexity control supports the view that SGD with momentum implicitly regularizes the Riemannian Dimension. Therefore, optimization is not merely as a mechanism for convergence; it is a primary driver of generalization through its systematic preference for low-complexity solutions. Riemannian Dimension provides a practical and theoretically grounded lens through which the implicit bias of optimizers in machine learning can be quantitatively assessed.

6 Conclusion

We have developed a pointwise, representation-aware foundation for DNN generalization in the nonlinear feature learning regime. Meeting this challenge required several technical innovations: a pointwise generalization framework, a non-perturbative calculus for network mappings, a hierarchical covering scheme, and an ellipsoidal entropy theory for the Grassmannian—yielding structural insights into cross-weight correlations and global geometric organization. The results strengthen the case that deep-learning generalization admits rigorous explanation and motivate a broader program in finite-scale geometric analysis of strongly correlated learning systems. Our experiments support the theory: the proposed Riemannian Dimension consistently tracks benign overparameterization, feature learning, and the optimizer’s implicit bias.

Conceptually, the paper also separates two levels of observability. The unconditional theorem proves compressed pointwise generalization through the mixed empirical-population geometry of the learned features; the fully observable theorem asks for a finite-resolution subspace isomorphism that lets the same geometry be read from the training sample alone. This distinction is analogous in spirit to observability and uncertainty phenomena in physics, and it gives rise to a concrete open problem on subspace-level feature regularity. The main message remains unchanged: **pointwise feature compression is intelligence**, with feature compression made rigorous as layerwise learned-feature spectral compression and pointwise made rigorous as hypothesis-dependent finite-scale complexity.

Important directions ahead include integrating the analysis with optimization and algorithmic perspectives, extending it to modern architectures, translating the theory into concrete design principles for new deep-learning systems, and resolving the subspace-level isomorphism question posed in Section 4.

A Related Works and Experimental Setup

A.1 Related Works

Given the breadth of work on generalization and its empirical proxies, the mathematical grounding of our approach, and its conceptual relevance to vision and language practice, we streamline the exposition by concentrating on the most relevant prior results.

Theoretical Generalization Bounds for DNN. A substantial line of work anchors generalization bounds to various norms of the network weights, including path norms [Neyshabur et al., 2015a], Frobenius norms [Neyshabur et al., 2015b, Golowich et al., 2020], and spectral norms [Bartlett et al., 2017, Neyshabur et al., 2018, Arora et al., 2018]. We refer to Neyshabur et al. [2017] as an important early contribution to this broader program. While offering conceptual insights, these bounds, often derived from globally uniform complexity measures like covering numbers or Rademacher complexity, frequently suffer from exponential dependencies on depth or layer norms, rendering them vacuous for practical, deep architectures. Compelling empirical evidence [Farhang et al., 2022, Razin and Cohen, 2020] further suggests that norm-based bounds alone are insufficient to fully elucidate the generalization phenomenon in deep learning. The kernel perspective [Belkin et al., 2018], epitomized by NTK theory [Jacot et al., 2018, Arora et al., 2019, Golikov et al., 2022], yields sharp guarantees by linearizing a network around its initialization—effectively casting training as kernel ridge regression with a fixed kernel. Within this linear/lazy regime, precise calculations explain both double descent [Belkin et al., 2019] and benign overfitting [Bartlett et al., 2020], and an eigenspace-projection viewpoint provides dimension-reduction and feature-compression insights [Bartlett et al., 2021]. Investigations beyond the lazy regime exist, but most analyses either study the two-layer infinite-width (mean-field) limit (e.g., [Mei et al., 2018, Chizat and Bach, 2018]) or remain in a neighborhood of initialization [Woodworth et al., 2020]. While insightful, these settings are idealized and struggle to capture the behavior of finite, deep networks (see Chapter 6 of [Misiakiewicz and Montanari, 2023]). More broadly, linear, lazy, or infinite-width approximations fail to reflect the feature learning that arises when parameters move far from initialization and representations evolve. This omission is widely viewed as a central bottleneck in current theory; indeed, the rich, representation-learning regime is often argued to be the key phenomenon distinguishing modern deep learning from long-standing frameworks (see, e.g., Bartlett et al. [2021], Misiakiewicz and Montanari [2023], Radhakrishnan et al. [2024], Wilson [2025]). Building on these directions, we establish—to our knowledge—the first pointwise generalization bounds for nonlinear DNN that are comparable in sharpness to prior linearization results and, crucially, remain valid in the practical feature-learning regime.

Other Theoretical Perspectives of Generalization. A growing line of work connects generalization to geometric notions of fractal dimension [Birdal et al., 2021, Dupuis et al., 2023, Simsekli et al., 2020, Andreeva et al., 2024, Camuto et al., 2021], typically through Hausdorff- or Minkowski-type dimensions of optimization trajectories or invariant measures. However, these fractal dimensions are globally uniform, infinitesimal-scale ($\varepsilon \rightarrow 0$) notions of complexity. In contrast, our theory is built on a pointwise, finite-scale notion of geometric dimension. Section 2.2 is precisely devoted to this distinction: we move from globally uniform to pointwise dimension and show that generalization is governed by the finite-scale pointwise dimension rather than its asymptotic limit. Several PAC-Bayesian approaches operate directly in parameter space W , en-

downing the weights with an explicit stochastic model and directly computing the KL divergence between a hand-designed prior and a posterior over W [Hinton and Van Camp, 1993, Dziugaite and Roy, 2017, Lotfi et al., 2022, 2024]; e.g., Gaussian distribution in Dziugaite and Roy [2017]. These parameter-space bounds are valuable for certifying that certain trained weight configurations admit nonvacuous PAC-Bayes guarantees, but they largely treat the network as a black box and do not directly capture how architecture and feature geometry control generalization. PAC-Bayes theory has also been used to formalize flatness/sharpness [Neyshabur et al., 2017] and spectral-norm bounds [Neyshabur et al., 2018], highlighting the central role of this methodology in deep learning generalization and providing strong motivation for pointwise generalization; see the mathematical background below. Alternative theoretical frameworks include algorithmic stability analyses, which are used primarily for one-hidden-layer networks and connected to the NTK/lazy-training viewpoint [Richards and Kuzborskij, 2021, Lei et al., 2022]; and VC-dimension methods [Bartlett et al., 2019], which has been discussed in Section 4.3.

Pointwise and Non-Perturbative Foundations. Our use of “pointwise” draws inspiration from several threads that emphasize hypothesis-specific complexity: the asymptotic pointwise dimension in fractal geometry [Falconer, 1997], PAC-Bayes analyses that tailor complexity to the chosen random posterior [McAllester, 1998, Dziugaite and Roy, 2017, Neyshabur et al., 2018], and the Fernique–Talagrand integral in the majorizing-measure formulation of generic chaining [Fernique, 1974, Talagrand, 1987, 2014]. Notably, Maurer [2010], Chu and Raginsky [2023] establish simultaneous pointwise envelopes for stochastic processes using variants of the Fernique–Talagrand integral, providing direct precedents for our pointwise-dimension viewpoint. The synthesis of PAC-Bayes bounds with generic chaining dates back to Audibert and Bousquet [2003], Audibert and Bousquet [2007], and mutual information based bounds have also been combined with chaining [Russo and Zou, 2016, Xu and Raginsky, 2017, Asadi et al., 2018, Liu, 2025]. To the best of our knowledge, this paper is the first work to establish a sharp pointwise bound for deterministic hypotheses in an uncountable class via localization to metric balls, explicitly connecting the result to the formal definition of pointwise dimension. A generic conversion from classical (subset-homogeneous) uniform convergence to pointwise generalization bounds, established in Xu and Zeevi [2020, 2025], serves as a guiding principle and plays a central role in our proof of Theorem 2. The adjective “non-perturbative,” borrowed from physics [nLab contributors, 2025a] and central to the study of strongly correlated systems [nLab contributors, 2025b], underscores that our theory remains valid far beyond infinitesimal neighborhoods of initialization—an essential property for deeply nonlinear, feature-learning DNN.

Connections to Differential Geometry and Lie Algebra. From a geometric perspective, Hausdorff dimension provides an asymptotic, covering-based notion of capacity (fundamental in geometric measure theory [Simon, 2018]), while differential and Riemannian geometry [Jost, 2008] develop the use of local charts and global atlases to analyze non-Euclidean manifolds. Our results motivate viewing generalization as a finite-scale problem in geometric analysis. The Grassmannian and families of orthogonal subspaces are traditionally studied via Lie groups; using differential-geometric tools, Szarek [1997], Pajor [1998] established finite-scale isotropic metric-entropy characterizations, which motivate our hierarchical covering viewpoint from local charts to a global atlas and our ellipsoidal entropy framework.

Predictor-Space Flatness and Regularization. Flatness has long been used as a geometric intuition for generalization and regularization [Hochreiter and Schmidhuber, 1997, Keskar et al., 2017, Foret et al., 2021], but parameter-space formulations are sensitive to reparametrization and scaling [Dinh et al., 2017]. The L_1 relaxation of the one-shot bound in Theorem 1 gives a coordinate-free version of this idea by defining flatness directly in predictor space. At scale ε , the bias term measures loss-level stability over the metric ball $B_{\varrho_{(S, \mathbb{P}), \ell}^{(1)}}(f, \varepsilon)$: nearby predictors have controlled loss deviation from f . The variance term, namely the pointwise dimension $\log \frac{1}{\pi \left(B_{\varrho_{(S, \mathbb{P}), \ell}^{(1)}}(f, \varepsilon) \right)}$,

measures the statistical size of the same stable neighborhood through its prior mass. Thus a favorable predictor is one whose loss is locally stable on a neighborhood that also has sufficiently large prior mass, and the optimal scale ε^* balances local approximation against local mass. This also suggests a regularization principle: prefer predictors with large stable neighborhoods in predictor space, rather than merely small sharpness in a chosen parameterization. Since the construction uses loss-metric balls in the function class $\mathcal{F}$, equivalent parameterizations of the same predictor induce the same ball and the same pointwise dimension. In this sense, (2.6) formalizes flatness as predictor-space stability plus local prior mass, avoiding the reparametrization pathologies of raw parameter-space sharpness.

Empirical Indicators of Generalization. Complementing theory, much research has focused on empirical indicators that explain the generalization of deep learning. Phenomena like *Neural Collapse* [Papayan et al., 2020, Parker et al., 2023, Kothapalli, 2022] reveals the emergence of low-rank geometric structures in last-layer features. Studies on *Intrinsic Dimension* [Li et al., 2018, Huh et al., 2021] similarly suggest that deeper models exhibit an inductive bias toward low-rank last-layer feature representations. A line of work focuses on *Dynamic NTK variants* [Atanasov et al., 2021, Baratin et al., 2021, Fort et al., 2020, Kopitkov and Indelman, 2020] or related feature-gradient kernels [Radhakrishnan et al., 2024], where the kernels evolve along optimization trajectories, has empirically shown that the dynamic kernel evolution is linked to generalization behaviour. Other probes, examining Fisher information [Karakida et al., 2019, Jastrzebski et al., 2021], Hessian spectral properties [Ghorbani et al., 2019, Rahaman et al., 2019], and output-input Jacobians [Novak et al., 2018], offer another lens. Collectively, existing empirical probes offer valuable, though often partial, insights—typically from a specific layer perspective, or through a constructed similarity analysis—without a unifying formalism and a theory foundation. Our proposed empirical indicator, rooted in a mathematically sharp theory, resonates with their goals (our theory is in fact supported by many of their experiments) while advancing them. It provides a principled, formal measure for studying pointwise generalization in deep neural networks.

Feature Compression in Deep Models for Vision and Language. Across vision and language, deep networks exhibit a robust layer-wise compression of representations. In computer vision, Ansuini et al. [2019] measure intrinsic dimensionality across convolutional layers and find early expansion followed by sharp reduction, with lower late-stage dimensionality correlating with stronger generalization; Feng et al. [2022] likewise show that feature matrices in CNNs and vision transformers become progressively low-rank with depth, at fixed width, indicating active compression of task-relevant information. Parallel trends appear in NLP: Cai et al. [2021] demonstrate that contextual embeddings (e.g., BERT) occupy narrow, anisotropic cones despite high nominal dimension, and Razzhigaev et al. [2024] document a two-phase training trajectory—initial expansion,

then sustained compression. A complementary line grounded in the Information Bottleneck [Tishby and Zaslavsky, 2015] interprets these findings as the selective removal of task-irrelevant variability: Shwartz-Ziv and Tishby [2017] observe that networks spend most of training compressing internal features toward a prediction-compression trade-off, while Patel and Shwartz-Ziv [2024] show gradient descent reduces the local rank of intermediate activations. Balzano et al. [2025] provide a complementary tutorial on low-rank structures arising during the training and adaptation of large models, emphasizing how gradient-descent dynamics and implicit regularization generate low-rank representations. Taken together, these phenomena motivate our investigation: compression is not merely qualitative, but admits precise, hypothesis-specific complexity that governs generalization.

A.2 Experimental Setup

We introduce detailed experimental setups. We evaluate the one-shot L_2 relaxation of the Riemannian Dimension bound derived from Theorem 6 on two standard architectures—Fully Connected Networks (FCNs) and ResNets, using two benchmark datasets—MNIST [LeCun et al., 1998] and CIFAR-10 [Krizhevsky, 2009], respectively. The architecture of FCNs: we consider a 9-hidden-layer FCN in which the first two hidden layers have width 2^{11} and the remaining seven hidden layers share a common width h , with $h \in \{2^6, 2^7, 2^8, 2^9, 2^{10}, 2^{11}, 2^{12}\}$. The output layer is a linear classifier mapping to 10 logits, and we use ReLU as the activation and use PyTorch’s default initialization (Kaiming uniform for ReLU). Increasing h monotonically enlarges both layer widths and the total parameter count, yielding a clean capacity sweep at fixed depth. The architecture of ResNets: we adopt the canonical ResNet architectures, ResNet-20, ResNet-32, ResNet-44, ResNet-56, ResNet-74, and ResNet-110, which differ only in the number of residual blocks per stage while maintaining the same overall architecture (three-stage, basic-block design) as introduced by [He et al., 2016]. Following the practice of [He et al., 2016], we apply BatchNorm and ReLU after each convolution, with shortcut connections added as needed, and a global average pooling layer precedes the final linear classifier. These ResNet architectures provides a clean capacity sweep via depth.

We adopt standard training pipelines widely used in the benchmarks. (1) The training Protocol of FCNs is: SGD with momentum optimizer where momentum = 0.9, learning rate = 0.01, and weight decay = 5×10^{-4} ; 200 epochs and 128 batch size; a step decay at epochs $\{100, 170\}$, where the learning rate is scaled by $\times 0.1$. (2) The training Protocol of ResNets is: SGD with momentum optimizer where momentum = 0.9, learning rate = 0.1, and weight decay = 5×10^{-4} ; 250 epochs and 128 batch size; a step decay at epochs $\{50, 150, 200\}$, where the learning rate is scaled by $\times 0.1$; Following practical training conditions, we apply standard data augmentation on CIFAR-10: random horizontal flips and 4-pixel random crops with zero-padding.

In the experiments of FCNs and ResNets, to enable layerwise analysis of the evolving feature representations and support our computation of Riemannian Dimension, we register forward hooks on all nonlinearity layers. For layers followed by pooling, we replace the last recorded ReLU activation with the corresponding pooled output. We also pre-register the input hook to capture the feature matrix of the data. These hooks ensure precise extraction of nonlinearity activations at each depth throughout training. We set the hyper-parameter ε via a one-dimensional ternary-search procedure: at the end of each training stage, we perform a 500-step ternary search for FCNs and a 50-step ternary search for ResNets over the admissible interval specified by the following finite-resolution search range: $[\sqrt{1/n}, \max_{l=1,\dots,L} \sqrt{\frac{2L\lambda_{\max}(F_{l-1}F_{l-1}^\top) \cdot \|W\|_F^2 \prod_{i>l} \|W_i\|_{\text{op}}^2}{n}}]$. The search selects the value of ε that minimizes the one-shot L_2 relaxation of the Riemannian Dimension-based

generalization bound in Theorem 6; it does not directly evaluate the gap-distance balls in (2.6). We note that tighter bounds could be achieved with more refined optimization procedures on ε . For FCNs, we compute full feature gram matrices. While for ResNets, the feature matrix F is formed by flattening each activation map into a vector of dimension $d = C \cdot H \cdot W$, where C, H, W are the channel, height, and width of the feature map respectively. To align with our theory, we simplify ResNets to fully connected (feed-forward) networks when computing our bound; we apply the same simplification to the associated VC-dimension and parameter-count calculations to maintain consistency. To avoid out-of-memory in computing full feature gram matrices in high-dimensional convolutional layers, we use the standard Gaussian sketching approximation, where each feature gram matrix uses a Gaussian sketch with parameter $r = \min(8192, \lfloor d/8 \rfloor)$ [Woodruff et al., 2014]. By standard subspace-embedding guarantees, such Gaussian sketches preserve Gram quadratic forms—and hence the spectra—of the feature matrices with high probability, introducing only negligible distortion and leaving our conclusions unchanged [Woodruff et al., 2014].

B Proofs for Pointwise Generalization Framework (Section 2)

Much of this section is devoted to a full proof of Theorem 2 (the integral upper bound). Conceptually, the pointwise–dimension principle already follows from elementary PAC-Bayes arguments—see Theorem 7 and the subsequent remark in Appendix B.2. We present the full derivation to make explicit structural properties (e.g., unified blueprint, subset homogeneity, mixed empirical-ghost comparison) that a rigorous proof requires.

B.1 The “Uniform Pointwise Convergence” Principle

In this section, we present a unified blueprint for establishing pointwise generalization bounds. We state necessary and sufficient conditions for pointwise generalization and show that, when applied carefully, the resulting pointwise bounds are no harder to obtain than classical uniform-convergence guarantees.

We begin by citing a general principle for converting *subset-homogeneous* uniform convergence guarantees—i.e., bounds in which the same pointwise complexity applies for every fixed subset $\mathcal{H} \subseteq \mathcal{F}$ —into pointwise generalization bounds. This conversion, introduced by the name “uniform localized convergence” principle in [Xu and Zeevi, 2020] (short conference version) and Xu and Zeevi [2025] (full journal version), identifies minimal conditions under which pointwise envelopes can be obtained, up to negligible double-logarithmic factors, while avoiding unrealistic Bernstein-type and sub-root assumptions. We state this result as “uniform pointwise convergence” principle. This subsection provides the necessary and sufficient justifications, together with discussions of data dependence.

Lemma 4 (“Uniform Pointwise Convergence” Principle) (Proposition 1 in Xu and Zeevi [2020, 2025]). *For a function class $\mathcal{F}$ and functional $d : \mathcal{F} \rightarrow [0, R]$, assume there is a function $\psi(r; \delta)$, which is non-decreasing with respect to r , non-increasing with respect to δ , and satisfies that $\forall \delta \in (0, 1), \forall r \in [0, R]$, with probability at least $1 - \delta$,*

$$\sup_{f \in \mathcal{F} : d(f) \leq r} (\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \psi(r; \delta). \quad (\text{B.1})$$

Then, given any $\delta \in (0, 1)$ and $r_0 \in (0, R]$, with probability at least $1 - \delta$, uniformly over all $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \psi \left(\max\{2d(f), r_0\}; \delta \left(\log_2 \frac{2R}{r_0} \right)^{-1} \right). \quad (\text{B.2})$$

This lemma provides a succinct proof that serves as a unifying principle to sharpen classical localization, building on Section 2 of [Xu and Zeevi \[2025\]](#). A key advantage of this framework is its level of abstraction: it establishes subset homogeneity as the necessary and sufficient condition for pointwise generalization when the complexity functional $d(\cdot)$ is data-independent, and likewise when $d(\cdot)$ is swap-invariant and depends on both the observed sample $S = \{z_i\}_{i=1}^n$ and an i.i.d. ghost sample $S' = \{z'_i\}_{i=1}^n$. It also provides a clean treatment of data-dependent functionals and their induced (random) sublevel sets $\{f \in \mathcal{F} : d(f) \leq r\}$, as outlined before Section 4 of [Xu and Zeevi \[2025\]](#). Crucially, this approach circumvents the circularities that often arise when combining symmetrization with localization or offset arguments.

B.1.1 Necessary and Sufficient Conditions for Pointwise Generalization

We leverage this “uniform pointwise convergence” principle to streamline the derivation of our bounds. Let $d(\cdot)$ denote a pointwise complexity functional, which we categorize into data-independent forms and data-dependent forms. Let $\psi(\cdot; \delta)$ be a non-decreasing function (typically $\psi(r; \delta) \asymp \sqrt{(r + \log(1/\delta))/n}$). We provide a clean characterization of pointwise generalization.

Necessary Condition: Subset Homogeneity. A valid pointwise generalization guarantee (i.e., (2.1)) necessitates *subset homogeneity*. That is, if the pointwise inequality

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \psi(d(f); \delta) \quad (\text{B.3})$$

holds with probability at least $1 - \delta$, then (B.3) must imply that for every fixed (i.e., data-independent) subset $\mathcal{H} \subseteq \mathcal{F}$,

$$\sup_{f \in \mathcal{H}} (\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \sup_{f \in \mathcal{H}} \psi(d(f); \delta).$$

Crucially, the complexity evaluation $d(f)$ must *not* depend on the chosen subset $\mathcal{H}$. For instance, for the pointwise dimension $\log \frac{1}{\pi(B_\varrho(f, \varepsilon))}$, the prior π (in particular, its support) should be independent of $\mathcal{H}$. This contrasts with classical empirical-process techniques—e.g., naive uses of Rademacher complexity and generic chaining—where the *pre-specified* index sets dictate the proxy $d(\cdot)$ via the chosen Rademacher expectation, admissible tree construction, or prior.

Subset homogeneity is thus the primary eligibility check for any candidate pointwise complexity functional. In Appendix B.3, we complete this check by establishing that the pointwise dimension is *ambiently equivalent*: using a prior $\pi \in \Delta(\mathcal{F})$ or its restriction $\pi \in \Delta(\mathcal{H})$ produces complexities that agree in order (up to absolute constants).

Sufficient Condition: Subset Homogeneity + Data-Independent (or Symmetrized) $d(\cdot)$. Assuming the following subset-homogeneity uniform convergence condition: for every fixed (i.e., data-independent) subset $\mathcal{H} \subseteq \mathcal{F}$ and $\delta \in (0, 1)$, with probability at least $1 - \delta$,

$$\sup_{f \in \mathcal{H}} (\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \sup_{f \in \mathcal{H}} \psi(d(f); \delta). \quad (\text{B.4})$$

By taking the sublevel set

$$\mathcal{H} = \{f \in \mathcal{F} : d(f) \leq r\},$$

the condition (B.4) (applied to this fixed sublevel set) directly implies the surrogate conditions (B.1) in Lemma 4, and hence the pointwise bound (B.2). Thus, subset homogeneity is a necessary and sufficient condition for a data-independent $d(\cdot)$ to imply a pointwise generalization bound.

Likewise, in Appendix B.4, we show that when the complexity $d(\cdot)$ may depend on both the observed sample $S = \{z_i\}_{i=1}^n$ and an i.i.d. ghost sample $S' = \{z'_i\}_{i=1}^n$, *provided it is swap-invariant in (S, S') (i.e., invariant under any exchange $z_i \leftrightarrow z'_i$)*, subset homogeneity suffices to yield a pointwise generalization bound via a final swap-symmetrization argument. This establishes Theorem 8: a pointwise generalization result in which the complexity is evaluated using both the observed sample S and the ghost sample S' .

Toward pointwise bounds using only observed sample. If one seeks bounds that are fully computable from the observed data $\{z_i\}_{i=1}^n$ alone, without ghost sample or sample splitting, the analysis is more involved. A practical route is two-step: (i) first derive a symmetrized pointwise bound using a complexity functional based on (S, S') or $(S, \mathbb{P})$ (which is already valid and sharp); (ii) then prove or assume an isomorphism between the $L_2(\mathbb{P}_S)$ - and $L_2(\mathbb{P})$ -induced pointwise complexities so as to replace population or ghost-dependent terms by empirical ones, yielding a fully data-dependent bound.

B.2 The PAC-Bayes Optimization Problem

We illustrate why pointwise dimension is a natural consequence of *best* PAC-Bayes optimization.

Lemma 5 (PAC-Bayes Bound [Catoni, 2003]; see also Theorems 2.1 and 2.4 in Alquier [2024])
Let π be a prior on a hypothesis class $\mathcal{F}$ independent of the data, and let $\ell: \mathcal{F} \times \mathcal{Z} \rightarrow [0, 1]$ be a bounded loss. Fix confidence $\delta \in (0, 1)$ and sample size n . For each fixed $\eta > 0$, with probability at least $1 - \delta$ over $S \sim \mathbb{P}^{\otimes n}$, simultaneously for every posterior μ ,

$$(\mathbb{P} - \mathbb{P}_n)\langle \mu, \ell \rangle \leq \frac{\text{KL}(\mu, \pi) + \log(1/\delta)}{\eta n} + \frac{\eta}{8}. \quad (\text{B.5})$$

Define

$$J_n := \lceil \log_2(4\sqrt{n}) \rceil, \quad M_n := J_n + 1.$$

There is also a single event of probability at least $1 - \delta$ on which, simultaneously for every μ ,

$$(\mathbb{P} - \mathbb{P}_n)\langle \mu, \ell \rangle \leq \frac{5}{4} \sqrt{\frac{\text{KL}(\mu, \pi) + \log(M_n/\delta)}{2n}}. \quad (\text{B.6})$$

On this same event, for every μ and every $\eta > 0$,

$$(\mathbb{P} - \mathbb{P}_n)\langle \mu, \ell \rangle \leq \frac{5}{4} \left\{ \frac{\eta}{8} + \frac{\text{KL}(\mu, \pi) + \log(M_n/\delta)}{\eta n} \right\}. \quad (\text{B.7})$$

Proof of the uniform optimization. Apply the fixed- η bound at confidence δ/M_n to each $\eta \in \{4 \cdot 2^{-j} : 0 \leq j \leq J_n\}$ and take a union bound. Write $K = \text{KL}(\mu, \pi) + \log(M_n/\delta)$. If $K \leq 2n$, then $\eta^* = \sqrt{8K/n}$ lies between the smallest grid point and 4. Choose a grid point $\eta \in [\eta^*, 2\eta^*]$. For $t = \eta/\eta^* \in [1, 2]$,

$$\frac{\eta}{8} + \frac{K}{\eta n} = \frac{t + t^{-1}}{2} \sqrt{\frac{K}{2n}} \leq \frac{5}{4} \sqrt{\frac{K}{2n}}.$$

If $K > 2n$, the same conclusion follows from the bound $(\mathbb{P} - \mathbb{P}_n)\langle \mu, \ell \rangle \leq 1$. Finally, the arithmetic-geometric mean inequality gives (B.7) for all $\eta > 0$ on the same event. $\square$

We now use the PAC-Bayes bound (which holds uniformly for every random posterior μ) to approximate a deterministic hypothesis f . On the event in (B.7), uniformly over every $\mu \in \Delta(\mathcal{F})$, $f \in \mathcal{F}$, and $\eta > 0$,

$$\begin{aligned} (\mathbb{P} - \mathbb{P}_n)\ell(f; z) &= \langle \mu, (\mathbb{P} - \mathbb{P}_n)\ell(\cdot; z) \rangle \\ &\quad + \langle \mu, (\mathbb{P} - \mathbb{P}_n)(\ell(f; z) - \ell(\cdot; z)) \rangle \\ &\leq \frac{5}{4} \left\{ \frac{\eta}{8} + \frac{\text{KL}(\mu, \pi) + \log(M_n/\delta)}{\eta n} \right\} + \langle \mu, \varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}(\cdot, f) \rangle, \end{aligned} \quad (\text{B.8})$$

where

$$\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}(f', f) := |(\mathbb{P} - \mathbb{P}_n)(\ell(f'; z) - \ell(f; z))| \leq \varrho_{(S, \mathbb{P}), \ell}^{(1)}(f', f). \quad (\text{B.9})$$

In (B.8), Lemma 5 bounds the variance term, and taking the absolute value of the signed gap difference bounds the bias term. Among symmetric nonnegative costs d satisfying $(\mathbb{P} - \mathbb{P}_n)(\ell(f; z) - \ell(g; z)) \leq d(f, g)$ for every ordered pair, $\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}$ is pointwise minimal: reversing the pair gives $|(\mathbb{P} - \mathbb{P}_n)(\ell(f; z) - \ell(g; z))| \leq d(f, g)$. This does not establish optimality among all PAC-Bayes bounds; retaining the signed posterior bias can be sharper.

Motivated by the above bias-variance optimization (B.8) via PAC-Bayes, for a given prior π , metric ϱ , and confidence $\delta \in (0, 1)$ we define the *PAC-Bayes optimization objective*

$$V(\mu, \eta, f, \varrho) := \underbrace{\frac{\eta}{8} + \frac{\text{KL}(\mu, \pi) + \log \frac{M_n}{\delta}}{\eta n}}_{\text{Variance}} + \underbrace{\langle \mu, \varrho(\cdot, f) \rangle}_{\text{Bias}}, \quad (\text{B.10})$$

where $\eta > 0$, n is the sample size, μ is a posterior over hypotheses, and $M_n = \lceil \log_2(4\sqrt{n}) \rceil + 1$ is the grid size for uniform optimization. Since the bias is nonnegative, (B.8) implies, on the same event,

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \frac{5}{4} V(\mu, \eta, f, \varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}) \quad \text{for all } f \in \mathcal{F}, \mu \in \Delta(\mathcal{F}), \eta > 0,$$

Here, the “Variance” term arises from a PAC-Bayes bound (Lemma 5) applied to μ , and the “Bias” term $\langle \mu, \varrho(\cdot, f) \rangle := \mathbb{E}_{h \sim \mu}[\varrho(h, f)]$ measures how well the randomized μ approximates the target f .

Optimizing the Posterior μ for the Objective (B.10) The intuitive analysis (B.8) explains how the PAC-Bayesian optimization objective naturally bounds the generalization gap. We now minimize the posterior μ in (B.10). For fixed η , (B.10) is minimized over μ by the Gibbs posterior. To obtain a closed-form characterization of the optimized value, we proceed in two steps: (i) derive an explicit pointwise-dimension upper bound by taking μ to be the π -normalized density on the metric ball $B_\varrho(f, \varepsilon)$ (Theorem 7), and (ii) show that this choice is near-optimal at matched scales (Lemma 6).

B.2.1 Pointwise Dimension Bound via Metric Ball

Given any prior π on $\mathcal{F}$ and any $f \in \mathcal{F}$, take μ to be the π -normalized density on the metric ball $B_\varrho(f, \varepsilon)$, i.e.,

$$\mu(A) = \frac{\pi(A \cap B_\varrho(f, \varepsilon))}{\pi(B_\varrho(f, \varepsilon))} \quad \text{for all measurable } A \subseteq \mathcal{F}. \quad (\text{B.11})$$

This choice gives an explicit upper bound on the optimized objective; Lemma 6 supplies the corresponding lower bound.

Theorem 7 (Pointwise Dimension and Pointwise Generalization Upper Bound) *For the PAC-Bayes objective (B.10), let μ be the π -normalized density on $B_\varrho(f, \varepsilon)$, assuming that this ball has positive prior mass, i.e.*

$$\frac{d\mu}{d\pi}(h) = \begin{cases} \frac{1}{\pi(B_\varrho(f, \varepsilon))}, & h \in B_\varrho(f, \varepsilon), \\ 0, & h \notin B_\varrho(f, \varepsilon). \end{cases}$$

Then, with $\eta^* = \sqrt{8(\text{KL}(\mu, \pi) + \log(M_n/\delta))}/n$,

$$V(\mu, \eta^*, f, \varrho) \leq \sqrt{\frac{\text{KL}(\mu, \pi) + \log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{2n}} + \varepsilon = \sqrt{\frac{\log \frac{1}{\pi(B_\varrho(f, \varepsilon))} + \log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{2n}} + \varepsilon. \quad (\text{B.12})$$

Combining the upper bound (B.12) with (B.8) yields the pointwise generalization bound: for every $\delta \in (0, 1)$, with probability at least $1 - \delta$, uniformly over every $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_n)\ell(f; z) \leq \frac{5}{4} \inf_{\varepsilon > 0} \left\{ \varepsilon + \sqrt{\frac{\log \frac{1}{\pi(B_{\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}}(f, \varepsilon))}} + \log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{2n}} \right\}.$$

Here $\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}}$ is the generalization-gap metric in (B.9). This proves Theorem 1, after enlarging its absolute constant.

Remark (why this intuition matters). The generalization-gap metric gives the sharpest symmetric pointwise bound on the bias, but depends on the unknown gaps themselves. Its mixed L_1 relaxation retains the loss geometry, while the mixed L_2 metric used by chaining is

$$\varrho_{(S, \mathbb{P}), \ell}(f', f) := ((\mathbb{P}_S + \mathbb{P})[\ell(f'; \cdot) - \ell(f; \cdot)]^2)^{1/2}. \quad (\text{B.13})$$

The inequalities $\varrho_{(S, \mathbb{P}), \ell}^{\text{gap}} \leq \varrho_{(S, \mathbb{P}), \ell}^{(1)} \leq \sqrt{2} \varrho_{(S, \mathbb{P}), \ell}$ connect these one-shot bounds. Corollary 1 records the common L_2 consequence of the PAC-Bayes and chaining theorems. The comparison between the integral and one-shot objectives is made in that same metric; it does not assert that the mixed L_2 integral always dominates the mixed L_1 one-shot bound.

Proof of Theorem 7: For the choice (B.11),

$$\text{KL}(\mu, \pi) = \int_{\mathcal{F}} \log\left(\frac{d\mu}{d\pi}(h)\right) \mu(dh) = \int_{B_\varrho(f, \varepsilon)} \log\left(\frac{1}{\pi(B_\varrho(f, \varepsilon))}\right) \mu(dh) = \log \frac{1}{\pi(B_\varrho(f, \varepsilon))}. \quad (\text{B.14})$$

Moreover, by construction,

$$\langle \mu, \varrho(\cdot, f) \rangle = \int_{B_\varrho(f, \varepsilon)} \varrho(h, f) \mu(dh) \leq \varepsilon.$$

Plugging (B.14) into (B.10) and minimizing $\frac{\eta}{8} + \frac{\text{KL}(\mu, \pi) + \log \frac{M_n}{\delta}}{\eta n}$ over $\eta > 0$ gives $\sqrt{\frac{\text{KL}(\mu, \pi) + \log \frac{M_n}{\delta}}{2n}}$, which together with the bias bound $\langle \mu, \varrho(\cdot, f) \rangle \leq \varepsilon$ yields the claimed bound (B.12). $\square$

B.2.2 Lower Bound and Optimality of PAC-Bayes Optimization

The following lemma compares the ball-posterior upper bound with the optimized Gibbs objective. The localization terms are $a + \varepsilon$ above and $\min\{a, \varepsilon\}$ below, apart from the stated additive term, where $a = \log(1/\pi(B_\varrho(f, \varepsilon)))/(\eta n)$. They are comparable when a and ε are comparable.

Lemma 6 (Optimality of Pointwise Dimension in PAC-Bayes Optimization) *For the PAC-Bayes optimization objective $V(\mu, \eta, f, \varrho)$ defined in (B.10), we have that for every $f \in \mathcal{F}$, $\eta > 0$, and $\varepsilon > 0$,*

$$\inf_{\mu} V(\mu, \eta, f, \varrho) \geq \frac{\eta}{8} + \frac{\log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{\eta n} + \min\left\{\frac{1}{\eta n} \log \frac{1}{\pi(B_\varrho(f, \varepsilon))}, \varepsilon\right\} - \frac{\log 2}{\eta n}. \quad (\text{B.15})$$

Consequently, for every $f \in \mathcal{F}$, $\eta > 0$, and $\varepsilon > 0$,

$$\begin{aligned} & \frac{\eta}{8} + \frac{\log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{\eta n} + \min\left\{\frac{\log \frac{1}{\pi(B_\varrho(f, \varepsilon))}}{\eta n}, \varepsilon\right\} - \frac{\log 2}{\eta n} \\ & \leq \inf_{\mu} V(\mu, \eta, f, \varrho) \\ & \leq \frac{\eta}{8} + \frac{\log \frac{1}{\pi(B_\varrho(f, \varepsilon))} + \log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{\eta n} + \varepsilon. \end{aligned} \quad (\text{B.16})$$

Proof of Lemma 6 The upper bound in (B.16) is already proved in Theorem 7, so we only need to prove the lower bound (B.15). The Donsker-Varadhan variational identity states that for any measurable h ,

$$-\log \int e^h d\pi = \inf_{\mu} \left\{ \text{KL}(\mu, \pi) - \int h d\mu \right\}.$$

Apply it with $h = -\eta n \varrho(\cdot, f)$ to obtain

$$-\log \int e^{-\eta n \varrho(\cdot, f)} d\pi = \inf_{\mu} \left\{ \text{KL}(\mu, \pi) + \int \eta n \varrho(\cdot, f) d\mu \right\},$$

which implies that

$$\frac{\eta}{8} + \frac{\log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{\eta n} - \frac{1}{\eta n} \log \int e^{-\eta n \varrho(\cdot, f)} d\pi = \inf_{\mu} \left\{ \frac{\eta}{8} + \frac{\text{KL}(\mu, \pi) + \log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{\eta n} + \langle \mu, \varrho(\cdot, f) \rangle \right\}. \quad (\text{B.17})$$

By splitting the dual integral,

$$\begin{aligned} \int e^{-\eta n \varrho(\cdot, f)} d\pi &= \int_{B_{\varrho}(f, \varepsilon)} e^{-\eta n \varrho(\cdot, f)} d\pi + \int_{B_{\varrho}(f, \varepsilon)^c} e^{-\eta n \varrho(\cdot, f)} d\pi \\ &\leq \pi(B_{\varrho}(f, \varepsilon)) + e^{-\eta n \varepsilon} (1 - \pi(B_{\varrho}(f, \varepsilon))) \\ &\leq \pi(B_{\varrho}(f, \varepsilon)) + e^{-\eta n \varepsilon}, \end{aligned}$$

where $B_{\varrho}(f, \varepsilon)^c$ is complement of $B_{\varrho}(f, \varepsilon)$; and we have used $e^{-\eta n \varrho(\cdot, f)} \leq 1$ on $B_{\varrho}(f, \varepsilon)$ and $e^{-\eta n \varrho(\cdot, f)} \leq e^{-\eta n \varepsilon}$ on $B_{\varrho}(f, \varepsilon)^c$. Hence

$$\inf_{\mu} V(\mu, \eta, f, \varrho) \geq \frac{\eta}{8} + \frac{\log \frac{\lceil \log_2(4\sqrt{n}) \rceil + 1}{\delta}}{\eta n} - \frac{1}{\eta n} \log \left(\pi(B_{\varrho}(f, \varepsilon)) + e^{-\eta n \varepsilon} \right). \quad (\text{B.18})$$

The simplified form (B.15) follows from $a + b \leq 2 \max\{a, b\}$ or equivalently $-\log(a + b) \geq -\log 2 + \min\{-\log a, -\log b\}$ on (B.18). Combining (B.10), (B.14) and (B.15) yields the sandwich (B.16). $\square$

B.3 Subset Homogeneity of Pointwise Dimension

We show that, for any $f \in \mathcal{H} \subseteq \mathcal{F}$, the pointwise–dimension functional defined with a prior π is unchanged in order (up to absolute constants) whether π is supported on $\mathcal{H}$ or on the ambient class $\mathcal{F}$. Hence one may take $\pi \in \Delta(\mathcal{F})$ without restricting it to any particular subset, which suffices to meet the subset–homogeneity condition in Appendix B.1.1.

Lemma 7 (Ambient Equivalence of Pointwise Dimension) *Let $(\mathcal{F}, \varrho)$ be a metric space and let $\mathcal{H} \subseteq \mathcal{F}$ be a subset. Consider a nearest-point selector $p : \mathcal{F} \rightarrow \mathcal{H}$ satisfying $\varrho(f, p(f)) = \min_{h \in \mathcal{H}} \varrho(f, h)$ for all $f \in \mathcal{F}$. For any ambient prior $\pi \in \Delta(\mathcal{F})$, let $\pi_{\mathcal{H}} := p_{\#} \pi$ be the pushforward measure induced by p , defined by*

$$\pi_{\mathcal{H}}(A) := \pi(\{g \in \mathcal{F} : p(g) \in A\}), \quad A \subseteq \mathcal{H} \text{ measurable.}$$

Equivalently, in the discrete case, $\pi_{\mathcal{H}}(h) = \sum_{g \in \mathcal{F} : p(g)=h} \pi(g)$, that is, $\pi_{\mathcal{H}}$ collects all ambient prior mass whose nearest point in $\mathcal{H}$ is h . Then for every $\varepsilon > 0$ we have

$$\pi_{\mathcal{H}}(B_{\varrho}(f, 2\varepsilon)) \geq \pi(B_{\varrho}(f, \varepsilon)), \quad \log \frac{1}{\pi_{\mathcal{H}}(B_{\varrho}(f, 2\varepsilon))} \leq \log \frac{1}{\pi(B_{\varrho}(f, \varepsilon))}.$$

Consequently, for any ambient prior $\pi \in \Delta(\mathcal{F})$ and any point $f \in \mathcal{F}$, define the majorizing measure integral

$$I(\pi, f, \varrho) := \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{+\infty} \sqrt{\log \frac{1}{\pi(B_{\varrho}(f, \varepsilon))}} d\varepsilon \right\}.$$

Then we have

$$\frac{1}{2} \inf_{\mu \in \Delta(\mathcal{H})} \sup_{f \in \mathcal{H}} I(\mu, f, \varrho) \leq \inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{H}} I(\pi, f, \varrho) \leq \inf_{\mu \in \Delta(\mathcal{H})} \sup_{f \in \mathcal{H}} I(\mu, f, \varrho). \quad (\text{B.19})$$

Proof of Lemma 7: The upper bound in (B.19) is immediate since $\Delta(\mathcal{H}) \subseteq \Delta(\mathcal{F})$: taking μ supported on $\mathcal{H}$ gives $\inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{H}} I(\pi, f, \varrho, r) \leq \inf_{\mu \in \Delta(\mathcal{H})} \sup_{f \in \mathcal{H}} I(\mu, f, \varrho, r)$.

For the lower bound in (B.19), take $\pi_{\mathcal{H}}$ to be the pushforward induced by the nearest-point selector. For any $f \in \mathcal{H}$ and $\varepsilon > 0$, if $f' \in B_{\varrho}(f, \varepsilon)$ then

$$\varrho(p(f'), f) \leq \varrho(p(f'), f') + \varrho(f', f) = \min_{h \in \mathcal{H}} \varrho(f', h) + \varrho(f', f) \leq 2\varepsilon,$$

hence $p(f') \in B_{\varrho}(f, 2\varepsilon)$ and

$$\pi_{\mathcal{H}}(B_{\varrho}(f, 2\varepsilon)) \geq \pi(B_{\varrho}(f, \varepsilon)), \quad \log \frac{1}{\pi_{\mathcal{H}}(B_{\varrho}(f, 2\varepsilon))} \leq \log \frac{1}{\pi(B_{\varrho}(f, \varepsilon))}. \quad (\text{B.20})$$

Therefore,

$$\begin{aligned} I(\pi, f, \varrho) &= \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{+\infty} \sqrt{\log \frac{1}{\pi(B_{\varrho}(f, \varepsilon))}} d\varepsilon \right\} \\ &\geq \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{+\infty} \sqrt{\log \frac{1}{\pi_{\mathcal{H}}(B_{\varrho}(f, 2\varepsilon))}} d\varepsilon \right\} \\ &= \frac{1}{2} \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{+\infty} \sqrt{\log \frac{1}{\pi_{\mathcal{H}}(B_{\varrho}(f, \varepsilon))}} d\varepsilon \right\} \\ &= \frac{1}{2} I(\pi_{\mathcal{H}}, f, \varrho), \end{aligned}$$

where the first inequality is by (B.20); the second equality is by the change of variables. Taking $\sup_{f \in \mathcal{H}}$ and then $\inf_{\pi \in \Delta(\mathcal{F})}$, $\inf_{\mu \in \Delta(\mathcal{H})}$ yields the desired lower bound. $\square$

Relationship to Fractional and Ordinary Covering Numbers As with ordinary covering numbers in Definition 5, we distinguish between internal and external fractional covering numbers. For $\mathcal{H} \subseteq \mathcal{F}$, define

$$N'(\mathcal{H}, \varrho, \varepsilon) := \inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{H}} \frac{1}{\pi(B_{\varrho}(f, \varepsilon))}.$$

This is the *external fractional covering number* relative to the ambient class $\mathcal{F}$; restricting the infimum to $\Delta(\mathcal{H})$ gives the internal version. See Chen et al. [2024, Section 4.1] for the general formulation and its connections to information-theoretic upper and lower bounds.

Let $N(\mathcal{H}, \varrho, \varepsilon)$ denote the internal covering number from Definition 5. A standard packing argument gives the following comparison; see Chen et al. [2024, Lemma 14] and, in the batch $L_2(\mathbb{P}_n)$ setting, Block et al. [2021, Lemma 8]:

$$\begin{aligned} \log N(\mathcal{H}, \varrho, 2\varepsilon) &\leq \log N'(\mathcal{H}, \varrho, \varepsilon) \\ &= \inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{H}} \log \frac{1}{\pi(B_{\varrho}(f, \varepsilon))} \leq \log N(\mathcal{H}, \varrho, \varepsilon). \end{aligned} \quad (\text{B.21})$$

The cited proofs establish this class-wide comparison through packing, whereas Lemma 7 provides a novel *pointwise* comparison of ball probabilities simultaneously across scales.

Finally, taking the ambient class to be the target class in the comparison above and applying Lemma 15 shows that the internal fractional covering number in (2.7) is equivalent to the ordinary external covering number from Definition 5, up to a constant rescaling of the radius. Together, these arguments establish the equivalence claimed in the main text and explicitly show that external and internal fractional covering numbers are equivalent up to a constant rescaling of the radius. This comparison is not stated separately by Block et al. [2021], Chen et al. [2024].

The equivalence of internal and external ordinary covering numbers established in Lemma 15 of Appendix D.6, together with the corresponding equivalence for fractional covering numbers derived above, helps clarify the contribution of our Lemma 7. The equivalence of internal and external covering numbers suggests that ambient equivalence should also hold for class-wide integral functionals, even when the truncation level is optimized separately for each hypothesis. The novelty of Lemma 7 lies in establishing a *pointwise* comparison that applies to *pointwise* dimensions as well as to these class-wide functionals.

Collapsing the Distinction between Chaining and Generic Chaining. A simple illustration of the strength of our pointwise blueprint is the multi-dimensional setting. Let $(d^{(1)}, \dots, d^{(k)}) : \mathcal{F} \rightarrow (0, R]^k$ be coordinatewise, data-independent complexity functions; and ψ be coordinatewise nondecreasing in $(d^{(1)}, \dots, d^{(k)})$ and nonincreasing in δ . Our blueprint makes no essential distinction between the two uniform forms

$$\begin{aligned} (\text{sup-inside}) \quad \sup_{f \in \mathcal{H}} (\mathbb{P} - \mathbb{P}_n) \ell(f; z) &\leq \psi \left(\sup_{f \in \mathcal{H}} d^{(1)}(f), \dots, \sup_{f \in \mathcal{H}} d^{(k)}(f); \delta \right), \\ (\text{sup-outside}) \quad \sup_{f \in \mathcal{H}} (\mathbb{P} - \mathbb{P}_n) \ell(f; z) &\leq \sup_{f \in \mathcal{H}} \psi \left(d^{(1)}(f), \dots, d^{(k)}(f); \delta \right), \end{aligned}$$

in the sense that *either* one leads to the same pointwise conclusion after peeling.

More precisely, fix a base scale $r_0 \in (0, R]$. Then with probability at least $1 - \delta$, for every $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_n) \ell(f; z) \leq \psi \left(\left(\dots, \max \left\{ 2d^{(j)}(f), r_0 \right\}, \dots \right); \delta \left(\log_2 \frac{2R}{r_0} \right)^{-k} \right). \quad (\text{B.22})$$

The most straightforward proof uses essentially the same peeling argument as in Lemma 4, with the only change that we use a grid of size $(\log_2(2R/r_0))^k$ (partition each coordinate into $\log_2(2R/r_0)$ dyadic scales); see the short proof of Proposition 1 in Xu and Zeevi [2025]. Alternatively, this can be proved by applying Lemma 4 for k times, where at each step we remove one dimension functional and divided confidence by $\log_2(2R/r_0)$. Moreover, the multi-dimensional pointwise bound (B.22) shows that its right-hand side, viewed as a *scalar* complexity, yields an equally tight pointwise bound. Hence the multi-dimensional formulation does not improve the best-achievable rates beyond a suitably defined one-dimensional complexity (as in generic chaining).

Conceptually, this shows that the apparent gap between classical chaining (entropy integral; sup-inside), generic chaining (majorizing measures; sup-outside), and our pointwise generic-chaining bound (Theorem 2) disappears within the blueprint: each is just a subset-homogeneous uniform statement that implies the same pointwise bound up to absolute constants and logarithms. The peeling argument itself does not compare empirical and mixed metrics. For the worst-case family in Theorem 3, the corresponding class-wide functionals are equivalent in expectation, including pointwise optimized truncation; Appendix B.6 proves this comparison and explains why it does not extend to arbitrary classes.

B.4 Pointwise Generalization Bound via Ghost Sample

In this section, we prove an easier variant of Theorem 2 that permits *swap-invariant* randomized priors depending on both the observed sample and its ghost counterpart. This setting subsumes—and strengthens—the conditional mutual information (CMI) framework of Steinke and Zakyntinou [2020].

Let $S = (z_1, \dots, z_n)$ and $S' = (z'_1, \dots, z'_n)$ be two i.i.d. samples drawn from $\mathbb{P}^{\otimes n}$, independent of each other. For each index $i \in \{1, \dots, n\}$, define the coordinate-swap map

$$\tau_i(S, S') := ((z_1, \dots, z_{i-1}, z'_i, z_{i+1}, \dots, z_n), (z'_1, \dots, z'_{i-1}, z_i, z'_{i+1}, \dots, z'_n)).$$

A randomized, data-dependent prior is a mapping $\pi_{(\cdot, \cdot)} : \mathcal{Z}^{2n} \rightarrow \Delta(\mathcal{F})$; we write $\pi_{(S, S')} \in \Delta(\mathcal{F})$ for the realized prior over $\mathcal{F}$ (a distribution on $\mathcal{F}$ that may depend on (S, S')). We say that π is *swap-invariant* on (S, S') , if

$$\pi_{(S, S')} = \pi_{\tau_i(S, S')} \quad \text{for all } i = 1, \dots, n \text{ and for } \mathbb{P}^{\otimes 2n}\text{-a.e. } (S, S').$$

Equivalently, π depends only on the unordered multiset $\{(z_i, z'_i)\}_{i=1}^n$ and not on which element of each pair is designated as “observed” versus “ghost.”

Connection to CMI. This notion covers the conditional-mutual-information (CMI) framework of Steinke and Zakyntinou [2020]. In the CMI setup, one draws paired data $Z = ((Z_i^{(0)}, Z_i^{(1)}))_{i=1}^n \stackrel{\text{i.i.d.}}{\sim} (\mathbb{P} \times \mathbb{P})^{\otimes n}$ and an independent selector $U \in \{0, 1\}^n$. The training and ghost sets are $S_U = (Z_1^{(U_1)}, \dots, Z_n^{(U_n)})$ and $S_{\bar{U}} = (Z_1^{(1-U_1)}, \dots, Z_n^{(1-U_n)})$. Any prior π that is a function of Z only (independent of U) is swap-invariant, since flipping U_i implements τ_i . Conversely, swap-invariance for all i is equivalent to invariance under all coordinatewise flips of U , hence independence from U .

Throughout, let $S = \{z_i\}_{i=1}^n$ be the observed sample and $S' = \{z'_i\}_{i=1}^n$ an i.i.d. ghost sample, independent of S . We write $\mathbb{P}_S$ for the empirical measure $\mathbb{P}_n$ based on S , and $\varrho_{S, \ell}$ for the metric $\varrho_{n, \ell}$ from the main paper. Let $\mathbb{P}_{S'}$ denote the empirical measure based on S' . For any integrable function $g : \mathcal{Z} \rightarrow \mathbb{R}$ (we write $g(z)$ when convenient; e.g., $g(z) = \ell(f; z)$), define the empirical averaging operators

$$\mathbb{P}_S g := \frac{1}{n} \sum_{i=1}^n g(z_i), \quad \mathbb{P}_{S'} g := \frac{1}{n} \sum_{i=1}^n g(z'_i).$$

We use the shorthand $(\mathbb{P}_S \pm \mathbb{P}_{S'})g := \mathbb{P}_S g \pm \mathbb{P}_{S'} g$ for the sum/difference of the two sample-average operators, and the same notation when $\mathbb{P}_S \pm \mathbb{P}_{S'}$ appear inside norms or distances.

Theorem 8 (Pointwise Generalization via Ghost Sample) *Let $\ell(f; z) \in [0, 1]$. There exists an absolute constant $C > 0$ such that for any swap-invariant prior $\pi_{(\cdot, \cdot)}$ on (S, S') , and any $\delta \in (0, 1)$, with probability at least $1 - \delta$ over (S, S') , uniformly in $f \in \mathcal{F}$,*

$$\begin{aligned} & (\mathbb{P}_{S'} - \mathbb{P}_S) \ell(f; z) \\ & \leq C \left(\inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi_{(S, S')} (B_{\varrho_{(S, S'), \ell}(f, \varepsilon))}}} d\varepsilon \right\} + \sqrt{\frac{\log(\log(2n)/\delta)}{n}} \right), \end{aligned}$$

where $\varrho_{(S, S'), \ell}(f_1, f_2) = \sqrt{(\mathbb{P}_S + \mathbb{P}_{S'}) (\ell(f_1; z) - \ell(f_2; z))^2}$.

Proof of Theorem 8: The proof of the upper bound in Theorem 8 consists of three steps: 1. Subset-Homogeneous Uniform Convergence; 2. Generic Conversion to Pointwise Generalization Bound; 3. High-Probability Symmetrization.

Step 1: Subset-Homogeneous Uniform Convergence. Let $S = \{z_i\}_{i=1}^n$ be the observed sample, and $S' = \{z'_i\}_{i=1}^n$ be an i.i.d. ghost sample. We consider the symmetrized loss

$$\tilde{\ell}(f; (z, z')) = \ell(f; z') - \ell(f; z). \quad (\text{B.23})$$

Since $\ell(f; z)$ is uniformly bounded in $[0, 1]$, $\tilde{\ell}(f; (z, z'))$ is uniformly bounded in $[-1, 1]$. We adopt the notation

$$\varrho_{S, \ell}(f_1, f_2) = \varrho_{n, \ell}(f_1, f_2) = \sqrt{\mathbb{P}_S(\ell(f_1; z) - \ell(f_2; z))^2}.$$

Furthermore, we define the loss-induced L_2 metrics $\varrho_{S', \ell}$ and $\varrho_{(S, S'), \ell}$ by

$$\begin{aligned} \varrho_{S', \ell}(f_1, f_2) &= \sqrt{\mathbb{P}_{S'}(\ell(f_1; z) - \ell(f_2; z))^2}, \\ \varrho_{(S, S'), \ell}(f_1, f_2) &= \sqrt{(\mathbb{P}_S + \mathbb{P}_{S'}) (\ell(f_1; z) - \ell(f_2; z))^2}. \end{aligned}$$

By Minkowski's inequality (see, e.g., [Wikipedia contributors \[2025c\]](#)) and $\sqrt{a} + \sqrt{b} \leq \sqrt{2(a+b)}$, we have

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\tilde{\ell}(f_1; (z_i, z'_i)) - \tilde{\ell}(f_2; (z_i, z'_i)))^2} \leq \varrho_{S, \ell}(f_1, f_2) + \varrho_{S', \ell}(f_1, f_2) \leq \sqrt{2} \varrho_{(S, S'), \ell}(f_1, f_2). \quad (\text{B.24})$$

Now applying the truncated integral bound (Lemma 13) to the empirical Rademacher complexity: let $\{\xi_i\}_{i=1}^n$ be i.i.d. Rademacher variables, then conditioned on (S, S') , given any subset $\mathcal{H} \subseteq \mathcal{F}$, we have that for all $\delta \in (0, 1)$, with probability at least $1 - \delta$ (the randomness all comes from $\{\xi_i\}_{i=1}^n$),

$$\begin{aligned} \sup_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \xi_i \tilde{\ell}(f; (z_i, z'_i)) &\leq \mathbb{E}_{\xi} \left[\sup_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \xi_i \tilde{\ell}(f; (z_i, z'_i)) \right] + \sqrt{\frac{2 \log \frac{1}{\delta}}{n}} \\ &\leq C_0 \inf_{\mu \in \Delta(\mathcal{H})} \sup_{f \in \mathcal{H}} \inf_{0 \leq \alpha \leq 2} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^2 \sqrt{\log \frac{1}{\mu(B_{\tilde{\varrho}}(f, \varepsilon))}} d\varepsilon \right\} + \sqrt{\frac{2 \log \frac{1}{\delta}}{n}}, \end{aligned}$$

where $C_0 > 0$ is an absolute constant, and $\tilde{\varrho}(f_1, f_2) := \sqrt{\frac{1}{n} \sum_{i=1}^n (\tilde{\ell}(f_1; (z_i, z'_i)) - \tilde{\ell}(f_2; (z_i, z'_i)))^2}$. Here, the first inequality is by McDiarmid's inequality (Lemma 16); and the second inequality is by Lemma 13; and the integral is capped at 2 because

$$\sup_{f_1 \in \mathcal{H}, f_2 \in \mathcal{H}} \tilde{\varrho}(f_1, f_2) \leq \sup_{f_1 \in \mathcal{H}, f_2 \in \mathcal{H}} \sqrt{2} \varrho_{(S, S'), \ell}(f_1, f_2) \leq 2,$$

where we have used (B.24). By the ambient-equivalence of the pointwise-dimension functional (Lemma 7), we have (note that we take the support of π to be $\mathcal{F}$ rather than $\mathcal{H}$)

$$\begin{aligned} &\inf_{\mu \in \Delta(\mathcal{H})} \sup_{f \in \mathcal{H}} \inf_{0 \leq \alpha \leq 2} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^2 \sqrt{\log \frac{1}{\mu(B_{\tilde{\varrho}}(f, \varepsilon))}} d\varepsilon \right\} \\ &\leq 2 \inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{H}} \inf_{0 \leq \alpha \leq 2} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^2 \sqrt{\log \frac{1}{\pi(B_{\tilde{\varrho}}(f, \varepsilon))}} d\varepsilon \right\}. \end{aligned}$$

By (B.24) and the fact that pointwise dimension is monotone in the underlying metric (Lemma 18), we have that for any $\pi \in \Delta(\mathcal{F})$,

$$\int_{\alpha}^2 \sqrt{\log \frac{1}{\pi(B_{\tilde{\varrho}}(f, \varepsilon))}} d\varepsilon \leq \int_{\alpha}^2 \sqrt{\log \frac{1}{\pi(B_{\sqrt{2}\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon = \sqrt{2} \int_{\alpha/\sqrt{2}}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon,$$

where the equality follows by a change of variables. Combining the above three inequalities, we prove the following *subset-homogeneous* uniform convergence argument when choosing an arbitrary $\pi \in \Delta(\mathcal{F})$: conditioned on (S, S') , for any $\mathcal{H} \subseteq \mathcal{F}$ and $\delta \in (0, 1)$, with probability at least $1 - \delta$ (the randomness all comes from $\{\xi_i\}_{i=1}^n$),

$$\sup_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \xi_i(\ell(f; z'_i) - \ell(f; z_i)) \leq \sup_{f \in \mathcal{H}} C_1 \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon \right\} + \sqrt{\frac{2 \log \frac{1}{\delta}}{n}}, \quad (\text{B.25})$$

where $C_1 = 2\sqrt{2}C_0 > 0$ is an absolute constant.

Conditioned on (S, S') , for fixed $\pi_{(S, S')} \in \Delta(\mathcal{F})$ that is independent with $\{\xi_i\}_{i=1}^n$, define the pointwise complexity

$$d_{S, S'}(f) := \left(\inf_{\alpha \geq 0} \left\{ \sqrt{n} \alpha + \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi_{(S, S')}(B_{\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon \right\} \right)^2. \quad (\text{B.26})$$

Then, by (B.25), conditioned on (S, S') , for any $\mathcal{H} \subseteq \mathcal{F}_R$ and any $\delta \in (0, 1)$, with probability at least $1 - \delta$ (the randomness all comes from $\{\xi_i\}_{i=1}^n$),

$$\sup_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \xi_i(\ell(f; z'_i) - \ell(f; z_i)) \leq \sup_{f \in \mathcal{H}} \left(C_1 \sqrt{\frac{d_{S, S'}(f)}{n}} + \sqrt{\frac{2 \log \frac{2}{\delta}}{n}} \right). \quad (\text{B.27})$$

As discussed in Appendix B.1.1, this condition is both necessary and sufficient to establish pointwise convergence when the complexity functional is the $\{\xi_i\}_{i=1}^n$ -independent $d_{S, S'}(\cdot)$ when conditioned on (S, S') .

Step 2: Generic Conversion to Pointwise Generalization Bound. All the analysis in this step is condition on (S, S') , thus all the randomness discussed here comes from $\{\xi_i\}_{i=1}^n$. For every $r \in [0, 2n]$, we take the subset

$$\mathcal{H} = \{f \in \mathcal{F} : d_{S, S'}(f) \leq r\},$$

which, by (B.27), implies that $\forall \delta \in (0, 1)$ and $\forall r \in [0, 2n]$, with probability at least $1 - \delta$

$$\sup_{f: d_{S, S'}(f) \leq r} \frac{1}{n} \sum_{i=1}^n \xi_i(\ell(f; z'_i) - \ell(f; z_i)) \leq C_1 \sqrt{\frac{r}{n}} + \sqrt{\frac{2 \log \frac{2}{\delta}}{n}}, \quad (\text{B.28})$$

where C_1 is an absolute constant. The inequality (B.28) is precisely the condition (B.1) in the generic conversion provided in Lemma 4 (here, the expectation (equal to 0) and the empirical

average are taken for $\{\xi_i\}_{i=1}^n$). Thus applying Lemma 4 we have the pointwise generalization bound: conditioned on (S, S') , for any $\delta \in (0, 1)$, by taking $r_0 = 1/n$, with probability at least $1 - \delta$, uniformly over all $f \in \mathcal{F}$,

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n \xi_i(\ell(f; z'_i) - \ell(f; z_i)) &\leq \frac{C_1}{\sqrt{n}} \sqrt{\max \left\{ 2d_{S, S'}(f), \frac{1}{n} \right\}} + \sqrt{\frac{2 \log \frac{2 \log_2(4n^2)}{\delta}}{n}} \\ &\leq C_1 \sqrt{\frac{2d_{S, S'}(f)}{n}} + \frac{C_1}{n} + \sqrt{\frac{2 \log \frac{4 \log_2(2n)}{\delta}}{n}}. \\ &\leq C_2 \left(\sqrt{\frac{d_{S, S'}(f)}{n}} + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right), \end{aligned} \quad (\text{B.29})$$

where $C_2 > 0$ is an absolute constant, where the second inequality is because there exists $C_2 \geq \sqrt{2}C_1$ such that for all positive integer n ,

$$\frac{C_1}{n} + \sqrt{\frac{2 \left(\log \frac{1}{\delta} + \log(\log 2 + \log n) + \log \frac{4}{\log 2} \right)}{n}} \leq C_2 \sqrt{\frac{\log \frac{1}{\delta} + \log(2(\log 2 + \log n))}{n}}.$$

Thus we prove the pointwise generalization bound (B.29) for the complexity functional $d_{S, S'}(\cdot)$ defined in (B.26), under the randomness of $\{\xi_i\}_{i=1}^n$: for all $\pi \in \Delta(\mathcal{F})$, conditioned on (S, S') , for any $\delta \in (0, 1)$, with probability at least $1 - \delta$, uniformly over all $f \in \mathcal{F}$,

$$\begin{aligned} &\frac{1}{n} \sum_{i=1}^n \xi_i(\ell(f; z'_i) - \ell(f; z_i)) \\ &\leq C_2 \left(\inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi_{(S, S')}(B_{\varrho_{(S, S'), \ell}(f, \varepsilon))}}} d\varepsilon \right\} + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right). \end{aligned} \quad (\text{B.30})$$

Step 3: High-Probability Symmetrization. Recall that $S = \{z_i\}_{i=1}^n$ and $S' = \{z'_i\}_{i=1}^n$ are i.i.d. samples, independent of each other, and $\{\xi_i\}_{i=1}^n$ are i.i.d. Rademacher signs, independent of (S, S') . The mixed (ghost) metric

$$\varrho_{(S, S'), \ell}(f_1, f_2) = \sqrt{(\mathbb{P}_S + \mathbb{P}_{S'}) (\ell(f_1; z) - \ell(f_2; z))^2},$$

is swap-invariant to the pair (z_i, z'_i) for each $i = 1, \dots, n$. By the definition of swap-invariant prior before Theorem 8, the prior $\pi_{(S, S')} \in \Delta(\mathcal{F})$ is also swap-invariant to the pair (z_i, z'_i) .

Denote the functionals

$$\begin{aligned} &X(f; S, S'; \delta) \\ &:= \frac{1}{n} \sum_{i=1}^n (\ell(f; z'_i) - \ell(f; z_i)) \\ &\quad - C_2 \left(\inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi_{(S, S')}(B_{\varrho_{(S, S'), \ell}(f, \varepsilon))}}} d\varepsilon \right\} + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right), \end{aligned}$$

and

$$\begin{aligned}
& Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \\
& := \frac{1}{n} \sum_{i=1}^n \xi_i (\ell(f; z'_i) - \ell(f; z_i)) \\
& \quad - C_2 \left(\inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi_{(S, S')}(\mathcal{B}_{\varrho_{(S, S'), \ell}(f, \varepsilon))}} d\varepsilon} \right\} + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right).
\end{aligned}$$

Symmetry Argument. We write $\stackrel{d}{=}$ to denote equality in distribution (i.e., the random variables have the same law, equivalently the same cumulative distribution function). For each $i \in \{1, \dots, n\}$, let $\tau_i(S, S')$ be the pair obtained by swapping z_i and z'_i . Since $\varrho_{(S, S'), \ell}$ and $\pi_{(S, S')}$ are invariant under $(S, S') \mapsto \tau_i(S, S')$ and $\tau_i(S, S') \stackrel{d}{=} (S, S')$, we have, for all $t \in \mathbb{R}$,

$$\begin{aligned}
& \Pr(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq t) \\
& = \frac{1}{2} \Pr(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq t | \xi_i = 1) + \frac{1}{2} \Pr(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq t | \xi_i = -1) \\
& = \Pr(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq t | \xi_i = 1),
\end{aligned}$$

i.e., $Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \stackrel{d}{=} Y(f; S, S', \{\xi_1, \dots, \xi_{i-1}, 1, \xi_{i+1}, \dots, \xi_n\}; \delta)$. In the second equality, we use the following joint-law symmetry. The map

$$(S, S', \{\xi_j\}_{j=1}^n) \mapsto (\tau_i(S, S'), \{\xi_1, \dots, \xi_{i-1}, -\xi_i, \xi_{i+1}, \dots, \xi_n\})$$

preserves the joint law of $(S, S', \{\xi_j\}_{j=1}^n)$, because (z_i, z'_i) are i.i.d. and ξ_i is symmetric. Moreover, by the swap-invariance of $\varrho_{(S, S'), \ell}$ and $\pi_{(S, S')}$, this transformation leaves the value of the process Y unchanged. Hence the two distributions coincide, and

$$\Pr(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq t | \xi_i = 1) = \Pr(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq t | \xi_i = -1).$$

Iterating over all indices $i = 1, \dots, n$, we obtain that

$$Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \stackrel{d}{=} Y(f; S, S', \{1, \dots, 1\}; \delta) = X(f; S, S'; \delta).$$

By the conclusion (B.30) in Step 2 and the tower property, we have

$$\Pr_{S, S', \xi} \left(Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta) \leq 0 \text{ for all } f \in \mathcal{F} \right) \geq 1 - \delta.$$

Moreover, the symmetry argument above gives the process-level equality in distribution

$$\{Y(f; S, S', \{\xi_i\}_{i=1}^n; \delta)\}_{f \in \mathcal{F}} \stackrel{d}{=} \{X(f; S, S'; \delta)\}_{f \in \mathcal{F}}.$$

Therefore,

$$\Pr_{S, S'} \left(X(f; S, S'; \delta) \leq 0 \text{ for all } f \in \mathcal{F} \right) \geq 1 - \delta.$$

Hence, with probability at least $1 - \delta$ over the draw of (S, S') , we have, uniformly over all $f \in \mathcal{F}$,

$$\begin{aligned} & \frac{1}{n} \sum_{i=1}^n (\ell(f; z'_i) - \ell(f; z_i)) \\ & \leq C_2 \left(\inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi_{(S, S')}(B_{\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon \right\} + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right), \end{aligned}$$

where $\varrho_{(S, S'), \ell}(f_1, f_2) = \sqrt{(\mathbb{P}_S + \mathbb{P}_{S'}) (\ell(f_1; z) - \ell(f_2; z))^2}$, and $C_2 > 0$ is an absolute constant. $\square$

B.5 Proof of Theorem 2

We first retain the bound with a conditional ghost expectation as an intermediate lemma.

Lemma 8 (Conditional ghost bound) *Let $\ell(f; z) \in [0, 1]$, and let π be a prior on $\mathcal{F}$ independent of both S and S' . There is an absolute constant $C > 0$ such that, for every $n \geq 1$ and $\delta \in (0, 1)$, with probability at least $1 - \delta$ over S , uniformly over $f \in \mathcal{F}$,*

$$\begin{aligned} (\mathbb{P} - \mathbb{P}_n) \ell(f; z) & \leq C \left(\mathbb{E}_{S'} \left[\inf_{0 \leq \alpha \leq \sqrt{2}} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon \right\} \mid S \right] \right. \\ & \quad \left. + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right), \end{aligned}$$

where $\varrho_{(S, S'), \ell}$ is the mixed empirical-ghost metric from Theorem 8.

Proof of Lemma 8. We consider data-independent prior π that is independent to both S and S' . Define

$$\Psi_{S, S'}(f) := C_0 \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f, \varepsilon))}} d\varepsilon \right\},$$

where $C_0 > 0$ is the absolute constant from Theorem 8. Finally, set

$$X(S, S') := \sup_{f \in \mathcal{F}} ((\mathbb{P}_{S'} - \mathbb{P}_S) \ell(f; z) - \Psi_{S, S'}(f))$$

and

$$Z(S) := \sup_{f \in \mathcal{F}} ((\mathbb{P} - \mathbb{P}_S) \ell(f; z) - \mathbb{E}_{S'} [\Psi_{S, S'}(f) \mid S]).$$

Core Poof Idea. We briefly overview the essential proof idea. The quantity $X(S, S')$ is the one-sided *ghost-sample fluctuation* controlled by Theorem 8, while $Z(S)$ is the one-sided fluctuation we ultimately want. The key point is that, conditional on S , the population mean $\mathbb{P}$ is exactly the conditional mean of the ghost empirical mean $\mathbb{P}_{S'}$. Therefore $Z(S)$ is dominated by the conditional expectation of $X(S, S')$:

$$Z(S) \leq \mathbb{E}_{S'} [X(S, S') \mid S].$$

Since we only require a one-sided upper-tail bound, we may transfer the ghost-sample tail bound to $Z(S)$ by applying Jensen's inequality to the exponential map, namely, to the moment generating function (MGF). This direct tail-to-MGF argument avoids both the unnecessary intermediate tail-to-expectation conversion and a separate McDiarmid step, which is designed for two-sided bounded-differences control and could lead to a worse-rate bound here unless additional techniques are introduced.

Step 1: Comparison between the Target Process and the Ghost Process. For each fixed $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_S)\ell(f; z) = \mathbb{E}_{S'}[(\mathbb{P}_{S'} - \mathbb{P}_S)\ell(f; z) \mid S].$$

Hence

$$(\mathbb{P} - \mathbb{P}_S)\ell(f; z) - \mathbb{E}_{S'}[\Psi_{S,S'}(f) \mid S] = \mathbb{E}_{S'}[(\mathbb{P}_{S'} - \mathbb{P}_S)\ell(f; z) - \Psi_{S,S'}(f) \mid S].$$

Taking the supremum over $f \in \mathcal{F}$ and using the generic comparison

$$\sup_{f \in \mathcal{F}} \mathbb{E}_{S'}[A_f \mid S] \leq \mathbb{E}_{S'} \left[\sup_{f \in \mathcal{F}} A_f \mid S \right],$$

we obtain

$$Z(S) \leq \mathbb{E}_{S'}[X(S, S') \mid S]. \quad (\text{B.31})$$

Step 2: Ghost-Sample Upper Tail. Because the prior π is independent of both S and S' , Theorem 8 yields that for every $\delta \in (0, 1)$,

$$\Pr \left(X(S, S') > C_0 \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right) \leq \delta.$$

Then for every $t \geq 0$,

$$\Pr \left(X(S, S') - C_0 \sqrt{\frac{\log(2 \log(2n))}{n}} > t \right) \leq \Pr \left(X(S, S') > \sqrt{C_0^2 \frac{\log(2 \log(2n))}{n}} + t^2 \right) \leq \exp \left(-\frac{nt^2}{C_0^2} \right).$$

Therefore, with

$$Y := (X(S, S') - b_n)_+, \quad b_n := C_0 \sqrt{\frac{\log(2 \log(2n))}{n}},$$

we have the sub-Gaussian tail bound

$$\Pr(Y > t) \leq \exp \left(-\frac{nt^2}{C_0^2} \right), \quad t \geq 0. \quad (\text{B.32})$$

Step 3: Exponential-Moment Transfer. Let $\lambda > 0$. Since $x \mapsto e^{\lambda x}$ is convex, Jensen's inequality gives

$$e^{\lambda Z(S)} \leq e^{\lambda \mathbb{E}_{S'}[X(S, S') \mid S]} \leq \mathbb{E}_{S'}[e^{\lambda X(S, S')} \mid S],$$

where the first inequality is due to (B.31). Taking expectation over S ,

$$\mathbb{E}[e^{\lambda Z(S)}] \leq \mathbb{E}[e^{\lambda X(S, S')}].$$

Since $X(S, S') \leq b_n + Y$, it follows that

$$\mathbb{E}[e^{\lambda Z(S)}] \leq e^{\lambda b_n} \mathbb{E}[e^{\lambda Y}].$$

We now bound the moment generating function of Y . Using the tail-integral formula for a nonnegative random variable,

$$\mathbb{E}[e^{\lambda Y}] = 1 + \lambda \int_0^\infty e^{\lambda t} \Pr(Y > t) dt. \quad (\text{B.33})$$

Integrating the tail bound (B.32) on Y to (B.33), we have the moment generating function bound

$$\mathbb{E}[e^{\lambda Y}] \leq 1 + \lambda \int_0^\infty \exp\left(\lambda t - \frac{nt^2}{C_0^2}\right) dt.$$

Completing the square,

$$\lambda t - \frac{nt^2}{C_0^2} = -\frac{n}{C_0^2} \left(t - \frac{C_0^2 \lambda}{2n}\right)^2 + \frac{C_0^2 \lambda^2}{4n},$$

so

$$\mathbb{E}[e^{\lambda Y}] \leq 1 + \lambda e^{C_0^2 \lambda^2 / (4n)} \int_{-\infty}^\infty \exp\left(-\frac{ns^2}{C_0^2}\right) ds.$$

The Gaussian integral equals $C_0 \sqrt{\pi/n}$, and therefore

$$\mathbb{E}[e^{\lambda Y}] \leq 1 + C_1 \frac{\lambda}{\sqrt{n}} \exp\left(\frac{C_0^2 \lambda^2}{4n}\right)$$

for some absolute constant $C_1 > 0$. Using the elementary bound

$$x \leq e^{x^2/2}, \quad x \geq 0,$$

with $x = \lambda/\sqrt{n}$, we obtain

$$\frac{\lambda}{\sqrt{n}} \leq \exp\left(\frac{\lambda^2}{2n}\right).$$

Hence

$$\mathbb{E}[e^{\lambda Y}] \leq 1 + C_1 \exp\left(C_2 \frac{\lambda^2}{n}\right) \leq C_3 \exp\left(C_2 \frac{\lambda^2}{n}\right)$$

for some absolute constants $C_2, C_3 > 0$. Consequently,

$$\mathbb{E}[e^{\lambda Z(S)}] \leq C_3 \exp\left(\lambda b_n + C_2 \frac{\lambda^2}{n}\right).$$

Step 4: Chernoff Bound to Optimize λ . For any $t > 0$ and any $\lambda > 0$,

$$\Pr(Z(S) > b_n + t) \leq e^{-\lambda(b_n + t)} \mathbb{E}[e^{\lambda Z(S)}] \leq C_3 \exp\left(-\lambda t + C_2 \frac{\lambda^2}{n}\right).$$

Optimizing at

$$\lambda = \frac{nt}{2C_2},$$

we obtain

$$\Pr(Z(S) > b_n + t) \leq C_3 \exp\left(-\frac{nt^2}{4C_2}\right).$$

Thus, for every $\delta \in (0, 1)$, choosing

$$t = \sqrt{\frac{4C_2}{n} \log \frac{C_3}{\delta}},$$

we get

$$\Pr\left(Z(S) > b_n + \sqrt{\frac{4C_2}{n} \log \frac{C_3}{\delta}}\right) \leq \delta.$$

Finally, since

$$b_n = C_0 \sqrt{\frac{\log(2 \log(2n))}{n}},$$

there exists an absolute constant $C > 0$ such that

$$b_n + \sqrt{\frac{4C_2}{n} \log \frac{C_3}{\delta}} \leq C \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}}.$$

Therefore, with probability at least $1 - \delta$,

$$Z(S) = \sup_{f \in \mathcal{F}} \left((\mathbb{P} - \mathbb{P}_S) \ell_f - \mathbb{E}_{S'} [\Psi_{S, S'}(f) \mid S] \right) \leq C \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}}.$$

Equivalently, uniformly over all $f \in \mathcal{F}$,

$$(\mathbb{P} - \mathbb{P}_n) \ell(f; z) \leq \mathbb{E}_{S'} \left[C_0 \inf_{\alpha \geq 0} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho(S, S'), \ell}(f, \varepsilon))}} d\varepsilon \right\} \mid S \right] + C \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}}.$$

Hence we finish the proof of Lemma 8. □

Lemma 9 (Averaging a random metric) *Let $n \geq 1$, $D > 0$, and let ϱ_ω be a jointly measurable random semi-metric on $\mathcal{F}$ with diameter at most D almost surely. Fix a prior π on $\mathcal{F}$ independent of ω , and set*

$$\bar{\varrho}(f, g) := (\mathbb{E}_\omega \varrho_\omega(f, g)^2)^{1/2}.$$

For any semi-metric ϱ of diameter at most D , define

$$I_{\pi, \varrho}(f) := \inf_{0 \leq \alpha \leq D} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^D \sqrt{\log \frac{1}{\pi(B_{\varrho}(f, r))}} dr \right\}.$$

Then, for every $f \in \mathcal{F}$,

$$\mathbb{E}_\omega I_{\pi, \varrho_\omega}(f) \leq \sqrt{2} I_{\pi, \bar{\varrho}}(f) + D \sqrt{\frac{\log 2}{n}}. \tag{B.34}$$

Proof. Fix f and define the prior quantile of its distance by

$$Q_\varrho(u) := \inf\{r \in [0, D] : \pi(B_\varrho(f, r)) \geq u\}, \quad 0 < u \leq 1.$$

The integrand in $I_{\pi, \varrho}(f)$ is nonincreasing in the radius. Optimizing the cutoff and then integrating its level sets therefore give

$$\begin{aligned} I_{\pi, \varrho}(f) &= \int_0^D \min\left\{1, \sqrt{\frac{\log(1/\pi(B_\varrho(f, r)))}{n}}\right\} dr \\ &= \int_0^1 Q_\varrho(e^{-nt^2}) dt. \end{aligned} \tag{B.35}$$

For $0 < u < 1/2$, set $r_u = Q_{\bar{\varrho}}(2u)$, $G = \{g : \bar{\varrho}(f, g) \leq r_u\}$, and $p = \pi(G) \geq 2u$. Markov's inequality on G gives

$$Q_{\varrho_\omega}(u)^2 \leq \frac{1}{p-u} \int_G \varrho_\omega(f, g)^2 \pi(dg).$$

Taking expectation and using Cauchy-Schwarz yields

$$\mathbb{E}_\omega Q_{\varrho_\omega}(u) \leq \left(\frac{1}{p-u} \int_G \bar{\varrho}(f, g)^2 \pi(dg) \right)^{1/2} \leq \sqrt{\frac{p}{p-u}} r_u \leq \sqrt{2} Q_{\bar{\varrho}}(2u).$$

Let $t_0 = \sqrt{\log 2/n} < 1$. Splitting (B.35) at t_0 and using $Q_{\varrho_\omega} \leq D$, we obtain

$$\begin{aligned} \mathbb{E}_\omega I_{\pi, \varrho_\omega}(f) &\leq Dt_0 + \sqrt{2} \int_{t_0}^1 Q_{\bar{\varrho}}(2e^{-nt^2}) dt \\ &\leq Dt_0 + \sqrt{2} \int_0^1 Q_{\bar{\varrho}}(e^{-ns^2}) ds \\ &= D\sqrt{\frac{\log 2}{n}} + \sqrt{2} I_{\pi, \bar{\varrho}}(f). \end{aligned}$$

The second inequality uses $s = \sqrt{t^2 - t_0^2}$, for which $2e^{-nt^2} = e^{-ns^2}$ and $dt/ds \leq 1$. $\square$

Proof of Theorem 2. Condition on S and apply Lemma 9 with $\omega = S'$, $D = \sqrt{2}$, and $\varrho_\omega = \varrho_{(S, S'), \ell}$. Independence of S' and S gives, for every $f, g \in \mathcal{F}$,

$$\mathbb{E}_{S'} [\varrho_{(S, S'), \ell}(f, g)^2 \mid S] = (\mathbb{P}_S + \mathbb{P})(\ell(f; z) - \ell(g; z))^2 = \varrho_{(S, \mathbb{P}), \ell}(f, g)^2.$$

Consequently, for every S and f ,

$$\mathbb{E}_{S'} [I_{\pi, \varrho_{(S, S'), \ell}}(f) \mid S] \leq \sqrt{2} I_{\pi, \varrho_{(S, \mathbb{P}), \ell}}(f) + \sqrt{\frac{2 \log 2}{n}}.$$

Substitute this deterministic comparison into Lemma 8. Since $\log(2 \log(2n)/\delta) \geq \log(2 \log 2) > 0$ for $n \geq 1$ and $\delta \in (0, 1)$, the additional $n^{-1/2}$ term is absorbed into the confidence term by enlarging an absolute constant. Hence we finish the proof of Theorem 2. $\square$

B.6 Proof of Theorem 3

Throughout this proof, $A_\pi(f)$ is the infimum in (2.6), and $C_\pi^\rho(f)$ is the infimum in Theorem 2 with metric ρ ; thus $C_\pi(f) = C_\pi^{\varrho(S, \mathbb{P}), \ell}(f)$. All priors are on the displayed loss class. Infima over priors are evaluated samplewise; the priors used for the upper comparisons below are fixed before sampling.

One-shot gap-distance lower bound. For the one-shot comparison, take the truncation level equal to the diameter in the gap distance, so every ball has prior mass one. This gives

$$\sup_f A_\pi(f) \leq \sup_{f, g} |(\mathbb{P} - \mathbb{P}_n)(\ell(f; z) - \ell(g; z))| \leq 2 \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)|,$$

proving (2.9) for every sample and prior.

Chaining integral lower bound. Lemma 14, followed by subtracting the Gaussian contribution of the means, gives

$$\mathbb{E} \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| \geq \frac{c_1}{\sqrt{\log n}} \left\{ \mathbb{E}_{S, g} \sup_f \frac{1}{n} \sum_{i=1}^n g_i \ell(f; z_i) - \sqrt{\frac{2}{\pi n}} \sup_f \mathbb{P} \ell(f; z) \right\}. \quad (\text{B.36})$$

Conditional on S , the process $n^{-1} \sum_i g_i \ell(f; z_i)$ has canonical metric $\varrho_{n, \ell}/\sqrt{n}$. Lemma 11 therefore yields

$$\mathbb{E}_{S, g} \sup_f \frac{1}{n} \sum_{i=1}^n g_i \ell(f; z_i) \geq \frac{c_2}{\sqrt{n}} \mathbb{E} \inf_\pi \sup_f \int_0^1 \sqrt{\log \frac{1}{\pi(B_{\varrho_{n, \ell}}(f, \varepsilon))}} d\varepsilon. \quad (\text{B.37})$$

Combining these inequalities proves the empirical lower bound. Complementation makes the absolute and one-sided suprema equal.

A family on which both bounds are sharp. Let $\mathbb{P}$ be uniform on $\{1, \dots, d\}$, where $2 \leq d \leq n$, and take all losses $\ell(f_A; j) = \mathbf{1}_{\{j \in A\}}$, $A \subseteq \{1, \dots, d\}$. Let U be uniform on these 2^d hypotheses, $N_j = \sum_i \mathbf{1}_{\{z_i = j\}}$, and $\lambda = n/d$. Put $a_j = 1/d - N_j/n$ and $Q^2 = \sum_j a_j^2$. Then

$$\begin{aligned} \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| &= \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| = \frac{1}{2} \sum_j |a_j|, \\ \mathbb{E} Q^2 &= \frac{1 - 1/d}{n}, \quad \mathbb{E} \left(\sum_j |a_j| \right)^2 \leq \frac{d}{n}. \end{aligned}$$

For $Y = N_1 - \lambda$, we have $\mathbb{E} Y^2 \geq \lambda/2$ and $\mathbb{E} Y^4 \leq 3\lambda^2 + \lambda \leq 4\lambda^2$. Hölder's inequality gives $\mathbb{E}|Y| \geq (\mathbb{E} Y^2)^{3/2}/(\mathbb{E} Y^4)^{1/2}$; hence

$$\frac{1}{8\sqrt{2}} \sqrt{\frac{d}{n}} \leq \mathbb{E} \sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| = \frac{d}{2n} \mathbb{E}|Y| \leq \frac{1}{2} \sqrt{\frac{d}{n}}. \quad (\text{B.38})$$

The untruncated comparison follows directly from Gaussian widths. For weights $w_j \geq 0$, the metric $\rho_w(f_A, f_B)^2 = \sum_{j \in A \Delta B} w_j$ is induced by $X_A = \sum_{j \in A} \sqrt{w_j} g_j$, and $\mathbb{E}_g \sup_A X_A =$

$(2\pi)^{-1/2} \sum_j \sqrt{w_j}$. The optimized majorizing-measure integral therefore lies between constant multiples of $\sum_j \sqrt{w_j}$. The empirical and mixed weights are N_j/n and $N_j/n + 1/d$, respectively. The mixed sum lies between $\sqrt{d}$ and $\sqrt{2d}$. Since $\mathbb{E}N_j^2 \leq 2\lambda^2$, Paley–Zygmund gives $\mathbb{P}\{N_j \geq \lambda/2\} \geq 1/8$. Jensen’s inequality supplies the upper bound, proving

$$c\sqrt{d} \leq \mathbb{E} \inf_{\pi} \sup_f \int_0^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\rho}(f, \varepsilon))}} d\varepsilon \leq C\sqrt{d} \quad (\text{B.39})$$

for either $\rho = \varrho_{n,\ell}$ or $\rho = \varrho_{(S,\mathbb{P}),\ell}$.

The individually truncated chaining objective. The following ball argument also applies to the rare-atom construction below. Suppose k coordinates of w are at least $b > 0$ and set $r = \sqrt{bk}/2$. A ball of radius r differs from its center on at most $k/4$ of those coordinates, so Hoeffding’s inequality gives $U(B_{\rho_w}(f, r)) \leq e^{-k/8}$ for every f . For any prior π , symmetry gives

$$\mathbb{E}_{f \sim U} \pi(B_{\rho_w}(f, r)) = \mathbb{E}_{g \sim \pi} U(B_{\rho_w}(g, r)) \leq e^{-k/8}.$$

Thus a single center f has ball mass at most $e^{-k/8}$. At this center, for every α , the chaining objective is at least $\alpha + (r - \alpha)_+ \sqrt{k/(8n)}$. Whenever $r \leq \sqrt{2}$, this proves

$$\inf_{\pi} \sup_f C_{\pi}^{\rho_w}(f) \geq \frac{\sqrt{bk}}{2} \min \left\{ 1, \sqrt{\frac{k}{8n}} \right\}. \quad (\text{B.40})$$

The center is fixed before optimizing α at that center; no interchange of infimum and supremum is involved. For the mixed metric, take $k = d$, $b = 1/d$. For the empirical metric, take $K = \#\{j : N_j \geq \lambda/2\}$, $k = K$, $b = 1/(2d)$; then $\inf_{\pi} \sup_f C_{\pi}^{\varrho_{n,\ell}}(f) \geq K/(8\sqrt{dn})$ and $\mathbb{E}K \geq d/8$. The case $K = 0$ gives the trivial lower bound zero. Conversely, singleton mass 2^{-d} and $\alpha = 0$ give $\sup_f C_U^{\rho}(f) \leq \sqrt{2d \log 2/n}$ for either metric. Together with $\varrho_{n,\ell} \leq \varrho_{(S,\mathbb{P}),\ell}$, this gives

$$c\sqrt{\frac{d}{n}} \leq \mathbb{E} \inf_{\pi} \sup_f C_{\pi}^{\varrho_{n,\ell}}(f) \leq \mathbb{E} \inf_{\pi} \sup_f C_{\pi}(f) \leq \mathbb{E} \sup_f C_U(f) \leq C\sqrt{\frac{d}{n}}. \quad (\text{B.41})$$

The gap-distance one-shot objective. Let f_* maximize $(\mathbb{P} - \mathbb{P}_n)\ell(f; z)$. Under U , its gap distance to a random hypothesis has the distribution $\sum_j |a_j| B_j$, where the B_j are independent Bernoulli variables of mean $1/2$. If $Q > 0$, Hoeffding’s inequality gives

$$U(B_{\varrho_{(S,\mathbb{P}),\ell}^{\text{gap}}}(f_*, \varepsilon)) \leq \exp \left(-\frac{(\sum_j |a_j|)^2}{8Q^2} \right) \quad \text{for } \varepsilon \leq \frac{1}{4} \sum_j |a_j|.$$

Splitting the optimization at $\frac{1}{4} \sum_j |a_j|$ yields

$$\sup_f A_U(f) \geq A_U(f_*) \geq \frac{\sum_j |a_j|}{4(1 + \sqrt{n}Q)}.$$

This also holds when $Q = 0$. Cauchy–Schwarz and the moment bounds above now give

$$\mathbb{E} \sup_f A_U(f) \geq \frac{(\mathbb{E} \sum_j |a_j|)^2}{4\{\mathbb{E} \sum_j |a_j| + \sqrt{n} \mathbb{E}[Q \sum_j |a_j|]\}} \geq c\sqrt{\frac{d}{n}}.$$

The reverse bound follows from singleton mass: $\sup_f A_U(f) \leq \sqrt{d \log 2/n}$. This proves sharpness of both complexity terms under the same sample-independent prior.

Rare atoms separate the two bounds. Give a background atom probability $1 - p$ and each of $d \geq 8n$ other atoms probability p/d . Retain all subset indicators on the rare atoms and set every loss to zero at the background atom. Let K count rare observations, so $\mathbb{E}K = np$, and keep the uniform prior U . The indicator of the unobserved rare atoms has gap at least $7p/8$, while $\sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| \leq p + K/n$. Thus $7p/8 \leq \mathbb{E}\sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)| \leq 2p$. The mixed weights are $N_j/n + p/d$; (B.40), with $k = d$ and $b = p/d$, gives $\inf_\pi \sup_f C_\pi(f) \geq \sqrt{p}/2$. Choosing the truncation level equal to the mixed diameter gives $\sup_f C_U(f) \leq \sqrt{p + K/n}$, whose expectation is at most $\sqrt{2p}$.

The diameter in the gap distance is at most $p + K/n$, so the same radius choice as in (2.9) gives $\mathbb{E}\sup_f A_U(f) \leq 2p$. For the reverse bound, let f_* indicate the unobserved rare atoms. A gap ball around f_* of radius at most $p/8$ contains only hypotheses with gap at least $3p/4$, which must include at least $3d/4$ unobserved rare atoms, since observed coordinates have negative gap coefficients $p/d - N_j/n$. Its U -mass is at most $e^{-d/8}$ by Hoeffding's inequality. Thus $A_U(f_*) \geq \min\{p/8, \sqrt{d/(8n)}\} = p/8$. This proves (2.11) and shows that truncation does not restore a universal mixed-metric lower bound.

Scope of tightness. The one-shot comparison (2.9) depends on the unknown population gaps; it is not an observable guarantee. Neither objective must match the gap at each hypothesis. In the first construction, $(\mathbb{P} - \mathbb{P}_n)\ell(f_\emptyset; z) = 0$, while $C_U(f_\emptyset) > 0$ for every sample. For $d = 2$,

$$A_U(f_\emptyset) = \min \left\{ \left| \frac{N_1}{n} - \frac{1}{2} \right|, \sqrt{\frac{\log 2}{n}} \right\},$$

which is positive whenever the counts are unequal. Finally, truncation and nontruncation are not universally equivalent for the mixed metric: for the same indicator class with n fixed and unrestricted d , the optimized untruncated mixed integral divided by $\sqrt{n}$ is at least $c\sqrt{d/n}$, whereas every truncated objective is at most $\sqrt{2}$. The sharpness assertions concern the complexity terms. At fixed confidence, the expected full upper bounds have the same rate as $\mathbb{E}\sup_f |(\mathbb{P} - \mathbb{P}_n)\ell(f; z)|$ on the first family when $d \geq \log(2 \log(2n)/\delta)$. In the second family, taking $p = n^{-1/4}$ and $d = 8n$ makes the confidence remainder negligible: the expected one-shot bound and deviation have order $n^{-1/4}$, while the expected chaining bound has order $n^{-1/8}$. $\square$

B.7 Background on Gaussian and Empirical Processes

It is now well understood that the supremum of Gaussian process can be tightly characterized by the majorizing measure integral via matching upper and lower bounds up to absolute constants [Fernique, 1974, Talagrand, 1987]; the goal of this section is to extend this characterization to (1) bounded empirical processes and (2) a truncated form of integral.

Background on Gaussian Processes. We begin by recalling several key results from a series of seminal papers by Talagrand, Fernique, and others, which introduced the majorizing-measure formulation of the generic chaining framework [Fernique, 1974, Talagrand, 1987]. Note that generic chaining have several equivalent formulations [Talagrand, 2014], and the one closest to our purpose is through majorizing measure.

A *centered Gaussian random variable* X is a real-valued measurable function on the outcome space such that the law of X has density

$$(2\pi\sigma^2)^{-1/2} \exp\left(-\frac{x^2}{2\sigma^2}\right).$$

The law of X is thus determined by $\sigma = (\mathbb{E}[X^2])^{1/2}$. If $\sigma = 1$, X is called *standard normal*.

A *Gaussian process* is a family $\{X_t\}_{t \in T}$ of random variables indexed by some set T , such that every finite linear combination $\sum_{j=1}^k \alpha_j X_{t_j}$ is Gaussian. On the index set T , consider the semi-metric ϱ given by

$$\varrho(u, v) = \sqrt{\mathbb{E}[(X_u - X_v)^2]}. \quad (\text{B.42})$$

Gaussian processes are thus a very rigid class of stochastic processes, with exceptionally nice properties that have been fully developed in the literature.

[Fernique \[1974\]](#) proved the following integral upper bound.

Lemma 10 (Upper Bound of Gaussian Processes via Majorizing Measure, [Fernique \[1974\]](#))

Given a Gaussian process $(X_t)_{t \in T}$ with its metric ϱ defined by (B.42), we have

$$\mathbb{E} \left[\sup_{t \in T} X_t \right] \leq C \inf_{\pi \in \Delta(T)} \sup_{t \in T} \int_0^\infty \sqrt{\log \frac{1}{\pi(B_\varrho(t, \varepsilon))}} d\varepsilon,$$

where $C > 0$ is an absolute constant.

A prior π that makes the right hand side in Lemma 10 finite is called a *majorizing measure*. Fernique conjectured as early as 1974 that the existence of majorizing measures might characterize the boundedness of Gaussian processes. He proved a number of important partial results, and his determination eventually motivated the Talagrand to attack the problem in 1987. [Talagrand \[1987\]](#) proved that the integral in Lemma 10 is tight up to absolute constants. The pointwise integral on the right-hand side of Lemma 10 is known as the Fernique-Talagrand (majorizing measure) integral; see [Talagrand \[2014, Section 6.2\]](#) for a standard treatment and [Chu and Raginsky \[2023\]](#) for a recent reference. A useful variance-adaptive refinement, whose integration range depends on the pointwise standard deviation $\sigma(t)$, is developed in [Xu \[2026\]](#).

Lemma 11 (Lower Bound of Gaussian Processes via Majorizing Measure, [Talagrand \[1987\]](#))

Given a Gaussian process $(X_t)_{t \in T}$ with its metric ϱ defined by (B.42), we have

$$\mathbb{E} \left[\sup_{t \in T} X_t \right] \geq c \inf_{\pi \in \Delta(T)} \sup_{t \in T} \int_0^\infty \sqrt{\log \frac{1}{\pi(B_\varrho(t, \varepsilon))}} d\varepsilon,$$

where $c > 0$ is an absolute constant.

Thus, the globally optimized Fernique-Talagrand integral, with the infimum over measures and the supremum over index points, characterizes the expected supremum of a Gaussian process up to universal constants.

Background on Empirical Processes. We now give several results on upper and lower bounding empirical process by Rademacher and Gaussian complexities [Giné and Zinn \[1984\]](#), [Bartlett and Mendelson \[2002\]](#).

Definition 4 (Rademacher and Gaussian complexities) *For a function class $\mathcal{F}$ that consists of mappings from $\mathcal{Z}$ to $\mathbb{R}$, define the Rademacher complexity of $\mathcal{F}$ as*

$$R_n(\mathcal{F}) := \mathbb{E}_{z, \xi} \left[\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \xi_i f(z_i) \right],$$

where $\{\xi_i\}_{i=1}^n$ are i.i.d. Rademacher variables; and define the Gaussian complexity of $\mathcal{F}$ as

$$G_n(\mathcal{F}) := \mathbb{E}_{z, g} \left[\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n g_i f(z_i) \right],$$

where $\{g_i\}_{i=1}^n$ are i.i.d. standard Gaussian variables.

It is well-known that Rademacher and Gaussian complexities are upper bounds of empirical processes (see, e.g., Lemma 7.4 in [Van Handel \[2014\]](#)):

Lemma 12 (Upper Bounds with Rademacher and Gaussian Complexities) *For any function class $\mathcal{F}$ that consists of mappings from $\mathcal{Z}$ to $\mathbb{R}$, we have*

$$\mathbb{E} \left[\sup_{f \in \mathcal{F}} (\mathbb{P} - \mathbb{P}_n) f(z) \right] \leq 2R_n(\mathcal{F}) \leq \sqrt{2\pi} G_n(\mathcal{F}),$$

where $R_n(\mathcal{F})$ and $G_n(\mathcal{F})$ are (expected) Rademacher and Gaussian complexities defined in Definition 4.

The following lemma strengthens Theorem 3 of [Block et al. \[2021\]](#) in the batch setting by allowing pointwise optimization of the truncation level inside the class-wide supremum. More broadly, this strengthening introduces a new type of truncated chaining or covering integrals for bounding suprema of empirical processes, with the truncation level optimized *separately* at each index. It also serves as a supporting step in the proof of Theorem 8.

Lemma 13 (Integral bound with separately optimized truncation level) *Let $\mathcal{F}$ be a nonempty separable class of functions from $\mathcal{Z}$ to $[0, 1]$, and fix $z_1, \dots, z_n \in \mathcal{Z}$. Define the empirical semi-metric*

$$\varrho_n(f_1, f_2) := \sqrt{\frac{1}{n} \sum_{i=1}^n (f_1(z_i) - f_2(z_i))^2}.$$

There exists an absolute constant $C > 0$ such that

$$\mathbb{E}_{\xi} \left[\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \xi_i f(z_i) \right] \leq C \inf_{\pi \in \Delta(\mathcal{F})} \sup_{f \in \mathcal{F}} \inf_{0 \leq \alpha \leq 1} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^1 \sqrt{\log \frac{1}{\pi(B_{\varrho_n}(f, \varepsilon))}} d\varepsilon \right\},$$

where $\{\xi_i\}_{i=1}^n$ are independent Rademacher variables.

Proof. Write $J_\pi(f, \alpha)$ for the expression in braces. Theorem 3 of Block et al. [2021], applied to trees whose values at each depth are independent of the path, gives

$$\mathbb{E}_\xi \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \xi_i f(z_i) \leq C' \inf_{\pi \in \Delta(\mathcal{F})} \inf_{0 \leq \alpha \leq 1} \sup_{f \in \mathcal{F}} J_\pi(f, \alpha),$$

where we have interchanged only the two infima. Fix π and put $M := \sup_f \inf_{0 \leq \alpha \leq 1} J_\pi(f, \alpha) \leq 1$. For $\eta > 0$, choose for each f a cutoff α_f such that $J_\pi(f, \alpha_f) \leq M + \eta$. Nonnegativity of the integral gives $\alpha_f \leq M + \eta$. Hence $a := \min\{1, M + \eta\} \geq \alpha_f$, and

$$J_\pi(f, a) \leq a + J_\pi(f, \alpha_f) - \alpha_f \leq 2(M + \eta).$$

Taking the supremum over f and letting $\eta \downarrow 0$ yields $\inf_\alpha \sup_f J_\pi(f, \alpha) \leq 2 \sup_f \inf_\alpha J_\pi(f, \alpha)$. The result follows by taking the infimum over π , with $C = 2C'$. $\square$

By translation and rescaling, the same bound holds for functions taking values in an interval of length $D > 0$, with the upper endpoint 1 replaced by D and $0 \leq \alpha \leq D$. In particular, $D = 2$ applies to the symmetrized loss in Theorem 8.

The following result illustrate that Gaussian and Rademacher complexities can also be used to lower bounding empirical processes.

Lemma 14 (Lower Bounds with Rademacher and Gaussian Complexities) *Let $n \geq 2$ and let $\mathcal{F}$ be a class of integrable real-valued functions on $\mathcal{Z}$, with centered class $\tilde{\mathcal{F}} = \{f - \mathbb{P}f : f \in \mathcal{F}\}$. Then*

$$\mathbb{E} \sup_{f \in \mathcal{F}} |(\mathbb{P} - \mathbb{P}_n)f| \geq \frac{1}{2} R_n(\tilde{\mathcal{F}}) \geq \frac{c}{\sqrt{\log n}} G_n(\tilde{\mathcal{F}}),$$

where $c > 0$ is an absolute constant. If $\tilde{\mathcal{F}}$ is symmetric, the absolute value on the left can be omitted.

Proof of Lemma 14. Let S' be an independent copy of S . Jensen's inequality and swapping each pair (Z_i, Z'_i) when $\xi_i = -1$ give

$$\begin{aligned} R_n(\tilde{\mathcal{F}}) &\leq \mathbb{E}_{S, S', \xi} \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \xi_i (f(Z_i) - f(Z'_i)) \\ &= \mathbb{E}_{S, S'} \sup_{f \in \mathcal{F}} (\mathbb{P}_S - \mathbb{P}_{S'})f \leq 2 \mathbb{E}_S \sup_{f \in \mathcal{F}} |(\mathbb{P}_S - \mathbb{P})f|. \end{aligned}$$

The Gaussian–Rademacher comparison gives the second inequality; see Problem 7.1 in Van Handel [2014]. Symmetry of the centered class makes the absolute and one-sided suprema equal. $\square$

Background on covering numbers Formally, we give the definition of covering number as follows.

Definition 5 (Covering numbers) *Let $(\mathcal{Y}, \varrho)$ be a metric space and let $\mathcal{Z} \subseteq \mathcal{Y}$. For $\varepsilon > 0$, a set $\mathcal{N} \subseteq \mathcal{Z}$ is an internal ε -cover of $\mathcal{Z}$ if for every $z \in \mathcal{Z}$ there exists $y \in \mathcal{N} \subseteq \mathcal{Z}$ with $\varrho(z, y) \leq \varepsilon$. The (internal) covering number is*

$$N(\mathcal{Z}, \varrho, \varepsilon) := \min\{m : \exists \text{ internal } \varepsilon\text{-cover of } \mathcal{Z} \text{ with size } m\}.$$

A set $\mathcal{N}_{\text{ext}} \subseteq \mathcal{Y}$ (not necessarily inside $\mathcal{Z}$) is an external ε -cover of $\mathcal{Z}$ if for every $z \in \mathcal{Z}$ there exists $y \in \mathcal{N}_{\text{ext}}$ with $\varrho(z, y) \leq \varepsilon$. The external covering number is

$$\mathbf{N}_{\text{ext}}(\mathcal{Z}, \varrho, \varepsilon) := \min\{m : \exists \text{ external } \varepsilon\text{-cover of } \mathcal{Z} \text{ with size } m\}.$$

Internal covering numbers depend only on the metric induced on $\mathcal{Z}$, while external covering numbers also depend on the ambient space $\mathcal{Y}$. Throughout the paper, “covering number” means the internal one unless otherwise stated.

We now relate the internal and external covering numbers, showing they are equivalent up to a constant factor in the radius—and thus interchangeable for our purposes.

Lemma 15 (Properties of External Covering Number) *For every $\varepsilon > 0$ and $\mathcal{Z} \subseteq \mathcal{Y}$,*

$$\mathbf{N}_{\text{ext}}(\mathcal{Z}, \varrho, \varepsilon) \leq \mathbf{N}(\mathcal{Z}, \varrho, \varepsilon) \leq \mathbf{N}_{\text{ext}}(\mathcal{Z}, \varrho, \varepsilon/2). \quad (\text{B.43})$$

And the external covering number enjoys monotonicity under set inclusion: if $\mathcal{Z}_1 \subseteq \mathcal{Z}_2$ then $\mathbf{N}_{\text{ext}}(\mathcal{Z}_1, \varrho, \varepsilon) \leq \mathbf{N}_{\text{ext}}(\mathcal{Z}_2, \varrho, \varepsilon)$.

Proof of Lemma 15: The left inequality in (B.43) is immediate since any internal cover is also an external cover. For the right inequality in (B.43), let $\{y_1, \dots, y_m\} \subseteq \mathcal{Y}$ be an external $(\varepsilon/2)$ -cover of $\mathcal{Z}$. For each i , define the (possibly empty) cell $V_i := \{z \in \mathcal{Z} : \varrho(z, y_i) \leq \varepsilon/2\}$. By the very definition of external $(\varepsilon/2)$ -cover, every $z \in \mathcal{Z}$ is within distance $\varepsilon/2$ of some y_i ; hence

$$\bigcup_{i=1}^m V_i = \mathcal{Z}.$$

If $V_i \neq \emptyset$, pick a representative $z_i \in V_i$. Then for any $z \in V_i$,

$$\varrho(z, z_i) \leq \varrho(z, y_i) + \varrho(y_i, z_i) \leq \varepsilon/2 + \varepsilon/2 = \varepsilon,$$

so the selected $\{z_i\} \subseteq \mathcal{Z}$ form an internal ε -cover. Hence $\mathbf{N}(\mathcal{Z}, \varrho, \varepsilon) \leq m = \mathbf{N}_{\text{ext}}(\mathcal{Z}, \varrho, \varepsilon/2)$. Lastly, the monotonicity under set inclusion for the external covering number is a straightforward consequence of its definition. □

Basic Concentration Inequalities. We state McDiarmid’s inequality, Hoeffding’s inequality, and Bernstein’s inequality.

Lemma 16 (McDiarmid’s inequality (bounded differences), [McDiarmid \[1998\]](#)) *Let $Z_1, \dots, Z_n$ be independent random variables with $Z_i \in \mathcal{Z}_i$. Let $h : \mathcal{Z}_1 \times \dots \times \mathcal{Z}_n \rightarrow \mathbb{R}$ be a measurable function satisfying the bounded difference property: there are constants $c_1, \dots, c_n \geq 0$ such that for all $i \in \{1, \dots, n\}$ and all $Z_1 \in \mathcal{Z}_1, \dots, Z_n \in \mathcal{Z}_n$,*

$$\sup_{Z'_i \in \mathcal{Z}_i} |h(Z_1, \dots, Z_{i-1}, Z_i, Z_{i+1}, \dots, Z_n) - h(Z_1, \dots, Z_{i-1}, Z'_i, Z_{i+1}, \dots, Z_n)| \leq c_i.$$

Then for every $t \geq 0$,

$$\Pr(h(Z_1, \dots, Z_n) - \mathbb{E}[h(Z_1, \dots, Z_n)] \geq t) \leq \exp\left(-\frac{2t^2}{\sum_{i=1}^n c_i^2}\right).$$

Lemma 17 (Hoeffding’s inequality, Chapter 2 in Vershynin [2018]) *Let $Z_1, \dots, Z_n$ be independent random variables with $a_i \leq Z_i \leq b_i$ almost surely. Then for every $t \geq 0$,*

$$\Pr\left(\sum_{i=1}^n Z_i - \mathbb{E}[Z] \geq t\right) \leq \exp\left(-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right).$$

C Proofs for Deep Neural Networks and Riemannian Dimension (Section 3)

C.1 Proof of Lemma 1 (Non-Perturbative Feature Expansion)

We start with the telescoping decomposition presented in the main paper, which serves as a non-perturbative replacement of conventional Taylor expansion, where in each summand the only difference lies in W'_l and W_l .

$$\begin{aligned} & F_L(W', X) - F_L(W, X) \\ &= \sum_{l=1}^L [\underbrace{\sigma_L(W'_L \cdots W'_{l+1})}_{\text{controlled by } M_{l \rightarrow L}} \underbrace{\sigma_l(W'_l)}_{\text{by 1}} \underbrace{F_{l-1}(W, X)}_{\text{learned feature}}) - \sigma_L(W'_L \cdots W'_{l+1}) \underbrace{\sigma_l(W_l)}_{\text{by 1}} \underbrace{F_{l-1}(W, X)}_{\text{learned feature}})]. \end{aligned}$$

Applying Cauchy-Schwarz inequality to the above identity, we have

$$\begin{aligned} & \|F_L(W', X) - F_L(W, X)\|_{\mathbf{F}}^2 \\ & \leq \sum_{l=1}^L L \|\sigma_L(W'_L \cdots W'_{l+1}) \sigma_l(W'_l F_{l-1}(W, X)) - \sigma_L(W'_L \cdots W'_{l+1}) \sigma_l(W_l F_{l-1}(W, X))\|_{\mathbf{F}}^2. \end{aligned} \quad (\text{C.1})$$

By the definition of local Lipschitz constant in Section 3, for all $W' \in B_{\varrho_n}(W, \varepsilon)$,

$$\begin{aligned} & \|\sigma_L(W'_L \cdots W'_{l+1}) \sigma_l(W'_l F_{l-1}(W, X)) - \sigma_L(W'_L \cdots W'_{l+1}) \sigma_l(W_l F_{l-1}(W, X))\|_{\mathbf{F}} \\ & \leq M_{l \rightarrow L}[W, \varepsilon] \|\sigma_l(W'_l F_{l-1}(W, X)) - \sigma_l(W_l F_{l-1}(W, X))\|_{\mathbf{F}}. \end{aligned} \quad (\text{C.2})$$

Because the activation function σ_l is 1-Lipschitz for each column, we have

$$\|\sigma_l(W'_l F_{l-1}(W, X)) - \sigma_l(W_l F_{l-1}(W, X))\|_{\mathbf{F}} \leq \|(W'_l - W_l) F_{l-1}(W, X)\|_{\mathbf{F}}. \quad (\text{C.3})$$

Combining (C.1) (C.2) and (C.3), we prove that

$$\|F_L(W', X) - F_L(W, X)\|_{\mathbf{F}}^2 \leq \sum_{l=1}^L L \cdot M_{l \rightarrow L}[W, \varepsilon]^2 \cdot \|(W'_l - W_l) F_{l-1}(W, X)\|_{\mathbf{F}}^2.$$

□

C.2 Metric Domination Lemma

Our non-perturbative expansion facilitates bounding the pointwise dimension of complex geometries via metric comparison. By constructing a simpler, dominating metric (i.e., one that is pointwise larger), we establish that the pointwise dimension of the original geometry is upper bounded by that of this new, more structured geometry. This “enlargement” for analytical tractability, a concept with roots in comparison geometry and majorization principles, is operationalized in Lemma 18.

Lemma 18 (Metric Domination Lemma) *For two metrics ϱ_1, ϱ_2 defined on $\mathcal{W}$, if $\varrho_1(W', W) \leq \varrho_2(W', W)$ for all $W' \in B_{\varrho_2}(W, \varepsilon)$, then for any prior $\pi \in \Delta(\mathcal{W})$ and any $\varepsilon > 0$, we have*

$$\log \frac{1}{\pi(B_{\varrho_1}(W, \varepsilon))} \leq \log \frac{1}{\pi(B_{\varrho_2}(W, \varepsilon))}.$$

Proof of Lemma 18: Because $\varrho_1(W', W) \leq \varrho_2(W', W)$ for all $W' \in B_{\varrho_2}(W, \varepsilon)$, we have that

$$B_{\varrho_1}(W, \varepsilon) \supseteq B_{\varrho_2}(W, \varepsilon).$$

So for any prior π on $\mathbb{R}^p$, monotonicity of measures gives

$$\pi(B_{\varrho_1}(W, \varepsilon)) \geq \pi(B_{\varrho_2}(W, \varepsilon)),$$

this implies

$$\log \frac{1}{\pi(B_{\varrho_1}(W, \varepsilon))} \leq \log \frac{1}{\pi(B_{\varrho_2}(W, \varepsilon))}.$$

□

We then state an extension of the metric domination lemma, which turns pointwise dimension in a high-dimensional space into a lower-dimensional subspace.

Lemma 19 (Subspace Metric Domination Lemma) *Given a metric ϱ_1 defined on $\mathbb{R}^p$ a subspace $\mathcal{V} \subseteq \mathbb{R}^p$, and a metric ϱ_2 defined on $\mathcal{V}$. Define the orthogonal projector to subspace $\mathcal{V}$ as $\mathcal{P}_{\mathcal{V}}(W) := \arg \min_{\tilde{W} \in \mathcal{V}} \|\tilde{W} - W\|_2$. If there exists $\varepsilon_1 \in (0, \varepsilon)$ such that for every $W' \in \mathcal{V}$,*

$$(\varrho_1(W', W))^2 \leq (\varrho_2(W', \mathcal{P}_{\mathcal{V}}(W)))^2 + \varepsilon_1^2, \quad (\text{C.4})$$

then for any prior $\pi \in \Delta(\mathcal{V})$, we have

$$\log \frac{1}{\pi(B_{\varrho_1}(W, \varepsilon))} \leq \log \frac{1}{\pi(B_{\varrho_2}(\mathcal{P}_{\mathcal{V}}(W), \sqrt{\varepsilon^2 - \varepsilon_1^2}))}. \quad (\text{C.5})$$

Proof of Lemma 19: By the condition (C.4), we know

$$B_{\varrho_1}(W, \varepsilon) \supseteq B_{\varrho_1}(W, \varepsilon) \cap \mathcal{V} \supseteq B_{\varrho_2}(\mathcal{P}_{\mathcal{V}}(W), \sqrt{\varepsilon^2 - \varepsilon_1^2}),$$

and this gives the desired conclusion (C.5) in Lemma 19.

□

C.3 Pointwise Dimension Bound with Reference Subspace

Set Up of Reference Effective Subspace Consider the weight space $B_2(R) \subset \mathbb{R}^p$ for vectorized weights W , where $B_2(R) := \{w \in \mathbb{R}^p : \|w\|_2 \leq R\}$. Given any fixed $p \times p$ PSD matrix $G(W)$, order the eigenvalues $\lambda_1(G(W)), \dots, \lambda_p(G(W))$ nonincreasingly. For notational convenience, we suppress the dependence on $G(W)$ and write simply λ_k when no confusion can arise. We denote $\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon)$ to be the *effective subspace*—the true top- r_{eff} eigenspace—of $G(W)$. For notational convenience, we use r_{eff} as the abbreviation of $r_{\text{eff}}(G(W), R, \varepsilon)$, and $\mathcal{V}$ as an abbreviation of $\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon)$ when no confusion can arise.

Assume there is another r -dimensional subspace $\bar{\mathcal{V}}$. We will show that if $\bar{\mathcal{V}}$ approximates $\mathcal{V}$, then using a prior supported on $\bar{\mathcal{V}}$ still yields a valid effective-dimension bound. This observation underpins the hierarchical covering argument in Theorem 4. For a self-contained introduction to subspaces (collectively known as the Grassmannian) and their frame parameterizations (the Stiefel manifold); see Section D.1, where we translate algebraic and differential-geometric insights into machine learning terminology.

Motivation of Approximate Effective Subspace. We can view the orthogonal projector to a subspace as a matrix (see the definition via the Stiefel parameterization in (D.6)), which is consistent with the earlier operator notation characterized by ℓ_2 -distance in Lemma 19. Now define the projected metric $\varrho_{G(W)}^{\bar{\mathcal{V}}}$ as

$$\begin{aligned}\varrho_{G(W)}^{\bar{\mathcal{V}}}(W_1, W_2) &= \sqrt{(\mathcal{P}_{\bar{\mathcal{V}}}(W_1) - \mathcal{P}_{\bar{\mathcal{V}}}(W_2))^\top G(W) (\mathcal{P}_{\bar{\mathcal{V}}}(W_1) - \mathcal{P}_{\bar{\mathcal{V}}}(W_2))} \\ &= \sqrt{(W_1 - W_2)^\top \mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}}(W_1 - W_2)}.\end{aligned}$$

By the subspace metric domination lemma (Lemma 19), if $\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}}$ approximates $G(W)$, we can use prior over $\bar{\mathcal{V}}$ to bound the pointwise dimension and achieve dimension reduction.

We will require the following approximation error condition:

$$\varrho_{\text{proj}, G(W)}(\mathcal{V}, \bar{\mathcal{V}}) = \|G(W)^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}} \leq \frac{\sqrt{n}\varepsilon}{4R}.$$

In Section D, we systematically study the ellipsoidal covering of Grassmannian, and establish that we can *always* find $\bar{\mathcal{V}}$ that approximates $\mathcal{V}$ to the desired precision, with an additional covering cost of the Grassmannian bound in the Riemannian Dimension. This generalizes the canonical projection metric between subspaces into ellipsoidal set-up.

Effective Dimension Bound for Approximate Effective Subspace. We now present the lemma that establishes effective dimension bound using prior supported on approximate effective subspace $\bar{\mathcal{V}}$ (not necessarily the true effective subspace $\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon)$). We state the main result of this subsection (Lemma 2 in the main paper).

Consider the weight space $B_2(R) \subset \mathbb{R}^p$ for vectorized weights, and a pointwise ellipsoidal metric defined via PSD $G(W)$. Let $\bar{\mathcal{V}} \subseteq \mathbb{R}^p$ be a fixed r -dimensional subspace. Define the prior $\pi_{\bar{\mathcal{V}}} = \text{Unif}(B_2(1.58R) \cap \bar{\mathcal{V}})$. Then, uniformly over all (W, ε) such that the top- r eigenspace $\mathcal{V}$ of $G(W)$ can be approximated by $\bar{\mathcal{V}}$ to precision

$$\varrho_{\text{proj}, G(W)}(\mathcal{V}, \bar{\mathcal{V}}) := \|G(W)^{1/2}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}} \leq \frac{\sqrt{n}\varepsilon}{4R}, \quad (\text{C.6})$$

we have

$$\log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\varrho_G(W)}(W, \sqrt{n}\varepsilon))} \leq \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}(G(W), R, \varepsilon)} \log \left(\frac{40R^2 \lambda_k(G(W))}{n\varepsilon^2} \right) = d_{\text{eff}}(G(W), \sqrt{5}R, \varepsilon).$$

Proof of Lemma 2: Given a fixed PSD matrix $G(W)$ with eigenvalues $\lambda_1 \geq \dots \geq \lambda_p$, denote $r_{\text{eff}} = r_{\text{eff}}(G(W), R, \varepsilon)$, and the projected metric $\varrho_{G(W)}^{\bar{\mathcal{V}}}$ on $\bar{\mathcal{V}}$:

$$\varrho_{G(W)}^{\bar{\mathcal{V}}}(W_1, W_2) = \sqrt{(W_1 - W_2)^\top \mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}}(W_1 - W_2)}.$$

Since $\mathcal{V}$ is the top- r_{eff} eigenspace of $G(W)$, by the elementary property of eigendecomposition we have that

$$\begin{aligned} G(W) &= \mathcal{P}_{\mathcal{V}}^\top G(W) \mathcal{P}_{\mathcal{V}} + \mathcal{P}_{\mathcal{V}_\perp}^\top G(W) \mathcal{P}_{\mathcal{V}_\perp} \\ &\preceq \mathcal{P}_{\mathcal{V}}^\top G(W) \mathcal{P}_{\mathcal{V}} + \lambda_{r_{\text{eff}}+1} \cdot \mathcal{P}_{\mathcal{V}_\perp}^\top \mathcal{P}_{\mathcal{V}_\perp}, \end{aligned} \quad (\text{C.7})$$

where $\mathcal{V}_\perp$ is orthogonal complement of $\mathcal{V}$. It is also straightforward to see

$$\mathcal{P}_{\mathcal{V}}^\top G(W) \mathcal{P}_{\mathcal{V}} \preceq 2\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}} + 2(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})^\top G(W) (\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}}). \quad (\text{C.8})$$

Combining (C.7) and (C.8), we have the fundamental loewner order inequality

$$G(W) \preceq 2\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}} + 2(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})^\top G(W) (\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}}) + \lambda_{r_{\text{eff}}+1} \cdot \mathcal{P}_{\mathcal{V}_\perp}^\top \mathcal{P}_{\mathcal{V}_\perp}. \quad (\text{C.9})$$

In order to apply the subspace metric domination lemma (Lemma 19), we hope to bound $\|W' - W\|_2^2$ and apply that bound to the two last remainder terms in the right hand side of (C.9).

To bound $\|W' - W\|_2^2$, we firstly state the following lemma on the eigenvalue of $\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}}$, whose proof is deferred until after the current proof. Here, the metric tensor is a Riemannian-geometric way of describing a pointwise positive semidefinite, matrix-valued function $G(W)$; see also Appendix C.4. The corresponding projected metric tensor is given by $\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}}$.

Lemma 20 (Eigenvalue Bound for Projected Metric Tensor) *Assume $\mathcal{V}$ is the top- r eigenspace of a PSD matrix Σ with eigenvalues $\lambda_1 \geq \dots \geq \lambda_p$, then for a r -dimensional subspace $\bar{\mathcal{V}}$ we have that for $k = 1, 2, \dots, r$,*

$$\lambda_k \geq \lambda_k(\mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}}) \geq \lambda_k/2 - \|\Sigma^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}^2.$$

For every $W' \in B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n}\varepsilon/4)$, we have $\forall k = 1, \dots, r_{\text{eff}}$,

$$\begin{aligned} \|W' - \mathcal{P}_{\bar{\mathcal{V}}}(W)\|_2^2 &\leq \frac{(W' - \mathcal{P}_{\bar{\mathcal{V}}}(W))^\top \mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}}(W' - \mathcal{P}_{\bar{\mathcal{V}}}(W))}{\lambda_{r_{\text{eff}}}(\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}})} \leq \frac{n\varepsilon^2}{16\lambda_{r_{\text{eff}}}(\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W) \mathcal{P}_{\bar{\mathcal{V}}})} \\ &\leq \frac{n\varepsilon^2}{8\lambda_{r_{\text{eff}}} - 16\|G(W)^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}^2} \leq \frac{1}{3}R^2, \end{aligned} \quad (\text{C.10})$$

where the first inequality holds because if A is a symmetric positive definite matrix, then for all vectors x , we have $x^\top A x \geq \lambda_{\min}(A)\|x\|_2^2$; the second inequality used the condition of $W' \in B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n}\varepsilon/4)$; the third inequality uses Lemma 20; and the last inequality uses $\lambda_{r_{\text{eff}}} \geq \frac{n\varepsilon^2}{2R^2}$ (by definition (3.3) of effective rank) and the approximation error condition (C.6). On the other hand, we have that $\|W\|_2^2 \leq R^2$, so that for every $W' \in B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n}\varepsilon/4)$

$$\|W' - W\|_2^2 = \|W' - \mathcal{P}_{\bar{\mathcal{V}}}(W)\|_2^2 + \|\mathcal{P}_{\bar{\mathcal{V}}_\perp}(W)\|_2^2 \leq \frac{4}{3}R^2,$$

combined with (C.10).

From the fundamental loewner order inequality (C.9), we establish the desired metric domination condition: for all $W' \in B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4)$ and $W \in B_2(R)$,

$$\begin{aligned} & (W' - W)^\top G(W)(W' - W) \\ & \leq (W' - W)^\top (2\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W)\mathcal{P}_{\bar{\mathcal{V}}})(W' - W) + (2\|G(W)^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}^2 + \lambda_{r_{\text{eff}}+1})\|W' - W\|_2^2 \\ & \leq 2\varrho_{G(W)}^{\bar{\mathcal{V}}}(W', \mathcal{P}_{\bar{\mathcal{V}}}(W))^2 + \frac{5n\varepsilon^2}{6}, \end{aligned}$$

where the first inequality holds because of the loewner order inequality (C.9) and the property of operator norm: $x^\top Ax \leq \|A\|_{\text{op}} \cdot \|x\|_2^2$ (one could also apply Lemma 20 to validate $\|\mathcal{P}_{\mathcal{V}_\perp}^\top \mathcal{P}_{\mathcal{V}_\perp}\|_{\text{op}} \leq 1$); and the last inequality uses the fact $\lambda_{r_{\text{eff}}+1} < \frac{n\varepsilon^2}{2R^2}$ (by definition 3.3 of effective rank) and the approximation error condition (C.6). Now we can apply the subspace metric domination lemma (Lemma 19) and obtain: for any $\pi \in \Delta(\bar{\mathcal{V}})$,

$$\log \frac{1}{\pi(B_{\varrho_{G(W)}}(W, \sqrt{n\varepsilon}))} \leq \log \frac{1}{\pi(B_{\sqrt{2}\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/\sqrt{6}))} \leq \log \frac{1}{\pi(B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4))}. \quad (\text{C.11})$$

In particular, we choose π to be the uniform prior over $\bar{\mathcal{V}}$:

$$\pi_{\bar{\mathcal{V}}} = \text{Unif}(B_2(1.58R) \cap \bar{\mathcal{V}}).$$

Then we aim to prove that $B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4) \subseteq \bar{\mathcal{V}} \cap B_2(1.58R)$. This is true because: 1) for every $W' \in B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4)$, (C.10) suggests $\|W' - \mathcal{P}_{\bar{\mathcal{V}}}(W)\|_2^2 \leq \frac{1}{3}R^2$, and 2) for every $W \in B_2(R)$, we have $\|\mathcal{P}_{\bar{\mathcal{V}}}(W)\|_2 \leq \|W\|_2 \leq R$. Combining this and the above inequality we have

$$\|W'\|_2 \leq \|W' - \mathcal{P}_{\bar{\mathcal{V}}}(W)\|_2 + \|\mathcal{P}_{\bar{\mathcal{V}}}(W)\|_2 \leq (\sqrt{1/3} + 1)R < 1.58R.$$

This proves that $B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4) \subseteq \bar{\mathcal{V}} \cap B_2(1.58R)$, so we have

$$\log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4))} = \log \frac{\text{Vol}(\bar{\mathcal{V}} \cap B_2(1.58R))}{\text{Vol}(B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4))}. \quad (\text{C.12})$$

By the change-of-variables theorem in multivariate calculus [Wikipedia contributors, 2025a], the linear map $T = G(W)^{\frac{1}{2}}$ implies the volume formula for ellipsoid $E = B_{\varrho_{G(W)}^{\bar{\mathcal{V}}}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4)$ with dimension r_{eff} , eigenvalues $\{\lambda_k(\mathcal{P}_{\bar{\mathcal{V}}}^\top G(W)\mathcal{P}_{\bar{\mathcal{V}}})\}_{k=1}^{r_{\text{eff}}}$ and radius $\sqrt{n\varepsilon}/4$

$$\text{Vol}(E) = |\det T|^{-1} \text{Vol}(T(E)) = (\det G(W))^{-1/2} \text{Vol}(B_2(\sqrt{n\varepsilon}/4)) = \left(\prod_{k=1}^{r_{\text{eff}}} \lambda_k \right)^{-1/2} \text{Vol}(B_2(\sqrt{n\varepsilon}/4)),$$

Also by the change-of-variable theorem, we have that the volume of r_{eff} -dimensional isotropic ball $\bar{\mathcal{V}} \cap B_2(1.58R)$ is

$$\text{Vol}(\bar{\mathcal{V}} \cap B_2(1.58R)) = \left(\frac{1.58R}{\sqrt{n\varepsilon}/4} \right)^{r_{\text{eff}}} \text{Vol}(B_2(\sqrt{n\varepsilon}/4)).$$

Hence, applying (C.11) (C.12) and combining it with the two above volume equalities, we have

$$\begin{aligned}
\log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\varrho_{G(W)}}(W, \sqrt{n\varepsilon}))} &\leq \log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\varrho_{G(W)}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4))} = \log \frac{\text{Vol}(\bar{\mathcal{V}} \cap B_2(1.58R))}{\text{Vol}(B_{\varrho_{G(W)}}(\mathcal{P}_{\bar{\mathcal{V}}}(W), \sqrt{n\varepsilon}/4))} \\
&\leq \frac{1}{2} \log \frac{(1.58R)^{2r_{\text{eff}}} \prod_{k=1}^{r_{\text{eff}}} \lambda_k}{(\sqrt{n\varepsilon}/4)^{2r_{\text{eff}}}} \leq \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}} \log \frac{40R^2 \lambda_k}{n\varepsilon^2} \\
&\leq d_{\text{eff}}(G(W), \sqrt{5}R, \varepsilon).
\end{aligned}$$

Finally, since the prior construction $\pi_{\bar{\mathcal{V}}} = \text{Unif}(B_2(1.58R) \cap \bar{\mathcal{V}})$ only depends on $\bar{\mathcal{V}}$ rather than W and ε , we have that uniformly over all $(W, \varepsilon) \in B_2(R) \times [0, \infty)$ such that $\bar{\mathcal{V}}$ approximates $\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon)$ to the precision (C.6),

$$\log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\varrho_{G(W)}}(W, \sqrt{n\varepsilon}))} \leq d_{\text{eff}}(G(W), \sqrt{5}R, \varepsilon),$$

which is the claimed bound. $\square$

Proof of Lemma 20: The Courant–Fischer–Weyl max-min characterization [Wikipedia contributors, 2025b] states that for any Hermitian (i.e. symmetric for real matrices studying here) matrix,

$$\lambda_k(\Sigma) = \max_{\substack{S \subseteq \mathbb{R}^p \\ \dim S = k}} \min_{\substack{W \in S \\ W \neq 0}} \frac{W^\top \Sigma W}{\|W\|_2^2},$$

and we have that for any r -dimensional subspace $\bar{\mathcal{V}}$,

$$\lambda_k(\mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}}) = \max_{\substack{S \subseteq \bar{\mathcal{V}} \\ \dim S = k}} \min_{\substack{W \in S \\ W \neq 0}} \frac{W^\top \mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}} W}{\|W\|_2^2},$$

so we have $\lambda_k(\mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}}) \leq \lambda_k$ for $k = 1, 2, \dots, r$.

Moreover, by the elementary property of eigendecomposition, since $\mathcal{V}$ is the top- r eigenspace of Σ , we have

$$\lambda_k = \lambda_k(\mathcal{P}_{\mathcal{V}}^\top \Sigma \mathcal{P}_{\mathcal{V}}), \quad k = 1, \dots, r.$$

Using the completing-square inequality (C.8), we have the Loewner-order bound

$$\mathcal{P}_{\mathcal{V}}^\top \Sigma \mathcal{P}_{\mathcal{V}} \preceq 2\mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}} + 2(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})^\top \Sigma (\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}}).$$

Furthermore,

$$(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})^\top \Sigma (\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}}) \preceq \|\Sigma^{1/2}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}^2 I.$$

Therefore,

$$\mathcal{P}_{\mathcal{V}}^\top \Sigma \mathcal{P}_{\mathcal{V}} \preceq 2\mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}} + 2\|\Sigma^{1/2}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}^2 I.$$

By the monotonicity of eigenvalues under the Loewner order, which follows immediately from the classical Courant–Fischer–Weyl min-max characterization [Wikipedia contributors, 2025b], we have

$$\lambda_k(\mathcal{P}_{\mathcal{V}}^\top \Sigma \mathcal{P}_{\mathcal{V}}) \leq 2\lambda_k(\mathcal{P}_{\bar{\mathcal{V}}}^\top \Sigma \mathcal{P}_{\bar{\mathcal{V}}}) + 2\|\Sigma^{1/2}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}^2.$$

Since $\lambda_k = \lambda_k(\mathcal{P}_V^\top \Sigma \mathcal{P}_V)$, rearranging gives

$$\lambda_k(\mathcal{P}_V^\top \Sigma \mathcal{P}_V) \geq \frac{\lambda_k}{2} - \|\Sigma^{1/2}(\mathcal{P}_V - \mathcal{P}_{\bar{V}})\|_{\text{op}}^2.$$

Combining this lower bound with the upper bound proved above yields the claim. $\square$

C.4 Proof of Riemannian Dimension Bound for DNN (Theorem 4)

In the language of Riemannian geometry [Jost, 2008], we regard a pointwise PSD, matrix-valued function $G(W)$ as a (possibly degenerate) *metric tensor*; such a $G(W)$ endows the parameter space $\mathbb{R}^{\sum_{l=1}^L d_{l-1}d_l}$ with a (semi-)Riemannian manifold structure. The pointwise ellipsoidal metric in (3.2) belongs to the following family of block-decomposable metric tensors.

Definition 6 (Metric Tensor of NN-surrogate Type) *A metric tensor $G(W)$ (pointwise PSD matrix-valued function of size $\sum_{l=1}^L d_{l-1}d_l \times \sum_{l=1}^L d_{l-1}d_l$) is of “NN-surrogate” type if $G(W)$ is in the form*

$$G(W) = \text{blockdiag}(A_1(W) \otimes I_{d_1}, \dots, A_l(W) \otimes I_{d_l}, \dots, A_L(W) \otimes I_{d_L})$$

where $A_l(W) \in \mathbb{R}^{d_{l-1} \times d_{l-1}}$.

By Lemma 1, the non-perturbative feature expansion gives rise to the metric tensor $G_{\text{NP}}(W)$ defined in (3.2); $G_{\text{NP}}(W)$ belongs to the “NN-surrogate” class. We first record some elementary decomposition properties for this family of NN-surrogate metric tensors, and then prove Theorem 4.

C.4.1 Decomposition Properties of NN-surrogate Metric Tensor

The NN-surrogate metric tensor $G(W)$ in Definition 6 has decomposition properties described by the next lemma.

Lemma 21 (Decomposition Properties of NN-surrogate Metric Tensor) *Given a NN-surrogate metric tensor $G(W)$ defined in Definition 6, for every W , we have the following decomposition properties: First, the effective rank and dimension decompose to*

$$\begin{aligned} r_{\text{eff}}(G(W), R, \varepsilon) &= \sum_{l=1}^L d_l \cdot r_{\text{eff}}(A_l(W), R, \varepsilon); \\ d_{\text{eff}}(G(W), R, \varepsilon) &= \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), R, \varepsilon). \end{aligned}$$

Second, denote $\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon)$ the effective subspace (i.e., the top- $r_{\text{eff}}(A_l(W), R, \varepsilon)$ eigenspace) of $A_l(W)$. Then the effective subspace of $G(W)$ is

$$\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon) = \mathcal{V}_{\text{eff}}(A_1(W), R, \varepsilon)^{d_1} \times \dots \times \mathcal{V}_{\text{eff}}(A_L(W), R, \varepsilon)^{d_L}.$$

Proof of Lemma 21. It is straightforward to see that, first, the effective rank of the fixed matrix $G(W)$ is

$$\begin{aligned}
& r_{\text{eff}}(G(W), R, \varepsilon) \\
&= \max\{k : 2\lambda_k(G(W))R^2 \geq n\varepsilon^2\} \\
&= \sum_{l=1}^L \max\{k : 2\lambda_k(A_l(W) \otimes I_{d_l})R^2 \geq n\varepsilon^2\} \\
&= \sum_{l=1}^L d_l \max\{k : 2\lambda_k(A_l(W))R^2 \geq n\varepsilon^2\} \\
&= \sum_{l=1}^L d_l \cdot r_{\text{eff}}(A_l(W), R, \varepsilon);
\end{aligned}$$

and the effective dimension of the fixed matrix $G(W)$ is

$$\begin{aligned}
& d_{\text{eff}}(G(W), R, \varepsilon) \\
&= \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}(G(W), R, \varepsilon)} \log \left(\frac{8R^2 \lambda_k(G(W))}{n\varepsilon^2} \right) \\
&= \sum_{l=1}^L \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}(A_l(W) \otimes I_{d_l}, R, \varepsilon)} \log \left(\frac{8R^2 \lambda_k(A_l(W) \otimes I_{d_l})}{n\varepsilon^2} \right) \\
&= \sum_{l=1}^L d_l \cdot \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}(A_l(W), R, \varepsilon)} \log \left(\frac{8R^2 \lambda_k(A_l(W))}{n\varepsilon^2} \right) \\
&= \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), R, \varepsilon).
\end{aligned}$$

Second, as the effective subspace of the matrix tensor product $A_l(W) \otimes I_{d_l}$ is subspace tensor product $\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon)^{d_l}$, the effective subspace for NN-surrogate metric tensor $G(W) = \text{blockdiag}(\dots; A_l(W) \otimes I_{d_l}; \dots)$ is

$$\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon) := \mathcal{V}_{\text{eff}}(A_1(W), R, \varepsilon)^{d_1} \times \dots \times \mathcal{V}_{\text{eff}}(A_L(W), R, \varepsilon)^{d_L}.$$

□

C.4.2 Proof of Theorem 4

We firstly prove the following result, which is almost Theorem 4, with the only difference being that the radius in the effective dimension depends on the global radius R rather than the pointwise Frobenious norm $\|W\|_{\mathbf{F}}$. Extending this result to Theorem 4 can be achieved via a simple application of the “uniform pointwise convergence” principle [Xu and Zeevi, 2025] illustrated in Lemma 4.

Lemma 22 (Riemannian Dimension for NN-surrogate Metric Tensor—Global Radius Version)

Consider the NN-surrogate metric tensor in Definition 6, and the weight space $B_{\mathbf{F}}(R)$. Then we have that the pointwise dimension is bounded by the pointwise Riemannian Dimension as the following: there exists a prior π such that uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\log \frac{1}{\pi(B_{\ell_G(W)}(W, \sqrt{n}\varepsilon))} \leq \sum_{l=1}^L \left(\underbrace{d_l \cdot d_{\text{eff}}(A_l(W), CR, \varepsilon)}_{\text{“must pay” cost at each } W} + \underbrace{d_{l-1} \cdot d_{\text{eff}}(A_l(W), CR, \varepsilon)}_{\text{covering cost of Grassmannian}} + \underbrace{\log(d_{l-1})}_{\text{covering cost of } r_{\text{eff}} \in [d_{l-1}]} \right),$$

where $C > 0$ is an absolute constant.

Proof of Lemma 22: The proof has two key steps: 1. Hierarchical covering argument, and 2. Bound covering Cost of the Grassmannian. A crucial lemma about the ellipsoidal covering of the Grassmannian, which is new even in the pure mathematics context, is deferred to Section D.

Step 1: Hierarchical Covering. As explained in the main paper, the major difficulty is that the prior measure $\pi_{\mathcal{V}}$ it constructed, is defined over the effective subspace $\mathcal{V}$, which itself encodes information of the point W and $\varepsilon > 0$. The goal of our proof is to construct a “universal” prior π that does not depend on $\mathcal{V}$. This is achieved via a hierarchical covering argument (3.6), which we make rigorous below.

The key idea of hierarchical covering is as follows: Firstly, for all W , we search for subspace $\bar{\mathcal{V}}$ that approximates the true effective subspace (top- r_{eff} eigenspace) $\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon)$ to the precision required by (C.6):

$$\|G(W)^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}} \leq \frac{\sqrt{n}\varepsilon}{4R}, \quad (\text{C.13})$$

where $G(W)^{\frac{1}{2}}$ is the unique square root of PSD matrix $G(W)$ (see, e.g, [Wikipedia contributors, 2025e]). Then by Lemma 2 (Pointwise Dimension Bound for Nonlinear Manifold with Approximate Effective Subspace), for every $(W, \varepsilon) \in B_2(R) \times [0, \infty)$ such that $\bar{\mathcal{V}}$ approximates $\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon)$ to the precision (C.13), the prior $\pi_{\bar{\mathcal{V}}} = \text{Unif}(B_2(1.58R) \cap \bar{\mathcal{V}})$ satisfies

$$\log \frac{1}{\pi_{\bar{\mathcal{V}}}(B_{\ell_G(W)}(W, \sqrt{n}\varepsilon))} \leq d_{\text{eff}}(G(W), \sqrt{5}R, \varepsilon) = \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), \sqrt{5}R, \varepsilon), \quad (\text{C.14})$$

where the first inequality is by Lemma 2 (see definition (3.4) of effective dimension); and the last equality is by the decomposition property of NN-surrogate metric tensor (Lemma 21).

Secondly, we put a prior μ over all possible subspaces $\mathcal{V}$ and construct the “universal” prior

$$\pi(W) = \sum_{\mathcal{V}} \mu(\mathcal{V}) \times \pi_{\mathcal{V}}(W), \quad (\text{C.15})$$

which implies that uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned}
& \log \frac{1}{\pi(B_{\varrho_{G(W)}}(W, \sqrt{n}\varepsilon))} \\
&= \log \frac{1}{\sum_{\mathcal{V}} \mu(\mathcal{V}) \pi_{\mathcal{V}}(B_{\varrho_{G(W)}}(W, \sqrt{n}\varepsilon))} \\
&\leq \log \frac{1}{\mu(\bar{\mathcal{V}} : \bar{\mathcal{V}} \text{ satisfies (C.13)}) \inf_{\bar{\mathcal{V}} \text{ satisfies (C.13)}} \pi_{\bar{\mathcal{V}}}(B_{\varrho_{G(W)}}(W, \sqrt{n}\varepsilon))} \\
&\leq \underbrace{\log \frac{1}{\mu(\bar{\mathcal{V}} : \bar{\mathcal{V}} \text{ satisfies (C.13)})}}_{\text{covering cost of the Grassmannian}} + \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), \sqrt{5}R, \varepsilon), \tag{C.16}
\end{aligned}$$

where the first equality is by definition (C.15) of the “universal” prior π ; the first inequality is straightforward; and the last inequality is by (C.14) (the result of the “must pay” part in the hierarchical covering) and the equivalence between $B_2(R)$ and $B_{\mathbf{F}}(R)$.

The above hierarchical covering argument successfully gives a valid Riemannian Dimension, with the cost of the additional covering cost given by the subspace prior μ . This explains our basic proof idea. The remaining proof executes this basic proof idea.

Step 2: Bounding Covering Cost of the Grassmannian. Section D provides a systematic study to the ellipsoidal metric entropy of Grassmannian manifold, which we detail the conclusion below.

Define

$$\text{Gr}(d, r) := \{r\text{-dimensional linear subspaces of } \mathbb{R}^d\}$$

as the *Grassmannian manifold*.

Given a $d \times d$ PSD Σ , define the anisometric projection metric between two subspaces by (labeled as Definition 7 in Section D)

$$\varrho_{\text{proj}, \Sigma}(\mathcal{V}, \bar{\mathcal{V}}) = \|\Sigma^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}, \tag{C.17}$$

where $\Sigma^{\frac{1}{2}}$ is the square root of the PSD matrix Σ (see, e.g., [Wikipedia contributors, 2025e]).

Lemma 3 states that (note that we use ε_1 and C_0 here instead of ε and C in the original statement of Lemma 3), given a Grassmannian $\text{Gr}(d, r)$, for uniform prior $\mu = \text{Unif}(\text{Gr}(d, r))$, we have that for every $\mathcal{V} \in \text{Gr}(d, r)$, every $\varepsilon_1 > 0$ and PSD matrix $\Sigma \in \mathbb{R}^{d \times d}$ with eigenvalues $\lambda_1 \geq \dots \lambda_d \geq 0$, we have the pointwise dimension bound

$$\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon_1))} \leq \frac{d-r}{2} \sum_{k=1}^r \log \frac{C_0 \max\{\lambda_k, \varepsilon_1^2\}}{\varepsilon_1^2} + \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{C_0 \max\{\lambda_k, \varepsilon_1^2\}}{\varepsilon_1^2}, \tag{C.18}$$

where $C_0 > 0$ is an absolute constant. We will use the result (C.18) and (C.16) to prove Theorem 4.

For a particular layer l , $d_{l-1} \times d_{l-1}$ PSD matrix $A_l(W)$, and a fixed rank r_l denote $\text{Gr}(d_{l-1}, r_l)$ as a Grassmannian (the collection of all r_l -dimensional in $\mathbb{R}^{d_{l-1}}$). By (C.18) we have that there exists a prior μ_l over $\text{Gr}(d_{l-1}, r_l)$ such that for every (W, ε_1) such that $r_{\text{eff}}(A_l(W), R, \varepsilon_1) = r_l$, and

$\lambda_{r_l+1}(A_l(W)) \leq c\varepsilon_1^2 \leq \lambda_{r_l}(A_l(W))$ where $c \geq 1$ can be any absolute constants no smaller than 1 (later we will specialize to $c = 8$),

$$\log \frac{1}{\mu_l(\bar{\mathcal{V}} : \varrho_{\text{proj}, A_l(W)}(\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon), \bar{\mathcal{V}}) \leq \varepsilon_1)} \leq \frac{d_{l-1}}{2} \sum_{k=1}^{r_l} \log \frac{C_1 \lambda_k(A_l(W))}{\varepsilon_1^2}, \quad (\text{C.19})$$

where $C_1 = c \max\{C_0, 1\} \geq 1$ is an absolute constant depending only on the absolute constant c (later we take $c = 8$ so $C_1 = 8 \max\{C_0, 1\}$ is indeed an absolute constant). (C.19) is because: 1) all eigenvalues with index at least $r_l + 1$ (each no larger than $c\varepsilon_1^2$) contribute only through the second term in (C.18). Their cumulative effect is at most

$$\mathbb{1}\{d_{l-1} - r_l > r_l\} \cdot \frac{r_l}{2} \sum_{k=r_l+1}^{d_{l-1}-r_l} \log \frac{C_0 c \varepsilon_1^2}{\varepsilon_1^2} = \frac{r_l \max\{d_{l-1} - 2r_l, 0\}}{2} \log C_0 c \leq \frac{r_l(d_{l-1} - r_l)}{2} \log C_0 c$$

unaffected to the spectrum, and we absorb this into the absolute constant C_1 . And 2) all eigenvalues with index at most r_l 's contribution leads to at most

$$\frac{d_{l-1} - r_l}{2} \sum_{k=1}^{r_l} \log \frac{C_0 \lambda_k(A_l(W))}{\varepsilon_1^2} + \frac{r_l}{2} \sum_{k=1}^{\max\{r_l, d_{l-1}-r_l\}} \log \frac{C_0 \lambda_k(A_l(W))}{\varepsilon_1^2} \leq \frac{d_{l-1}}{2} \sum_{k=1}^{r_l} \log \frac{\max\{C_0, 1\} \lambda_k(A_l(W))}{\varepsilon_1^2}.$$

Summing up the contributions two parts of the spectrum together, we get the right hand side of (C.19).

By the subspace decomposition property in Lemma 21, we have that for $\bar{\mathcal{V}} = (\cdots, \underbrace{\bar{\mathcal{V}}_l, \cdots, \bar{\mathcal{V}}_l}_{\text{repeat } d_l \text{ times}}, \cdots)$,

$$\begin{aligned} & \varrho_{\text{proj}, G(W)}(\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon), \bar{\mathcal{V}}) \\ &= \varrho_{\text{proj}, G(W)}\left(\prod_{l=1}^L \mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon)^{d_l}, \prod_{l=1}^L \bar{\mathcal{V}}_l^{d_l}\right) \\ &= \max_l \varrho_{\text{proj}, A_l(W)}(\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon), \bar{\mathcal{V}}_l), \end{aligned} \quad (\text{C.20})$$

where the first equality is by Lemma 21, and the second equality is by the properties of the spectral norm: $\|\text{blockdiag}(A, B)\|_{\text{op}} = \max\{\|A\|_{\text{op}}, \|B\|_{\text{op}}\}$ and $\|A \otimes I_d\|_{\text{op}} = \|A\|_{\text{op}}$.

Taking $\varepsilon_1 = \frac{\sqrt{n\varepsilon}}{4R}$, by definition (3.3) on the threshold to determine effective rank, we obtain $\lambda_{r_l+1}(A_l(W)) \leq 8\varepsilon_1^2 = n\varepsilon^2/(2R^2) \leq \lambda_{r_l}(A_l(W))$, thus this particular choice satisfies the required eigenvalue condition to establish (C.19) with $c = 8$. Then for all layers $l = 1, \cdots, L$, given a fixed $\{r_1, \cdots, r_L\}$, by (C.19), we have that there exists a prior

$$\mu_{\{r_l\}_{l=1}^L} = \mu_1^{d_1} \otimes \cdots \otimes \mu_L^{d_L} \quad (\text{C.21})$$

over the product Grassmannian $\text{Gr}(d_0, r_1)^{d_1} \times \cdots \times \text{Gr}(d_{L-1}, r_L)^{d_L}$ such that uniformly over all $W \in B_{\mathbf{F}}(R)$ such that $r_{\text{eff}}(A_l(W), R, \varepsilon) = r_l, \forall l \in [L]$ (here $[L]$ is the notation of $\{1, 2, \cdots, L\}$), the

“Grassmannian covering cost” term in (C.16) is bounded by

$$\begin{aligned}
& \log \frac{1}{\mu(\bar{\mathcal{V}} : \bar{\mathcal{V}} \text{ satisfies (C.13)})} \\
&= \log \frac{1}{\mu_{\{r_l\}_{l=1}^L}(\{\bar{\mathcal{V}} : \varrho_{\text{proj}, G(W)}(\mathcal{V}_{\text{eff}}(G(W), R, \varepsilon), \bar{\mathcal{V}}) \leq \frac{\sqrt{n}\varepsilon}{4R} = \varepsilon_1\})} \\
&\leq \log \frac{1}{\mu_{\{r_l\}_{l=1}^L}(\{(\cdots, \underbrace{\bar{\mathcal{V}}_l, \cdots, \bar{\mathcal{V}}_l}_{d_l \text{ times}}, \cdots) : \varrho_{\text{proj}, A_l(W)}(\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon), \bar{\mathcal{V}}_l) \leq \varepsilon_1, \quad \forall l \in [L]\})} \\
&= \sum_{l=1}^L \log \frac{1}{\mu_{\{r_l\}, l}(\{\bar{\mathcal{V}}_l : \varrho_{\text{proj}, A_l(W)}(\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon), \bar{\mathcal{V}}_l) \leq \varepsilon_1\})} \\
&\leq \sum_{l=1}^L \frac{d_{l-1}}{2} \sum_{k=1}^{r_l} \log \frac{C_1 \lambda_k(A_l(W))}{\varepsilon_1^2} \\
&\leq \sum_{l=1}^L d_{l-1} d_{\text{eff}}(A_l(W), \sqrt{2C_1}R, \varepsilon), \tag{C.22}
\end{aligned}$$

where the first inequality follows by restricting the reference subspace $\bar{\mathcal{V}}$ to the repeated-product form $\prod_{l=1}^L \bar{\mathcal{V}}_l^{d_l}$ and using (C.20): for such subspaces, the global projection error is the maximum of the layerwise projection errors, so requiring every layerwise error to be at most ε_1 implies the global condition (C.13). The second equality follows from the layer-wise product structure of the prior in (C.21). More explicitly, if $E_l := \{\bar{\mathcal{V}}_l : \varrho_{\text{proj}, A_l(W)}(\mathcal{V}_{\text{eff}}(A_l(W), R, \varepsilon), \bar{\mathcal{V}}_l) \leq \varepsilon_1\}$, then the repeated-product event is $\bigcap_{l=1}^L E_l$, and the product prior gives $\mu_{\{r_l\}_{l=1}^L}(\bigcap_{l=1}^L E_l) = \prod_{l=1}^L \mu_{l, r_l}(E_l)$. Here $\mu_l^{d_l}$ in (C.21) denotes the pushforward of μ_l under the repeated-copy map $\bar{\mathcal{V}}_l \mapsto \bar{\mathcal{V}}_l^{d_l}$, not the d_l -fold independent product measure. Thus one samples a single reference subspace $\bar{\mathcal{V}}_l \sim \mu_l$ at layer l and repeats this same subspace d_l times, so no extra d_l factor appears in the atlas cost. The second inequality is by the layer-wise covering bound (C.19); and the last inequality follows from the choice $\varepsilon_1 = \sqrt{n}\varepsilon/(4R)$ and the definition of effective dimension in (3.4).

Note that (C.22) is uniformly over all $W \in B_{\mathbf{F}}(R)$ such that $r_{\text{eff}}(A_l(W), R, \varepsilon) = r_l$, $\forall l \in [L]$, not uniformly over all $W \in B_{\mathbf{F}}(R)$. We would like to extend (C.22) to all $W \in B_{\mathbf{F}}(R)$ over uniform prior over possible integer values of r_l . Now assign uniform prior over $[d_{l-1}] = \{1, \dots, d_{l-1}\}$ for r_l , we obtain the “universal” prior π (as we have pursued in our hierarchical covering argument (C.15)) defined by

$$\begin{aligned}
\mu(\mathcal{V}) &= \prod_{l=1}^L \underbrace{\text{Unif}([d_{l-1}])}_{\text{prior of } r_l} \otimes \underbrace{\mu_{\{r_k\}_{k=1}^L}}_{\text{prior over product Grassmannian in (C.21)}}, \\
\pi(W) &= \sum_{\mathcal{V}} \underbrace{\mu(\mathcal{V})}_{\text{prior over subspaces defined above}} \otimes \underbrace{\text{Unif}(B_2(1.58R) \cap \bar{\mathcal{V}})}_{\text{uniform prior constrained in subspace}}. \tag{C.23}
\end{aligned}$$

Then we have that uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned}
& \log \frac{1}{\pi(B_{\varrho_G(W)}(W, \sqrt{n\varepsilon}))} \\
& \leq \log \frac{1}{\mu(\bar{\mathcal{V}} : \bar{\mathcal{V}} \text{ satisfies (C.13)})} + \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), \sqrt{5}R, \varepsilon) \\
& \leq \sum_{l=1}^L \log d_{l-1} + \log \frac{1}{\mu_{\{r_k\}_{k=1}^L}(\bar{\mathcal{V}} : \bar{\mathcal{V}} \text{ satisfies (C.13)})} + \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), \sqrt{5}R, \varepsilon) \\
& \leq \sum_{l=1}^L \log d_{l-1} + \sum_{l=1}^L d_{l-1} \cdot d_{\text{eff}}(A_l(W), \sqrt{2C_1}R, \varepsilon) + \sum_{l=1}^L d_l \cdot d_{\text{eff}}(A_l(W), \sqrt{5}R, \varepsilon),
\end{aligned}$$

where $C_1 > 0$ is an absolute constant. Here the first inequality is by the hierarchical covering argument (C.16); the second inequality is by the prior construction (C.23); and the third inequality is by the Grassmannian covering bound (C.22) for fixed $\{r_k\}_{k=1}^L$. This shows that for NN-surrogate metric tensor $G(W)$, the pointwise dimension is bounded by the Riemannian Dimension as the following:

$$\log \frac{1}{\pi(B_{\varrho_G(W)}(W, \sqrt{n\varepsilon}))} \leq \sum_{l=1}^L (d_l + d_{l-1}) \cdot d_{\text{eff}}(A_l(W), CR, \varepsilon) + \log(d_{l-1}),$$

where C is a positive absolute constant. This finishes the proof of Lemma 22 with R in effective dimension being a global upper bound of $\|W\|_{\mathbf{F}}$. $\square$

Proof of Theorem 4: Motivated by the “uniform pointwise convergence” principle (proposed in Xu and Zeevi [2025] and illustrated in Lemma 4), we apply a peeling argument to adapt the Riemannian Dimension to $\|W\|_{\mathbf{F}}$. Given any $R_0 \in (0, R]$, we take $R_k = 2^k R_0$ for $k = 0, 1, \dots, \log_2 \lceil R/R_0 \rceil$. Taking a uniform prior on these R_k , and set

$$\tilde{\pi} = \underbrace{\text{Unif}(\{R_0, \dots, 2^{\log_2 \lceil R/R_0 \rceil} R_0\})}_{\text{prior over upper bound } \tilde{R} \text{ of } \|W\|_{\mathbf{F}}} \otimes \underbrace{\pi_{\tilde{R}}}_{\text{prior defined via (C.23)}},$$

where $\pi_{\tilde{R}}$ is the prior defined via (C.23) in the proof of Lemma 22. Then for every $W \in B_{\mathbf{F}}(R)$ where $\|W\|_{\mathbf{F}} > R_0$, denote $k(W)$ to be the integer such that $2^{k(W)} R_0 < \|W\|_{\mathbf{F}} \leq 2^{k(W)+1} R_0$, then

$$\begin{aligned}
& \log \frac{1}{\tilde{\pi}(B_{\varrho_G(W)}(W, \sqrt{n\varepsilon}))} \\
& \leq \underbrace{\log \log_2 \lceil R/R_0 \rceil}_{\text{density of } 2^{k(W)+1} R_0} + \underbrace{\log \frac{1}{\pi_{2^{k(W)+1} R_0}(B_{\varrho_G(W)}(W, \sqrt{n\varepsilon}))}}_{\pi \text{ is constructed via (C.23), with global radius taken to be } 2^{k(W)+1} R_0} \\
& \leq \log \log_2 \lceil R/R_0 \rceil + \sum_{l=1}^L ((d_l + d_{l-1}) \cdot d_{\text{eff}}(A_l(W), C_1 2^{k(W)+1} R_0, \varepsilon) + \log d_{l-1}) \\
& \leq \log \log_2 \lceil R/R_0 \rceil + \sum_{l=1}^L ((d_l + d_{l-1}) \cdot d_{\text{eff}}(A_l(W), C_1 \cdot 2\|W\|_{\mathbf{F}}, \varepsilon) + \log d_{l-1}),
\end{aligned}$$

where the first inequality is due to the product construction of $\tilde{\pi}$; the second inequality is due to Lemma 22, with $C_1 > 0$ being an absolute constant; and the last inequality uses the fact $\|W\|_{\mathbf{F}} \leq 2^{k(W)+1} R_0 \leq 2\|W\|_{\mathbf{F}}$, with $C_1 > 0$.

The above bound assumes $\|W\|_{\mathbf{F}} > R_0$. When $\|W\|_{\mathbf{F}} \leq R_0$, we directly apply Lemma 22 and obtain

$$\begin{aligned} & \log \frac{1}{\tilde{\pi}(B_{\varrho_{G(W)}}(W, \sqrt{n}\varepsilon))} \\ & \leq \underbrace{\log \log_2[R/R_0]}_{\text{density of } R_0} + \underbrace{\log \frac{1}{\pi_{R_0}(B_{\varrho_{G(W)}}(W, \sqrt{n}\varepsilon))}}_{\pi \text{ is constructed via (C.23), with global radius taken to be } R_0} \\ & \leq \log \log_2[R/R_0] + \sum_{l=1}^L ((d_l + d_{l-1}) \cdot d_{\text{eff}}(A_l(W), C_1 \cdot R_0, \varepsilon) + \log d_{l-1}). \end{aligned}$$

Combining the two cases discussed above, we conclude that the pointwise dimension for NN-surrogate metric tensor $G(W)$ in Definition 6 is bounded by the Riemmanin Dimension

$$\begin{aligned} & \log \frac{1}{\tilde{\pi}(B_{\varrho_{G(W)}}(W, \sqrt{n}\varepsilon))} \leq d_{\text{R}}(W, \varepsilon) \\ & = \sum_{l=1}^L ((d_l + d_{l-1}) \cdot d_{\text{eff}}(A_l(W), C \max\{\|W\|_{\mathbf{F}}, R_0\}, \varepsilon) + \log(d_{l-1} \log_2[R/R_0])), \end{aligned}$$

where $C = 2C_1$ is a positive absolute constant. The term $\log_2[R/R_0]$ is enlarged by an additional factor of L .

Finally, by the sentence below (3.2) (which is a straightforward result from non-perturbative feature expansion for DNN (Lemma 1) and the metric domination lemma (Lemma 18)), we know that there exists a prior $\tilde{\pi}$ such that uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned} & \log \frac{1}{\tilde{\pi}(B_{\varrho_n}(f(W, \cdot), \varepsilon))} \leq \log \frac{1}{\tilde{\pi}(B_{\varrho_{G_{\text{NP}}(W)}}(W, \sqrt{n}\varepsilon))} \\ & \leq d_{\text{R}}(W, \varepsilon) = \sum_{l=1}^L ((d_l + d_{l-1}) \cdot d_{\text{eff}}(A_l(W), C \max\{\|W\|_{\mathbf{F}}, R_0\}, \varepsilon) + \log(d_{l-1} \log_2[R/R_0])), \end{aligned}$$

where $A_l(W) = LM_{l \rightarrow L}^2(W, \varepsilon) \cdot F_{l-1}(W, X) F_{l-1}^\top(W, X)$ when taking $G(W)$ to be $G_{\text{NP}}(W)$ defined in (3.2). Taking $R_0 = R/2^n$ proves Theorem 4. □

D Ellipsoidal Covering of the Grassmannian (Lemma 3)

The central goal of this section is to prove the following result on the ellipsoidal metric entropy of the Grassmannian manifold. The definition for Gr (Grassmannian manifold), St (Stiefel parameterization manifold) are temporarily deferred to Section D.1.

Definition 7 (Ellipsoidal Projection Metric) For two subspaces $\mathcal{V}, \bar{\mathcal{V}} \in \text{Gr}(d, r)$, and a positive semidefinite matrix Σ , define the ellipsoidal projection metric $\varrho_{\text{proj}, \Sigma}$ by

$$\varrho_{\text{proj}, \Sigma}(\mathcal{V}, \bar{\mathcal{V}}) = \|\Sigma^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}},$$

where $\mathcal{P}_{\mathcal{V}}$ and $\mathcal{P}_{\bar{\mathcal{V}}}$ are orthogonal projectors to subspace $\mathcal{V}$ and $\bar{\mathcal{V}}$, respectively.

We view orthogonal projectors as matrices (see the definition via the Stiefel parameterization in (D.6)), consistent with the earlier operator notation characterized by ℓ_2 -distance in Lemma 19. In the isotropic case $\Sigma = I_d$, the ellipsoidal projection metric reduces to the standard isotropic projection metric

$$\varrho_{\text{proj}}(\mathcal{V}, \bar{\mathcal{V}}) = \|\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}}\|_{\text{op}}.$$

We now state our main result in this section (Lemma 3 in the main paper).

Consider the Grassmannian $\text{Gr}(d, r)$ and the uniform prior $\mu = \text{Unif}(\text{Gr}(d, r))$, then for every $\mathcal{V} \in \text{Gr}(d, r)$, every $\varepsilon > 0$ and every PSD matrix Σ with eigenvalues $\lambda_1 \geq \dots \geq \lambda_d \geq 0$, we have

$$\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} \leq \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{C \max\{\lambda_k, \varepsilon^2\}}{\varepsilon^2} + \frac{d-r}{2} \sum_{k=1}^r \log \frac{C \max\{\lambda_k, \varepsilon^2\}}{\varepsilon^2}, \quad (\text{D.1})$$

where $C > 0$ is an absolute constant.

Recall that the traditional covering number bound for the Grassmannian manifold states that

$$\left(\frac{C}{\varepsilon}\right)^{r(d-r)} \leq N(\text{Gr}(d, r), \varrho_{\text{proj}}, \varepsilon) \leq \left(\frac{C}{\varepsilon}\right)^{r(d-r)}. \quad (\text{D.2})$$

Here $N(\mathcal{F}, \varrho, \varepsilon)$ is the standard covering number—the smallest size of an ε -net that covers $\mathcal{F}$ under the metric ϱ ; see Definition 5 for details. In comparison, Lemma 3 is much more challenging than proving classical isotropic covering number bounds (D.2) because

- 1) we consider ellipsoidal metric;
- 2) we require the prior μ to be independent with Σ and ε .

We need to firstly understand how such classical results are proved, and then proceed to generalized them. This suggests that deep mathematical insights are necessary for the purpose to study neural networks generalization, as we will introduce below.

From Pure Mathematics to Machine Learning Language. Understanding the classical proof for the Grassmannian and generalizing them to prove Lemma 3 necessitate the a deep dive in to the geometry and algebra of subspaces and Grassmannians. In fact, traditional treatments to study Grassmannian manifold often invoke advanced machinery—ranging from differential geometry [Bendokat et al., 2024] and Lie-group theory [Szarek, 1997] to algebraic geometry [Devriendt et al., 2024], and the seminal covering number proof [Szarek, 1997] is particularly stated in Lie-algebra and differential-geometry language.

Motivated by the subsequent covering number proof [Pajor, 1998] that uses relatively more elementary language, we give an exposition that is elementary and entirely self-contained, relying only on matrix-analysis and learning-theoretic techniques familiar from machine learning. In

particular, every “advanced” fact—for example, the group theory of continuous symmetries traditionally handled via Lie groups—is derived by elementary means (explicit matrix parameterizations, principal-angle/cosine-sine representations, and basic spectral arguments) while preserving the high-level geometric intuition. We hope that this versatile framework—and our novel contributions (e.g., Definition 7 and Lemma 3), which are new even in a pure-mathematics setting—will establish subspaces, the Grassmannian, and their underlying algebraic structures as powerful tools for future machine learning applications.

Effective Rank vs. Full-Spectrum Complexity. Consider a covariance matrix Σ with eigenvalues $\lambda_1 \geq \dots \lambda_d \geq 0$. By Definition 7, the ellipsoidal metric satisfies

$$\varrho_{\text{proj}, \Sigma}(\mathcal{V}, \bar{\mathcal{V}}) \leq \lambda_1^{\frac{1}{2}} \varrho_{\text{proj}}(\mathcal{V}, \bar{\mathcal{V}}).$$

If one is willing to accept a coarser complexity scaling, then one could invoke existing Grassmannian covering results under the canonical isotropic metric (D.2) (taking $\mu = \text{Unif}(\text{Gr}(d, r))$) and obtain

$$\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} \leq \log \frac{1}{\mu(B_{\varrho_{\text{proj}}}(\mathcal{V}, \varepsilon/\sqrt{\lambda_1}))} \leq (d - r_{\text{eff}}(\Sigma, R, \varepsilon)) r_{\text{eff}}(\Sigma, R, \varepsilon) \log \frac{C\lambda_1}{\varepsilon^2}. \quad (\text{D.3})$$

However, this makes the *global atlas* cost dominate the *local chart* cost, yielding a suboptimal bound than the full-spectrum effective dimension in (D.1). The refined analysis in this section—also simplifying and strengthening the isotropic route—establishes the correct structural principle: the global-atlas cost must be balanced by the local-chart cost. Thus, while the effective-rank bound (D.3) serves as a useful sanity check, the full-spectrum treatment is what delivers the sharpened complexities required for our main results.

D.1 Grassmannian Manifold, Stiefel Parameterization, and Orthogonal Groups

Fix integers $r \leq d$. Define

$$\text{Gr}(d, r) := \{r\text{-dimensional linear subspaces of } \mathbb{R}^d\}$$

as the *Grassmann manifold*. Write

$$\text{St}(d, r) := \{V \in \mathbb{R}^{d \times r} : V^\top V = I_r\}$$

for the *Stiefel manifold* of r orthonormal columns in $\mathbb{R}^d$. $\text{St}(d, r)$ is a convenient *parameterization* of that class $\text{Gr}(d, r)$.

If for subspace $\mathcal{V} \in \text{Gr}(d, r)$ and matrix $V \in \text{St}(d, r)$ we have $\mathcal{V} = \text{span}(V)$, then we say V is a *parameterization matrix* of $\mathcal{V}$. Though such parameterization is not unique, the associated orthogonal projector and projection metric are both unique. Moreover, the anisometric projection we define in Definition 7 is also unique. We will prove these shortly.

Write

$$O(r) := \{Q \in \mathbb{R}^{r \times r} : Q^\top Q = QQ^\top = I_r\}$$

to be the *orthogonal group*. Optionally, we also state that (in the real setting)

$$\text{Gr}(d, r) \cong O(d) / (O(r) \times O(d-r)) \cong \text{Gr}(d, d-r), \quad (\text{D.4})$$

where “/” denotes the *quotient* and “ $\cong$ ” denotes a canonical *isomorphism* (indeed, a *diffeomorphism* of smooth manifolds or a *homeomorphism* of topological manifolds; see, e.g., Chapter 1.5 in [Awodey, 2010]). Moreover, $\text{Gr}(d, r)$ can be regarded as a standard *algebraic variety* [Devriendt et al., 2024]. We do not aim to explain these notions in detail, but merely note that:

1. The geometric properties of $\text{Gr}(d, r)$ coincide with those of $\text{Gr}(d, d-r)$ under this isomorphism (geometric equivalence).
2. The number of degrees of freedom of $\text{Gr}(d, r)$ is

$$\underbrace{\frac{d(d-1)}{2}}_{\dim O(d)} - \underbrace{\frac{r(r-1)}{2}}_{\dim O(r)} - \underbrace{\frac{(d-r)(d-r-1)}{2}}_{\dim O(d-r)} = r(d-r), \quad (\text{D.5})$$

which also appears as the dimension factor in the precise covering-number bounds (D.2).

We now define the orthogonal projector and the projection metric on the Grassmannian manifold.

Definition of Orthogonal Projector. For $V \in \text{St}(d, r)$ and its column-space $\mathcal{V} = \text{span}(V)$, define the rank- r orthogonal projector⁴

$$\mathcal{P}_{\mathcal{V}} := VV^{\top} \in \mathbb{R}^{d \times d}. \quad (\text{D.6})$$

Then $\mathcal{P}_{\mathcal{V}}$ depends *only* on the subspace $\mathcal{V}$. Indeed, if $Q \in O(r)$ then $(VQ)(VQ)^{\top} = VQQ^{\top}V^{\top} = VV^{\top}$, so V and VQ represent the same subspace. Hence the map

$$\Psi : \text{St}(d, r) \longrightarrow \text{Gr}(d, r), \quad V \mapsto \text{span}(V),$$

is an $O(r)$ -*quotient*: two frames give the same subspace iff they differ by a right orthogonal factor.

Ellipsoidal Projection Metric. Following Definition 7, for $\mathcal{V}, \bar{\mathcal{V}} \in \text{Gr}(d, r)$,

$$\varrho_{\text{proj}, \Sigma}(\mathcal{V}, \bar{\mathcal{V}}) := \|\Sigma^{\frac{1}{2}}(\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}})\|_{\text{op}}, \quad (\text{D.7})$$

where $\mathcal{P}_{\mathcal{V}} := VV^{\top}$ for *any* V such that $\text{span}(V) = \mathcal{V}$ (similarly $\mathcal{P}_{\bar{\mathcal{V}}}$). Because $\mathcal{P}_{\mathcal{V}}$ is unique for each subspace, $\varrho_{\text{proj}, \Sigma}$ is well defined (independent of the chosen V). The metric can be pulled back to $\text{St}(d, r)$:

$$\varrho_{\text{proj}, \Sigma}(V, \bar{V}) := \varrho_{\text{proj}, \Sigma}(\text{span}(V), \text{span}(\bar{V})) = \|\Sigma^{\frac{1}{2}}(VV^{\top} - \bar{V}\bar{V}^{\top})\|_{\text{op}}. \quad (\text{D.8})$$

⁴By elementary linear algebra, the matrix definition of the orthogonal projector $\mathcal{P}$ here coincides with the ℓ_2 -projection characterized in Lemma 19; thus the notation is consistent.

D.2 Principal Angles between Subspaces

We study how metrics and angles between images $\mathcal{V}$ and $\bar{\mathcal{V}}$ affect their spectral properties. We introduce principal angles and the cosine–sine (CS) decomposition—standard tools for analyzing subspaces (see, e.g., Chapter 6.4.3 in [Golub and Van Loan, 2013]).

Principle Angles and Cosine-Sine representation. Let U and $\bar{U}$ be two $d \times d$ orthogonal matrix, and V and $\bar{V}$ be the first r columns of U and $\bar{U}$, respectively. We are interested in studying the metrics and angles between r –dimensional subspaces $\mathcal{V} = \text{span}(V)$ and $\bar{\mathcal{V}} = \text{span}(\bar{V})$. Formally, denote

$$U, \bar{U} \in O(d), \quad U = [V \ V_\perp], \quad \bar{U} = [\bar{V} \ \bar{V}_\perp],$$

where

$$V, \bar{V} \in \mathbb{R}^{d \times r}, \quad V^\top V = I_r, \quad \bar{V}^\top \bar{V} = I_r,$$

and

$$V_\perp, \bar{V}_\perp \in \mathbb{R}^{d \times (d-r)}, \quad V_\perp^\top V_\perp = I_{d-r}, \quad \bar{V}_\perp^\top \bar{V}_\perp = I_{d-r}.$$

Since $U, \bar{U} \in O(d)$, their product $U^\top \bar{U}$ is itself orthogonal. Writing

$$U^\top \bar{U} = \begin{pmatrix} V^\top \\ V_\perp^\top \end{pmatrix} [\bar{V} \ \bar{V}_\perp] = \begin{pmatrix} V^\top \bar{V} & V^\top \bar{V}_\perp \\ V_\perp^\top \bar{V} & V_\perp^\top \bar{V}_\perp \end{pmatrix},$$

define the four blocks

$$\underbrace{C}_{r \times r} = V^\top \bar{V}, \quad \underbrace{C_\perp}_{r \times (d-r)} = V^\top \bar{V}_\perp, \quad (\text{D.9})$$

$$\underbrace{S}_{(d-r) \times r} = V_\perp^\top \bar{V}, \quad \underbrace{S_\perp}_{(d-r) \times (d-r)} = V_\perp^\top \bar{V}_\perp. \quad (\text{D.10})$$

Thus

$$U^\top \bar{U} = \begin{pmatrix} C & C_\perp \\ S & S_\perp \end{pmatrix} \in O(d).$$

Now we introduce principal angles between $\mathcal{V} = \text{span}(V)$ and $\bar{\mathcal{V}} = \text{span}(\bar{V})$ by writing

$$C = V^\top \bar{V} = Q_1 \text{diag}(\cos \theta_1, \dots, \cos \theta_r) W_1^\top, \quad Q_1, W_1 \in O(r), \quad (\text{D.11})$$

where

$$0 \leq \theta_1 \leq \theta_2 \leq \dots \leq \theta_r \leq \pi/2$$

are called the principle angles between subspaces $\mathcal{V}$ and $\bar{\mathcal{V}}$; and where $\{\cos \theta_1, \dots, \cos \theta_r\}$ are the singular values of C . Simultaneously, we have that the eigenvalues of $S, C_\perp, S_\perp$ are (notation sepc

means spectrum, the set of singular values)

$$\begin{aligned}
\text{spec}(S) &= \{\sin \theta_1, \dots, \sin \theta_{\min\{r, d-r\}}, \underbrace{0, \dots, 0}_{\max\{d-2r, 0\}}\}, \\
\text{spec}(C_\perp) &= \{\sin \theta_1, \dots, \sin \theta_{\min\{r, d-r\}}, \underbrace{0, \dots, 0}_{\max\{d-2r, 0\}}\} \\
\text{spec}(S_\perp) &= \{\cos \theta_1, \dots, \cos \theta_{\min\{r, d-r\}}, \underbrace{1, \dots, 1}_{\max\{d-2r, 0\}}\}.
\end{aligned} \tag{D.12}$$

The above representation in (D.11) and (D.12) are without loss of generality: if $r \leq d - r$, then all the four spectrum contain all r principal angles; if $r > d - r$, then only first $d - r$ principal angles $\{\theta_k\}_{k=1}^{d-r}$ can be smaller than $\pi/2$ and $\theta_k = 0$ for all $d - r + 1 \leq k \leq r$.

The cosine–sine representation of the eigenvalues in (D.11) and (D.12) motivates our notation C and S when defining block matrices in (D.9) and (D.10). This representation is an immediate consequence of the classical CS decomposition for orthogonal matrices [Paige and Wei, 1994, Golub and Van Loan, 2013], and we henceforth regard the resulting eigenvalue characterization as given.

Projection Metric via Principal Angles. For subspaces $\mathcal{V}$ and $\bar{\mathcal{V}}$, recall that for orthogonal projectors

$$\mathcal{P}_{\mathcal{V}} = VV^\top, \quad \mathcal{P}_{\bar{\mathcal{V}}} = \bar{V}\bar{V}^\top.$$

It is known that the projection metric defined in (D.7) and (D.8) are equal to $\sin \theta_r$, sine of the largest principal angle between the two subspaces. Formally, there is the fact (see, e.g., the last equation in Section 6.4.3 in [Golub and Van Loan, 2013])

$$\varrho_{\text{proj}} = \|\mathcal{P}_{\mathcal{V}} - \mathcal{P}_{\bar{\mathcal{V}}}\|_{\text{op}} = \max_{1 \leq k \leq r} \sin \theta_k = \sin \theta_r. \tag{D.13}$$

Here θ_i is the i -th principal-angle between $\mathcal{V}$ and $\bar{\mathcal{V}}$, and the spectral norm of the difference of two projectors equals the largest of these sines.

D.3 Local Charts of the Grassmannian

In differential geometry, a *chart* is a single local coordinate map. An *atlas* is the whole collection of charts that covers the manifold. We introduce a useful atlas that consists of finite graph charts, which only rely on elementary linear algebra and avoid more advanced Lie algebra and exponential map techniques in Szarek [1997].

Choose a reference subspace $\bar{\mathcal{V}} \in \text{Gr}(d, r)$ and its parameterization matrix $\bar{V} \in \text{St}(d, r)$. Denote $X \in \mathbb{R}^{(d-r) \times r}$ to be mappings from r -dimensional subspace $\bar{\mathcal{V}}$ to $(d - r)$ -dimensional subspace $\bar{\mathcal{V}}_\perp$. Every r -dimensional subspace close to $\bar{\mathcal{V}}$ can be written as the *graph*

$$\mathcal{V}(X) := \text{span}\left\{[\bar{V}\bar{V}_\perp] \begin{pmatrix} I_r \\ X \end{pmatrix}\right\}, \quad X \in \mathbb{R}^{(d-r) \times r}, \tag{D.14}$$

where $\mathcal{V}(X)$ is the subspace spanned by the columns of $[\bar{V} \bar{V}_\perp] \begin{pmatrix} I_r \\ X \end{pmatrix}$ (the matrix multiplication).

Given the reference subspace $\bar{\mathcal{V}}$, define the local *graph chart* from $\mathbb{R}^{(d-r) \times r}$ to $\text{Gr}(d, r)$ by

$$\phi_{\bar{\mathcal{V}}} : X \mapsto \mathcal{V}(X) \in \text{Gr}(d, r). \tag{D.15}$$

Note that for the $(d - r) \times r$ zero matrix (denoted as 0), we have $\phi_{\bar{\mathcal{V}}}(0) = \bar{\mathcal{V}}$.

Intuition for the graph chart. If a subspace $\mathcal{V}$ is close to $\bar{\mathcal{V}}$ —specifically, $\varrho_{\text{proj}}(\mathcal{V}, \bar{\mathcal{V}}) = \sin \theta_r < 1$ —then all principal angles between $\mathcal{V}$ and $\bar{\mathcal{V}}$ satisfy $\theta_i < \pi/2$. Equivalently, the orthogonal projection $\mathcal{P}_{\bar{\mathcal{V}}}$ restricted to $\mathcal{V}$ is a bijection $\mathcal{P}_{\bar{\mathcal{V}}}|\mathcal{V} : \mathcal{V} \rightarrow \bar{\mathcal{V}}$. In the orthonormal basis $[\bar{V} \ \bar{V}_\perp]$, this means every $v \in \mathcal{V}$ can be written uniquely as

$$v = [\bar{V} \ \bar{V}_\perp] \begin{pmatrix} \bar{v} \\ X \bar{v} \end{pmatrix}, \quad \begin{pmatrix} \bar{v} \\ 0 \end{pmatrix} \in \text{span} \left\{ \begin{pmatrix} I_r \\ 0 \end{pmatrix} \right\},$$

for a linear map $X \in \mathbb{R}^{(d-r) \times r}$. Thus, locally around $\bar{\mathcal{V}}$ (all principal angles $< \pi/2$), every r –plane admits—and is uniquely determined by—its graph parameter X . We call X the *graph parameterization* of $\mathcal{V}(X)$ in this image. This is formalized as the following lemma.

Lemma 23 (Local Bijection of Graph Chart) *Fix an orthonormal decomposition $\mathbb{R}^d = \bar{\mathcal{V}} \oplus \bar{\mathcal{V}}_\perp$ with basis $[\bar{V} \ \bar{V}_\perp]$. Then every r –dimensional subspace $\mathcal{V}$ such that $\varrho_{\text{proj}}(\mathcal{V}, \bar{\mathcal{V}}) < 1$ (i.e., all principal angles $< \pi/2$) can be written uniquely as a graph*

$$\mathcal{V} = \phi_{\bar{\mathcal{V}}}(X) = \text{span} \left\{ [\bar{V} \ \bar{V}_\perp] \begin{pmatrix} I_r \\ X \end{pmatrix} \right\}, \quad X \in \mathbb{R}^{(d-r) \times r}.$$

Proof of Lemma 23: If $V \in \text{St}(d, r)$ spans $\mathcal{V}$, block it in the $[\bar{V} \ \bar{V}_\perp]$ basis: denote

$$\begin{pmatrix} A \\ B \end{pmatrix} := \begin{pmatrix} \bar{V}^\top \\ \bar{V}_\perp^\top \end{pmatrix} V \quad (A \in \mathbb{R}^{r \times r}, B \in \mathbb{R}^{(d-r) \times r}).$$

Then by the principal angle representation (D.11), $A = \bar{V}^\top V$ is invertible iff all principal angles $< \pi/2$, and choosing

$$X = B A^{-1}$$

leads to

$$\mathcal{V} = \text{span}(V) = \text{span} \left\{ [\bar{V} \ \bar{V}_\perp] \begin{pmatrix} A \\ B \end{pmatrix} \right\} = \text{span} \left\{ [\bar{V} \ \bar{V}_\perp] \begin{pmatrix} I_r \\ X \end{pmatrix} \right\},$$

where the last equality is because for invertible A one always have $\text{span}(ZA) = \text{span}(Z)$ for any matrix Z .

We have already shown existence. For uniqueness, assuming there are two different X_1, X_2 such that $\phi_{\bar{\mathcal{V}}}(X_1) = \phi_{\bar{\mathcal{V}}}(X_2)$. Because two bases of the same r –dimensional subspace differ by an invertible change of coordinates, so there exists an invertible $r \times r$ matrix Y such that

$$[\bar{V} \ \bar{V}_\perp] \begin{pmatrix} I_r \\ X_1 \end{pmatrix} Y = [\bar{V} \ \bar{V}_\perp] \begin{pmatrix} I_r \\ X_2 \end{pmatrix},$$

which results in $Y = I_r$ and $X_1 = X_2$. Thus the parameterization X of $\mathcal{V}$ is unique. □

Sine-tangent Relationship in Graph Chart. We will show that there is a sine-tangent relationship between $\varrho_{\text{proj}}(\mathcal{V}, \bar{\mathcal{V}})$ and $\|X\|_{\text{op}}$. To be specific, we have the following lemma.

Lemma 24 (Sine-Tangent Relationship in Graph Chart) *Denote θ_r is the maximal principal angle between the subspaces $\mathcal{V}(X)$ and $\bar{\mathcal{V}}$, defined in (D.11). For the graph chart (D.15), we have*

$$\varrho_{\text{proj}}(\mathcal{V}(X), \bar{\mathcal{V}}) = \sin \theta_r, \quad \|X\|_{\text{op}} = \tan \theta_r.$$

The above relationship immediately implies that

$$\varrho_{\text{proj}}(\mathcal{V}(X), \bar{\mathcal{V}}) = \|X\|_{\text{op}} / \sqrt{1 + \|X\|_{\text{op}}^2}.$$

Proof of Lemma 24: Given the fact $\varrho_{\text{proj}}(\mathcal{V}(X), \bar{\mathcal{V}}) = \sin \theta_r$ (which is already shown in (D.13)), where θ_r is the largest principal angle between the subspaces $\mathcal{V}(X)$ and the reference subspace $\bar{\mathcal{V}}$, we want to show $\|X\|_{\text{op}} = \tan \theta_r$.

Step 1: Setup and Simplification. The projection metric is invariant under orthogonal transformations of the ambient space $\mathbb{R}^d$. We can therefore choose a coordinate system that simplifies the calculations without loss of generality. We choose a basis such that the reference frame $\bar{V}$ and its orthogonal complement $\bar{V}_\perp$ are represented as:

$$\bar{V} = \begin{pmatrix} I_r \\ 0 \end{pmatrix} \in \text{St}(d, r), \quad \bar{V}_\perp = \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} \in \text{St}(d, d-r). \quad (\text{D.16})$$

In this basis, the reference subspace is $\bar{\mathcal{V}} = \text{span}(\bar{V})$. The parameterization matrix (orthonormal basis) $V(X)$ for the subspace $\mathcal{V}(X)$ simplifies to (here $(I_r + X^\top X)^{-1/2}$ normalize $V(X)$ to be an orthogonal matrix):

$$V(X) = [\bar{V} \ \bar{V}_\perp] \begin{pmatrix} I_r \\ X \end{pmatrix} (I_r + X^\top X)^{-1/2} = I_d \begin{pmatrix} I_r \\ X \end{pmatrix} (I_r + X^\top X)^{-1/2} = \begin{pmatrix} I_r \\ X \end{pmatrix} (I_r + X^\top X)^{-1/2}, \quad (\text{D.17})$$

where the second equality follows from our choice of basis without loss of generality: the reference frame $\bar{V}$ and its complement $\bar{V}_\perp$ are represented as block identity matrices as in (D.16).

Step 2: Projection Metric and Principal Angles. A fundamental result in matrix analysis, our equation (D.11), states that the cosines of the principal angles, $\cos \theta_i$, between two subspaces spanned by orthonormal bases V and $\bar{V}$ are the singular values of $V^\top \bar{V}$. In our case, the principal angles between $\mathcal{V}(X)$ and $\bar{\mathcal{V}}$ are determined by the singular values of $V(X)^\top \bar{V}$ —which are, equivalently, the singular values of $\bar{V}^\top V(X)$.

Step 3: Calculation of $\cos \theta_i$. Let's compute the matrix product $\bar{V}^\top V(X)$ using our simplified forms:

$$\begin{aligned}\bar{V}^\top V(X) &= (I_r \ 0) \left[\begin{pmatrix} I_r \\ X \end{pmatrix} (I_r + X^\top X)^{-1/2} \right] \\ &= \left((I_r \ 0) \begin{pmatrix} I_r \\ X \end{pmatrix} \right) (I_r + X^\top X)^{-1/2} \\ &= I_r \cdot (I_r + X^\top X)^{-1/2} \\ &= (I_r + X^\top X)^{-1/2}.\end{aligned}$$

To find the singular values of this matrix, we use the Singular Value Decomposition (SVD) of X . Let $X = U\Sigma W^\top$, where $U \in \mathbb{R}^{(d-r) \times (d-r)}$ and $W \in \mathbb{R}^{r \times r}$ are orthogonal, and $\Sigma \in \mathbb{R}^{(d-r) \times r}$ is a rectangular diagonal matrix with the singular values $\lambda_1 \geq \lambda_2 \geq \dots \geq 0$ on its diagonal. The spectral norm is $\|X\|_{\text{op}} = \lambda_1$.

Then, $X^\top X = (U\Sigma W^\top)^\top (U\Sigma W^\top) = W\Sigma^\top U^\top U\Sigma W^\top = W\Sigma_r^2 W^\top$, where Σ_r^2 is the $r \times r$ diagonal matrix with entries λ_i^2 . So, the matrix $I_r + X^\top X = W(I_r + \Sigma_r^2)W^\top$. Its inverse square root is: $(I_r + X^\top X)^{-1/2} = W(I_r + \Sigma_r^2)^{-1/2}W^\top$.

The singular values of $\bar{V}^\top V(X)$ are the diagonal entries of $(I_r + \Sigma_r^2)^{-1/2}$, which are: $s_i = \frac{1}{\sqrt{1 + \lambda_i^2}}$. These singular values are the values of $\cos \theta_i$. The largest principal angle, θ_r , corresponds to the smallest cosine value. This occurs when the singular value λ_i is largest, i.e., for $\lambda_1 = \|X\|_{\text{op}}$. Thus,

$$\cos \theta_r = \frac{1}{\sqrt{1 + \|X\|_{\text{op}}^2}}.$$

Step 4: Deriving $\tan \theta_r$. Using the fundamental trigonometric identity $\sin^2 \theta + \cos^2 \theta = 1$ and the fact that principal angles lie in $[0, \pi/2)$, we have:

$$\tan \theta_r = \|X\|_{\text{op}}.$$

We have shown that for graph charts, there is the relationship $\varrho_{\text{proj}}(\mathcal{V}(X), \bar{\mathcal{V}}) = \sin \theta_r$ and $\|X\|_{\text{op}} = \tan \theta_r$. This suggests

$$\varrho_{\text{proj}}(\mathcal{V}(X), \bar{\mathcal{V}}) = \frac{\|X\|_{\text{op}}}{\sqrt{1 + \|X\|_{\text{op}}^2}}.$$

□

D.4 Global Atlas of Graph Charts

For the Grassmannian $\text{Gr}(d, r)$ we have that for all $\varepsilon > 0$, we have the coarse covering number bound $N(\text{Gr}(d, r), \varrho_{\text{proj}}, \varepsilon) \leq C^{\frac{r(d-r)}{\varepsilon}}$, where $C > 0$ is an absolute constant. This is a coarse bound—its dependence is exponential in $1/\varepsilon$ (hence not rate-optimal; the optimal dependence is polynomial)—and we use it only as a preliminary supporting estimate. This coarse estimate suggests that, a finite $O(e^{r(d-r)})$ number of graph charts are sufficient to cover the entire $\text{Gr}(d, r)$ such that every subspace $\mathcal{V} \in \text{Gr}(d, r)$ is contained in the image of a graph chart with its graph parameterization X satisfies $\|X\|_{\text{op}} \leq 1$. From this intuition, we have the following lemma.

Lemma 25 (Pointwise Dimension Consequence of Finite Global Atlas) *The uniform prior $\mu = \text{Unif}(\text{Gr}(d, r))$ satisfies that for every $\mathcal{V} \in \text{Gr}(d, r)$, every PSD matrix Σ and every $\varepsilon > 0$,*

$$\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} \leq C_1 r(d - r) + \sup_{X \in \mathcal{X}} \log \frac{1}{\text{Unif}(\bar{\mathcal{X}})\{X' \in \bar{\mathcal{X}} : \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \leq \varepsilon\}},$$

where $\mathcal{X} = \{X \in \mathbb{R}^{(d-r)r} : \|X\|_{\text{op}} \leq 1\}$ and $\bar{\mathcal{X}} = \{X \in \mathbb{R}^{(d-r)r} : \|X\|_{\text{op}} \leq 2\}$ (we make $\bar{\mathcal{X}}$ slightly larger than $\mathcal{X}$ for later technical derivation), $\text{Unif}(\bar{\mathcal{X}})\{\cdot\}$ is the uniform measure over $\bar{\mathcal{X}}$, and $C_1 > 0$ is an absolute constant.

Proof of Lemma 25: Proposition 6 in [Pajor, 1998] prove a coarse covering number bound

$$N(\text{Gr}(d, r), \varrho_{\text{proj}}, \varepsilon) \leq C^{\frac{r(d-r)}{\varepsilon}}$$

where $C > 0$ is an absolute constant; this coarse estimate is exponential rather than polynomial in ε , so it is used only for preliminary supporting purposes. For every $\mathcal{V} \in \text{Gr}(d, r)$, by the homogeneity of the Grassmannian (under the action of $O(d)$), the ϱ_{proj} -ball $B_{\text{proj}}(\mathcal{V}, \varepsilon)$ has volume independent of its center. We therefore refer to this common value as the volume of an ε - ϱ_{proj} ball, written as $\text{Vol}(\varepsilon - \varrho_{\text{proj}} \text{ ball})$. By the definition of covering number (see Definition 5 and the subsequent inequality for background), we have that

$$N(\text{Gr}(d, r), \varrho_{\text{proj}}, \varepsilon) \cdot \text{Vol}(\varepsilon - \varrho_{\text{proj}} \text{ ball}) \geq \text{Vol}(\text{Gr}(d, r)),$$

then for the uniform prior $\nu = \text{Unif}(\text{Gr}(d, r))$, we have that for every $\bar{\mathcal{V}} \in \text{Gr}(d, r)$,

$$\log \frac{1}{\nu(B_{\varrho_{\text{proj}}}(\bar{\mathcal{V}}, \varepsilon))} = \log \frac{\text{Vol}(\text{Gr}(d, r))}{\text{Vol}(\varepsilon - \varrho_{\text{proj}} \text{ ball})} \leq r(d - r) \frac{\log C}{\varepsilon}.$$

Note that ϱ_{proj} is not the target metric; our goal is the ellipsoidal metric $\varrho_{\text{proj}, \Sigma}$. Taking $\varepsilon = 1/\sqrt{2}$, we obtain:

$$\log \frac{1}{\nu(B_{\varrho_{\text{proj}}}(\bar{\mathcal{V}}, 1/\sqrt{2}))} \leq C_1 r(d - r), \quad (\text{D.18})$$

where $C_1 > 0$ is an absolute constant. By Lemma 24, we have that inside the ball $B_{\varrho_{\text{proj}}}(\bar{\mathcal{V}}, 1/\sqrt{2})$, by choosing $\bar{\mathcal{V}}$ as the reference subspace, the graph parameterization X of $\mathcal{V}$ satisfies

$$\|X\|_{\text{op}} \leq 1,$$

which follows from that if $\varrho_{\text{proj}}(\mathcal{V}(X), \bar{\mathcal{V}}) \leq 1/\sqrt{2}$ (i.e., $\sin \theta_r \leq 1/\sqrt{2}$), we have $\|X\|_{\text{op}} \leq 1$. See (D.14) for the definition of this graph chart parameterization; the existence and uniqueness of the parameterization X is by Lemma 23 (local bijection of graph chart). Furthermore, again by Lemma 23 and Lemma 24, $\mathcal{X} = \{X \in \mathbb{R}^{(d-r)r} : \|X\|_{\text{op}} \leq 1\}$ satisfies ($\cong$ means isomorphism/bijection)

$$B_{\varrho_{\text{proj}}}(\bar{\mathcal{V}}, 1/\sqrt{2}) \cong \mathcal{X} \subset \bar{\mathcal{X}} \cong B_{\varrho_{\text{proj}}}(\bar{\mathcal{V}}, 2/\sqrt{5}). \quad (\text{D.19})$$

Let

$$\mu_{\bar{\mathcal{V}}} = \text{Unif}(B_{\text{proj}}(\bar{\mathcal{V}}, 2/\sqrt{5})), \quad \mu(\mathcal{V}) = \int \nu(\bar{\mathcal{V}}) \mu_{\bar{\mathcal{V}}}(\mathcal{V}) d\bar{\mathcal{V}} = \text{Unif}(\text{Gr}(d, r)).$$

Then we have

$$\begin{aligned}
\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} &= \log \frac{1}{\int \nu(\bar{\mathcal{V}}) \mu_{\bar{\mathcal{V}}}(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon)) d\bar{\mathcal{V}}} \\
&= \log \frac{1}{\int \nu(\bar{\mathcal{V}}) \mu_{\bar{\mathcal{V}}}(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon) \cap B_{\text{proj}}(\bar{\mathcal{V}}, 2/\sqrt{5})) d\bar{\mathcal{V}}} \\
&\leq \log \frac{1}{\nu(B_{\varrho_{\text{proj}}}(\mathcal{V}, 1/\sqrt{2})) \min_{\bar{\mathcal{V}} \in B_{\varrho_{\text{proj}}}(\mathcal{V}, 1/\sqrt{2})} \mu_{\bar{\mathcal{V}}}(X' \in \bar{\mathcal{X}} : \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \leq \varepsilon)} \\
&\leq C_1 r(d-r) + \sup_{X \in \mathcal{X}} \log \frac{1}{\text{Unif}(\bar{\mathcal{X}})\{X' \in \bar{\mathcal{X}} : \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \leq \varepsilon\}},
\end{aligned}$$

where the first inequality is by restricting $\bar{\mathcal{V}}$ to $B_{\varrho_{\text{proj}}}(\mathcal{V}, 1/\sqrt{2})$; and the second inequality is by (D.18) as well as the bijection stated in (D.19) and Lemma 23. Note that we use different radius here than in $\mu_{\bar{\mathcal{V}}}$ to ensure that the set $\bar{\mathcal{X}}$ for X' , which is inside the uniform distribution in the final bound, to be larger than the domain $\mathcal{X}$ for X to take sup. This will help later technical derivation. $\square$

D.5 Decomposition and Lipschitz Properties inside Graph Chart

We apply a non-perturbative analysis to the ellipsoidal projection metric.

Lemma 26 (Non-Perturbative Decomposition of Projector Difference) *Let $X, X' \in \mathbb{R}^{(d-r) \times r}$ be two matrices. Given any reference subspace $\bar{\mathcal{V}}$, consider the graph chart $\phi_{\bar{\mathcal{V}}} : X \mapsto \mathcal{V}(X)$ defined in (D.14). Then the difference between two projectors $\mathcal{P}_{\mathcal{V}(X)}$, $\mathcal{P}_{\mathcal{V}(X')}$ be decomposed as follows:*

$$\begin{aligned}
&\mathcal{P}_{\mathcal{V}(X)} - \mathcal{P}_{\mathcal{V}(X')} \\
&= \mathcal{P}_{\mathcal{V}(X)_{\perp}} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') \begin{pmatrix} I_r & 0 \end{pmatrix} \mathcal{P}_{\mathcal{V}(X')} + \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^{\top} - X'^{\top}) \begin{pmatrix} 0 & I_{d-r} \end{pmatrix} \mathcal{P}_{\mathcal{V}(X')_{\perp}}.
\end{aligned}$$

Proof of Lemma 26: The projector is invariant under orthogonal transformations of the ambient space $\mathbb{R}^d$. We can therefore choose a coordinate system that simplifies the calculations without loss of generality. By the matrix representation (D.17) (which, without loss of generality, uses a convenient orthogonal basis specified by (D.16)), we denote

$$A(X) = \begin{pmatrix} I_r \\ X \end{pmatrix}, \quad M(X) = (I_r + X^{\top} X)^{-1},$$

and have the following facts:

$$\begin{aligned}
V(X) &= A(X)M(X)^{1/2}, \\
\mathcal{P}_{\mathcal{V}(X)} &= A(X)M(X)A(X)^{\top} = A(X)M(X) \begin{pmatrix} I_r & X^{\top} \end{pmatrix}
\end{aligned} \tag{D.20}$$

$$\begin{aligned}
\mathcal{P}_{\mathcal{V}(X)} - \mathcal{P}_{\mathcal{V}(X')} &= A(X)M(X)A(X)^{\top} - A(X')M(X')A(X')^{\top} \\
A(X)M(X) &= \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix}
\end{aligned} \tag{D.21}$$

$$A(X)M(X)X^{\top} = \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix}, \tag{D.22}$$

where (D.21) and (D.22) are straightforward consequences of (D.20).

We begin with a non-perturbative decomposition:

$$\begin{aligned}
& \mathcal{P}_{\mathcal{V}(X)} - \mathcal{P}_{\mathcal{V}(X')} \\
&= A(X)M(X)A(X)^\top - A(X')M(X')A(X')^\top \\
&= (A(X) - A(X'))M(X')A(X')^\top + A(X)(M(X) - M(X'))A(X')^\top + A(X)M(X)(A(X) - A(X'))^\top.
\end{aligned} \tag{D.23}$$

We continue to decompose each term non-perturbatively. First,

$$\begin{aligned}
& (A(X) - A(X'))M(X')A(X')^\top \\
&= \begin{pmatrix} 0 \\ X - X' \end{pmatrix} M(X')A(X')^\top \\
&= \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X')M(X')A(X')^\top \\
&= \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') \begin{pmatrix} I_r & 0 \end{pmatrix} \mathcal{P}_{\mathcal{V}(X')},
\end{aligned} \tag{D.24}$$

where the last equality uses the fact (D.21) and symmetry of $\mathcal{P}_{\mathcal{V}(X)}$.

Second, because we have the non-perturbative decomposition

$$\begin{aligned}
& M(X) - M(X') \\
&= (I_r + X^\top X)^{-1} \left((I_r + X'^\top X') - (I_r + X^\top X) \right) (I_r + X'^\top X')^{-1} \\
&= (I_r + X^\top X)^{-1} \left(X'^\top X' - X^\top X \right) (I_r + X'^\top X')^{-1} \\
&= (I_r + X^\top X)^{-1} \left(X^\top (X' - X) + (X'^\top - X^\top)X' \right) (I_r + X'^\top X')^{-1} \\
&= M(X)X^\top (X' - X)M(X') + M(X)(X'^\top - X^\top)X'M(X'),
\end{aligned}$$

we have

$$\begin{aligned}
& A(X)(M(X) - M(X'))A(X')^\top \\
&= A(X)M(X)X^\top (X' - X)M(X')A(X')^\top + A(X)M(X)(X'^\top - X^\top)X'M(X')A(X')^\top \\
&= -\mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') \begin{pmatrix} I_r & 0 \end{pmatrix} \mathcal{P}_{\mathcal{V}(X')} - \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) \begin{pmatrix} 0 & I_{d-r} \end{pmatrix} \mathcal{P}_{\mathcal{V}(X')},
\end{aligned} \tag{D.25}$$

where the last equality uses the fact (D.21) and the fact (D.22).

Third, we have

$$\begin{aligned}
& A(X)M(X)(A(X) - A(X'))^\top \\
&= A(X)M(X) \begin{pmatrix} 0 & X^\top - X'^\top \end{pmatrix} \\
&= \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) \begin{pmatrix} 0 & I_{d-r} \end{pmatrix},
\end{aligned} \tag{D.26}$$

where the last equality uses the fact (D.21).

Substituting (D.24), (D.25), (D.26) back into (D.23), we have

$$\begin{aligned}
& \mathcal{P}_{\mathcal{V}(X)} - \mathcal{P}_{\mathcal{V}(X')} \\
&= \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') (I_r \ 0) \mathcal{P}_{\mathcal{V}(X')} \\
&\quad - \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') (I_r \ 0) \mathcal{P}_{\mathcal{V}(X')} - \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) (0 \ I_{d-r}) \mathcal{P}_{\mathcal{V}(X')} \\
&\quad + \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) (0 \ I_{d-r}) \\
&= \mathcal{P}_{\mathcal{V}(X)^\perp} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') (I_r \ 0) \mathcal{P}_{\mathcal{V}(X')} + \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) (0 \ I_{d-r}) \mathcal{P}_{\mathcal{V}(X')^\perp},
\end{aligned}$$

where the last equality uses $I_d - \mathcal{P}_{\mathcal{V}(X)} = \mathcal{P}_{\mathcal{V}(X)^\perp}$ and $I_d - \mathcal{P}_{\mathcal{V}(X')} = \mathcal{P}_{\mathcal{V}(X')^\perp}$. $\square$

Building upon the non-perturbative decomposition in Lemma 26, we have the following Lipschitz property of graph chart.

Lemma 27 (Lipschitz of Graph Chart) *Let $X, X' \in \mathbb{R}^{(d-r) \times r}$ be two matrices. Given any reference subspace $\bar{\mathcal{V}}$, consider the graph chart defined in (D.17). Then the ellipsoidal projection metric is Lipschitz to ellipsoidal spectral metrics as follows: for every rank- r PSD $\Sigma \in \mathbb{R}^{d \times d}$,*

$$\begin{aligned}
& \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \\
&\leq \left\| \left((0 \ I_{d-r}) \mathcal{P}_{\mathcal{V}(X)^\perp}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)^\perp} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} \right)^{\frac{1}{2}} (X - X') \right\|_{\text{op}} + \left\| \left((I_r \ 0) \mathcal{P}_{\mathcal{V}(X)}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} \right)^{\frac{1}{2}} (X^\top - X'^\top) \right\|_{\text{op}}.
\end{aligned}$$

Proof of Lemma 27: By Lemma 26, we have

$$\begin{aligned}
& \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) = \left\| \Sigma^{\frac{1}{2}} (\mathcal{P}_{\mathcal{V}(X)} - \mathcal{P}_{\mathcal{V}(X')}) \right\|_{\text{op}} \\
&= \left\| \Sigma^{\frac{1}{2}} \mathcal{P}_{\mathcal{V}(X)^\perp} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') (I_r \ 0) \mathcal{P}_{\mathcal{V}(X')} + \Sigma^{\frac{1}{2}} \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) (0 \ I_{d-r}) \mathcal{P}_{\mathcal{V}(X')^\perp} \right\|_{\text{op}} \\
&\leq \left\| \Sigma^{\frac{1}{2}} \mathcal{P}_{\mathcal{V}(X)^\perp} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} (X - X') \right\|_{\text{op}} + \left\| \Sigma^{\frac{1}{2}} \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} (X^\top - X'^\top) \right\|_{\text{op}} \\
&= \left\| \left((0 \ I_{d-r}) \mathcal{P}_{\mathcal{V}(X)^\perp}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)^\perp} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix} \right)^{\frac{1}{2}} (X - X') \right\|_{\text{op}} + \left\| \left((I_r \ 0) \mathcal{P}_{\mathcal{V}(X)}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix} \right)^{\frac{1}{2}} (X^\top - X'^\top) \right\|_{\text{op}}.
\end{aligned}$$

where the inequality follows from the triangle inequality and the facts that the spectral norms of $\mathcal{P}_{\mathcal{V}(X')}$, $\mathcal{P}_{\mathcal{V}(X')^\perp}$, and the two block-identity matrices are all at most 1 (the fact that spectral norms of projectors are at most 1 can be proved via the first inequality in Lemma 20); and the last equality is because for any matrices A, B we have

$$\|\Sigma^{\frac{1}{2}} AB\|_{\text{op}} = \sqrt{\|B^\top A^\top \Sigma AB\|_{\text{op}}} = \|(A^\top \Sigma A)^{\frac{1}{2}} B\|_{\text{op}}.$$

□

We continue to present the following lemma, which implies that the projectors and the block-identity matrices in Lemma 27 only reduces the effective dimensions of the ellipsoidal map, and does not increase the eigenvalues (up to absolute constants).

Lemma 28 (Spectral domination under contractions) *Let $\Sigma \succeq 0$ be a $d \times d$ PSD matrix with ordered eigenvalues $\lambda_1(\Sigma) \geq \dots \geq \lambda_d(\Sigma)$. Let $A \in \mathbb{R}^{d \times m}$ for some $m \leq d$ and write $s := \|A\|_{\text{op}}$. Denote by $\mu_1 \geq \dots \geq \mu_m$ the eigenvalues of $A^\top \Sigma A$. Then, for every $k = 1, \dots, m$,*

$$\mu_m \leq s^2 \lambda_m(\Sigma).$$

Proof of Lemma 28: By the Courant–Fischer–Weyl max-min characterization (see, e.g., [Wikipedia contributors, 2025b]), we have

$$\begin{aligned} \lambda_k(A^\top \Sigma A) &= \min_{\substack{S \subset \mathbb{R}^d \\ \dim S = d-k+1}} \sup\{\|A^\top \Sigma^{\frac{1}{2}} x\|_2^2 : x \in S, \|x\|_2 = 1\} \\ &\leq s^2 \cdot \min_{\substack{S \subset \mathbb{R}^d \\ \dim S = d-k+1}} \sup\{\|\Sigma^{1/2} x\|_2^2 : x \in S, \|x\|_2 = 1\} \\ &= s^2 \lambda_k(\Sigma). \end{aligned}$$

□

D.6 Proof of the Main Result

From Lemma 25, to cover $\text{Gr}(d, r)$ it suffices to cover the unit ball of $(d-r) \times r$ matrices under the ellipsoidal spectral metric. We are now ready to prove Lemma 3, the main result for ellipsoidal Grassmannian covering.

Proof of Lemma 3: We present the proof in multiple parts.

Part 1: Applying Lemma 25. Define $\mathcal{X} = \{X \in \mathbb{R}^{(d-r) \times r} : \|X\|_{\text{op}} \leq 1\}$ and further $\bar{\mathcal{X}} = \{X \in \mathbb{R}^{(d-r) \times r} : \|X\|_{\text{op}} \leq 2\}$. By Lemma 25 (Pointwise Dimension Consequence of Finite Global Atlas), for $\mu = \text{Unif}(\text{Gr}(d, r))$, we have that for all $\mathcal{V} \in \text{Gr}(d, r)$ and all $\varepsilon > 0$,

$$\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} \leq C_1 r(d-r) + \sup_{X \in \mathcal{X}} \log \frac{1}{\text{Unif}(\bar{\mathcal{X}})\{X' \in \bar{\mathcal{X}} : \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \leq \varepsilon\}}, \quad (\text{D.27})$$

where $C_1 > 0$ is an absolute constant.

Define the $(d-r) \times (d-r)$ positive definite matrices $H_1(X)$ and the $r \times r$ positive definite matrix $H_2(X)$ as the following

$$\begin{aligned} H_1(X) &= \begin{pmatrix} 0 & I_{d-r} \end{pmatrix} \mathcal{P}_{\mathcal{V}(X)^\perp}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)^\perp} \begin{pmatrix} 0 \\ I_{d-r} \end{pmatrix}, \\ H_2(X) &= \begin{pmatrix} I_r & 0 \end{pmatrix} \mathcal{P}_{\mathcal{V}(X)}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)} \begin{pmatrix} I_r \\ 0 \end{pmatrix}. \end{aligned}$$

By Lemma 27 (Lipschitz of Graph Chart), we have that

$$\varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \leq \|H_1(X)^{\frac{1}{2}}(X' - X)\|_{\text{op}} + \|H_2(X)^{\frac{1}{2}}(X' - X)^\top\|_{\text{op}}.$$

Part 2: Volumetric Arguments. We analyze the log density complexity in (D.27) via volumetric arguments.

A technical step: ball inclusion via thresholding In order to compute the log density complexity with the uniform prior, one needs the operator norm ball to be included in the support of the prior. Given a PSD matrix $H \in \mathbb{R}^{m \times m}$ and an eigenvalue threshold α , assume its eigendecomposition is $H = U \text{diag}(\beta_1, \dots, \beta_m) U^\top$, define the thresholding function T_α by

$$T_\alpha(H) = U \text{diag}(\max\{\beta_1, \alpha\}, \dots, \max\{\beta_m, \alpha\}) U^\top.$$

Clearly this function only increases the metric. We further define the following two ellipsoidal metrics:

$$\begin{aligned} \varrho_1^2(X, X') &= \|(X' - X)^\top \bar{H}_1(X)(X' - X)\|_{\text{op}}, \quad \bar{H}_1(X) = T_{\varepsilon^2}(H_1(X)) \\ \varrho_2^2(X, X') &= \|(X' - X) \bar{H}_2(X)(X' - X)^\top\|_{\text{op}}, \quad \bar{H}_2(X) = T_{\varepsilon^2}(H_2(X)) \end{aligned}$$

We note that the two balls $B_{\varrho_1}(X, \varepsilon)$, $B_{\varrho_2}(X, \varepsilon)$ are contained in $\bar{\mathcal{X}}$, as we have applied the thresholding function to ensure this inclusion. For example, for the first ball, from

$$X' - X = (\bar{H}_1(X))^{-1/2} \underbrace{(\bar{H}_1(X))^{\frac{1}{2}}(X' - X)}_{\text{spectral norm} \leq \varepsilon \text{ for } X' \in B_{\varrho_1}(X, \varepsilon)},$$

we have (by using the ε estimate from the second underbraced term above, and combining it with the thresholding guarantee $\lambda_{\min}(\bar{H}_1(X)) \geq \varepsilon^2$)

$$\|X' - X\|_{\text{op}} \leq \lambda_{\min}(\bar{H}_1(X))^{-1/2} \cdot \varepsilon \leq 1,$$

which resulting in $\|X'\|_{\text{op}} \leq \|X' - X\|_{\text{op}} + \|X\|_{\text{op}} \leq 2$ and thus $B_{\varrho_1}(X, \varepsilon) \subseteq \bar{\mathcal{X}}$. Similarly, we can show $B_{\varrho_2}(X, \varepsilon) \subseteq \bar{\mathcal{X}}$. this gives us the auxiliary ball-inclusion result:

$$B_{\varrho_1 + \varrho_2}(X, \varepsilon) \subseteq B_{\varrho_1}(X, \varepsilon) \cap B_{\varrho_2}(X, \varepsilon) \subseteq B_{\varrho_1}(X, \varepsilon) \cup B_{\varrho_2}(X, \varepsilon) \subseteq \bar{\mathcal{X}}. \quad (\text{D.28})$$

Now we are ready to proceed with the main part of the proof. By Lemma 27 (Lipschitz of Graph Chart) and the fact that thresholding only increases the spectral norm, the ellipsoidal projection metric is bounded by $\varrho_1 + \varrho_2$, so for any $X \in \mathcal{X}$,

$$\begin{aligned} & \log \frac{1}{\text{Unif}(\bar{\mathcal{X}})\{X' \in \bar{\mathcal{X}} : \varrho_{\text{proj}, \Sigma}(\mathcal{V}(X), \mathcal{V}(X')) \leq \varepsilon\}} \\ & \leq \log \frac{1}{\text{Unif}(\bar{\mathcal{X}})\{X' \in \bar{\mathcal{X}} : \varrho_1(X, X') + \varrho_2(X, X') \leq \varepsilon\}} \\ & = \log \frac{1}{\text{Unif}(\bar{\mathcal{X}})\{B_{\varrho_1 + \varrho_2}(X, \varepsilon)\}} \end{aligned} \quad (\text{D.29})$$

$$= \frac{\text{Vol}(\bar{\mathcal{X}})}{\text{Vol}(B_{\varrho_1 + \varrho_2}(X, \varepsilon))}, \quad (\text{D.30})$$

where the first equality uses the ball-inclusion result (D.28).

Covering number in normed vector space. Definition 5 is stated for a general metric space. In the special case of a normed vector space (see, e.g., [Wikipedia contributors \[2025d\]](#)), however, the covering number admits a more explicit characterization via standard volume-ratio arguments, up to absolute-constant factors in the radius. This simplification arises from the strong homogeneity of the metric induced by the norm, in particular the absolute homogeneity property

$$\varrho(\lambda y, \lambda z) = |\lambda| \varrho(y, z), \quad \forall \lambda \in \mathbb{R}. \quad (\text{D.31})$$

Classical volume-ratio arguments give the following results on the covering number of balls in general normed vector space $\mathcal{Y}$. For a p -dimensional normed vector space equipped with the metric associated to its norm $\|\cdot\|$, we denote by $B(y, R)$ the ball in $\mathcal{Y}$ centered at $y \in \mathcal{Y}$ with radius R , and by $N(\mathcal{Z}, \|\cdot\|, \varepsilon)$ the covering number of a subset $\mathcal{Z} \subseteq \mathcal{Y}$.

Proposition 4.2.10 in [Vershynin \[2018\]](#) (the proof is elementary and clearly holds true for general metric in a normed vector space) states that for $\mathcal{Z} \subseteq \mathcal{Y}$ and general metric $\|\cdot\|$, we have that for any $y \in \mathcal{Y}$,

$$\frac{\text{Vol}(\mathcal{Z})}{\text{Vol}(B(y, \varepsilon))} \leq N(\mathcal{Z}, \|\cdot\|, \varepsilon) \leq \frac{\text{Vol}(\mathcal{Z} + B(y, \frac{\varepsilon}{2}))}{\text{Vol}(B(y, \frac{\varepsilon}{2}))},$$

where the set $\mathcal{A} + \mathcal{B} := \{a + b : a \in \mathcal{A}, b \in \mathcal{B}\}$. When $\mathcal{Z}$ is convex and $B(y, \varepsilon) \subseteq \mathcal{Z}$, we further have

$$\frac{\text{Vol}(\mathcal{Z})}{\text{Vol}(B(y, \varepsilon))} \leq N(\mathcal{Z}, \|\cdot\|, \varepsilon) \leq \frac{\text{Vol}(\mathcal{Z} + B(y, \frac{\varepsilon}{2}))}{\text{Vol}(B(y, \frac{\varepsilon}{2}))} \leq \frac{\text{Vol}(\frac{3}{2}\mathcal{Z})}{\text{Vol}(B(y, \frac{\varepsilon}{2}))} = 3^p \frac{\text{Vol}(\mathcal{Z})}{\text{Vol}(B(y, \varepsilon))}, \quad (\text{D.32})$$

where $\lambda\mathcal{A} := \{\lambda a : a \in \mathcal{A}\}$ for $\lambda > 0$. Lastly, when the normed space $\mathcal{Y}$ is p -dimensional, for every $\varepsilon \in (0, R]$, setting $\mathcal{Z} = B(0, R)$ turns the above inequality (D.32) into the optimal covering number bound

$$\left(\frac{R}{\varepsilon}\right)^p \leq N(B(0, R), \|\cdot\|, \varepsilon) \leq \left(\frac{3R}{\varepsilon}\right)^p. \quad (\text{D.33})$$

Note that this result is for general normed space, not only for the ℓ_2 norm in Euclidean space (see, e.g., display (1) in [Pajor \[1998\]](#); see also [Milman and Schechtman \[1986\]](#), [Pisier \[1999\]](#)).

A technical step—lifting to product space. Consider the product space $\mathbb{R}^{(d-r) \times r} \times \mathbb{R}^{(d-r) \times r}$ (of dimension $2 \times (d-r) \times r$). Given any $(d-r) \times (d-r)$ positive definite matrix H_1 and $r \times r$ positive definite matrix H_2 , define the modified spectral norm by

$$\|(X_1, X_2) - (X'_1, X'_2)\|_{\text{op}, H_1, H_2} := \|H_1^{\frac{1}{2}}(X_1 - X'_1)\|_{\text{op}} + \|H_2^{\frac{1}{2}}(X_2^\top - X'^{\top}_2)\|_{\text{op}}.$$

Consider the constrained set

$$\mathcal{S} := \{(X_1, X_2) \in \mathbb{R}^{(d-r) \times r} \times \mathbb{R}^{(d-r) \times r} : X_1 = X_2\} = \{(X, X) : X \in \mathbb{R}^{(d-r) \times r}\},$$

which is a normed space with dimension $(d-r) \times r$ (isomorphic to $\mathbb{R}^{(d-r) \times r}$), equipped with the modified spectral norm

$$\|(X, X) - (X', X')\|_{\text{op}, H_1, H_2} = \|H_1^{\frac{1}{2}}(X - X')\|_{\text{op}} + \|H_2^{\frac{1}{2}}(X^\top - X'^\top)\|_{\text{op}}.$$

Denote $B_{\text{op}, H_1, H_2}^{\mathcal{S}}((X, X), R) = \{(X', X') \in \mathcal{S} : \|(X', X') - (X, X)\|_{\text{op}, H_1, H_2} \leq R\}$ (the ball constrained in $\mathcal{S}$). Because there is a bijective, distance-preserving (isometric) map between $B_{\varrho_1 + \varrho_2}(X, \varepsilon)$ and $B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}^{\mathcal{S}}((X, X), \varepsilon)$, and likewise $B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4)$ and $\bar{\mathcal{X}}$ (here 0 denotes the $(d-r) \times r$ 0 matrix), we obtain

$$\frac{\text{Vol}(\bar{\mathcal{X}})}{\text{Vol}(B_{\varrho_1 + \varrho_2}(X, \varepsilon))} = \frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4))}{\text{Vol}(B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}^{\mathcal{S}}((X, X), \varepsilon))}, \quad (\text{D.34})$$

where the volume on $\mathcal{S}$ is defined via the surface area measure. (D.34) is exactly the objective we need to bound in (D.29).

Given $\varepsilon > 0$, by the property (D.32) of covering number, we have that for every $X \in \mathcal{X}$ and $\varepsilon > 0$,

$$\frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4))}{\text{Vol}(B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}^{\mathcal{S}}((X, X), \varepsilon))} \leq N(B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon). \quad (\text{D.35})$$

Remark on why lifting to product space double the degree of freedom. We now lift the $\mathcal{S}$ -constrained ball $B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4)$ to the product space $\bar{\mathcal{X}} \times \bar{\mathcal{X}}$, using the covering number of the lifted product space to bound the covering number of the original space, in order to obtain an upper bound on (D.35) and (D.34). This is the reason why our final bound will scale (in the isotropic case) in the order $O((d-r)r \log \frac{1}{\varepsilon^2}) = O(2(d-r)r \log \frac{1}{\varepsilon})$ rather than the classical optimal order $\Theta((d-r)r \log \frac{1}{\varepsilon})$ —the lifting to product space increase the number of freedom by a multiplicative factor of 2. Nevertheless, such difference is negligible in our theory.

For every $(X_1, X_2) \in \mathbb{R}^{(d-r) \times r} \times \mathbb{R}^{(d-r) \times r}$, every $(d-r) \times (d-r)$ matrix $H_1 \succ 0$, and every $r \times r$ matrix $H_2 \succ 0$, and radius R , denote $B_{\text{op}, H_1, H_2}((X_1, X_2), R)$ to be the unconstrained ball in $\mathbb{R}^{(d-r) \times r} \times \mathbb{R}^{(d-r) \times r}$:

$$B_{\text{op}, H_1, H_2}((X_1, X_2), R) := \{(X'_1, X'_2) \in \mathbb{R}^{(d-r) \times r} \times \mathbb{R}^{(d-r) \times r} : \|(X_1, X_2) - (X'_1, X'_2)\|_{\text{op}, H_1, H_2} \leq R\}.$$

Lifting to the product space can only increase the external covering number (monotonicity under set inclusion), and the external covering number is equivalent to the internal covering number up to a constant factor in the radius. To be specific, by Lemma 15, we have

$$\begin{aligned} & N(B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon) \\ & \leq N_{\text{ext}}(B_{\text{op}, I_{d-r}, I_r}^{\mathcal{S}}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon/2) \\ & \leq N_{\text{ext}}(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon/2) \\ & \leq N(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon/2). \end{aligned} \quad (\text{D.36})$$

For every $X \in \mathcal{X}$, the ball-inclusion argument (D.28) is strong enough to imply that the unconstrained ball $B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon) \subseteq \mathbb{R}^{(d-r) \times r} \times \mathbb{R}^{(d-r) \times r}$ is also included in the lifted ball $B_{\text{op}, I_{d-r}, I_r}((0, 0), 4)$, which gives that

$$B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon/2) \subset B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon) \subseteq B_{\text{op}, I_{d-r}, I_r}((0, 0), 4).$$

This satisfies the inclusion condition required to establish (D.32), and we have

$$\begin{aligned} N(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon/2) &\leq 3^{2(d-r)r} \frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4))}{\text{Vol}(B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon/2))} \\ &= 6^{2(d-r)r} \frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4))}{\text{Vol}(B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon))}. \end{aligned} \quad (\text{D.37})$$

Part 3: Applying Change of Variable and Calculating the Jacobian Determinant.
Applying the standard change of variables

$$Y_1 = \bar{H}_1(X)^{1/2} X_1, \quad Y_2 = X_2 \bar{H}_2(X)^{1/2},$$

the map on vectorized variables is

$$\text{vec}(Y_1) = (I_r \otimes \bar{H}_1(X)^{1/2}) \text{vec}(X_1), \quad \text{vec}(Y_2) = (\bar{H}_2(X)^{\top 1/2} \otimes I_{d-r}) \text{vec}(X_2),$$

and the total Jacobian is

$$J(X) = \begin{pmatrix} I_r \otimes \bar{H}_1(X)^{1/2} & 0 \\ 0 & \bar{H}_2(X)^{\top 1/2} \otimes I_{d-r} \end{pmatrix}.$$

The two block-diagonal Jacobian determinants are

$$\begin{aligned} \left| \det(I_r \otimes \bar{H}_1(X)^{1/2}) \right| &= (\det \bar{H}_1(X)^{1/2})^r = \det(\bar{H}_1(X))^{r/2}, \\ \left| \det(\bar{H}_2(X)^{\top 1/2} \otimes I_{d-r}) \right| &= (\det \bar{H}_2(X)^{1/2})^{d-r} = \det(\bar{H}_2(X))^{(d-r)/2}. \end{aligned}$$

Multiplying the two factors, the total Jacobian of the linear change of variables is

$$\det(J(X)) = \det(\bar{H}_1(X))^{r/2} \det(\bar{H}_2(X))^{(d-r)/2}.$$

(We used $\det(B^\top) = \det(B)$ and that $\bar{H}_1(X), \bar{H}_2(X) \succ 0$, so determinants are positive.) By the change of variable formula in integration (see, e.g., [Wikipedia contributors \[2025a\]](#)), we have

$$\begin{aligned} &\text{Vol}(B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon)) \\ &= \text{Vol}(B_{\text{op}, I_{d-1}, I_r}((X, X), \varepsilon)) (\det(J(X)))^{-1} \\ &= \text{Vol}(B_{\text{op}, I_{d-1}, I_r}((X, X), \varepsilon)) \prod_{k=1}^{d-r} \lambda_k(\bar{H}_1(X))^{-r/2} \prod_{k=1}^r \lambda_k(\bar{H}_2(X))^{-(d-r)/2}, \end{aligned}$$

which implies

$$\frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4))}{\text{Vol}(B_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}((X, X), \varepsilon))} = \prod_{k=1}^{d-r} \lambda_k(\bar{H}_1(X))^{r/2} \prod_{k=1}^r \lambda_k(\bar{H}_2(X))^{(d-r)/2} \frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4))}{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((X, X), \varepsilon))}. \quad (\text{D.38})$$

Part 4: Proving the Final Bound. For all $X \in \mathcal{X}$ and $\varepsilon \leq 1$, we have that $B_{\text{op}, I_{d-r}, I_r}((X, X), \varepsilon) \subseteq B_{\text{op}, I_{d-r}, I_r}((0, 0), 4)$ and thus by (D.32) and (D.33), we have

$$\frac{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4))}{\text{Vol}(B_{\text{op}, I_{d-r}, I_r}((X, X), \varepsilon))} \leq \left(\frac{12}{\varepsilon}\right)^{2(d-r)r}. \quad (\text{D.39})$$

Combining the above inequality (D.39) with (D.37) and (D.38), we have

$$\begin{aligned} & \log N(B_{\text{op}, I_{d-r}, I_r}((0, 0), 4), \|\cdot\|_{\text{op}, \bar{H}_1(X), \bar{H}_2(X)}, \varepsilon/2) \\ & \leq 2(d-r)r \log \frac{72}{\varepsilon} + \frac{r}{2} \sum_{k=1}^{d-r} \log \lambda_k(\bar{H}_1(X)) + \frac{d-r}{2} \sum_{k=1}^r \log \lambda_k(\bar{H}_2(X)) \\ & = \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{72^2 \lambda_k(\bar{H}_1(X))}{\varepsilon^2} + \frac{d-r}{2} \sum_{k=1}^r \log \frac{72^2 \lambda_k(\bar{H}_2(X))}{\varepsilon^2}. \end{aligned} \quad (\text{D.40})$$

Combining the above inequality (D.40) with (D.34), (D.35) and (D.36), we have that for all $X \in \mathcal{X}$,

$$\log \frac{\text{Vol}(\bar{\mathcal{X}})}{\text{Vol}(B_{\varrho_1 + \varrho_2}(X, \varepsilon))} \leq \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{72^2 \lambda_k(\bar{H}_1(X))}{\varepsilon^2} + \frac{d-r}{2} \sum_{k=1}^r \log \frac{72^2 \lambda_k(\bar{H}_2(X))}{\varepsilon^2}. \quad (\text{D.41})$$

Finally, combine the above inequality (D.41) with (D.27) and (D.29), we prove that for $\mu = \text{Unif}(\text{Gr}(d, r))$, we have that for all $\mathcal{V} \in \text{Gr}(d, r)$ and all $\varepsilon > 0$,

$$\begin{aligned} \log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} & \leq C_1 r(d-r) + \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{72^2 \lambda_k(\bar{H}_1(X))}{\varepsilon^2} + \frac{d-r}{2} \sum_{k=1}^r \log \frac{72^2 \lambda_k(\bar{H}_2(X))}{\varepsilon^2} \\ & = \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{C \lambda_k(\bar{H}_1(X))}{\varepsilon^2} + \frac{d-r}{2} \sum_{k=1}^r \log \frac{C \lambda_k(\bar{H}_2(X))}{\varepsilon^2}, \end{aligned} \quad (\text{D.42})$$

where $C > 0$ is an absolute constant.

We end the proof by applying Lemma 28 and Lemma 20: since

$$\begin{aligned} \lambda_k(H_1(X)) & \leq \lambda_k(\mathcal{P}_{\mathcal{V}(X)^\perp}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)^\perp}) \leq \lambda_k, \quad k = 1, \dots, d-r; \\ \lambda_k(H_2(X)) & \leq \lambda_k(\mathcal{P}_{\mathcal{V}(X)}^\top \Sigma \mathcal{P}_{\mathcal{V}(X)}) \leq \lambda_k, \quad k = 1, \dots, r, \end{aligned}$$

we have

$$\begin{aligned} \lambda_k(\bar{H}_1(X)) & \leq \max\{\lambda_k, \varepsilon^2\}, \quad k = 1, \dots, d-r; \\ \lambda_k(\bar{H}_2(X)) & \leq \max\{\lambda_k, \varepsilon^2\}, \quad k = 1, \dots, r. \end{aligned}$$

Substituting this bound to (D.42), we prove that for $\mu = \text{Unif}(\text{Gr}(d, r))$, we have that for all $\mathcal{V} \in \text{Gr}(d, r)$ and all $\varepsilon > 0$,

$$\log \frac{1}{\mu(B_{\varrho_{\text{proj}, \Sigma}}(\mathcal{V}, \varepsilon))} \leq \frac{r}{2} \sum_{k=1}^{d-r} \log \frac{C \max\{\lambda_k, \varepsilon^2\}}{\varepsilon^2} + \frac{d-r}{2} \sum_{k=1}^r \log \frac{C \max\{\lambda_k, \varepsilon^2\}}{\varepsilon^2},$$

where $C > 0$ is an absolute constant.

□

E Proofs for Generalization Bounds and Comparison (Section 4)

E.1 Proof of Theorem 5

Let π be the data-independent hierarchical prior of Theorem 4. For fixed S, W and scale u , define

$$G_{\text{NP}}^{S, \mathbb{P}}(W, u) = \text{blockdiag}\left(\dots, L\bar{M}_{l \rightarrow L}^2(W, u)\Gamma_{l-1}^{S, \mathbb{P}}(W) \otimes I_{d_l}, \dots\right).$$

The non-perturbative expansion on the same bounded parameter charts applies in the mixed $L_2(\mathbb{P}_S + \mathbb{P})$ norm. Its quadratic terms satisfy

$$\mathbb{P}_S \|\Delta_l f_{l-1}(W, X)\|_2^2 + \mathbb{P} \|\Delta_l f_{l-1}(W, X)\|_2^2 = \frac{1}{n} \text{Tr}\left(\Delta_l \Gamma_{l-1}^{S, \mathbb{P}}(W) \Delta_l^\top\right).$$

The local chart and metric-domination argument used in Theorem 4, followed by the β -Lipschitz loss comparison, therefore gives

$$\log \frac{1}{\pi(B_{\varrho_{(S, \mathbb{P}), \ell}}(f(W, \cdot), \eta))} \leq d_{\text{R}}^{S, \mathbb{P}}(W, c\eta/\beta).$$

Here $\varrho_{(S, \mathbb{P}), \ell}$ is the mixed empirical-population loss metric in Theorem 2. The geometric calculation applies to every finite PSD feature matrix; it does not require a finite-column realization. In particular, it applies directly to $\Gamma^{S, \mathbb{P}}$. Equivalently, with ghost-independent common transfer bounds, the quadratic tensor is the conditional mean of its ghost counterpart. This observation concerns the quadratic form only: no expectation is passed through an effective rank or a truncated spectral sum.

Substitution into Theorem 2, followed by $\eta = \beta\varepsilon$ and the same absolute-constant endpoint rescaling as below, proves the integral bound. Expanding each block by the definition of d_{eff} gives (4.1). This proves Theorem 5. $\square$

E.2 Proof of Theorem 6

Work on the event in Definition 3. Put $K_j = n\Sigma_j^{\mathbb{P}}(W)$, $P = P_j^S(W, \varepsilon)$, $Q = I - P$, and $\Delta_l = U_l - W_l$. Positivity of K_j and the two assumed subspace inequalities give

$$\text{Tr}(\Delta_l K_j \Delta_l^\top) \leq 2\kappa \text{Tr}(\Delta_l P \Gamma_j^S(W) P \Delta_l^\top) + 2b_{\text{sub}} \vartheta_j(W, \varepsilon) \|\Delta_l\|_{\mathbf{F}}^2.$$

Thus the finite-scale domination argument proved in detail in Appendix E.4 applies with $\Gamma_j^{S'}$ replaced by K_j and the empirical-ghost loss metric replaced by $\varrho_{(S, \mathbb{P}), \ell}$. On the same bounded local support, $\sum_l \|\Delta_l\|_{\mathbf{F}}^2 \leq CR_W^2$, while

$$L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \vartheta_{l-1}(W, \varepsilon) = \frac{n\varepsilon^2}{2R_W^2}.$$

The inactive contribution is therefore at most $Cn\varepsilon^2$. Choosing c_{sub} sufficiently small gives the same finite-scale ball inclusion and hence

$$\log \frac{1}{\pi(B_{\varrho_{(S, \mathbb{P}), \ell}}(f(W, \cdot), \eta))} \leq d_{\text{R}}^S(W, c_{\text{sub}}\eta/\beta).$$

Apply Theorem 2 and use a union bound with the population-subspace event. This proves (4.7). No ghost exceptional event is used in this argument. $\square$

The population condition and the conditional ghost condition below are alternative sufficient conditions. With positive ghost failure probability, averaging the latter need not imply the former, because the feature covariance on the exceptional event need not be controlled. The conditional ghost statement retains its original ζ loss term.

E.3 Ghost formulation and proof of the DNN bound

For the proofs and refinements below, retain the original conditional ghost formulation. Write $F_l^T(W) = F_l(W, X_T)$ and $\Gamma_l^T(W) = F_l^T(W)F_l^T(W)^\top$ for $T \in \{S, S'\}$, and

$$\Gamma_l^{S,S'}(W) = F_l^S(W)F_l^S(W)^\top + F_l^{S'}(W)F_l^{S'}(W)^\top.$$

Use common outer transfer bounds valid for the observed and ghost samples and their concatenation. They may be chosen to be the uniform bounds over inputs used in the main theorem. The following result needs no population feature moment assumption.

Lemma 29 (Generalization Bound for DNN; mixed empirical-ghost form) *Let the loss $\ell(f(W, x), y)$ be bounded in $[0, 1]$ and β -Lipschitz with respect to $f(W, x)$. For every $\delta \in (0, 1)$, with probability at least $1 - \delta$ over the observed sample S , uniformly over all $W \in B_{\mathbf{F}}(R)$,*

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C_1 \left(\beta \mathbb{E}_{S'} \left[\inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\mathbf{R}}^{S,S'}(W, c\varepsilon)} d\varepsilon \right\} \mid S \right] + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right), \end{aligned}$$

where $c, C_1 > 0$ are absolute constants and

$$\begin{aligned} d_{\mathbf{R}}^{S,S'}(W, \varepsilon) = & \frac{1}{2} \sum_{l=1}^L \left((d_l + d_{l-1}) \sum_{k=1}^{r_{\text{eff}}^{S,S'}[W, l]} \underbrace{\log \frac{8C_2^2 \lambda_k(\Gamma_{l-1}^{S,S'}(W))}{n\varepsilon^2}}_{\text{spectrum of inner layers } 1:l-1} \right. \\ & \left. + (d_l + d_{l-1}) r_{\text{eff}}^{S,S'}[W, l] \cdot \underbrace{\log \left(\bar{M}_{l \rightarrow L}^2(W, \varepsilon) L \max\{\|W\|_{\mathbf{F}}^2, R^2/4^n\} \right)}_{\text{spectrum of outer layers } l+1:L} + 2 \log(d_{l-1}n) \right), \end{aligned} \tag{E.1}$$

with

$$r_{\text{eff}}^{S,S'}[W, l] := r_{\text{eff}} \left(L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) \Gamma_{l-1}^{S,S'}(W), C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon \right),$$

and $C_2 > 0$ is an absolute constant.

Proof of Lemma 29. The proof consists of two steps. First, we combine Lemma 8 with the Riemannian Dimension obtained from Theorem 4, applied to the mixed sample; this yields an integral bound in terms of the mixed empirical-ghost Riemannian Dimension. Second, we expand this mixed Riemannian Dimension into the layerwise expression displayed in the theorem.

Step 1: Obtaining the Integral Bound on Generalization Gap. Let π be the data-independent hierarchical prior from Theorem 4. Lemma 8 gives, with probability at least $1 - \delta$ over the observed sample S , uniformly over all $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C_0 \left(\mathbb{E}_{S'} \left[\inf_{0 \leq \alpha \leq \sqrt{2}} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f(W, \cdot), \eta)))} d\eta} \right\} \mid S \right] + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right), \end{aligned} \quad (\text{E.2})$$

where $C_0 > 0$ is the absolute constant in Lemma 8. We now compare the mixed empirical-ghost loss metric with the mixed non-perturbative NN-surrogate metric. As presented in (3.2), we construct the mixed metric tensor for the mixed sample (S, S') ,

$$\begin{aligned} G_{\text{NP}}^{S, S'}(W, u) &:= \text{blockdiag} \left(\cdots, L\bar{M}_{l \rightarrow L}^2(W, u) \cdot \Gamma_{l-1}^{S, S'}(W) \otimes I_{d_l}, \cdots \right), \\ \varrho_{G_{\text{NP}}^{S, S'}(W, u)}(W, W')^2 &= \text{vec}(W' - W)^\top G_{\text{NP}}^{S, S'}(W, u) \text{vec}(W' - W), \end{aligned}$$

where $G_{\text{NP}}^{S, S'}$ is obtained by replacing each empirical feature Gram matrix by the mixed feature Gram matrix $\Gamma_{l-1}^{S, S'}(W)$, and $\varrho_{G_{\text{NP}}^{S, S'}(W, u)}$ is the corresponding NN-surrogate metric. By Lipschitz property of the loss function we have

$$\varrho_{(S, S'), \ell}(f(U, \cdot), f(W, \cdot)) \leq \beta \left((\mathbb{P}_S + \mathbb{P}_{S'}) \|f(U, x) - f(W, x)\|_2^2 \right)^{1/2} = \frac{\beta}{\sqrt{n}} \|F_L^{S, S'}(U) - F_L^{S, S'}(W)\|_{\mathbf{F}},$$

where $F_l^{S, S'}(W)$ is defined as

$$F_l^{S, S'}(W) := [F_l^S(W), F_l^{S'}(W)].$$

The non-perturbative expansion used in Lemma 1 gives

$$\|F_L^{S, S'}(U) - F_L^{S, S'}(W)\|_{\mathbf{F}}^2 \leq \varrho_{G_{\text{NP}}^{S, S'}(W, u)}(U, W)^2.$$

Combining the above two inequalities and taking $u = \eta/\beta$ we have

$$\varrho_{(S, S'), \ell}(f(U, \cdot), f(W, \cdot)) \leq \frac{\beta}{\sqrt{n}} \varrho_{G_{\text{NP}}^{S, S'}(W, \eta/\beta)}(U, W).$$

In particular, we have the ball inclusion

$$B_{\varrho_{G_{\text{NP}}^{S, S'}(W, \eta/\beta)}} \left(W, \frac{\sqrt{n}\eta}{\beta} \right) \subseteq B_{\varrho_{(S, S'), \ell}}(f(W, \cdot), \eta).$$

By the metric domination lemma (Lemma 18), we have the pointwise dimension bound: for every $W \in B_{\mathbf{F}}(R)$,

$$\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f(W, \cdot), \eta))} \leq \log \frac{1}{\pi(B_{\varrho_{G_{\text{NP}}^{S, S'}(W, \eta/\beta)}}(W, \sqrt{n}\eta/\beta))}.$$

Applying Theorem 4 (Riemannian Dimension Bound for DNN) to this mixed metric, we have that there exists a prior π such that uniformly over every $W \in B_{\mathbf{F}}(R)$,

$$\log \frac{1}{\pi(B_{\varrho_{(S,S')},\ell}(f(W,\cdot),\eta))} \leq d_{\mathbf{R}}^{S,S'}(W, c\eta/\beta)$$

with absolute-constant changes, where $c > 0$ is an absolute constant. Substituting this into (E.2) and changing variables $\eta = \beta\varepsilon$ gives

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C_0 \left(\beta \mathbb{E}_{S'} \left[\inf_{0 \leq \alpha \leq \sqrt{2}/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}/\beta} \sqrt{d_{\mathbf{R}}^{S,S'}(W, c\varepsilon)} d\varepsilon \right\} \mid S \right] + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right). \end{aligned}$$

The substitution $\varepsilon = \sqrt{2}u$, $\alpha = \sqrt{2}a$ changes the upper endpoint to $1/\beta$; absorb the resulting factors of $\sqrt{2}$ into the absolute constants C_1 and c . Thus,

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C_1 \left(\beta \mathbb{E}_{S'} \left[\inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\mathbf{R}}^{S,S'}(W, c\varepsilon)} d\varepsilon \right\} \mid S \right] + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} \right). \quad (\text{E.3}) \end{aligned}$$

This proves the integral part of Lemma 29.

Step 2: Obtaining the Expression of Riemannian Dimension. It remains to expand the mixed Riemannian Dimension $d_{\mathbf{R}}^{S,S'}$ by Theorem 4. By definition of the mixed feature Gram matrix,

$$\begin{aligned} d_{\mathbf{R}}^{S,S'}(W, \varepsilon) &= \sum_{l=1}^L \left((d_l + d_{l-1}) \cdot d_{\text{eff}}(L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \cdot \Gamma_{l-1}^{S,S'}(W), C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon) \right. \\ & \quad \left. + \log(d_{l-1}n) \right), \end{aligned} \quad (\text{E.4})$$

where $R = \sup_{\mathcal{W}} \|W\|_{\mathbf{F}}$ and C_2 is an absolute constant. The effective dimension is

$$\begin{aligned} & d_{\text{eff}}(L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \cdot \Gamma_{l-1}^{S,S'}(W), C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon) \\ &= \frac{1}{2} \sum_{k=1}^{r_{\text{eff}}^{S,S'}[W,l]} \log \frac{8C_2^2 \max\{\|W\|_{\mathbf{F}}^2, R^2/4^n\} L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \lambda_k(\Gamma_{l-1}^{S,S'}(W))}{n\varepsilon^2}, \end{aligned} \quad (\text{E.5})$$

where $r_{\text{eff}}^{S,S'}[W, l]$ abbreviates $r_{\text{eff}}(L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \Gamma_{l-1}^{S,S'}(W), C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}, \varepsilon)$. Combining (E.4) and (E.5) gives the displayed expression (E.1). Combining the integral upper bound (E.3) with this expression concludes the proof of Lemma 29. $\square$

E.4 Empirical-Ghost Subspace Isomorphism for the Pointwise Ellipsoidal Metric

Definition 8 (Conditional ghost subspace isomorphism) Fix constants $\kappa \geq 1$, $b_{\text{sub}} \geq 1$, and a failure level $\zeta \in [0, 1]$. We say that the observed sample S satisfies the conditional layer-wise

finite-resolution subspace isomorphism if, conditionally on S , with probability at least $1 - \zeta$ over an independent ghost sample S' , the following two estimates hold simultaneously for every feature layer $j = 0, \dots, L - 1$, every center $W \in B_{\mathbf{F}}(R)$, and every scale $\varepsilon \in (0, 1/\beta]$:

$$P_j^S(W, \varepsilon) \Gamma_j^{S'}(W) P_j^S(W, \varepsilon) \preceq \kappa P_j^S(W, \varepsilon) \Gamma_j^S(W) P_j^S(W, \varepsilon), \quad (\text{E.6})$$

$$\|Q_j^S(W, \varepsilon) \Gamma_j^{S'}(W) Q_j^S(W, \varepsilon)\|_{\text{op}} \leq b_{\text{sub}} \vartheta_j(W, \varepsilon). \quad (\text{E.7})$$

Lemma 30 (Fully empirical DNN bound under subspace isomorphism) *Assume the setting of Lemma 29. Suppose that, with probability at least $1 - \delta_{\text{iso}}$ over S , the conditional layer-wise finite-resolution subspace isomorphism in Definition 8 holds. Then there is a constant $c_{\text{sub}} > 0$, depending only on (κ, b_{sub}) , such that for every $\delta \in (0, 1)$, with probability at least $1 - \delta - \delta_{\text{iso}}$ over S , uniformly over all $W \in B_{\mathbf{F}}(R)$,*

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n) \ell(f(W, x), y) \\ & \leq C \left(\beta \inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\text{R}}^S(W, c_{\text{sub}} \varepsilon)} d\varepsilon \right\} + \sqrt{\frac{\log(2 \log(2n)/\delta)}{n}} + \zeta \right), \end{aligned} \quad (\text{E.8})$$

where $C > 0$ is an absolute constant. In particular, if $\kappa, b_{\text{sub}} = O(1)$, then c_{sub} is an absolute constant. If moreover $\zeta \leq \sqrt{\log(2 \log(2n)/\delta)}/n$, the last term is absorbed into the displayed concentration term and the leading complexity is the clean observed-sample Riemannian Dimension d_{R}^S .

This subsection proves Lemma 30. The proof uses only the finite-resolution subspace isomorphism in Definition 8. Its role is to show that the mixed empirical-ghost metric in Lemma 8 is dominated, at the finite scale used in the pointwise dimension, by the clean observed-sample ellipsoidal metric. Unlike a full Loewner isomorphism, the argument does not insert a fixed isotropic ridge into the observed spectrum.

For a sample T of size n , write X_T for its input matrix and

$$F_{l-1}^T(W) := F_{l-1}(W, X_T), \quad A_{l,T}(W, u) := L \bar{M}_{l \rightarrow L}^2(W, u) F_{l-1}^T(W) F_{l-1}^T(W)^\top.$$

The empirical non-perturbative metric tensor is

$$G_{\text{NP}}^T(W, u) := \text{blockdiag}(\dots, A_{l,T}(W, u) \otimes I_{d_l}, \dots).$$

The block-diagonal form is essential: the local-chart contribution is d_l copies of the ellipsoid generated by $A_{l,T}(W, u)$, while the Grassmannian atlas cost is governed by the same spectrum and contributes the d_{l-1} factor.

Subspace isomorphism converts the mixed metric to the empirical ellipsoid. We now prove the fully empirical statement under Definition 8. For $j = 0, \dots, L - 1$, let $P_j^S(W, \varepsilon)$ and $Q_j^S(W, \varepsilon)$ be the active and inactive projectors below (4.4). The next lemma is the precise point at which the subspace isomorphism is used.

Lemma 31 (Subspace isomorphism implies finite-scale metric domination) *Assume that the event in Definition 8 holds for a pair (S, S') . Then there is a constant $c_{\text{sub}} > 0$, depending only*

on (κ, b_{sub}) , such that the following holds. For every center $W \in B_{\mathbf{F}}(R)$ and every scale $\eta \in (0, 1]$, the empirical ellipsoidal ball at radius $c_{\text{sub}}\sqrt{n}\eta/\beta$, restricted to the same Euclidean shell used in the proof of Theorem 4, is contained in the mixed loss ball:

$$B_{\varrho_{G_{\text{NP}}^S(W, c_{\text{sub}}\eta/\beta)}}(W, c_{\text{sub}}\sqrt{n}\eta/\beta) \subseteq B_{\varrho_{(S, S'), \ell}}(f(W, \cdot), \eta). \quad (\text{E.9})$$

Proof of Lemma 31. Recall that the proof of Theorem 4 uses a dyadic peeling over the Frobenius norm scale. For the dyadic component containing the center W , the corresponding local support has Euclidean radius comparable to

$$R_W := C_2 \max\{\|W\|_{\mathbf{F}}, R/2^n\}.$$

Throughout this lemma we work on this same local support. Put $\varepsilon := c_{\text{sub}}\eta/\beta$ and write $\Delta_l := U_l - W_l$. It suffices to show that any

$$U \in B_{\varrho_{G_{\text{NP}}^S(W, \varepsilon)}}(W, \sqrt{n}\varepsilon)$$

inside the same Euclidean shell used in the proof of Theorem 4 satisfies

$$\varrho_{(S, S'), \ell}(f(U, \cdot), f(W, \cdot)) \leq \eta.$$

Hence, for every U in the relevant local support,

$$\sum_{l=1}^L \|\Delta_l\|_{\mathbf{F}}^2 = \|U - W\|_{\mathbf{F}}^2 \leq CR_W^2. \quad (\text{E.10})$$

By the non-perturbative expansion in Lemma 1, applied to S and to S' on the same local-Lipschitz event,

$$\begin{aligned} \|F_L(U, X_S) - F_L(W, X_S)\|_{\mathbf{F}}^2 &\leq \sum_{l=1}^L L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \|\Delta_l F_{l-1}^S(W)\|_{\mathbf{F}}^2, \\ \|F_L(U, X_{S'}) - F_L(W, X_{S'})\|_{\mathbf{F}}^2 &\leq \sum_{l=1}^L L\bar{M}_{l \rightarrow L}^2(W, \varepsilon) \|\Delta_l F_{l-1}^{S'}(W)\|_{\mathbf{F}}^2. \end{aligned}$$

Fix the summand l , set $j = l - 1$, and write

$$P = P_j^S(W, \varepsilon), \quad Q = Q_j^S(W, \varepsilon) = I - P.$$

Since P, Q are spectral projectors of

$$\Gamma_j^S(W) = F_j^S(W)F_j^S(W)^\top,$$

the observed covariance has no P - Q cross term:

$$P\Gamma_j^S(W)Q = 0, \quad Q\Gamma_j^S(W)P = 0.$$

Therefore,

$$\Gamma_j^S(W) = P\Gamma_j^S(W)P + Q\Gamma_j^S(W)Q,$$

and hence

$$\|\Delta_l F_j^S(W)\|_{\mathbf{F}}^2 = \text{Tr}(\Delta_l \Gamma_j^S(W) \Delta_l^\top) = \text{Tr}(\Delta_l P \Gamma_j^S(W) P \Delta_l^\top) + \text{Tr}(\Delta_l Q \Gamma_j^S(W) Q \Delta_l^\top).$$

For the ghost term, P, Q are not spectral projectors of $\Gamma_j^{S'}(W)$, so cross terms may appear. Since $\Gamma_j^{S'}(W) \succeq 0$, for any row vector a ,

$$a \Gamma_j^{S'}(W) a^\top \leq 2(aP) \Gamma_j^{S'}(W) (aP)^\top + 2(aQ) \Gamma_j^{S'}(W) (aQ)^\top.$$

Applying this row-wise to the rows of Δ_l and summing gives

$$\text{Tr}(\Delta_l \Gamma_j^{S'}(W) \Delta_l^\top) \leq 2 \text{Tr}(\Delta_l P \Gamma_j^{S'}(W) P \Delta_l^\top) + 2 \text{Tr}(\Delta_l Q \Gamma_j^{S'}(W) Q \Delta_l^\top).$$

By Definition 8, on the active subspace,

$$P \Gamma_j^{S'}(W) P \preceq \kappa P \Gamma_j^S(W) P,$$

while on the inactive subspace,

$$Q \Gamma_j^{S'}(W) Q \preceq b_{\text{sub}} \vartheta_j(W, \varepsilon) Q.$$

Consequently,

$$\begin{aligned} \|\Delta_l F_j^{S'}(W)\|_{\mathbf{F}}^2 &= \text{Tr}(\Delta_l \Gamma_j^{S'}(W) \Delta_l^\top) \leq 2\kappa \text{Tr}(\Delta_l P \Gamma_j^S(W) P \Delta_l^\top) + 2b_{\text{sub}} \vartheta_j(W, \varepsilon) \|\Delta_l Q\|_{\mathbf{F}}^2 \\ &\leq 2\kappa \text{Tr}(\Delta_l P \Gamma_j^S(W) P \Delta_l^\top) + 2b_{\text{sub}} \vartheta_j(W, \varepsilon) \|\Delta_l\|_{\mathbf{F}}^2. \end{aligned}$$

Moreover, by the definition of Q as the inactive spectral projector of $\Gamma_j^S(W)$,

$$Q \Gamma_j^S(W) Q \preceq \vartheta_j(W, \varepsilon) Q.$$

Combining the observed decomposition with the preceding ghost estimate gives

$$\|\Delta_l F_j^S(W)\|_{\mathbf{F}}^2 + \|\Delta_l F_j^{S'}(W)\|_{\mathbf{F}}^2 \leq C_\kappa \|\Delta_l F_j^S(W)\|_{\mathbf{F}}^2 + C_b \vartheta_j(W, \varepsilon) \|\Delta_l\|_{\mathbf{F}}^2.$$

Substituting this estimate into the two non-perturbative expansions and summing over l , we obtain

$$\begin{aligned} &\|F_L(U, X_S) - F_L(W, X_S)\|_{\mathbf{F}}^2 + \|F_L(U, X_{S'}) - F_L(W, X_{S'})\|_{\mathbf{F}}^2 \\ &\leq C_\kappa \sum_{l=1}^L L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) \|\Delta_l F_{l-1}^S(W)\|_{\mathbf{F}}^2 + C_b \sum_{l=1}^L L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) \vartheta_{l-1}(W, \varepsilon) \|\Delta_l\|_{\mathbf{F}}^2. \end{aligned}$$

The first sum is exactly the observed empirical ellipsoidal energy,

$$\sum_{l=1}^L L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) \|\Delta_l F_{l-1}^S(W)\|_{\mathbf{F}}^2 = \varrho_{G_{\text{NP}}^S(W, \varepsilon)}(U, W)^2.$$

For the second sum, using the definition of the finite resolution in (4.4),

$$\vartheta_{l-1}(W, \varepsilon) = \frac{n\varepsilon^2}{2L \bar{M}_{l \rightarrow L}^2(W, \varepsilon) R_W^2},$$

we have

$$L\bar{M}_{l \rightarrow L}^2(W, \varepsilon)\vartheta_{l-1}(W, \varepsilon) = \frac{n\varepsilon^2}{2R_W^2}.$$

Therefore, by (E.10),

$$\sum_{l=1}^L L\bar{M}_{l \rightarrow L}^2(W, \varepsilon)\vartheta_{l-1}(W, \varepsilon)\|\Delta_l\|_{\mathbf{F}}^2 \leq Cn\varepsilon^2.$$

Thus

$$\begin{aligned} & \|F_L(U, X_S) - F_L(W, X_S)\|_{\mathbf{F}}^2 + \|F_L(U, X_{S'}) - F_L(W, X_{S'})\|_{\mathbf{F}}^2 \\ & \leq C_{\kappa, b} \varrho_{G_{\text{NP}}^S(W, \varepsilon)}(U, W)^2 + C_{\kappa, b} n\varepsilon^2. \end{aligned}$$

If U belongs to the empirical ellipsoidal ball on the left-hand side of (E.9), then

$$\varrho_{G_{\text{NP}}^S(W, \varepsilon)}(U, W) \leq \sqrt{n\varepsilon}.$$

Hence the right-hand side above is bounded by $C_{\kappa, b} n\varepsilon^2$. Since the loss is β -Lipschitz,

$$\begin{aligned} \varrho_{(S, S'), \ell}(f(U, \cdot), f(W, \cdot))^2 & \leq \frac{\beta^2}{n} \left(\|F_L(U, X_S) - F_L(W, X_S)\|_{\mathbf{F}}^2 + \|F_L(U, X_{S'}) - F_L(W, X_{S'})\|_{\mathbf{F}}^2 \right) \\ & \leq C_{\kappa, b} \beta^2 \varepsilon^2. \end{aligned}$$

Since $\varepsilon = c_{\text{sub}}\eta/\beta$, this becomes

$$\varrho_{(S, S'), \ell}(f(U, \cdot), f(W, \cdot))^2 \leq C_{\kappa, b} c_{\text{sub}}^2 \eta^2.$$

Choosing $c_{\text{sub}} > 0$ sufficiently small, depending only on (κ, b_{sub}) , gives

$$\varrho_{(S, S'), \ell}(f(U, \cdot), f(W, \cdot)) \leq \eta.$$

This proves (E.9). □

Lemma 32 (Empirical pointwise dimension under subspace isomorphism) *Assume that the event in Definition 8 holds for the pair (S, S') . Then the hierarchical prior π from Theorem 4 satisfies, uniformly over $W \in B_{\mathbf{F}}(R)$ and $\eta > 0$,*

$$\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f(W, \cdot), \eta))} \leq d_{\text{R}}^S\left(W, \frac{c_{\text{sub}}\eta}{\beta}\right). \quad (\text{E.11})$$

Proof of Lemma 32. Lemma 31 gives the finite-scale ball inclusion required to compare pointwise dimensions. By monotonicity of pointwise dimension under metric domination, Lemma 18,

$$\log \frac{1}{\pi(B_{\varrho_{(S, S'), \ell}}(f(W, \cdot), \eta))} \leq \log \frac{1}{\pi(B_{\varrho_{G_{\text{NP}}^S(W, c_{\text{sub}}\eta/\beta)}}(W, c_{\text{sub}}\sqrt{n}\eta/\beta))}.$$

The tensor G_{NP}^S has exactly the same NN-surrogate block-diagonal form treated in Theorem 4, with the feature Gram equal to the observed matrix $\Gamma_{l-1}^S(W)$. Thus the hierarchical covering theorem applies directly to this observed empirical metric tensor. Applying Theorem 4 with $\varepsilon = c_{\text{sub}}\eta/\beta$ gives

$$\log \frac{1}{\pi(B_{\varrho_{G_{\text{NP}}^S(W, c_{\text{sub}}\eta/\beta)}}(W, c_{\text{sub}}\sqrt{n}\eta/\beta))} \leq d_{\text{R}}^S\left(W, \frac{c_{\text{sub}}\eta}{\beta}\right).$$

Combining this with the preceding display proves (E.11). □

Proof of Lemma 30. Let $\Psi_{S,S'}(W)$ denote the generic-chaining functional inside Lemma 8,

$$\Psi_{S,S'}(W) := \inf_{0 \leq \alpha \leq \sqrt{2}} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{\sqrt{2}} \sqrt{\log \frac{1}{\pi(B_{\varrho(S,S'),\ell}(f(W,\cdot),\eta))}} d\eta \right\}.$$

Since the pointwise dimension is nonincreasing in the radius η , a change of variables $\eta = \sqrt{2}t$, followed by taking the infimum over α , allows us to replace the upper endpoint $\sqrt{2}$ by 1 at the cost of an absolute constant:

$$\Psi_{S,S'}(W) \leq C \inf_{0 \leq \alpha \leq 1} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^1 \sqrt{\log \frac{1}{\pi(B_{\varrho(S,S'),\ell}(f(W,\cdot),\eta))}} d\eta \right\},$$

where the constant is enlarged harmlessly. Conditionally on an observed sample S satisfying Definition 8, let $\mathcal{E}_{\text{iso}}(S, S')$ denote the event over the ghost sample S' on which the two estimates in Definition 8 hold. By definition,

$$\mathbb{P}_{S'}(\mathcal{E}_{\text{iso}}(S, S') \mid S) \geq 1 - \zeta.$$

On $\mathcal{E}_{\text{iso}}(S, S')$, Lemma 32 and the change of variables $\eta = \beta\varepsilon$ imply

$$\Psi_{S,S'}(W) \leq C\beta \inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_R^S(W, c_{\text{sub}}\varepsilon)} d\varepsilon \right\}.$$

On the complement $\mathcal{E}_{\text{iso}}(S, S')^c$, the trivial choice $\alpha = \sqrt{2}$ in the original definition of $\Psi_{S,S'}(W)$ gives $\Psi_{S,S'}(W) \leq \sqrt{2}$ for every W . Hence, conditionally on such an S ,

$$\begin{aligned} \mathbb{E}_{S'}[\Psi_{S,S'}(W) \mid S] &= \mathbb{E}_{S'}[\Psi_{S,S'}(W)\mathbf{1}_{\mathcal{E}_{\text{iso}}} \mid S] + \mathbb{E}_{S'}[\Psi_{S,S'}(W)\mathbf{1}_{\mathcal{E}_{\text{iso}}^c} \mid S] \\ &\leq \mathbb{E}_{S'}[\Psi_{S,S'}(W)\mathbf{1}_{\mathcal{E}_{\text{iso}}} \mid S] + \sqrt{2}\mathbb{P}_{S'}(\mathcal{E}_{\text{iso}}^c \mid S) \\ &\leq C\beta \inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_R^S(W, c_{\text{sub}}\varepsilon)} d\varepsilon \right\} + C\zeta, \end{aligned} \quad (\text{E.12})$$

where the last inequality absorbs $\sqrt{2}$ into the absolute constant C . Substituting (E.12) into Lemma 8 proves (E.8). The probability is at least $1 - \delta - \delta_{\text{iso}}$ by a union bound between the generic-chaining event and the event that Definition 8 holds for S . $\square$

E.5 How feature regularity can imply the subspace isomorphism

The two-sample verification below is retained for the conditional ghost result. The same relative covariance certificate restricted to $T = S$ also implies the population condition in Definition 3: rearrange it as $(1 - \gamma)\Sigma_j \preceq \widehat{\Sigma}_{j,S} + \omega_j I$, then restrict to the observed active and inactive subspaces. This gives constants $\kappa, b_{\text{sub}} \leq (1 + c_0)/(1 - \gamma)$ at the stated resolution.

Definition 8 is a finite-resolution, center-uniform condition on the learned feature class. It is weaker than full Loewner domination because it only asks for multiplicative transfer on the empirical active subspace and below-resolution leakage on its orthogonal complement. Let

$$\widehat{\Sigma}_{j,T}(W) := \frac{1}{n} F_j^T(W) F_j^T(W)^\top, \quad \Sigma_j(W) := \mathbb{P}[f_j(W, x) f_j(W, x)^\top], \quad \bar{\vartheta}_j(W, \varepsilon) := \frac{\vartheta_j(W, \varepsilon)}{n}.$$

A convenient verification certificate for Definition 8 is the following regularized relative covariance estimate: for some $\gamma \in (0, 1/4)$ and some $\omega_j(W, \varepsilon)$ satisfying $\omega_j(W, \varepsilon) \leq c_0 \bar{\vartheta}_j(W, \varepsilon)$ with a sufficiently small absolute constant c_0 , both samples $T \in \{S, S'\}$ obey, uniformly over all j, W, ε and all $v \in \mathbb{R}^{d_j}$,

$$\left| v^\top (\widehat{\Sigma}_{j,T}(W) - \Sigma_j(W)) v \right| \leq \gamma v^\top \Sigma_j(W) v + \omega_j(W, \varepsilon) \|v\|_2^2. \quad (\text{E.13})$$

Indeed, (E.13) implies Definition 8. Applying (E.13) with $T = S'$ gives

$$v^\top \widehat{\Sigma}_{j,S'}(W) v \leq (1 + \gamma) v^\top \Sigma_j(W) v + \omega_j(W, \varepsilon) \|v\|_2^2.$$

On the other hand, applying (E.13) with $T = S$ gives

$$v^\top \widehat{\Sigma}_{j,S}(W) v \geq (1 - \gamma) v^\top \Sigma_j(W) v - \omega_j(W, \varepsilon) \|v\|_2^2.$$

Combining the two inequalities yields

$$v^\top \widehat{\Sigma}_{j,S'}(W) v \leq \frac{1 + \gamma}{1 - \gamma} v^\top \widehat{\Sigma}_{j,S}(W) v + \left(1 + \frac{1 + \gamma}{1 - \gamma}\right) \omega_j(W, \varepsilon) \|v\|_2^2. \quad (\text{E.14})$$

For v in the active empirical subspace, $v^\top \widehat{\Sigma}_{j,S}(W) v \geq \bar{\vartheta}_j(W, \varepsilon) \|v\|_2^2$. Since $\omega_j(W, \varepsilon) \leq c_0 \bar{\vartheta}_j(W, \varepsilon)$, the preceding inequality (E.14) gives

$$v^\top \widehat{\Sigma}_{j,S'}(W) v \leq \kappa v^\top \widehat{\Sigma}_{j,S}(W) v, \quad \kappa := \frac{1 + \gamma}{1 - \gamma} + \left(1 + \frac{1 + \gamma}{1 - \gamma}\right) c_0.$$

For v in the inactive empirical subspace, $v^\top \widehat{\Sigma}_{j,S}(W) v \leq \bar{\vartheta}_j(W, \varepsilon) \|v\|_2^2$, so the same inequality (E.14) gives

$$v^\top \widehat{\Sigma}_{j,S'}(W) v \leq b_{\text{sub}} \bar{\vartheta}_j(W, \varepsilon) \|v\|_2^2, \quad b_{\text{sub}} := \frac{1 + \gamma}{1 - \gamma} + \left(1 + \frac{1 + \gamma}{1 - \gamma}\right) c_0.$$

Thus Definition 8 holds with constants that may be taken as

$$\kappa = b_{\text{sub}} := \frac{1 + \gamma}{1 - \gamma} + \left(1 + \frac{1 + \gamma}{1 - \gamma}\right) c_0.$$

Since $\gamma < 1/4$ and c_0 is an absolute constant chosen sufficiently small, $\kappa = b_{\text{sub}} = O(1)$. Since

$$\Gamma_j^T(W) = n \widehat{\Sigma}_{j,T}(W), \quad \vartheta_j(W, \varepsilon) = n \bar{\vartheta}_j(W, \varepsilon),$$

multiplying by n gives (E.6)–(E.7).

Thus, sub-Gaussian or small-ball assumptions at the subspace level (see, e.g., Mendelson [2015, 2021]) should be used to prove the relative certificate (E.13), rather than a fixed additive-ridge theorem. The relevant empirical process is the scalar quadratic class

$$\mathcal{Q}_j = \left\{ x \mapsto \langle v, f_j(W, x) \rangle^2 : W \in B_{\mathbf{F}}(R), v \in S^{d_j-1} \right\}. \quad (\text{E.15})$$

Consequently, the directional part of the verification scales with the vector-sphere dimension d_j , rather than with the matrix dimension d_j^2 ; the remaining cost is the complexity of the learned feature class as W varies. This vector-size saving is crucial for proving the isomorphism event, while the theorem above keeps the final generalization bound pointwise and compressed, since no below-resolution leakage is inserted into $d_{\mathbf{R}}^S$. Concretely, the open question in Section E.4 can be reframed in the following special case:

Can (E.13) be proved for the two-layer quadratic class (E.15) under sub-Gaussian or small-ball feature regularity assumptions at the subspace level?

We hope the discussion above provides a concrete route toward resolving this open question.

E.6 Proof for Regularized ERM in Section 4.4

Lemma 33 (Excess Risk Bound for Regularized ERM) *Assume we have high-probability point-wise generalization bound in the form of (2.1), and the loss $\ell(f; z)$ is uniformly bounded by $[0, 1]$. Then for the regularized ERM*

$$\hat{f} = \operatorname{argmin}_f \left\{ \mathbb{P}_n \ell(f; z) + C \sqrt{\frac{d(f) + \log(2/\delta)}{n}} \right\},$$

we have the excess risk bound against the population risk minimizer $f^ := \arg \min_{\mathcal{F}} \mathbb{P} \ell(f; z)$: with probability at least $1 - \delta$,*

$$\begin{aligned} \mathbb{P} \ell(\hat{f}; z) - \mathbb{P} \ell(f^*; z) &\leq \inf_{f \in \mathcal{F}} \left\{ \mathbb{P}_n \ell(f; z) + C \sqrt{\frac{d(f) + \log(2/\delta)}{n}} \right\} - \mathbb{P} \ell(f^*; z) \\ &\leq (C + \sqrt{1/2}) \sqrt{\frac{d(f^*) + \log(2/\delta)}{n}}. \end{aligned}$$

Proof of Lemma 33: By (2.1), for every $\delta \in (0, 1)$, take $\delta_1 = \delta_2 = \delta/2$, we have that with probability at least $1 - \delta_1 - \delta_2 = 1 - \delta$, we have

$$\begin{aligned} \mathbb{P} \ell(\hat{f}; z) &\leq \inf_{f \in \mathcal{F}} \left\{ \mathbb{P}_n \ell(f; z) + C \sqrt{\frac{d(f) + \log(1/\delta_1)}{n}} \right\} \\ &\leq \mathbb{P}_n \ell(f^*; z) + C \sqrt{\frac{d(f^*) + \log(1/\delta_1)}{n}} \\ &\leq \mathbb{P} \ell(f^*; z) + \sqrt{\frac{\log(1/\delta_2)}{2n}} + C \sqrt{\frac{d(f^*) + \log(1/\delta_1)}{n}} \\ &= \mathbb{P} \ell(f^*; z) + \sqrt{\frac{\log(2/\delta)}{2n}} + C \sqrt{\frac{d(f^*) + \log(2/\delta)}{n}} \\ &\leq \mathbb{P} \ell(f^*; z) + (C + \sqrt{1/2}) \sqrt{\frac{d(f^*) + \log(2/\delta)}{n}}, \end{aligned}$$

where the first inequality uses the bound of the form (2.1); the second inequality uses definition of $\hat{f}$; and the third inequality is an application of the Hoeffding's inequality (Lemma 17) at f^* ; the equality is by $\delta_1 = \delta_2 = \delta/2$; and the last inequality follows from the monotonicity of the square root function. Thus we have that the excess risk is bounded by

$$\begin{aligned} \mathbb{P} \ell(\hat{f}; z) - \mathbb{P} \ell(f^*; z) &\leq \inf_{f \in \mathcal{F}} \left\{ \mathbb{P}_n \ell(f; z) + C \sqrt{\frac{d(f) + \log(2/\delta)}{n}} \right\} - \mathbb{P} \ell(f^*; z) \\ &\leq (C + \sqrt{1/2}) \sqrt{\frac{d(f^*) + \log(2/\delta)}{n}}. \end{aligned}$$

□

E.7 Improvement over Norm Bounds in Section 4.3

E.7.1 Exponential Improvement to a Norm Bound and Comparison

This subsection records a norm-bound relaxation that follows directly from the mixed empirical-population bound in Theorem 5. It does not use the subspace-isomorphism condition in Definition 3, which is needed only for replacing the population feature contribution by observed feature information.

Recall that

$$\Gamma_l^{S,\mathbb{P}}(W) = F_l^S(W)F_l^S(W)^\top + n\mathbb{E}[f_l(W, X)f_l(W, X)^\top].$$

Invoking the elementary bound $\log x \leq \log(1+x) \leq x$ for $x > 0$, the effective-dimension part in Theorem 5 satisfies, for every layer l ,

$$\begin{aligned} & \sum_{k=1}^{r_{\text{eff}}^{S,\mathbb{P}}[W,l]} \log \left(\frac{8C_2^2 \lambda_k(\Gamma_{l-1}^{S,\mathbb{P}}(W)) R_n(W)^2 L \bar{M}_{l \rightarrow L}^2(W, \varepsilon)}{n\varepsilon^2} \right) \\ & \leq \frac{8C_2^2 \text{Tr}(\Gamma_{l-1}^{S,\mathbb{P}}(W)) R_n(W)^2 L \bar{M}_{l \rightarrow L}^2(W, \varepsilon)}{n\varepsilon^2}, \end{aligned} \quad (\text{E.16})$$

where $R_n(W) := \max\{\|W\|_{\mathbf{F}}, R/2^n\}$. Here

$$\text{Tr}(\Gamma_{l-1}^{S,\mathbb{P}}(W)) = \|F_{l-1}^S(W)\|_{\mathbf{F}}^2 + n\mathbb{E}\|f_{l-1}(W, X)\|_2^2.$$

This observation gives the following corollary.

Corollary 2 (Mixed empirical-population rank-free and spectral-norm relaxation) *Assume the hypotheses of Theorem 5. Define $\bar{M}_l(W) := \sup_{0 < \varepsilon \leq \max\{1, c\}/\beta} \bar{M}_{l \rightarrow L}(W, \varepsilon)$, where c is the radius constant in Theorem 5, and let*

$$\alpha_n := \min\{n^{-5}, (2\beta)^{-1}\}, \quad \Lambda_n := 1 + \log \frac{1}{\alpha_n \beta}.$$

There exists an absolute constant $C > 0$ such that, with probability at least $1 - \delta$ over S , uniformly over $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C \left(\frac{\beta \Lambda_n R_n(W)}{n} \sqrt{\sum_{l=1}^L L(d_l + d_{l-1}) \text{Tr}(\Gamma_{l-1}^{S,\mathbb{P}}(W)) \bar{M}_l(W)^2} \right. \\ & \quad \left. + \beta \alpha_n + \sqrt{\frac{\sum_{l=1}^L \log(d_{l-1}n) + \log \frac{2 \log(2n)}{\delta}}{n}} \right). \end{aligned} \quad (\text{E.17})$$

The bound is rank-free and contains no r_{eff} or d_{eff} .

If, in addition, the activations are 1-Lipschitz and satisfy $\sigma_l(0) = 0$, then

$$\text{Tr}(\Gamma_{l-1}^{S,\mathbb{P}}(W))^{1/2} \leq \left(\prod_{i < l} \|W_i\|_{\text{op}} \right) (\|X_S\|_{\mathbf{F}}^2 + n\mathbb{E}\|X\|_2^2)^{1/2}. \quad (\text{E.18})$$

If the common outer transfer bounds on the relevant weight charts also satisfy

$$\overline{M}_l(W) \leq C_0 \prod_{i>l} \|W_i\|_{\text{op}}$$

for an absolute constant C_0 , then

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq \tilde{O}\left(\frac{\beta R_n(W)}{n} (\|X_S\|_{\mathbf{F}}^2 + n\mathbb{E}\|X\|_2^2)^{1/2} \sqrt{\sum_{l=1}^L L(d_l + d_{l-1}) \prod_{i \neq l} \|W_i\|_{\text{op}}^2}\right. \\ & \quad \left. + \beta\alpha_n + \sqrt{\frac{\sum_{l=1}^L \log(d_{l-1}n) + \log \frac{2\log(2n)}{\delta}}{n}}\right), \end{aligned} \quad (\text{E.19})$$

where $\tilde{O}$ hides absolute constants and Λ_n . The outer-chart condition is additional: the spectral product at the center alone need not dominate transfer constants throughout a chart.

If the inputs are almost surely bounded, $\|X\|_2 \leq B_x$, then

$$(\|X_S\|_{\mathbf{F}}^2 + n\mathbb{E}\|X\|_2^2)^{1/2} \leq (\|X_S\|_{\mathbf{F}}^2 + nB_x^2)^{1/2} \leq \sqrt{2n} B_x. \quad (\text{E.20})$$

Hence (E.19) becomes the familiar $n^{-1/2}$ -order spectral-norm bound

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq \tilde{O}\left(\frac{\beta B_x R_n(W)}{\sqrt{n}} \sqrt{\sum_{l=1}^L L(d_l + d_{l-1}) \prod_{i \neq l} \|W_i\|_{\text{op}}^2}\right. \\ & \quad \left. + \beta\alpha_n + \sqrt{\frac{\sum_{l=1}^L \log(d_{l-1}n) + \log \frac{2\log(2n)}{\delta}}{n}}\right). \end{aligned} \quad (\text{E.21})$$

Discussion of Corollary 2: We proceed in three paragraphs of discussion. First, we show that the Riemannian Dimension bound in Theorem 5 can be *exponentially* tighter than the spectrally normalized bound in (E.19). Second, we offer a metric-tensor interpretation that clarifies the source of this improvement. Finally, we position (E.19) relative to the most representative spectrally normalized bounds (SNB) in the existing literature. Besides, the additional normalization $\sigma_l(0) = 0$ is used only for the spectral-norm relaxation (E.18). It is standard in spectrally normalized norm-bound analyses; for example, Bartlett et al. [2017].

I: Why the improvement can be exponential. The Riemannian-Dimension bound retains logarithmic ellipsoidal volumes, whereas $\log(1+x) \leq x$ replaces them by linear trace bounds. For a concrete comparison, take a scalar one-layer network with identity activation, $W = R = \beta = 1$, deterministic input $X = e^t$, and $n = \lceil e^{4t} \rceil$, for $t \geq 1$. The loss may be any bounded 1-Lipschitz loss, and the outer transfer constant is exactly one. Then $\Gamma_0^{S, \mathbb{P}} = 2ne^{2t}$, and

$$d_{\mathbf{R}}^{S, \mathbb{P}}(W, c\varepsilon) = O\left(t + \log \frac{1}{\varepsilon}\right), \quad 0 < \varepsilon \leq 1.$$

At fixed confidence, choosing $\alpha = 0$ in Theorem 5 gives $O(\sqrt{t}e^{-2t})$, whereas the leading spectral term in (E.19) is $\tilde{\Theta}(e^{-t})$. Thus the displayed bounds can differ exponentially in t , even when both vanish. This compares the bounds, not the actual generalization gap. Feature compression can give an additional improvement: replacing $\text{Tr}(\Gamma_{l-1}^{S,\mathbb{P}})^{1/2}$ by the right-hand side of (E.18) discards the learned feature spectrum. Both improvements are already available in Theorem 5, without the subspace-isomorphism condition.

II: Metric tensor interpretation. Under the outer-chart condition in Corollary 2, the spectral-norm relaxation replaces the mixed ellipsoidal metric tensor $G_{\text{NP}}^{S,\mathbb{P}}(W)$, up to an absolute factor, by

$$G_{\text{SNB}}^{S,\mathbb{P}}(W) = \text{blockdiag} \left(\dots, L(\|X_S\|_{\mathbf{F}}^2 + n\mathbb{E}\|X\|_2^2) \prod_{i \neq l} \|W_i\|_{\text{op}}^2 I_{d_l d_{l-1}}, \dots \right).$$

This tensor is isotropic inside each layer and forgets the spectrum of the learned feature Gram. The Riemannian-Dimension bound keeps the anisotropic tensor $L\bar{M}_{l \rightarrow L}^2 \Gamma_{l-1}^{S,\mathbb{P}}(W) \otimes I_{d_l}$, which is why it adapts to low rank and spectral decay in the learned representations.

III: Relation to existing spectrally normalized bounds. The bound in (E.19) is structurally close to the classical SNB results of Bartlett et al. [2017] and Neyshabur et al. [2018]; the three bounds differ only in the *global ball* used to constraint the hypothesis class.

- (a) Our bound (E.19) controls *all* layers simultaneously via the global Frobenius norm $\|W\|_{\mathbf{F}}$, hence the factor $R_n(W)$ in the numerator, which differs from $\|W\|_{\mathbf{F}}$ only below the discretization scale $R/2^n$.
- (b) Neyshabur et al. [2018] bounds each layer l separately by its Frobenius norm $\|W_l\|_{\mathbf{F}}$. Strengthening their argument with Dudley’s entropy integral (one-shot optimization in the original paper) gives

$$(\mathbb{P} - \mathbb{P}_n) \ell(f(W, x), y) \leq \tilde{O} \left(\frac{\beta \|X\|_{\mathbf{F}} \sqrt{\sum_{l=1}^L L^2 (d_l + d_{l-1}) \|W_l\|_{\mathbf{F}}^2 \prod_{i \neq l} \|W_i\|_{\text{op}}^2}}{n} + \sqrt{\frac{\log \frac{1}{\delta}}{n}} \right). \quad (\text{E.22})$$

Neither (E.19) nor (E.22) strictly dominates the other, since factors of the form $(\sum_l a_l)(\sum_l b_l)$ in (E.19) *vs.* factors of the form $L \sum_l a_l b_l$ in (E.22) can swap their relative order.

- (c) Bartlett et al. [2017] replaces each Frobenius norm by the $\|\cdot\|_{2,1}$ norm, obtaining the tighter

$$(\mathbb{P} - \mathbb{P}_n) \ell(f(W, x), y) \leq \tilde{O} \left(\frac{\beta \|X\|_{\mathbf{F}} (\sum_l \|W_l\|_{2,1}^{2/3} \sum_l (\prod_{i \neq l} \|W_i\|_{\text{op}})^{2/3})^{3/2}}{n} + \sqrt{\frac{\log \frac{1}{\delta}}{n}} \right), \quad (\text{E.23})$$

which improves on (E.19) and (E.22) thanks to the sharper 2,1 norm. Extending our Riemannian-dimension analysis to the 2,1 norm setting is an interesting direction for future work.

- (d) Size-independent SNB bounds (pioneered by Golowich et al. [2020]) remove all depth/width dependence at the price of a worse scaling in n ; incorporating their technique is left for future research.

- (e) [Pinto et al. \[2025\]](#) impose explicit per-layer rank constraints on the weight matrices, thereby replacing the width factors in (E.22) with the corresponding ranks while leaving the product of spectral norms unchanged. Their bound includes an additional C^L factor, which is subsequently removed by [Ledent et al. \[2025\]](#). Moreover, [Ledent et al. \[2025\]](#) seek to bridge the spectral-norm and parameter-count regimes by leveraging the Schatten- p framework of [Golowich et al. \[2020\]](#), which interpolates between the product-of-spectral-norm regime ($p \rightarrow \infty$) and layerwise low-rank scalings ($p \rightarrow 0$). In the extreme $p \rightarrow 0$ limit, a representative consequence (Theorem E.8 of [Ledent et al., 2025](#)) yields

$$(\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \leq \tilde{O} \left(\frac{\sup_i \|x_i\|_2^2}{\sqrt{n}} \sqrt{\sum_{l=1}^L L(d_l + d_{l-1}) \text{rank}(W_l)} \right).$$

Notably, the explicit dependence on the *ranks of the weight matrices*—rather than on spectrum-aware or *feature-rank* quantities—renders this result structurally similar to VC-dimension bounds (indeed, the proof proceeds via uniform covering numbers, and packing/VC dimensions for matrices are known to adapt to explicit rank constraints [[Srebro et al., 2004](#)]). As the authors acknowledge, this is a principal limitation: empirical evidence suggests that deep networks exhibit low rank in their *features* rather than their weights, a phenomenon this bound does not capture.

In any case, (E.19) is a representative SNB bound, and the key message in this subsection is that our Riemannian-Dimension result in Theorem 5 can be *exponentially* sharper than (E.19).

E.7.2 Proof of Corollary 2

We work directly with the empirical-population metric throughout. Theorem 5 gives, with probability at least $1 - \delta$, uniformly over $W \in B_{\mathbf{F}}(R)$,

$$\begin{aligned} & (\mathbb{P} - \mathbb{P}_n)\ell(f(W, x), y) \\ & \leq C_1 \left(\beta \inf_{0 \leq \alpha \leq 1/\beta} \left\{ \alpha + \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\mathbf{R}}^{S, \mathbb{P}}(W, c\varepsilon)} d\varepsilon \right\} + \sqrt{\frac{\log \frac{2 \log(2n)}{\delta}}{n}} \right). \end{aligned} \quad (\text{E.24})$$

The supremum defining $\overline{M}_l$ covers all scales $c\varepsilon$ used here; the fixed factor c is absorbed into the absolute constants below.

Step 1: The trace relaxation. Summing (E.16) over layers gives

$$d_{\mathbf{R}}^{S, \mathbb{P}}(W, c\varepsilon) \leq \frac{CR_n(W)^2}{n\varepsilon^2} \sum_{l=1}^L L(d_l + d_{l-1}) \text{Tr}(\Gamma_{l-1}^{S, \mathbb{P}}(W)) \overline{M}_l(W)^2 + \sum_{l=1}^L \log(d_{l-1}n). \quad (\text{E.25})$$

Here we used the identity

$$\sum_k \lambda_k(\Gamma_{l-1}^{S, \mathbb{P}}(W)) = \text{Tr}(\Gamma_{l-1}^{S, \mathbb{P}}(W)) = \|F_{l-1}^S(W)\|_{\mathbf{F}}^2 + n\mathbb{E}\|f_{l-1}(W, X)\|_2^2. \quad (\text{E.26})$$

Thus, for $0 < \alpha \leq 1/\beta$,

$$\begin{aligned} & \frac{1}{\sqrt{n}} \int_{\alpha}^{1/\beta} \sqrt{d_{\mathbf{R}}^{S, \mathbb{P}}(W, c\varepsilon)} d\varepsilon \\ & \leq \frac{CR_n(W)}{n} \sqrt{\sum_{l=1}^L L(d_l + d_{l-1}) \text{Tr}(\Gamma_{l-1}^{S, \mathbb{P}}(W)) \overline{M}_l(W)^2 \log \frac{1}{\alpha\beta} + \frac{1}{\beta} \sqrt{\frac{\sum_{l=1}^L \log(d_{l-1}n)}{n}}}. \end{aligned} \quad (\text{E.27})$$

Choose $\alpha = \alpha_n$ in (E.24) and use $\log(1/(\alpha_n\beta)) \leq \Lambda_n$. This proves (E.17). The truncation contributes only

$$\beta\alpha_n = \min\{\beta n^{-5}, 1/2\}. \quad (\text{E.28})$$

In particular, there is no need to integrate the trace bound down to zero.

Step 2: The spectral-norm relaxation. For every input x , the activation assumptions give

$$\|f_{l-1}(W, x)\|_2 \leq \left(\prod_{i < l} \|W_i\|_{\text{op}} \right) \|x\|_2.$$

Squaring and integrating against $n(\mathbb{P}_S + \mathbb{P})$ proves (E.18). Under the stated common outer-chart bound, substituting this inequality and $\overline{M}_l(W) \leq C_0 \prod_{i > l} \|W_i\|_{\text{op}}$ into (E.17) proves (E.19). The bound is simultaneous in W on the same event. Replacing $R_n(W)$ by $\|W\|_{\mathbf{F}}$, as in the main-text display, incurs at most the additional term

$$\frac{C\beta\Lambda_n R}{n2^n} (\|X_S\|_{\mathbf{F}}^2 + n\mathbb{E}\|X\|_2^2)^{1/2} \sqrt{\sum_{l=1}^L L(d_l + d_{l-1}) \prod_{i \neq l} \|W_i\|_{\text{op}}^2}. \quad (\text{E.29})$$

Finally, (E.20) yields (E.21). No subspace-isomorphism assumption or ghost-sample averaging is used. $\square$

References

- Pierre Alquier. User-friendly introduction to pac-bayes bounds. *Foundations and Trends® in Machine Learning*, 17(2):174–303, 2024.
- Rayna Andreeva, Benjamin Dupuis, Rik Sarkar, Tolga Birdal, and Umut Simsekli. Topological generalization bounds for discrete-time stochastic optimization algorithms. *Advances in Neural Information Processing Systems*, 37:4765–4818, 2024.
- Alessio Ansuini, Alessandro Laio, Jakob H Macke, and Davide Zoccolan. Intrinsic dimension of data representations in deep neural networks. *Advances in Neural Information Processing Systems*, 32, 2019.
- Sanjeev Arora, Rong Ge, Behnam Neyshabur, and Yi Zhang. Stronger generalization bounds for deep nets via a compression approach. In *International conference on machine learning*, pages 254–263. PMLR, 2018.
- Sanjeev Arora, Simon S Du, Wei Hu, Zhiyuan Li, Russ R Salakhutdinov, and Ruosong Wang. On exact computation with an infinitely wide neural net. *Advances in neural information processing systems*, 32, 2019.
- Amir Asadi, Emmanuel Abbe, and Sergio Verdú. Chaining mutual information and tightening generalization bounds. In *Advances in Neural Information Processing Systems 31*, pages 7245–7254, 2018.
- Alexander Atanasov, Blake Bordelon, and Cengiz Pehlevan. Neural networks as kernel learners: The silent alignment effect. In *International Conference on Learning Representations*, 2021.
- Jean-Yves Audibert and Olivier Bousquet. Combining pac-bayesian and generic chaining bounds. *Journal of Machine Learning Research*, 8(4), 2007.
- Jean-Yves Audibert and Olivier Bousquet. Pac-bayesian generic chaining. *Advances in neural information processing systems*, 16, 2003.
- Steve Awodey. *Category theory*, volume 52. OUP Oxford, 2010.
- Laura Balzano, Tianjiao Ding, Benjamin D Haeffele, Soo Min Kwon, Qing Qu, Peng Wang, Zhangyang Wang, and Can Yaras. An overview of low-rank structures in the training and adaptation of large models. *arXiv preprint arXiv:2503.19859*, 2025.
- Aristide Baratin, Thomas George, César Laurent, R Devon Hjelm, Guillaume Lajoie, Pascal Vincent, and Simon Lacoste-Julien. Implicit regularization via neural feature alignment. In *International Conference on Artificial Intelligence and Statistics*, pages 2269–2277. PMLR, 2021.
- Peter L Bartlett and Shahr Mendelson. Rademacher and gaussian complexities: Risk bounds and structural results. *Journal of Machine Learning Research*, 3(Nov):463–482, 2002.
- Peter L Bartlett, Dylan J Foster, and Matus J Telgarsky. Spectrally-normalized margin bounds for neural networks. *Advances in neural information processing systems*, 30, 2017.

- Peter L. Bartlett, Nick Harvey, Christopher Liaw, and Abbas Mehrabian. Nearly-tight vc-dimension and pseudodimension bounds for piecewise linear neural networks. *Journal of Machine Learning Research*, 20(63):1–17, 2019.
- Peter L Bartlett, Philip M Long, Gábor Lugosi, and Alexander Tsigler. Benign overfitting in linear regression. *Proceedings of the National Academy of Sciences*, 117(48):30063–30070, 2020.
- Peter L Bartlett, Andrea Montanari, and Alexander Rakhlin. Deep learning: a statistical viewpoint. *Acta numerica*, 30:87–201, 2021.
- Mikhail Belkin, Siyuan Ma, and Soumik Mandal. To understand deep learning we need to understand kernel learning. In *International Conference on Machine Learning*, pages 541–549. PMLR, 2018.
- Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal. Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32):15849–15854, 2019.
- Thomas Bendokat, Ralf Zimmermann, and P-A Absil. A grassmann manifold handbook: Basic geometry and computational aspects. *Advances in Computational Mathematics*, 50(1):6, 2024.
- Yoshua Bengio, Aaron Courville, and Pascal Vincent. Representation learning: A review and new perspectives. *IEEE transactions on pattern analysis and machine intelligence*, 35(8):1798–1828, 2013.
- Tolga Birdal, Aaron Lou, Leonidas J Guibas, and Umut Simsekli. Intrinsic dimension, persistent homology and generalization in neural networks. *Advances in neural information processing systems*, 34:6776–6789, 2021.
- Adam Block, Yuval Dagan, and Alexander Rakhlin. Majorizing measures, sequential complexities, and online learning. In *Conference on Learning Theory*, pages 587–590. PMLR, 2021.
- Xingyu Cai, Jiaji Huang, Yuchen Bian, and Kenneth Church. Isotropy in the contextual embedding space: Clusters and manifolds. In *International conference on learning representations*, 2021.
- Alexander Camuto, George Deligiannidis, Murat A Erdogdu, Mert Gurbuzbalaban, Umut Simsekli, and Lingjiong Zhu. Fractal structure and generalization properties of stochastic optimization algorithms. *Advances in neural information processing systems*, 34:18774–18788, 2021.
- Olivier Catoni. A pac-bayesian approach to adaptive classification. *preprint*, 840(2):6, 2003.
- Fan Chen, Dylan J Foster, Yanjun Han, Jian Qian, Alexander Rakhlin, and Yunbei Xu. Assouad, fano, and le cam with interaction: A unifying lower bound framework and characterization for bandit learnability. *Advances in Neural Information Processing Systems*, 37:75585–75641, 2024.
- Lenaïc Chizat and Francis Bach. On the global convergence of gradient descent for over-parameterized models using optimal transport. *Advances in neural information processing systems*, 31, 2018.
- Yifeng Chu and Maxim Raginsky. Majorizing measures, codes, and information. In *2023 IEEE International Symposium on Information Theory (ISIT)*, pages 660–665. IEEE, 2023.

- Karel Devriendt, Hannah Friedman, Bernhard Reinke, and Bernd Sturmfels. The two lives of the grassmannian. *arXiv preprint arXiv:2401.03684*, 2024.
- Laurent Dinh, Razvan Pascanu, Samy Bengio, and Yoshua Bengio. Sharp minima can generalize for deep nets. In *International Conference on Machine Learning*, pages 1019–1028. PMLR, 2017.
- Benjamin Dupuis, George Deligiannidis, and Umut Simsekli. Generalization bounds using data-dependent fractal dimensions. In *International Conference on Machine Learning*, pages 8922–8968. PMLR, 2023.
- Gintare Karolina Dziugaite and Daniel M Roy. Computing nonvacuous generalization bounds for deep (stochastic) neural networks with many more parameters than training data. *arXiv preprint arXiv:1703.11008*, 2017.
- Mathieu Even and Laurent Massoulié. Concentration of non-isotropic random tensors with applications to learning and empirical risk minimization. In *Conference on Learning Theory*, pages 1847–1886. PMLR, 2021.
- Kenneth J Falconer. *Techniques in fractal geometry*, volume 3. Wiley Chichester, 1997.
- Alexander R Farhang, Jeremy D Bernstein, Kushal Tirumala, Yang Liu, and Yisong Yue. Investigating generalization by controlling normalized margin. In *International Conference on Machine Learning*, pages 6324–6336. PMLR, 2022.
- Ruili Feng, Kecheng Zheng, Yukun Huang, Deli Zhao, Michael Jordan, and Zheng-Jun Zha. Rank diminishing in deep neural networks. *Advances in Neural Information Processing Systems*, 35: 33054–33065, 2022.
- X Fernique. Regularite des trajectoires des fonctions aleatoires gaussiennes. *Ecole d’Eté de Probabilités de Saint-Flour IV*, pages 1–96, 1974.
- Pierre Foret, Ariel Kleiner, Hossein Mobahi, and Behnam Neyshabur. Sharpness-aware minimization for efficiently improving generalization. In *International Conference on Learning Representations*, 2021.
- Stanislav Fort, Gintare Karolina Dziugaite, Mansheej Paul, Sepideh Kharaghani, Daniel M Roy, and Surya Ganguli. Deep learning versus kernel learning: an empirical study of loss landscape geometry and the time evolution of the neural tangent kernel. *Advances in Neural Information Processing Systems*, 33:5850–5861, 2020.
- Behrooz Ghorbani, Shankar Krishnan, and Ying Xiao. An investigation into neural net optimization via hessian eigenvalue density. In *International Conference on Machine Learning*, pages 2232–2241. PMLR, 2019.
- Evarist Giné and Joel Zinn. Some limit theorems for empirical processes. *The Annals of Probability*, pages 929–989, 1984.
- Eugene Golikov, Eduard Pokonechnyy, and Vladimir Korviakov. Neural tangent kernel: A survey. *arXiv preprint arXiv:2208.13614*, 2022.

- Noah Golowich, Alexander Rakhlin, and Ohad Shamir. Size-independent sample complexity of neural networks. *Information and Inference: A Journal of the IMA*, 9(2):473–504, 2020.
- Gene H Golub and Charles F Van Loan. *Matrix computations*. JHU press, 2013.
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*. MIT press Cambridge, 2016.
- Suriya Gunasekar, Jason Lee, Daniel Soudry, and Nathan Srebro. Characterizing implicit bias in terms of optimization geometry. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80, pages 1832–1841. PMLR, 2018.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- Geoffrey Hinton, Nitish Srivastava, and Kevin Swersky. Neural networks for machine learning lecture 6a overview of mini-batch gradient descent. *Cited on*, 14(8):2, 2012.
- Geoffrey E Hinton and Drew Van Camp. Keeping the neural networks simple by minimizing the description length of the weights. In *Proceedings of the sixth annual conference on Computational learning theory*, pages 5–13, 1993.
- Sepp Hochreiter and Jürgen Schmidhuber. Flat minima. *Neural computation*, 9(1):1–42, 1997.
- Minyoung Huh, Hossein Mobahi, Richard Zhang, Brian Cheung, Pulkit Agrawal, and Phillip Isola. The low-rank simplicity bias in deep networks. *Transactions on Machine Learning Research*, 2021.
- Arthur Jacot, Franck Gabriel, and Clément Hongler. Neural tangent kernel: Convergence and generalization in neural networks. *Advances in neural information processing systems*, 31, 2018.
- Stanislaw Jastrzebski, Devansh Arpit, Oliver Astrand, Giancarlo B Kerg, Huan Wang, Caiming Xiong, Richard Socher, Kyunghyun Cho, and Krzysztof J Geras. Catastrophic fisher explosion: Early phase fisher matrix impacts generalization. In *International Conference on Machine Learning*, pages 4772–4784. PMLR, 2021.
- Jürgen Jost. *Riemannian geometry and geometric analysis*, volume 42005. Springer, 2008.
- Ryo Karakida, Shotaro Akaho, and Shun-ichi Amari. Universal statistics of fisher information in deep neural networks: Mean field approach. In *The 22nd International Conference on Artificial Intelligence and Statistics*, pages 1032–1041. PMLR, 2019.
- Nitish Shirish Keskar, Dheevatsa Mudigere, Jorge Nocedal, Mikhail Smelyanskiy, and Ping Tak Peter Tang. On large-batch training for deep learning: Generalization gap and sharp minima. In *International Conference on Learning Representations*, 2017.
- Dmitry Kopitkov and Vadim Indelman. Neural spectrum alignment: Empirical study. In *International Conference on Artificial Neural Networks*, pages 168–179. Springer, 2020.

- Vignesh Kothapalli. Neural collapse: A review on modelling principles and generalization. *Transactions on Machine Learning Research*, 2022.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *nature*, 521(7553):436–444, 2015.
- Antoine Ledent, Rodrigo Alves, and Yunwen Lei. Generalization bounds for rank-sparse neural networks. *arXiv preprint arXiv:2510.21945*, 2025.
- Jaehoon Lee, Yasaman Bahri, Roman Novak, Samuel S. Schoenholz, Jeffrey Pennington, and Jascha Sohl-Dickstein. Deep neural networks as gaussian processes. In *International Conference on Learning Representations*, 2018.
- Yunwen Lei, Rong Jin, and Yiming Ying. Stability and generalization analysis of gradient methods for shallow neural networks. *Advances in neural information processing systems*, 35:38557–38570, 2022.
- Chunyuan Li, Heerad Farkhor, Rosanne Liu, and Jason Yosinski. Measuring the intrinsic dimension of objective landscapes. In *International Conference on Learning Representations*, 2018.
- Tengyuan Liang and Alexander Rakhlin. Just interpolate: Kernel “ridgeless” regression can generalize. *The Annals of Statistics*, 48(3):1329–1347, 2020.
- Jingbo Liu. Simple and sharp generalization bounds via lifting. *arXiv preprint arXiv:2508.18682*, 2025.
- Sanae Lotfi, Marc Finzi, Sanyam Kapoor, Andres Potapczynski, Micah Goldblum, and Andrew G Wilson. Pac-bayes compression bounds so tight that they can explain generalization. *Advances in Neural Information Processing Systems*, 35:31459–31473, 2022.
- Sanae Lotfi, Marc Finzi, Yilun Kuang, Tim GJ Rudner, Micah Goldblum, and Andrew Gordon Wilson. Non-vacuous generalization bounds for large language models. In *International Conference on Machine Learning*, pages 32801–32818. PMLR, 2024.
- Neil Lutz. A note on pointwise dimensions. *arXiv preprint arXiv:1612.05849*, 2016.
- Andreas Maurer. Majorizing codes and measures, 2010. URL <https://www.andreas-maurer.eu/mmnnotes.pdf>. Preprint.
- David A McAllester. Some pac-bayesian theorems. In *Proceedings of the eleventh annual conference on Computational learning theory*, pages 230–234, 1998.
- Colin McDiarmid. Concentration. In *Probabilistic methods for algorithmic discrete mathematics*, pages 195–248. Springer, 1998.

- Song Mei, Andrea Montanari, and Phan-Minh Nguyen. A mean field view of the landscape of two-layer neural networks. *Proceedings of the National Academy of Sciences*, 115(33):E7665–E7671, 2018.
- Shahar Mendelson. Learning without concentration. *Journal of the ACM*, 62(3):1–25, 2015.
- Shahar Mendelson. Extending the scope of the small-ball method. *Studia Mathematica*, 256(2):147–167, 2021.
- Vitali D Milman and Gideon Schechtman. *Asymptotic theory of finite dimensional normed spaces: Isoperimetric inequalities in riemannian manifolds*, volume 1200. Springer Science & Business Media, 1986.
- Theodor Misiakiewicz and Andrea Montanari. Six lectures on linearized neural networks. *arXiv preprint arXiv:2308.13431*, 2023.
- Vaishnavh Nagarajan and J. Zico Kolter. Uniform convergence may be unable to explain generalization in deep learning. *arXiv preprint arXiv:1902.04742*, 2019.
- Gergely Neu and Lorenzo Rosasco. Iterate averaging as regularization for stochastic gradient descent. In *Proceedings of the Conference on Learning Theory*, volume 75, pages 3222–3242. PMLR, 2018.
- Behnam Neyshabur, Russ R Salakhutdinov, and Nathan Srebro. Path-sgd: Path-normalized optimization in deep neural networks. *Advances in neural information processing systems*, 28, 2015a.
- Behnam Neyshabur, Ryota Tomioka, and Nathan Srebro. Norm-based capacity control in neural networks. In *Conference on learning theory*, pages 1376–1401. PMLR, 2015b.
- Behnam Neyshabur, Srinadh Bhojanapalli, David McAllester, and Nathan Srebro. Exploring generalization in deep learning. *Advances in neural information processing systems*, 30, 2017.
- Behnam Neyshabur, Srinadh Bhojanapalli, and Nathan Srebro. A pac-bayesian approach to spectrally-normalized margin bounds for neural networks. In *International Conference on Learning Representations*, 2018.
- nLab contributors. Non-perturbative quantum field theory. <https://ncatlab.org/nlab/show/non-perturbative+quantum+field+theory>, 2025a. nLab (The n -Category Café).
- nLab contributors. Strongly correlated system. <https://ncatlab.org/nlab/show/strongly+correlated+system>, 2025b. nLab (The n -Category Café).
- Roman Novak, Yasaman Bahri, Daniel A Abolafia, Jeffrey Pennington, and Jascha Sohl-Dickstein. Sensitivity and generalization in neural networks: an empirical study. In *International Conference on Learning Representations*, 2018.
- Christopher C Paige and Musheng Wei. History and generality of the cs decomposition. *Linear Algebra and its Applications*, 208:303–326, 1994.
- Alain Pajor. Metric entropy of the grassmann manifold. *Convex Geometric Analysis*, 34(181-188):0942–46013, 1998.

- Vardan Papyan, XY Han, and David L Donoho. Prevalence of neural collapse during the terminal phase of deep learning training. *Proceedings of the National Academy of Sciences*, 117(40):24652–24663, 2020.
- Liam Parker, Emre Onal, Anton Stengel, and Jake Intrater. Neural collapse in the intermediate hidden layers of classification neural networks. *arXiv preprint arXiv:2308.02760*, 2023.
- Niket Patel and Ravid Shwartz-Ziv. Learning to compress: Local rank and information compression in deep neural networks. *arXiv preprint arXiv:2410.07687*, 2024.
- Andrea Pinto, Akshay Rangamani, and Tomaso A. Poggio. On generalization bounds for neural networks with low rank layers. In *Proceedings of The 36th International Conference on Algorithmic Learning Theory*, volume 272, pages 921–936. PMLR, 2025.
- Gilles Pisier. *The volume of convex bodies and Banach space geometry*, volume 94. Cambridge University Press, 1999.
- Adityanarayanan Radhakrishnan, Daniel Beaglehole, Parthe Pandit, and Mikhail Belkin. Mechanism for feature learning in neural networks and backpropagation-free machine learning models. *Science*, 383(6690):1461–1467, 2024.
- Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred Hamprecht, Yoshua Bengio, and Aaron Courville. On the spectral bias of neural networks. In *International conference on machine learning*, pages 5301–5310. PMLR, 2019.
- Noam Razin and Nadav Cohen. Implicit regularization in deep learning may not be explainable by norms. *Advances in neural information processing systems*, 33:21174–21187, 2020.
- Anton Razzhigaev, Matvey Mikhalechuk, Elizaveta Goncharova, Ivan Oseledets, Denis Dimitrov, and Andrey Kuznetsov. The shape of learning: Anisotropy and intrinsic dimensions in transformer-based models. In *Findings of the Association for Computational Linguistics*, pages 868–874, 2024.
- Dominic Richards and Ilja Kuzborskij. Stability & generalisation of gradient descent for shallow neural networks without the neural tangent kernel. *Advances in neural information processing systems*, 34:8609–8621, 2021.
- Daniel Russo and James Zou. Controlling bias in adaptive data analysis using information theory. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics*, volume 51, pages 1232–1240. PMLR, 2016.
- Zhouxing Shi, Yihan Wang, Huan Zhang, J Zico Kolter, and Cho-Jui Hsieh. Efficiently computing local lipschitz constants of neural networks via bound propagation. *Advances in Neural Information Processing Systems*, 35:2350–2364, 2022.
- Ravid Shwartz-Ziv and Naftali Tishby. Opening the black box of deep neural networks via information. *arXiv preprint arXiv:1703.00810*, 2017.
- Leon Simon. Introduction to geometric measure theory. <https://math.stanford.edu/~lms/ntu-gmt-text.pdf>, 2018. NTU Lecture Notes, updated 7 Mar 2018.

- Umut Simsekli, Ozan Sener, George Deligiannidis, and Murat A Erdogdu. Hausdorff dimension, heavy tails, and generalization in neural networks. *Advances in Neural Information Processing Systems*, 33:5138–5151, 2020.
- Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, and Nathan Srebro. The implicit bias of gradient descent on separable data. *Journal of Machine Learning Research*, 19: 1–57, 2018.
- Nathan Srebro, Noga Alon, and Tommi Jaakkola. Generalization error bounds for collaborative prediction with low-rank matrices. *Advances In Neural Information Processing Systems*, 17, 2004.
- Thomas Steinke and Lydia Zakynthinou. Reasoning about generalization via conditional mutual information. In *Conference on Learning Theory*, pages 3437–3452. PMLR, 2020.
- Ilya Sutskever. An observation on generalization. <https://simons.berkeley.edu/talks/ilya-sutskever-openai-2023-08-14>, 2023. Workshop in Large Language Models and Transformers, Simons Institute for the Theory of Computing, recorded 14 Aug 2023.
- Stanislaw J Szarek. Metric entropy of homogeneous spaces. *arXiv preprint math/9701213*, 1997.
- Michel Talagrand. Regularity of gaussian processes. *Acta Mathematica*, 159:99–149, 1987.
- Michel Talagrand. *Upper and Lower Bounds for Stochastic Processes: Modern Methods and Classical Problems*, volume 60 of *Ergebnisse der Mathematik und ihrer Grenzgebiete. 3. Folge / A Series of Modern Surveys in Mathematics*. Springer, Berlin, Heidelberg, 2014.
- Naftali Tishby and Noga Zaslavsky. Deep learning and the information bottleneck principle. In *ieee information theory workshop*, pages 1–5. Ieee, 2015.
- Ramon Van Handel. Probability in high dimension. *Lecture Notes (Princeton University)*, 2(3): 2–3, 2014.
- Vladimir Vapnik. *The nature of statistical learning theory*. Springer science & business media, 2013.
- Gal Vardi. On the implicit bias in deep-learning algorithms. *Communications of the ACM*, 66(6): 86–93, 2023.
- Roman Vershynin. *High-dimensional probability: An introduction with applications in data science*, volume 47. Cambridge university press, 2018.
- Abraham Wald. Statistical decision functions which minimize the maximum risk. *Annals of Mathematics*, 46(2):265–280, 1945.
- Peng Wang, Xiao Li, Can Yaras, Zhihui Zhu, Laura Balzano, Wei Hu, and Qing Qu. Understanding deep representation learning via layerwise feature compression and discrimination. *arXiv preprint arXiv:2311.02960*, 2025.
- Wikipedia contributors. Change of variables. https://en.wikipedia.org/wiki/Change_of_variables, 2025a.

- Wikipedia contributors. Min-max theorem. https://en.wikipedia.org/wiki/Min-max_theorem, 2025b.
- Wikipedia contributors. Minkowski inequality. https://en.wikipedia.org/wiki/Minkowski_inequality, 2025c.
- Wikipedia contributors. Normed vector space. https://en.wikipedia.org/wiki/Normed_vector_space, 2025d.
- Wikipedia contributors. Square root of a matrix. https://https://en.wikipedia.org/wiki/Square_root_of_a_matrix, 2025e.
- Andrew Gordon Wilson. Deep learning is not so mysterious or different. *arXiv preprint arXiv:2503.02113*, 2025.
- David P Woodruff et al. Sketching as a tool for numerical linear algebra. *Foundations and Trends® in Theoretical Computer Science*, 10(1–2):1–157, 2014.
- Blake Woodworth, Suriya Gunasekar, Jason D Lee, Edward Moroshko, Pedro Savarese, Itay Golan, Daniel Soudry, and Nathan Srebro. Kernel and rich regimes in overparametrized models. In *Conference on Learning Theory*, pages 3635–3673. PMLR, 2020.
- Aolin Xu and Maxim Raginsky. Information-theoretic analysis of generalization capability of learning algorithms. In *Advances in Neural Information Processing Systems 30 (NeurIPS 2017)*, pages 2524–2533, 2017.
- Yunbei Xu. Pointwise complexity for Gaussian fields: Upper envelopes, algorithmic lower bounds, and separation. *arXiv preprint arXiv:2606.07931*, 2026.
- Yunbei Xu and Assaf Zeevi. Towards problem-dependent optimal learning rates. *Advances in Neural Information Processing Systems*, 33:2196–2206, 2020.
- Yunbei Xu and Assaf Zeevi. Towards optimal problem dependent generalization error bounds in statistical learning theory. *Mathematics of Operations Research*, 50(1):40–67, 2025.
- Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. Understanding deep learning (still) requires rethinking generalization. *Communications of the ACM*, 64(3):107–115, 2021.